

Image processing

370: Intro to Computer Vision

Subhransu Maji

February 27 & March 4

College of
INFORMATION AND
COMPUTER SCIENCES



Overview of the next two lectures

Digitizing an image

Image processing

- Example: Improving contrast

The screenshot shows the official website for scikit-image. At the top is a navigation bar with links for Download, Gallery, Documentation, Community Guidelines, Source, and a search bar. Below the navigation is a banner for the 'Stable (release notes)' version 0.18.1 from December 2020, featuring a 'Download' button. To the right of the banner is a section titled 'Image processing in Python' which describes the project as a collection of algorithms for image processing, available free of charge and under an open license, with contributions from a community of volunteers.

Convolution and filtering

- Mathematical model
- Implementation details

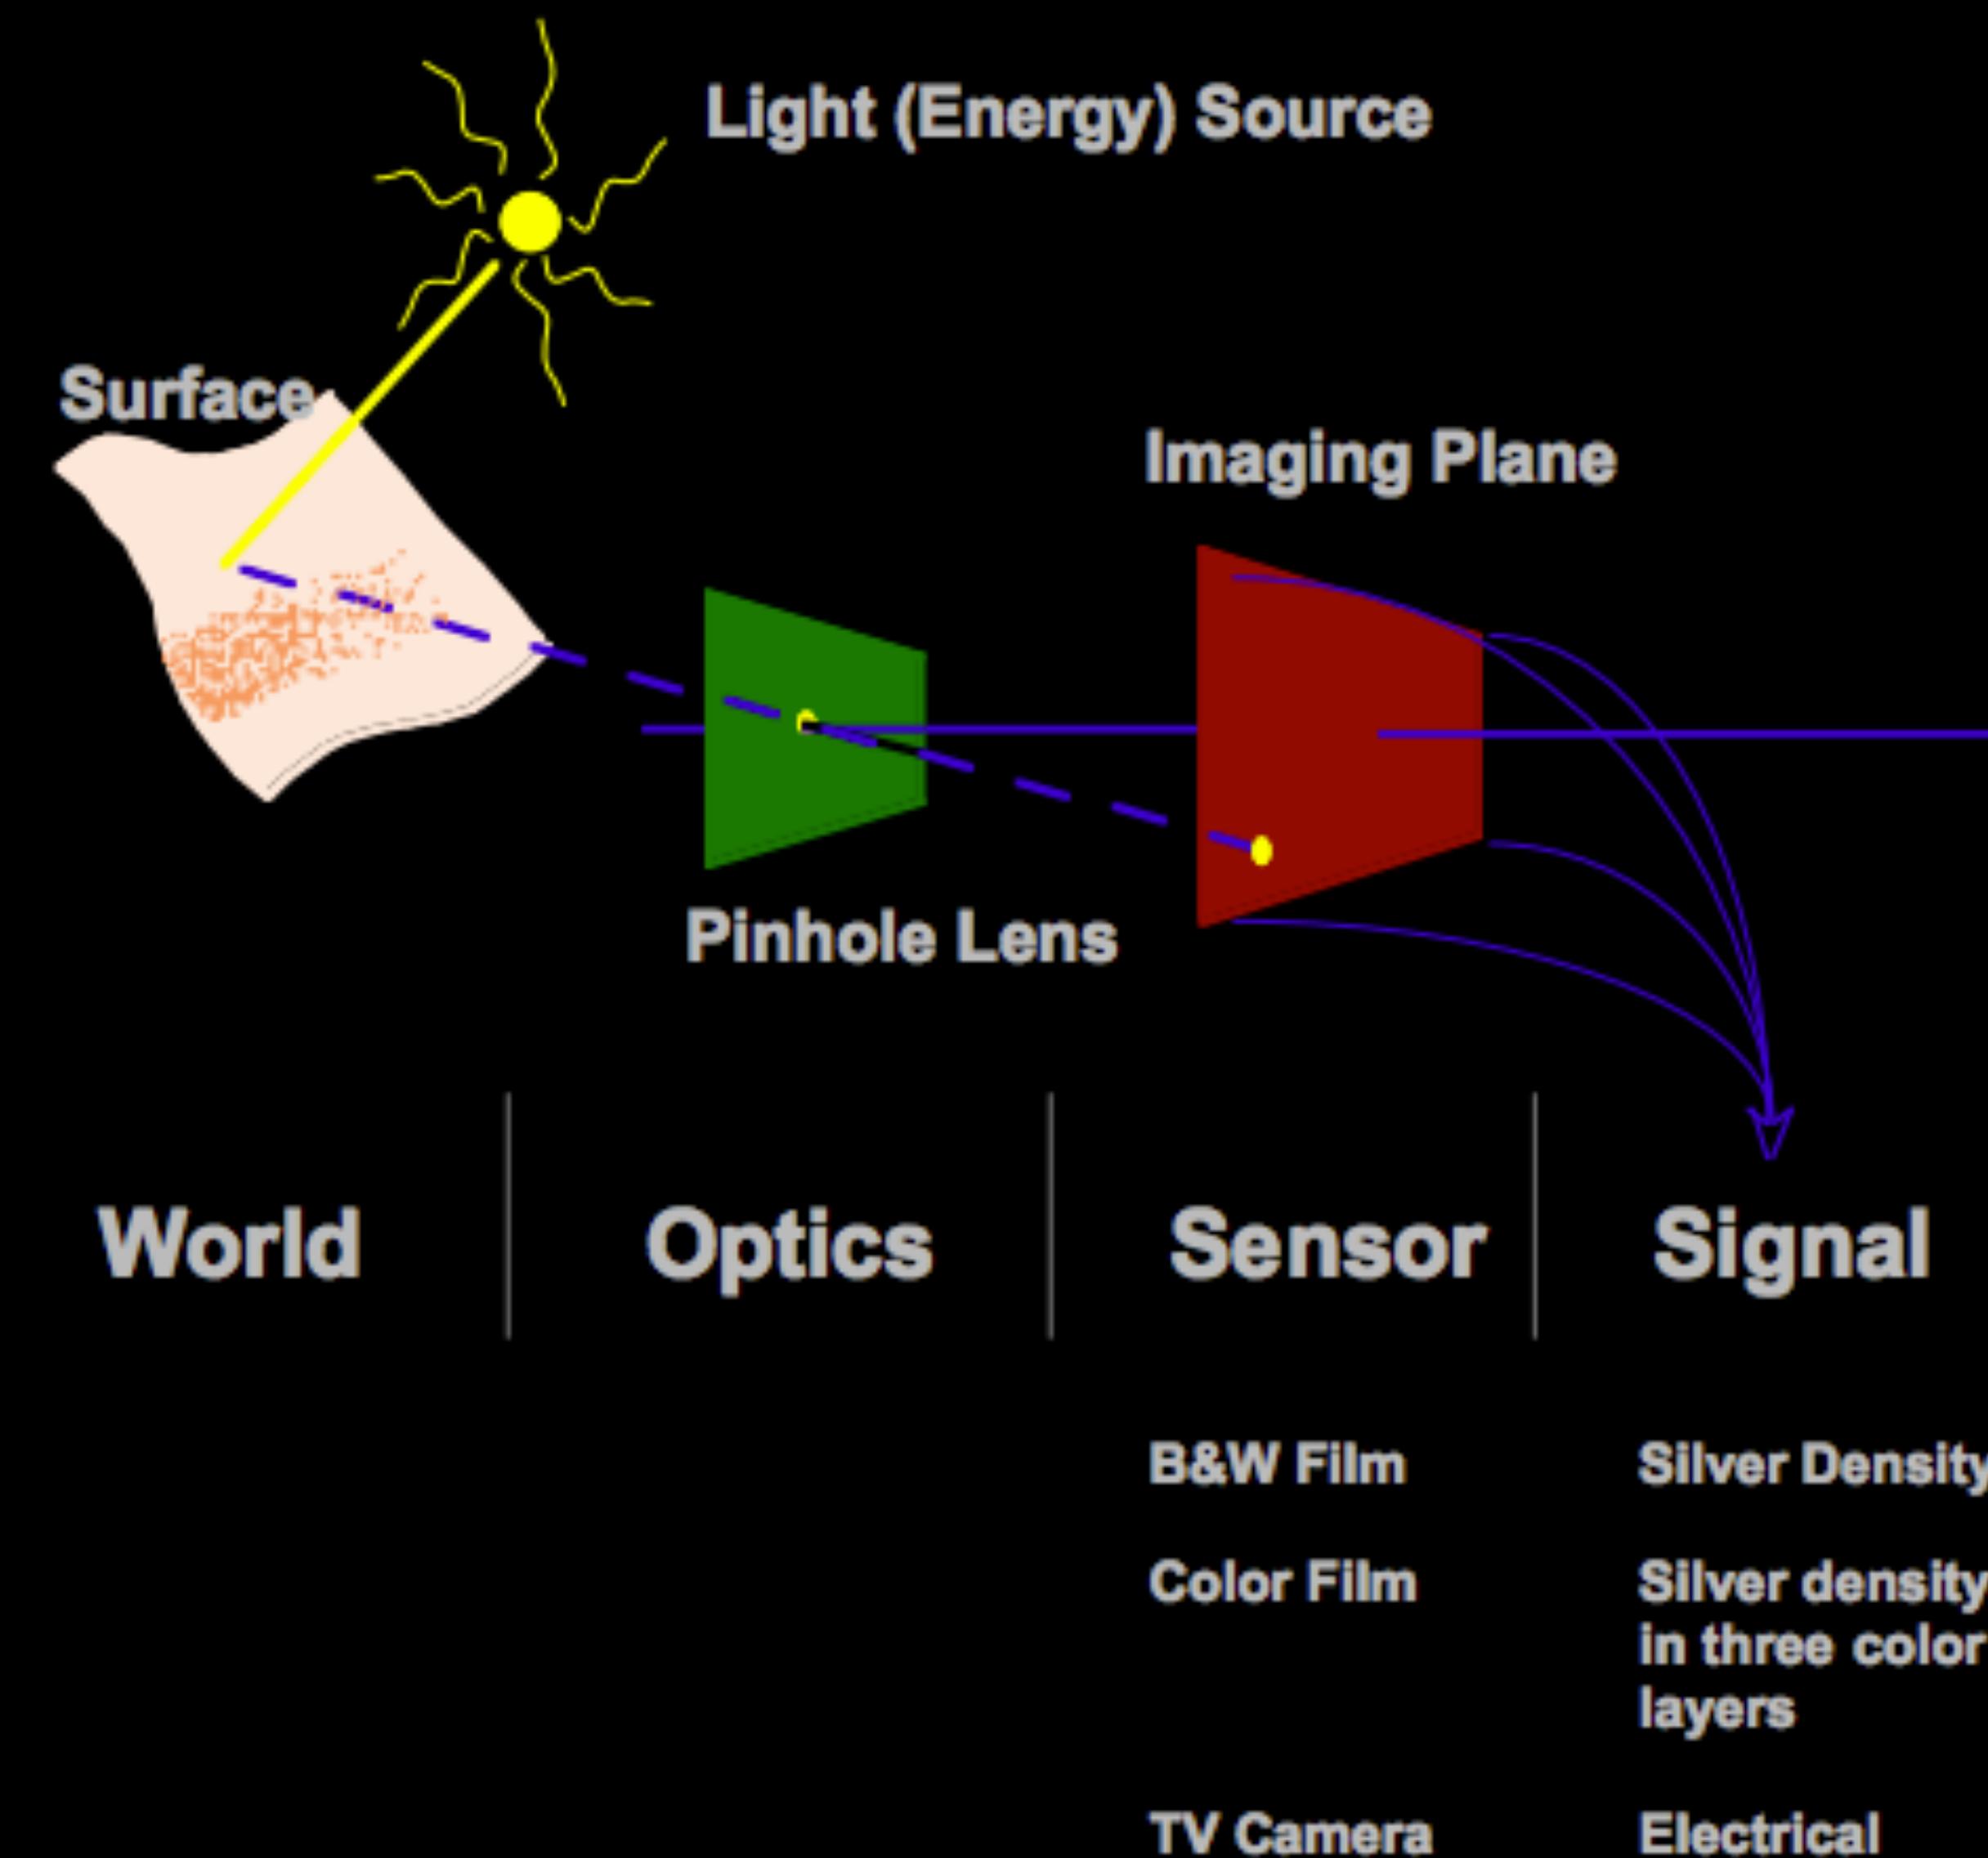
The screenshot shows the PyPI project page for 'opencv-python'. It features a large image of a 3D cube composed of colored blocks. The title 'opencv-python 4.5.1.48' is prominently displayed, along with a 'pip install opencv-python' button. On the right, there's a green 'Latest version' button and a note indicating it was released on Jan 2, 2021.

Applications

- Denoising
- Sharpening
- Edge detection

The screenshot shows the PyTorch homepage. The header includes links for PyTorch, Get Started, Ecosystem, Mobile, Blog, Tutorials, Docs, Resources, GitHub, and a search icon. The main visual is a large purple-to-red gradient background with the text 'FROM RESEARCH TO PRODUCTION' in white. Below this, a subtitle reads 'An open source machine learning framework that accelerates the path from research to production'.

Image formation



Pre-digitization image

What is an image before you digitize it?

- ▶ Continuous range of wavelengths
- ▶ 2-dimensional extent
- ▶ Continuous range of power at each point

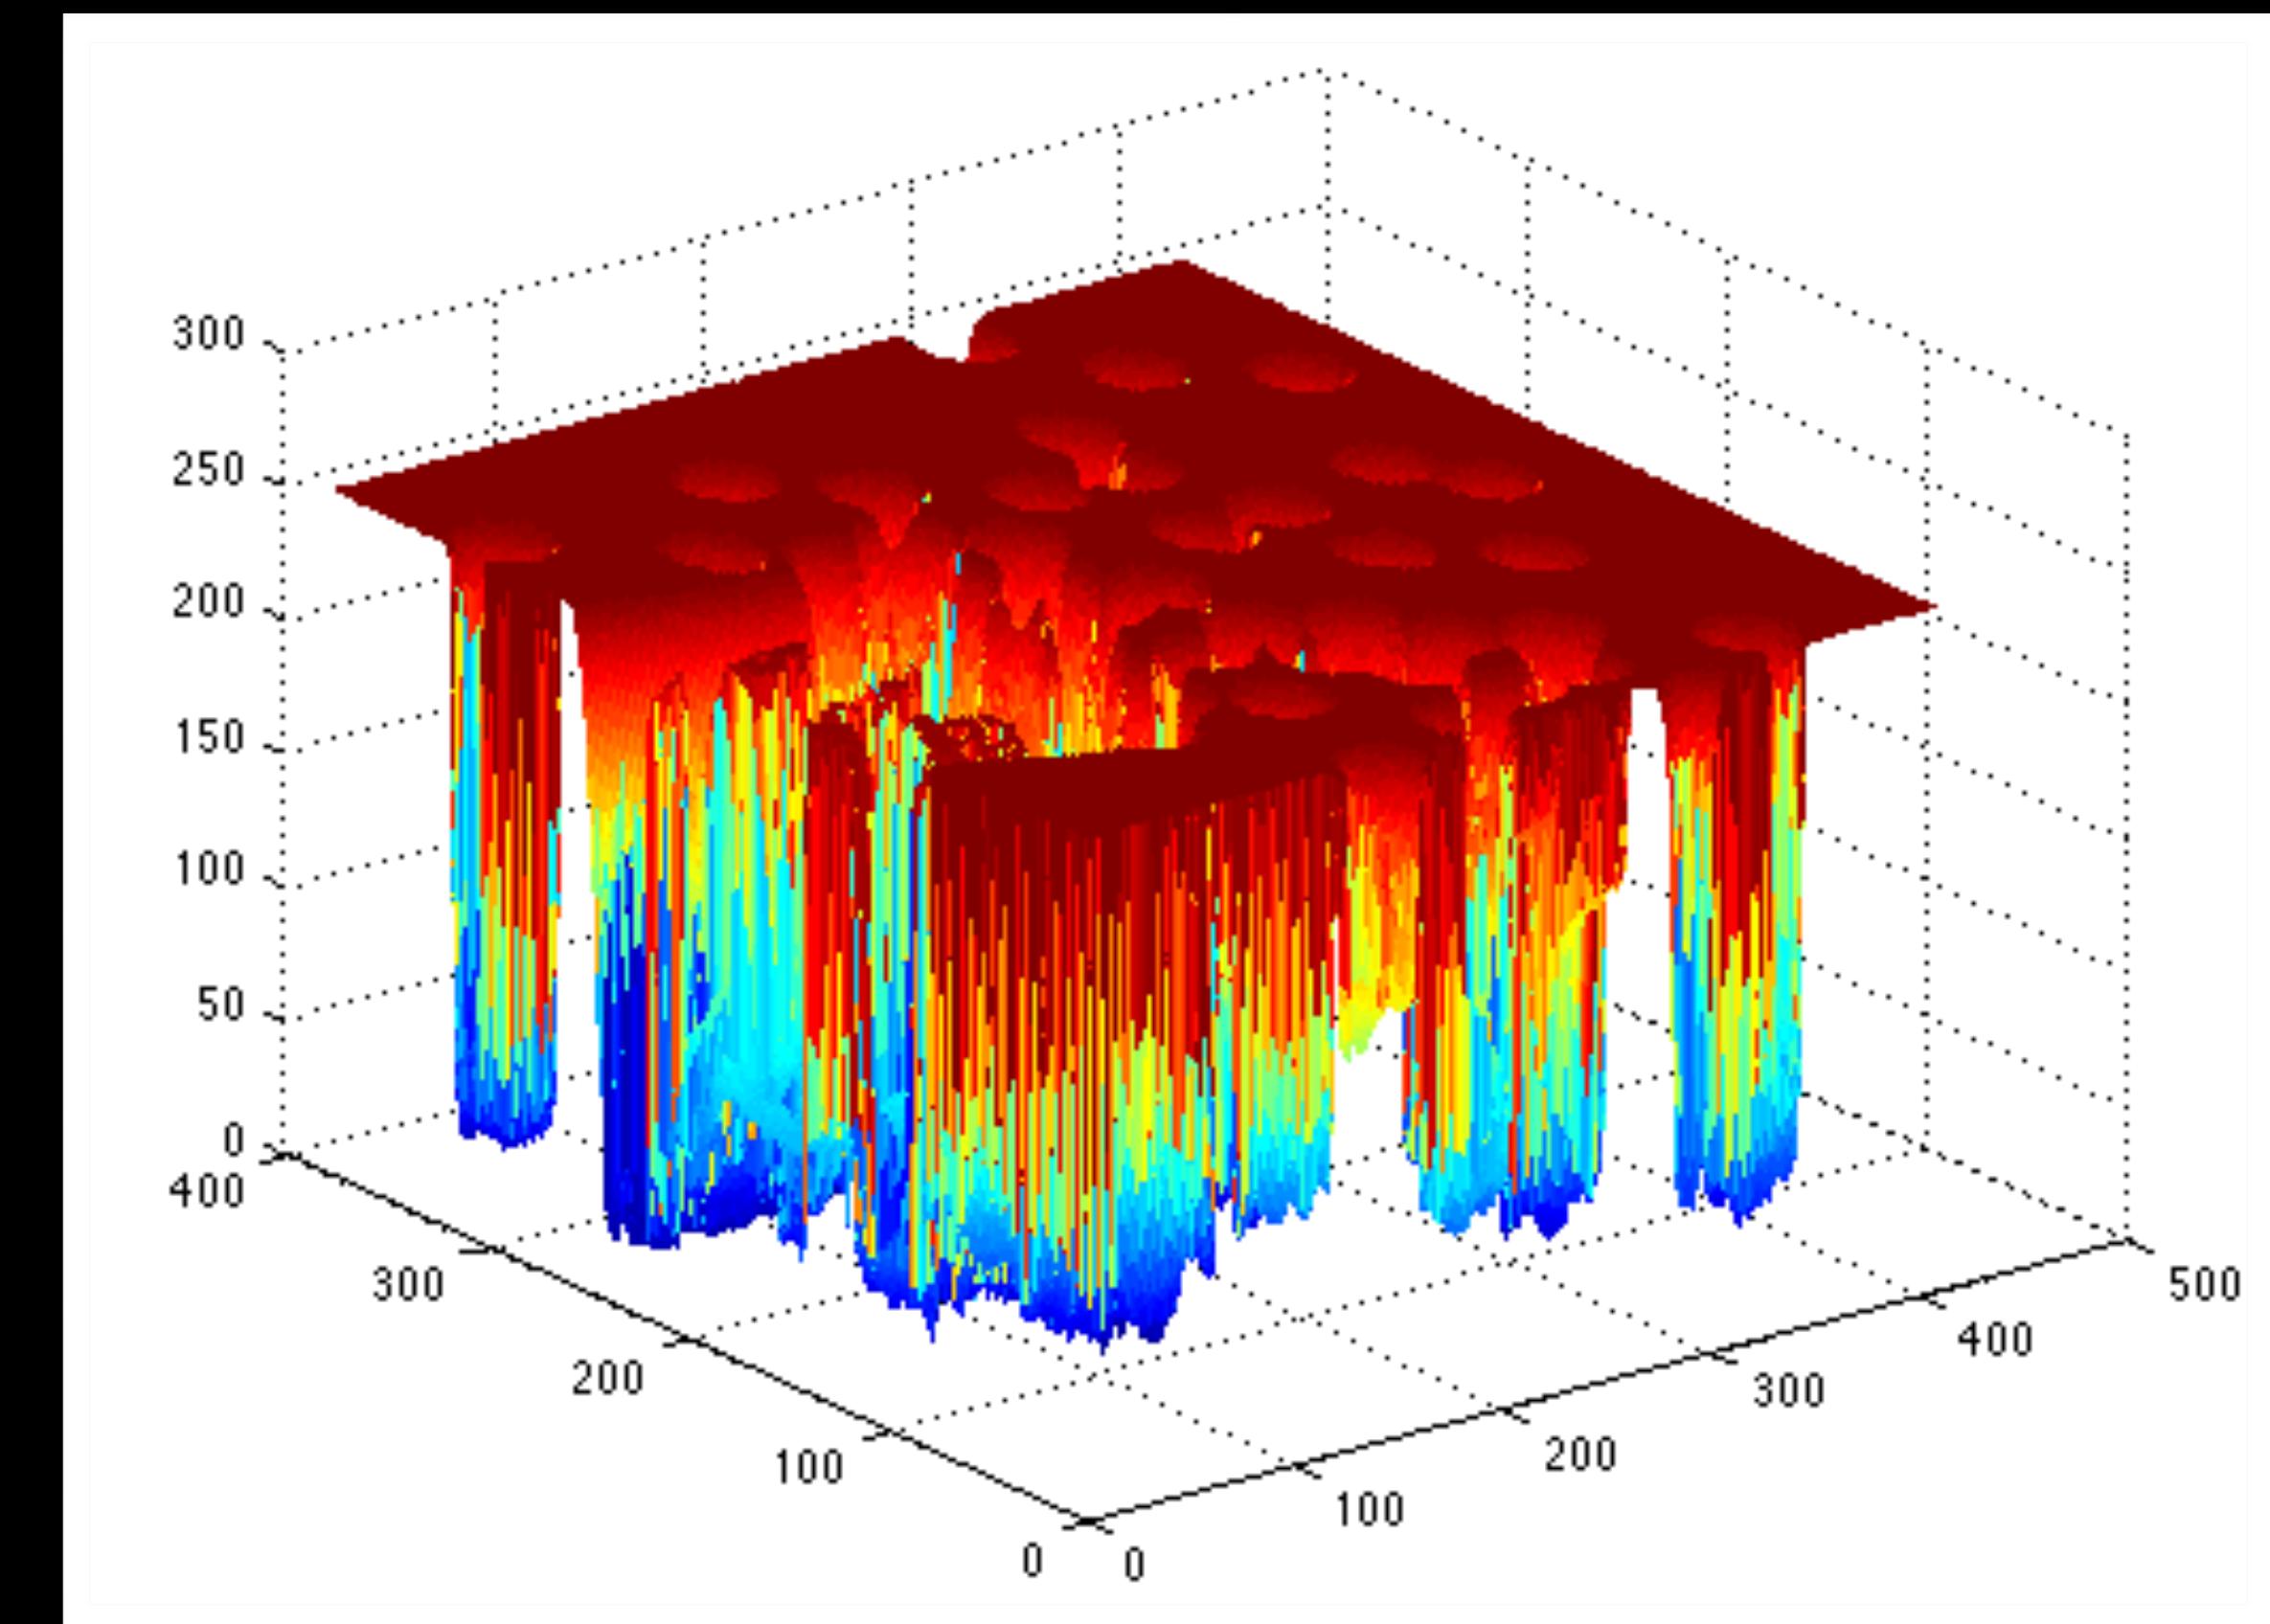
Brightness images

To simplify, consider only a brightness image

- ▶ Two-dimensional (continuous range of locations)
- ▶ Continuous range of brightness values

This is equivalent to a two-dimensional function over a plane

An image as a surface



How do we represent this continuous two dimensional surface efficiently?

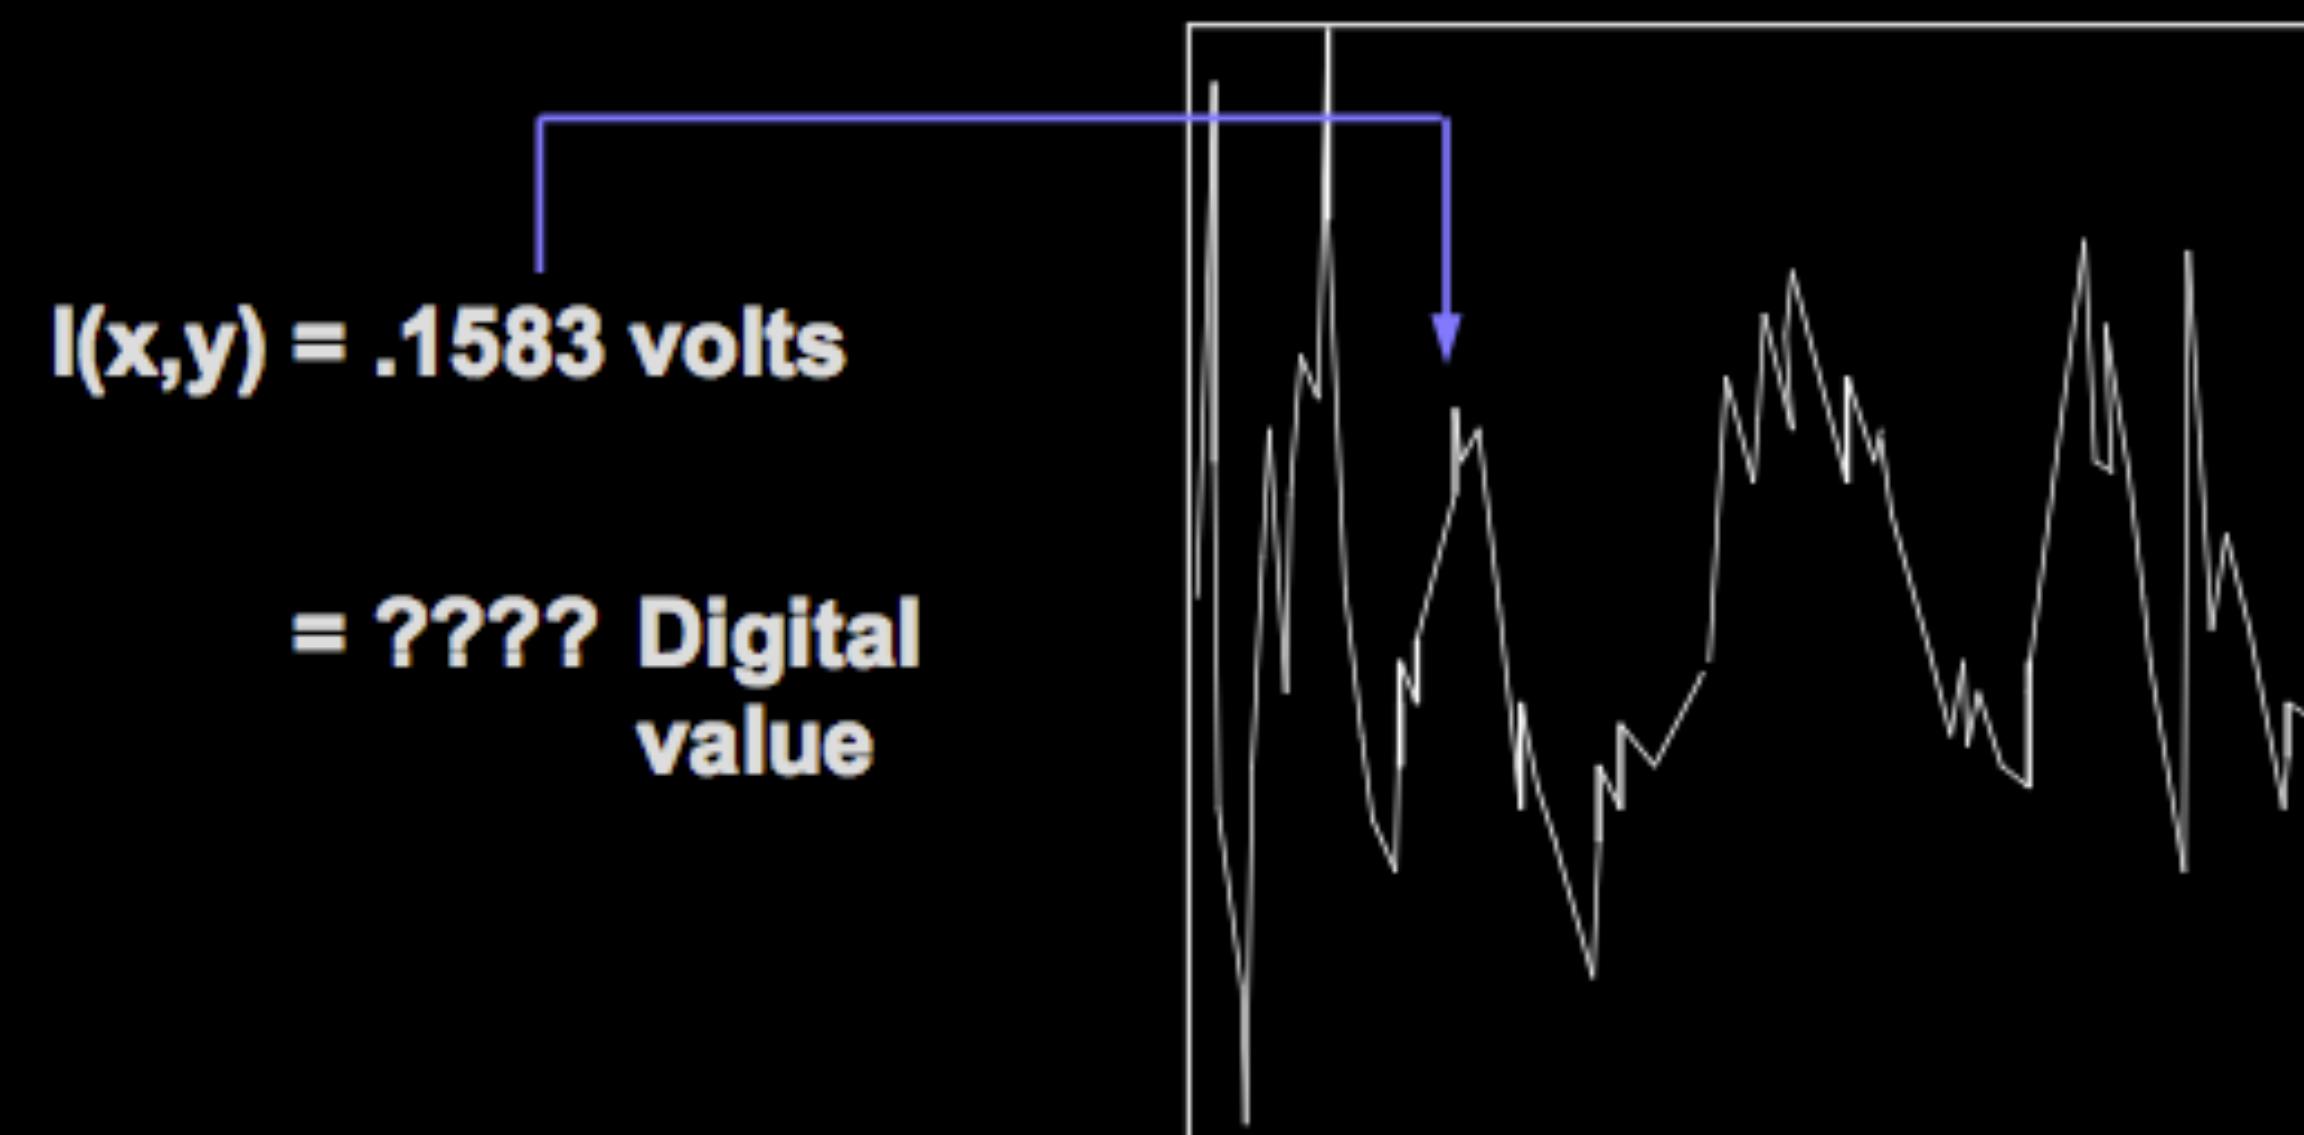
Discretization

Sampling strategies

- ▶ Spatial sampling
 - How many pixels?
 - What arrangement of pixels?
- ▶ Brightness sampling
 - How many brightness values?
 - Spacing of brightness values?
- ▶ For video, also the question of time sampling.

Signal quantization

Goal: determine a mapping from a continuous signal (e.g. analog video signal) to one of K discrete (digital) levels.



Quantization

$I(x,y)$ = continuous signal: $0 \leq I \leq M$

Want to quantize to K values $0, 1, \dots, K-1$

K usually chosen to be a power of 2:

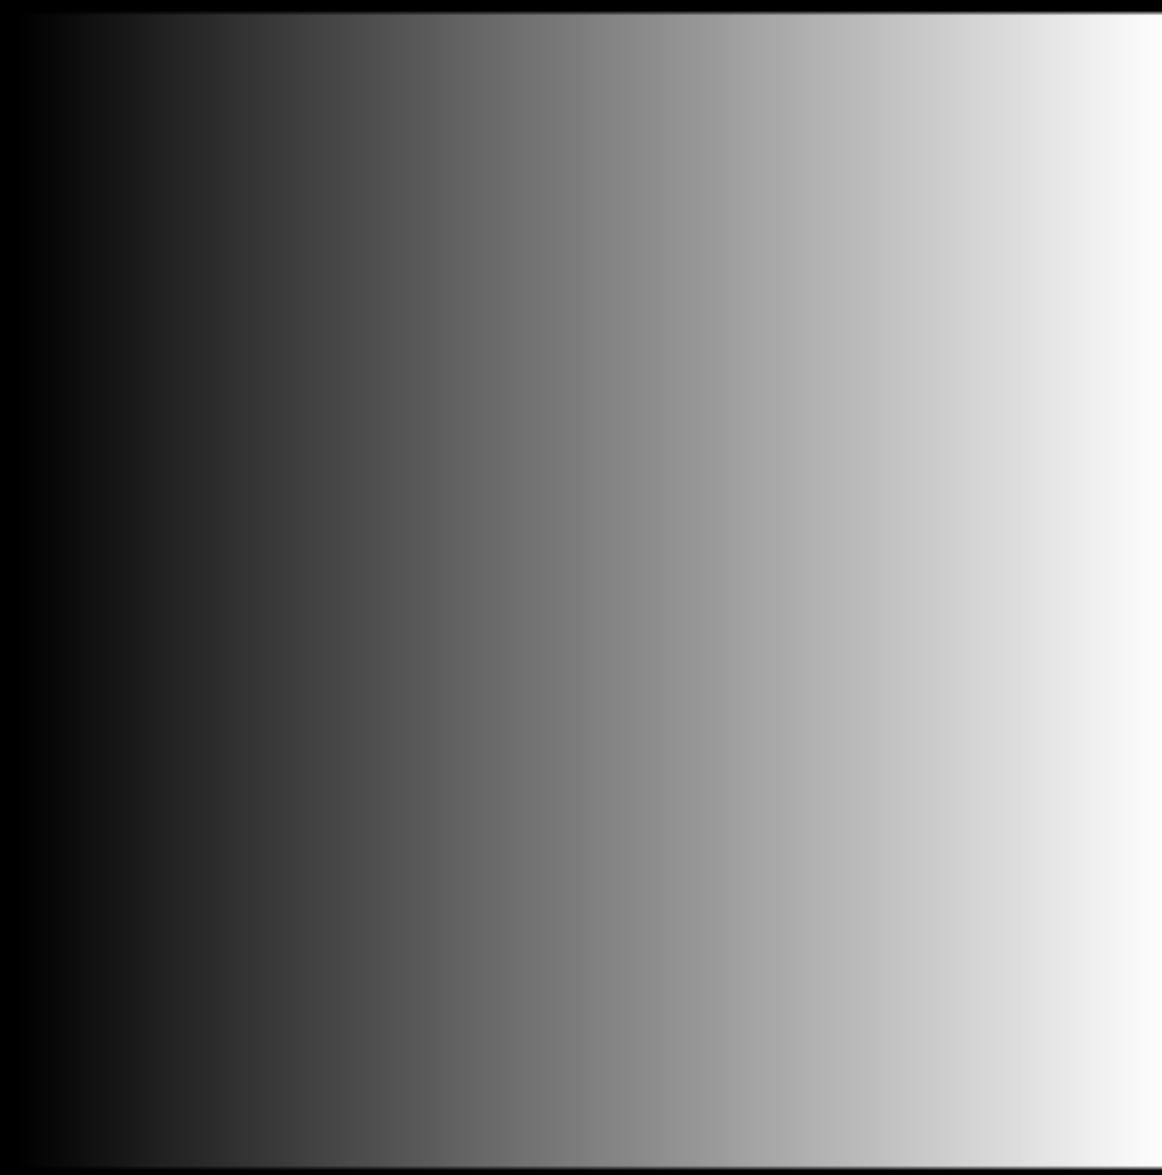
K: #Levels	#Bits
2	1
4	2
8	3
16	4
32	5
64	6
128	7
256	8

Mapping from input signal to output signal is to be determined.

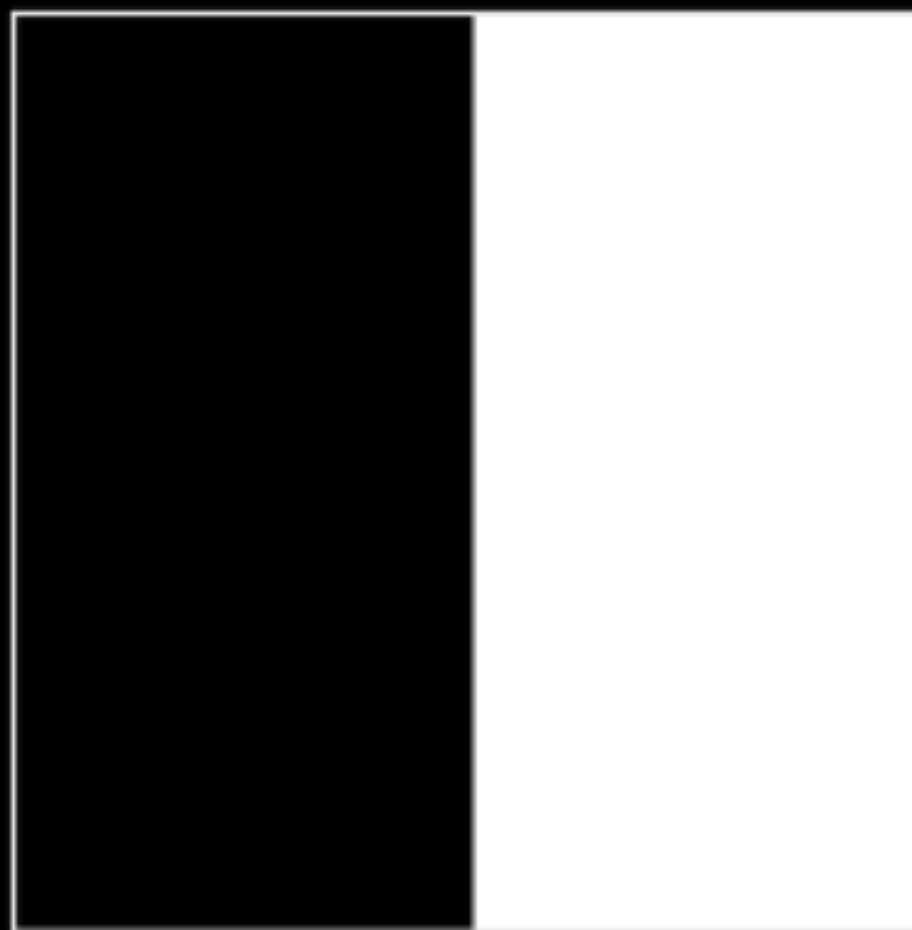
Several types of mappings: uniform, logarithmic, etc.

Choice of K

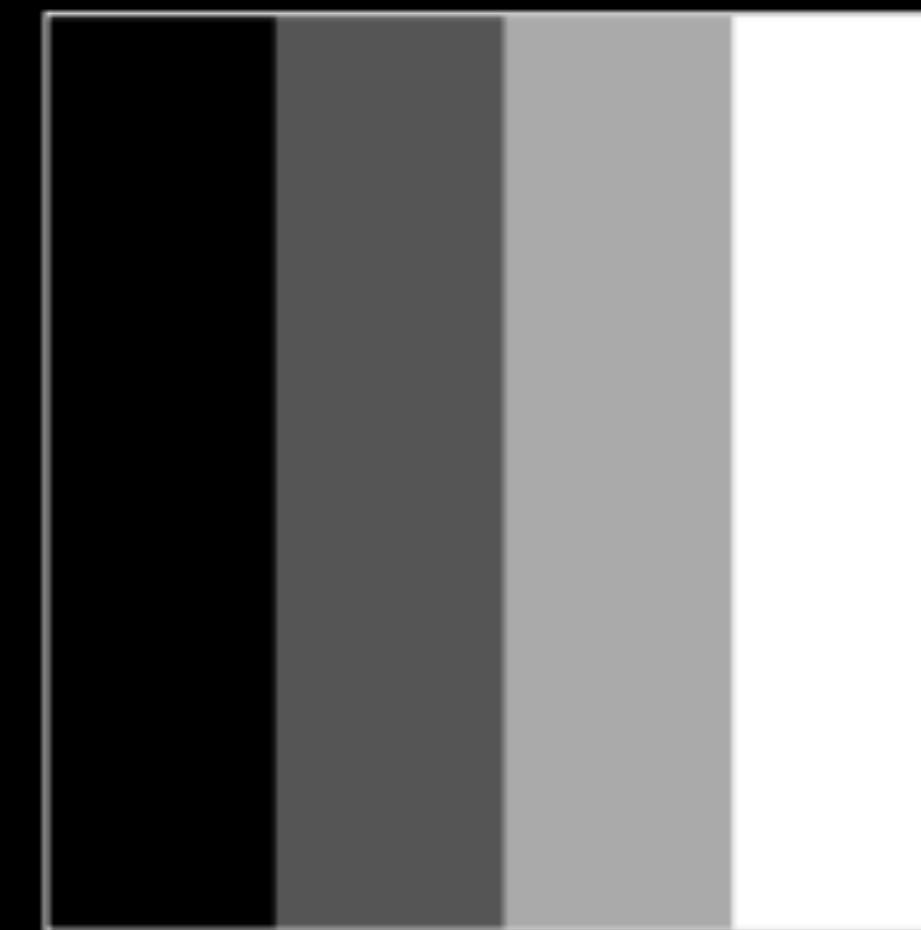
Original



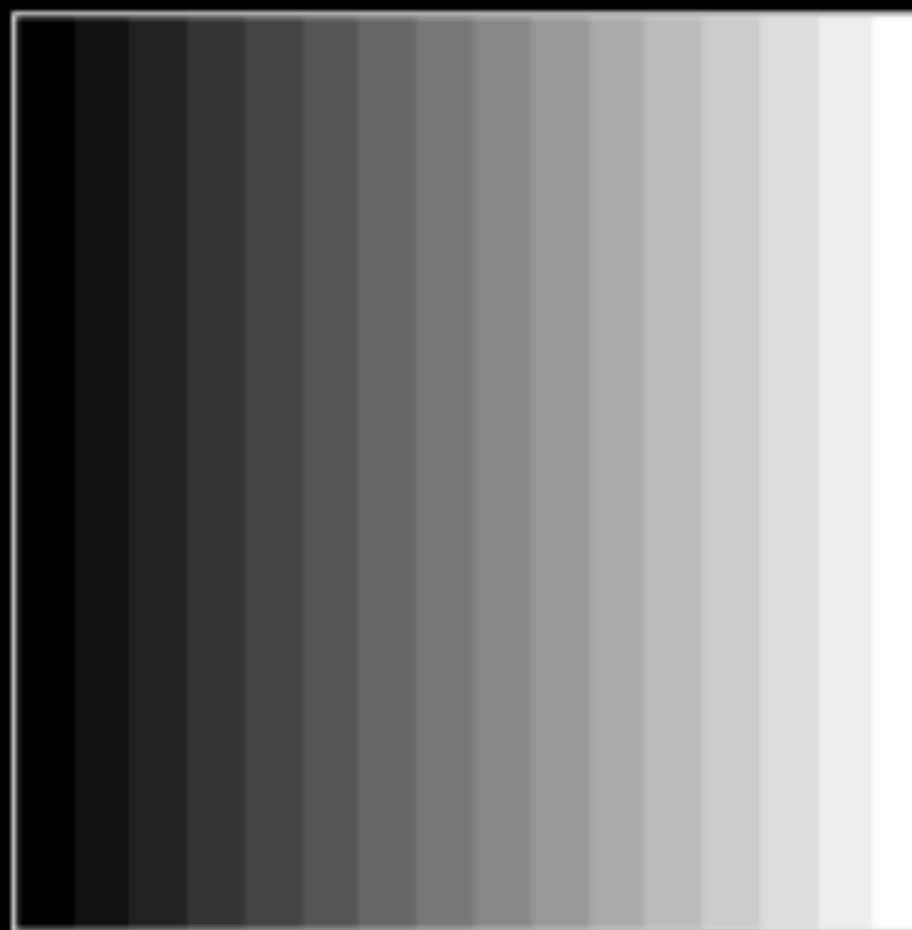
Linear Ramp



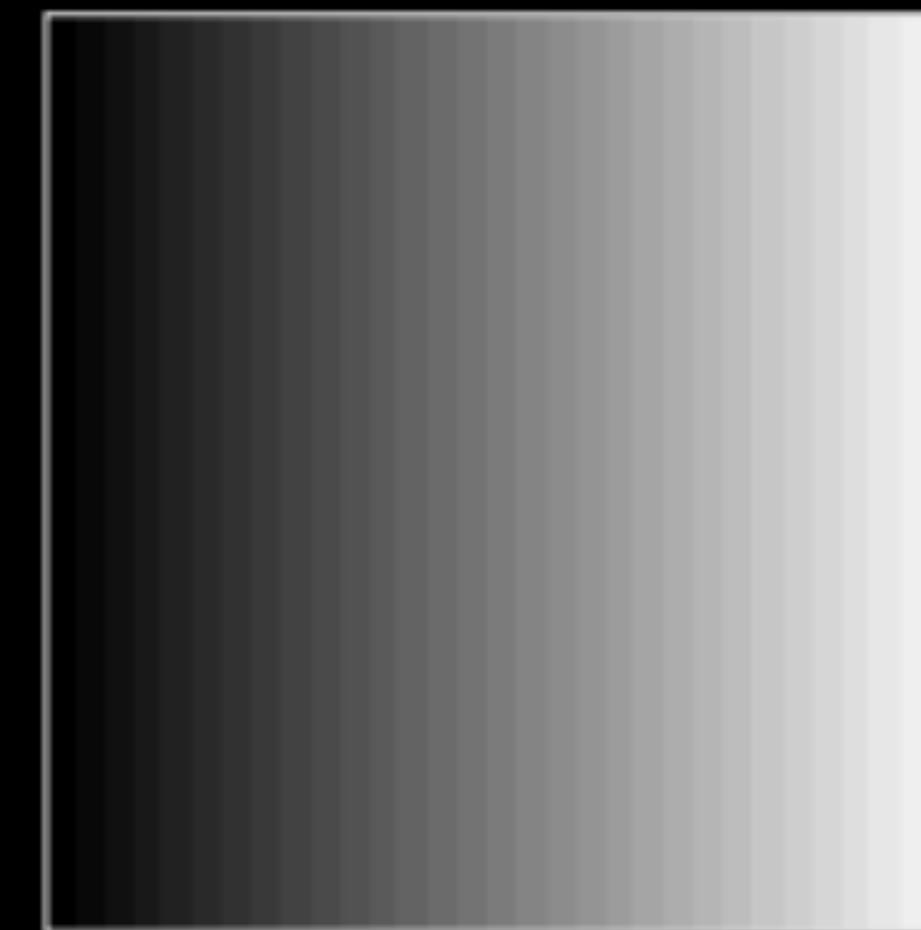
K=2



K=4

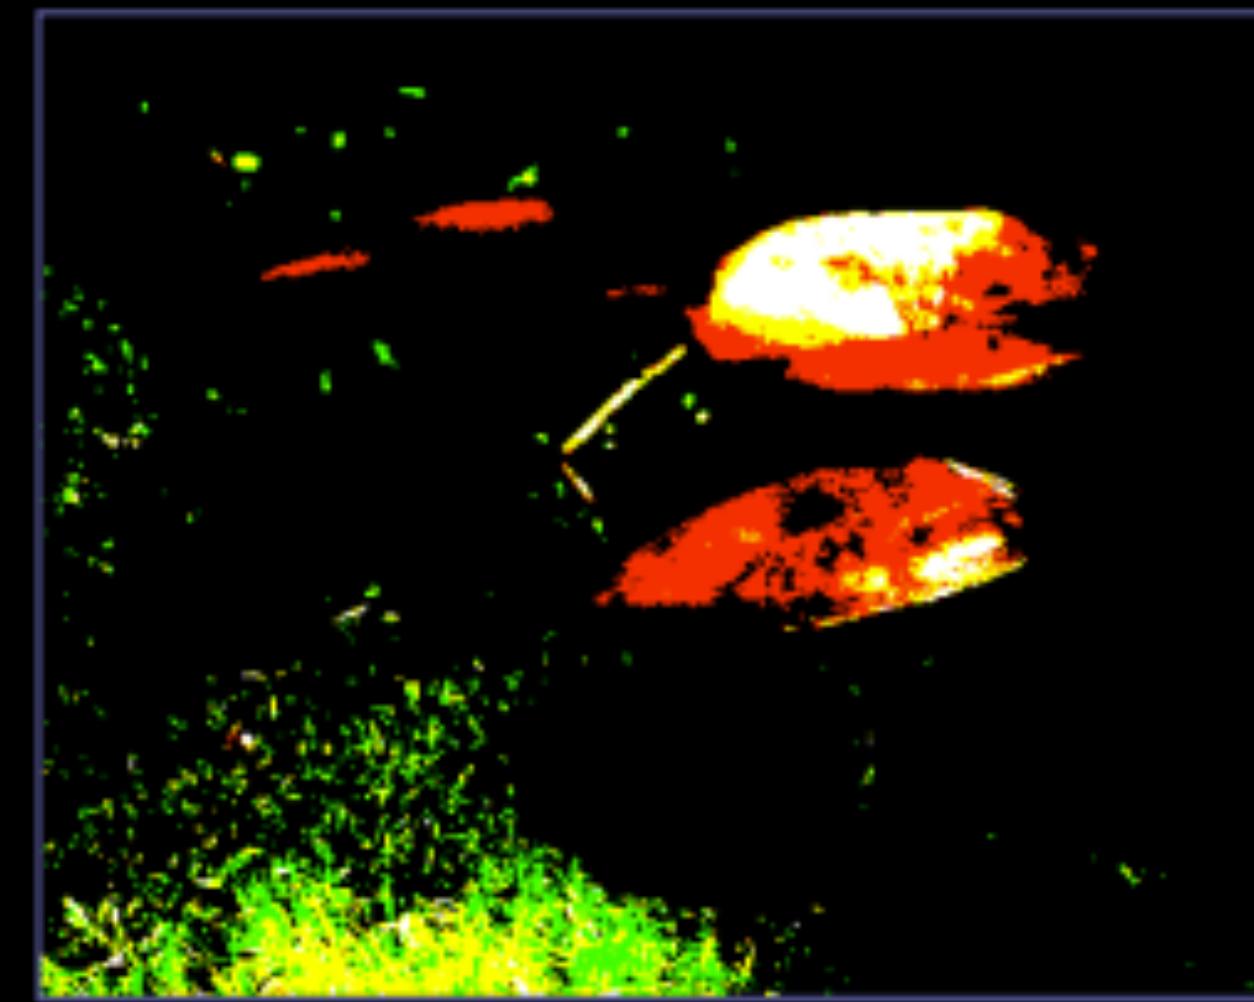


K=16



K=32

Choice of K



K=2 (each color)



K=4 (each color)

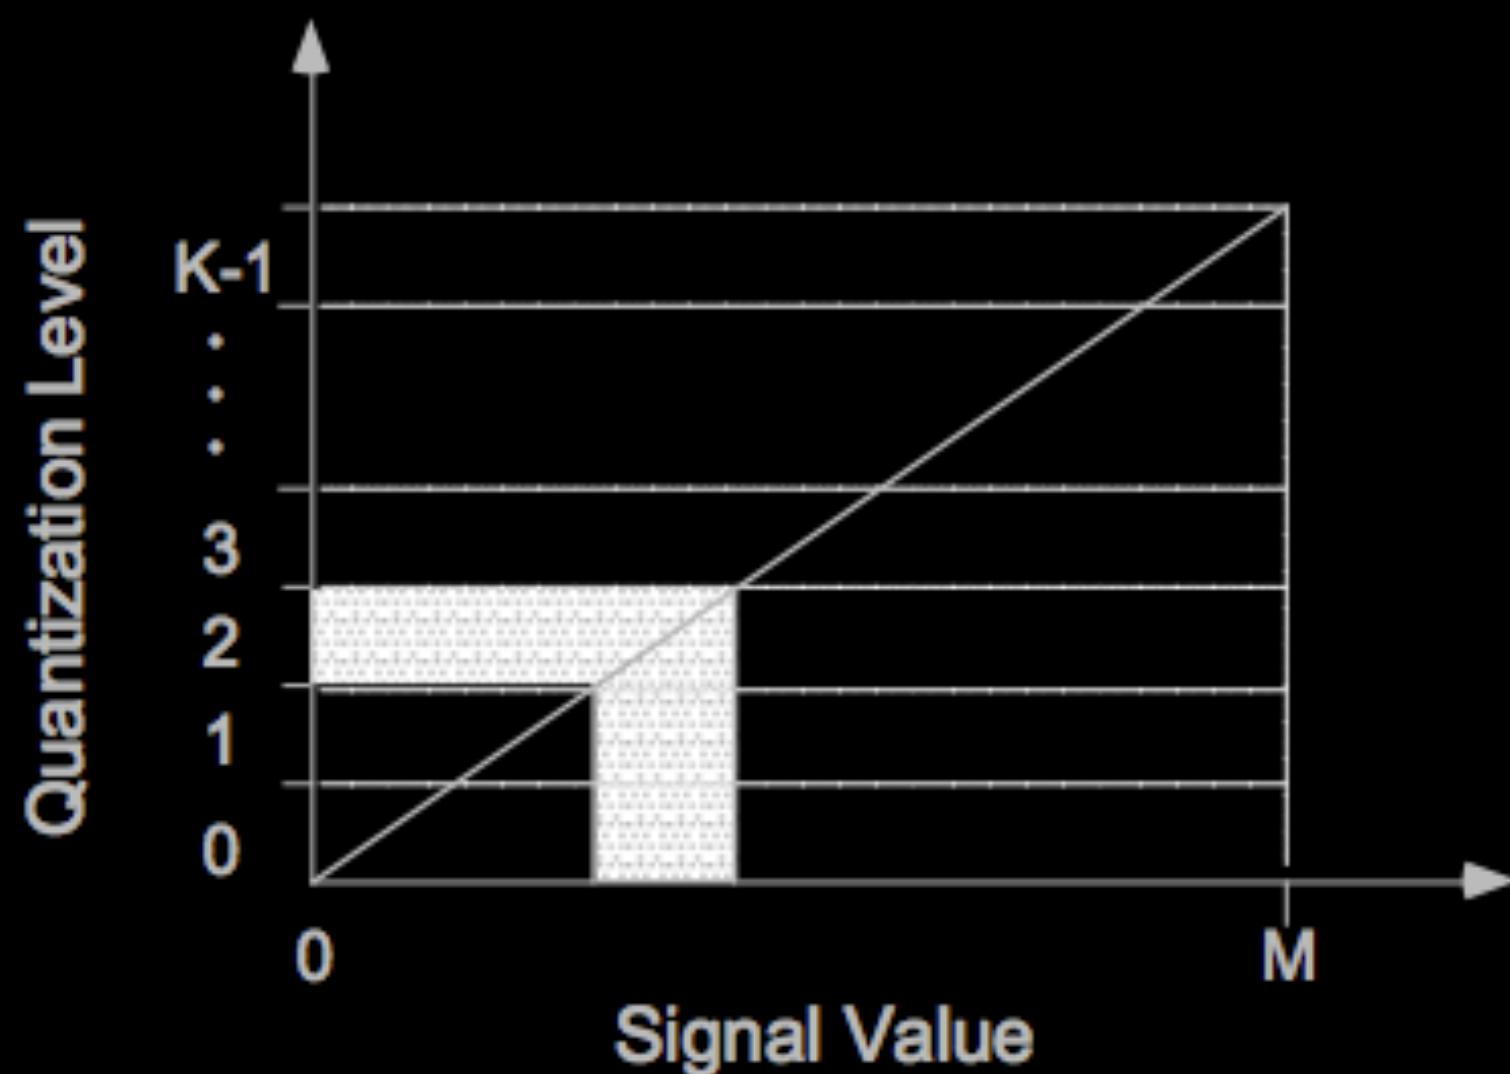
Choice of the function: uniform

Uniform sampling divides the signal range $[0-M]$ into K equal-sized intervals.

The integers $0, \dots, K-1$ are assigned to these intervals.

All signal values within an interval are represented by the associated integer value.

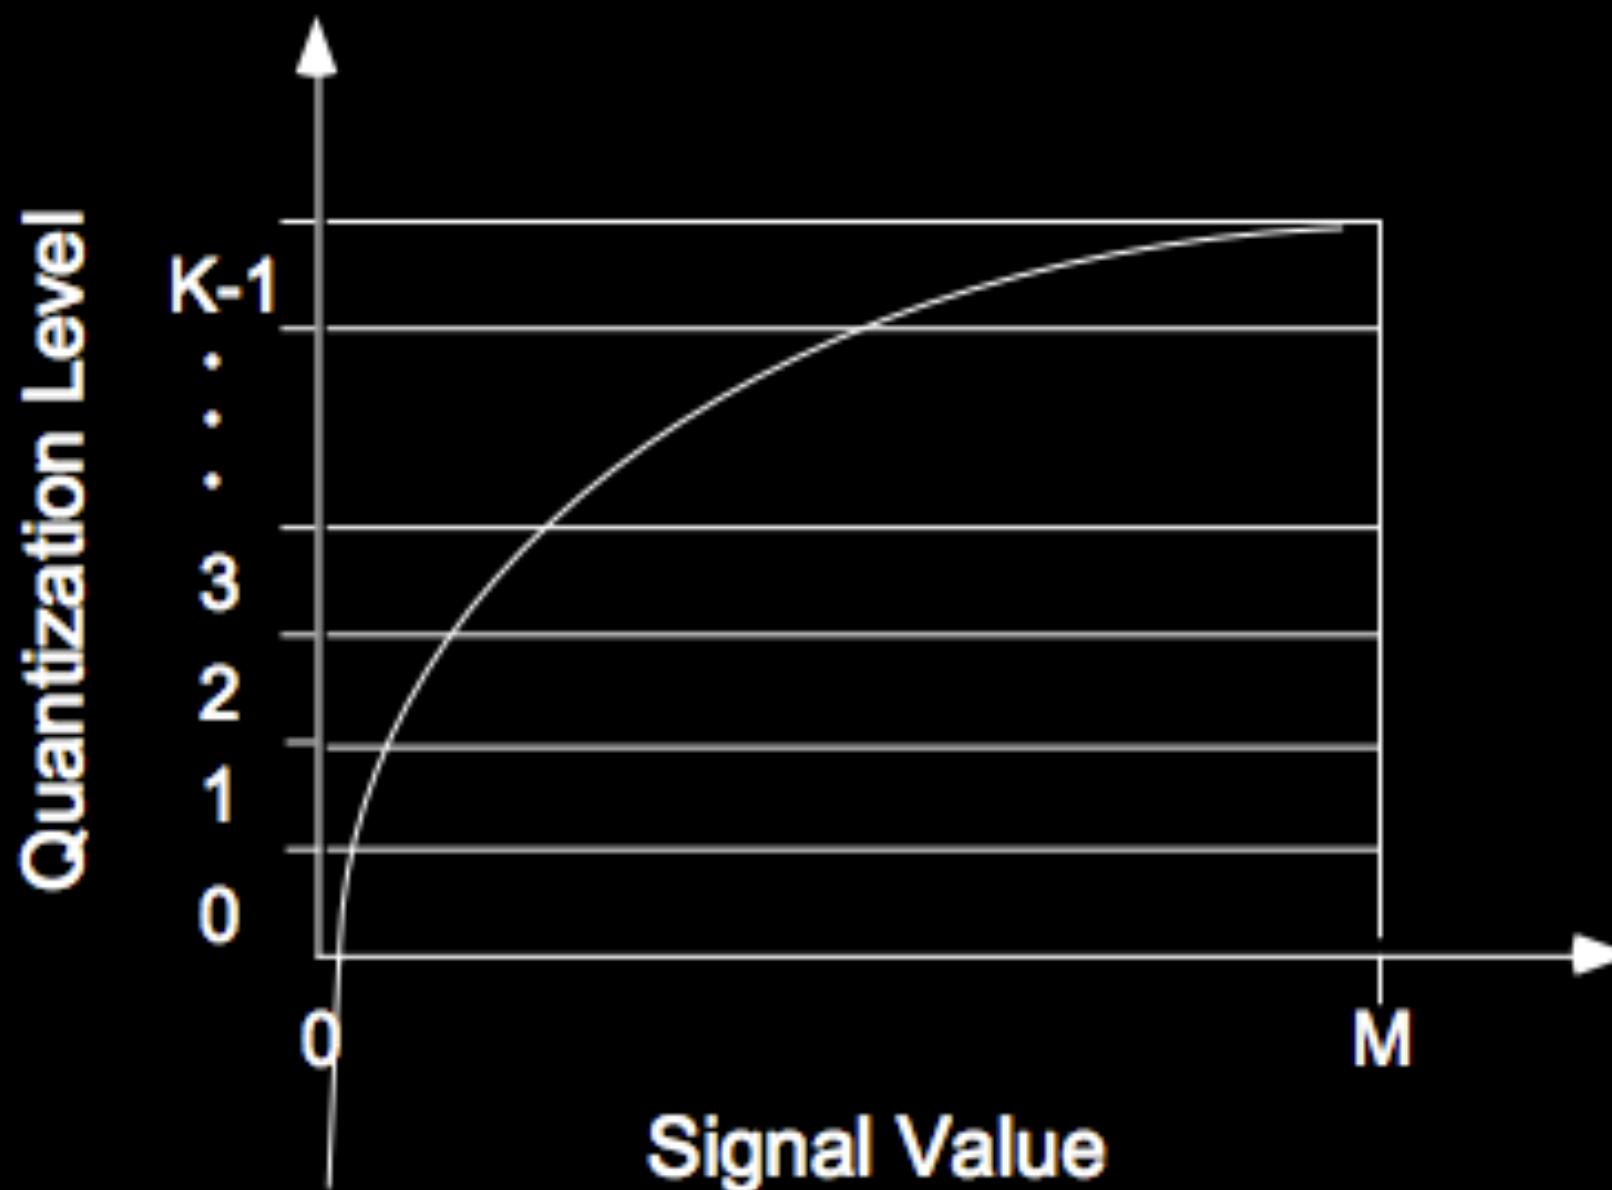
Defines a mapping:



Logarithmic quantization

Signal is: $\log I(x,y)$

Effect is:

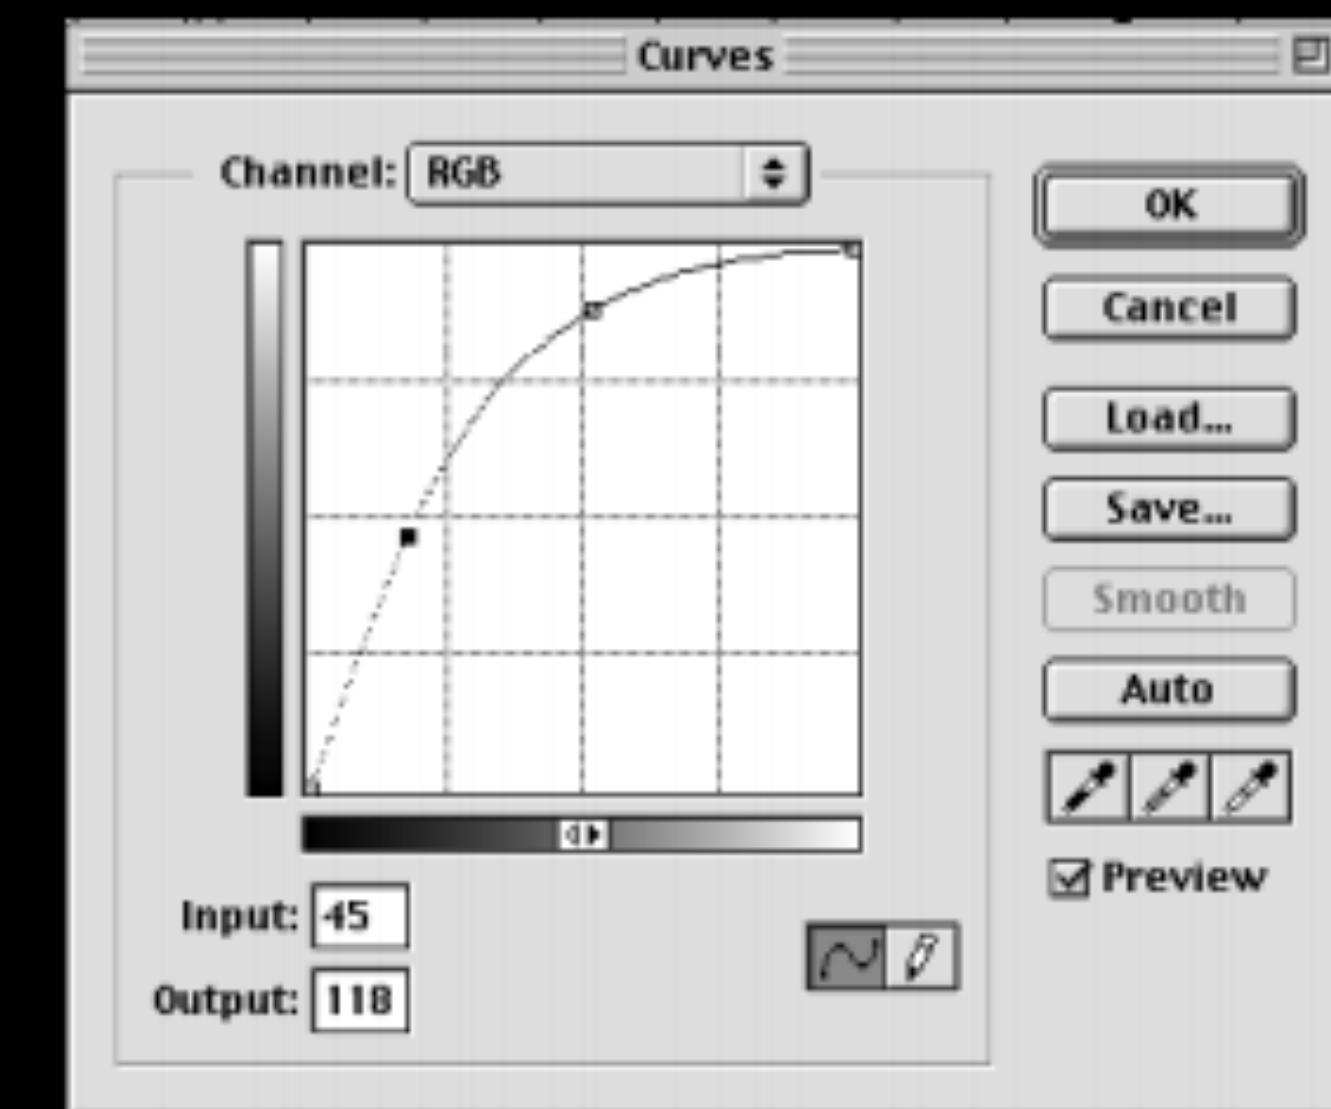


Detail enhanced in the low signal values at expense of detail in high signal values.

Logarithmic quantization



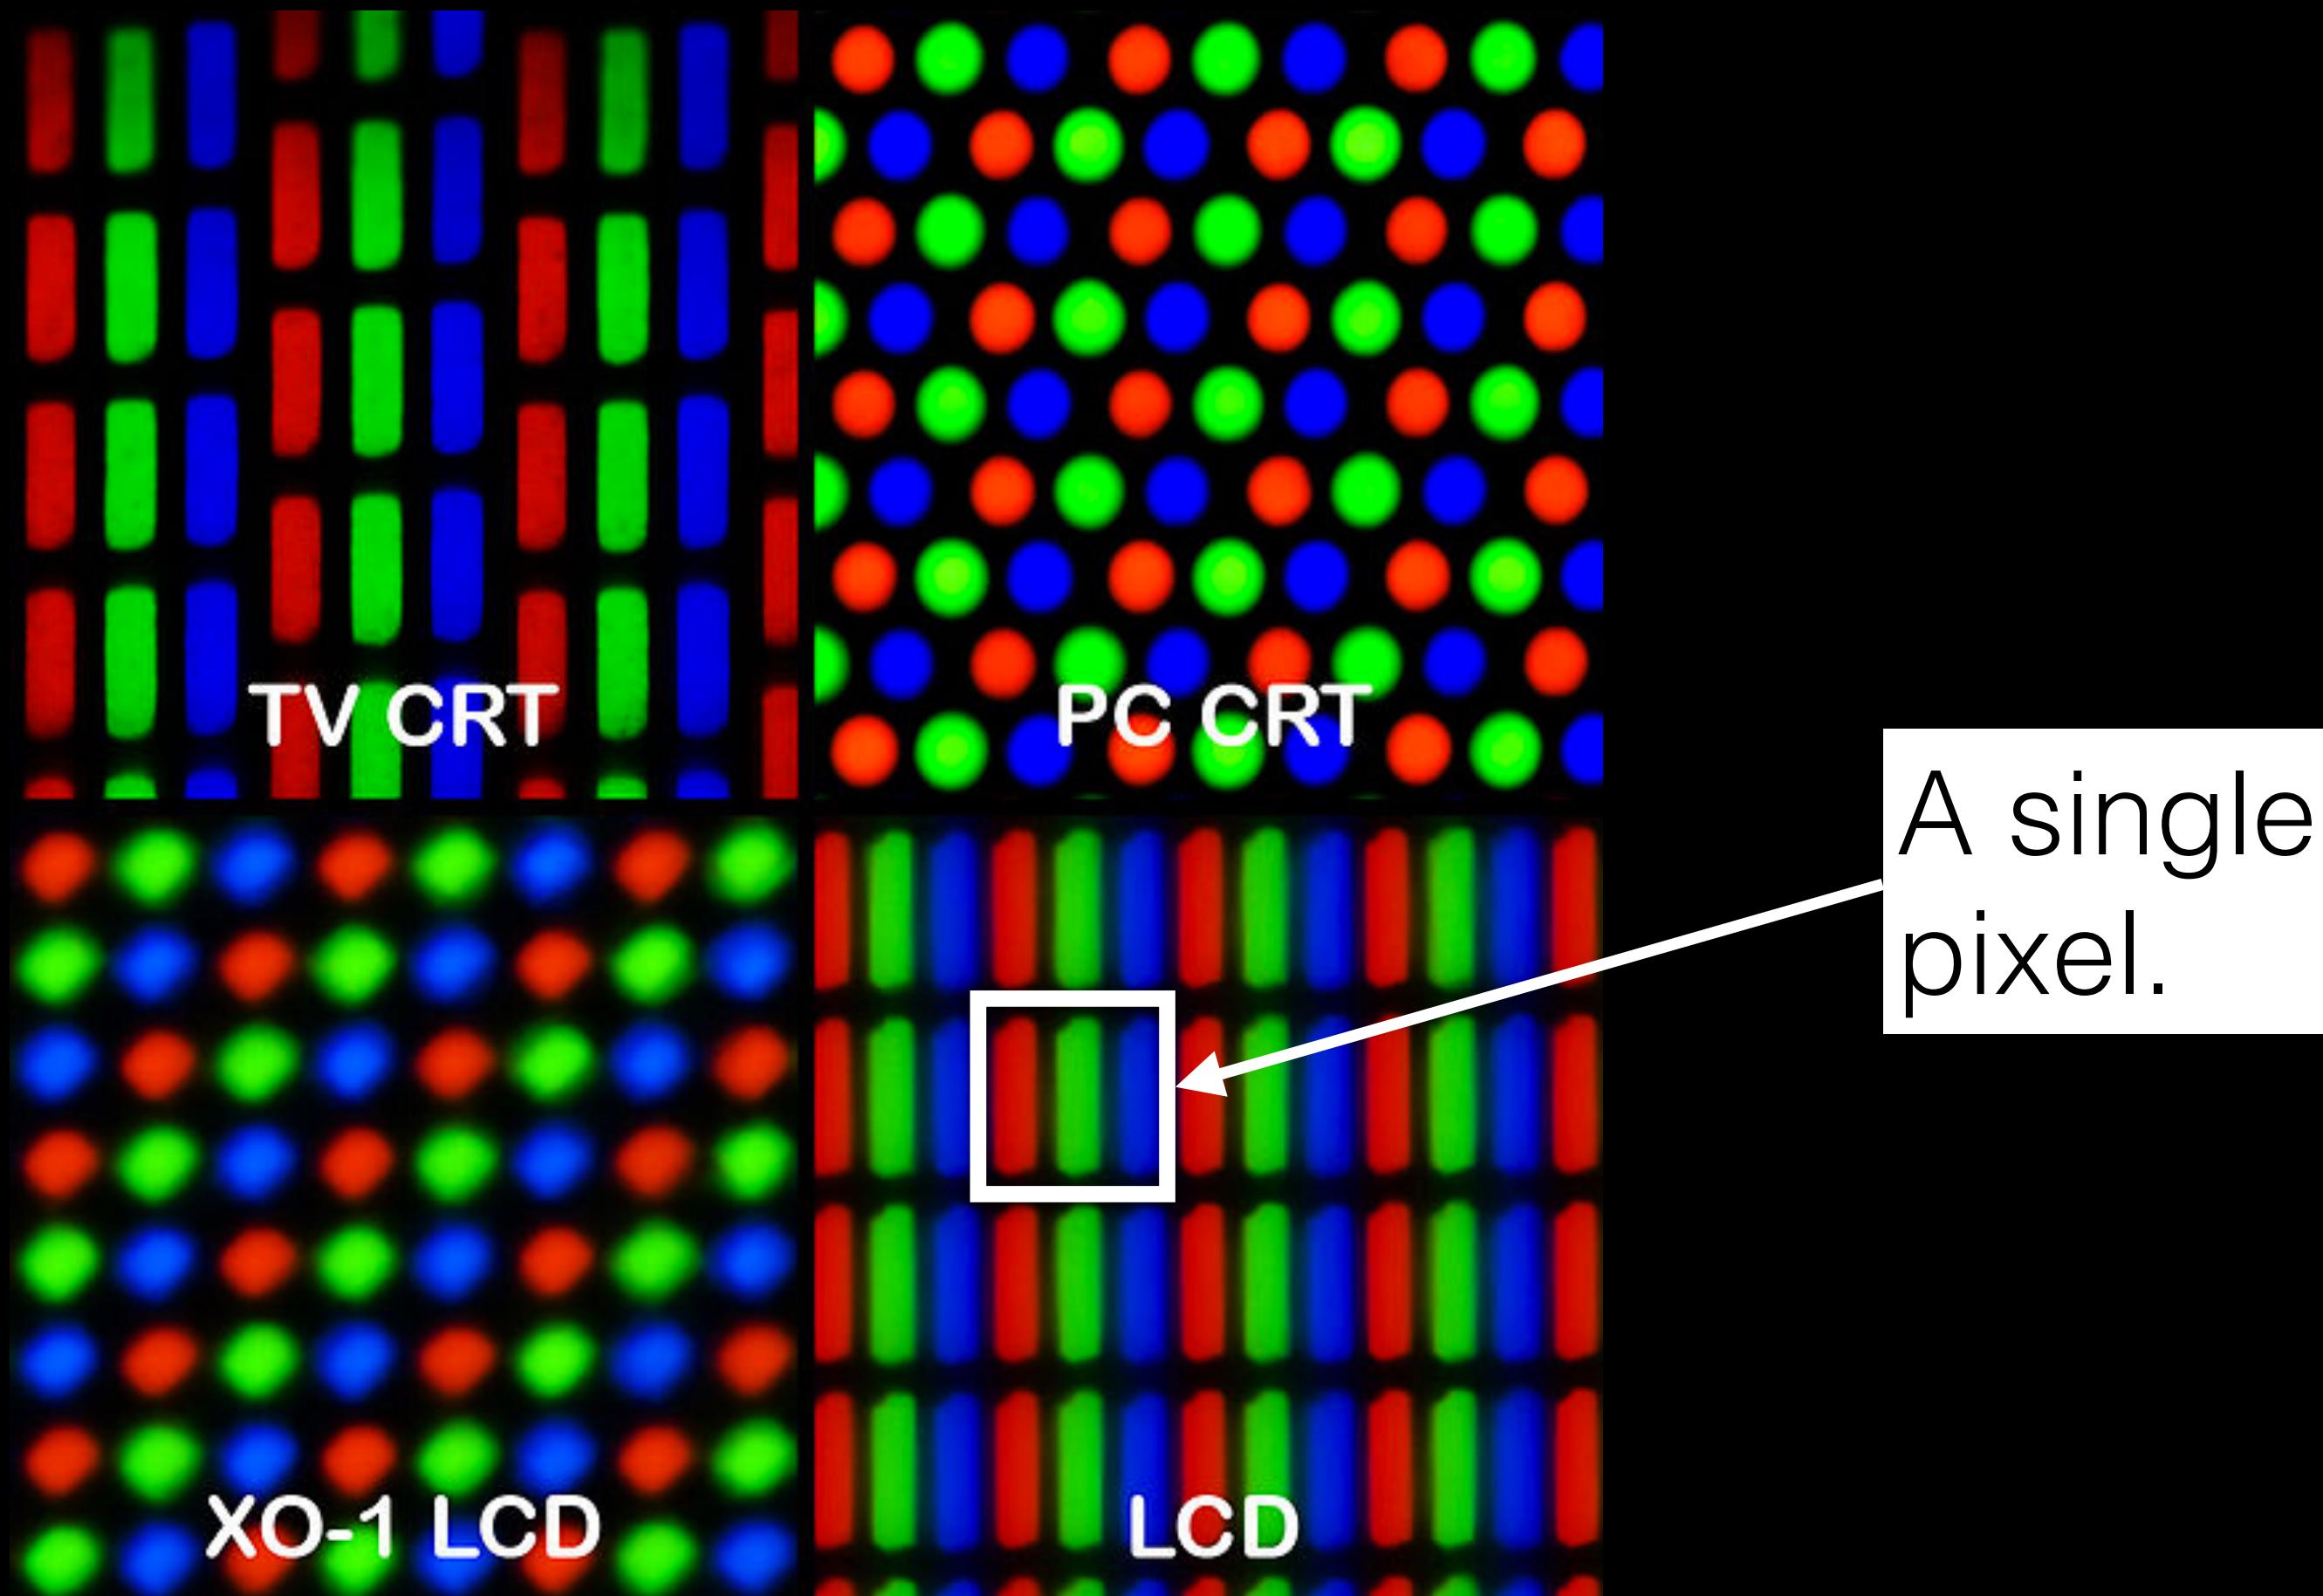
Quantization Curve



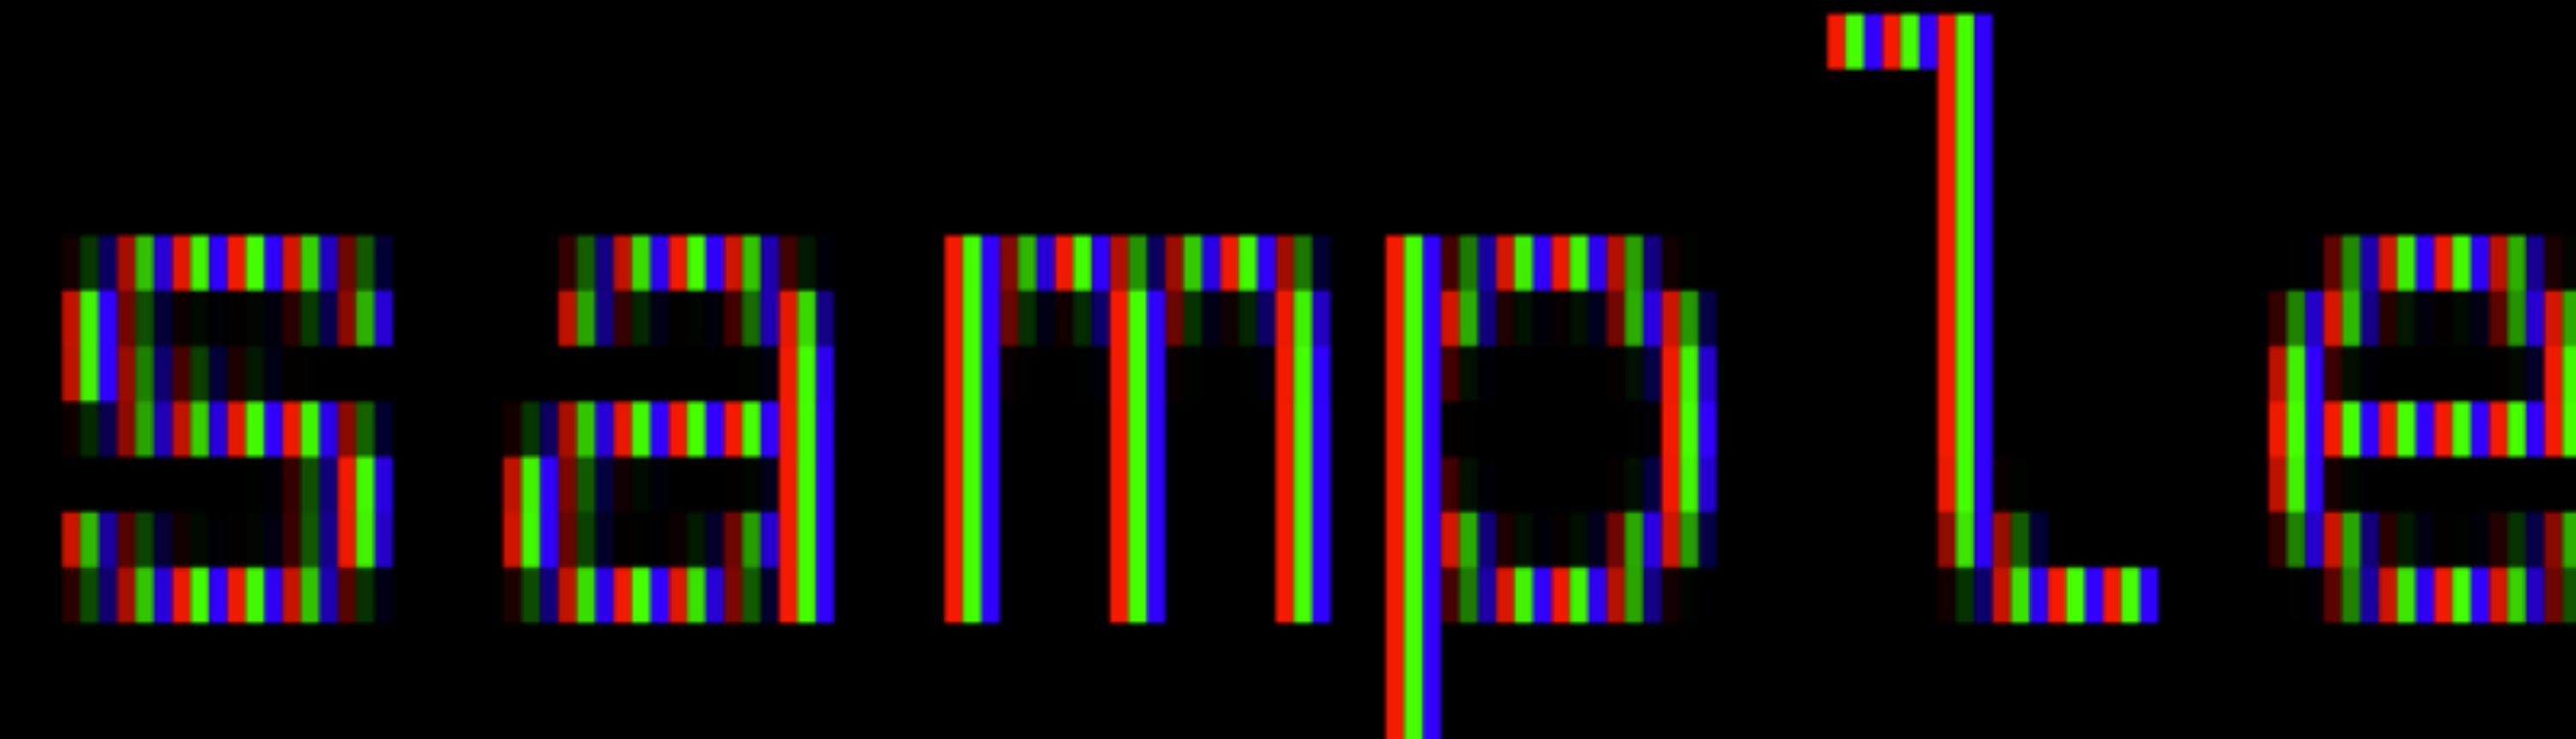
Color displays

Given a 24 bit color image (8 bits for R, G, B)

- ▶ Turn on 3 subpixels with power proportional to RGB values



“White” text on color display



sample

http://en.wikipedia.org/wiki/Subpixel_rendering

Lookup tables

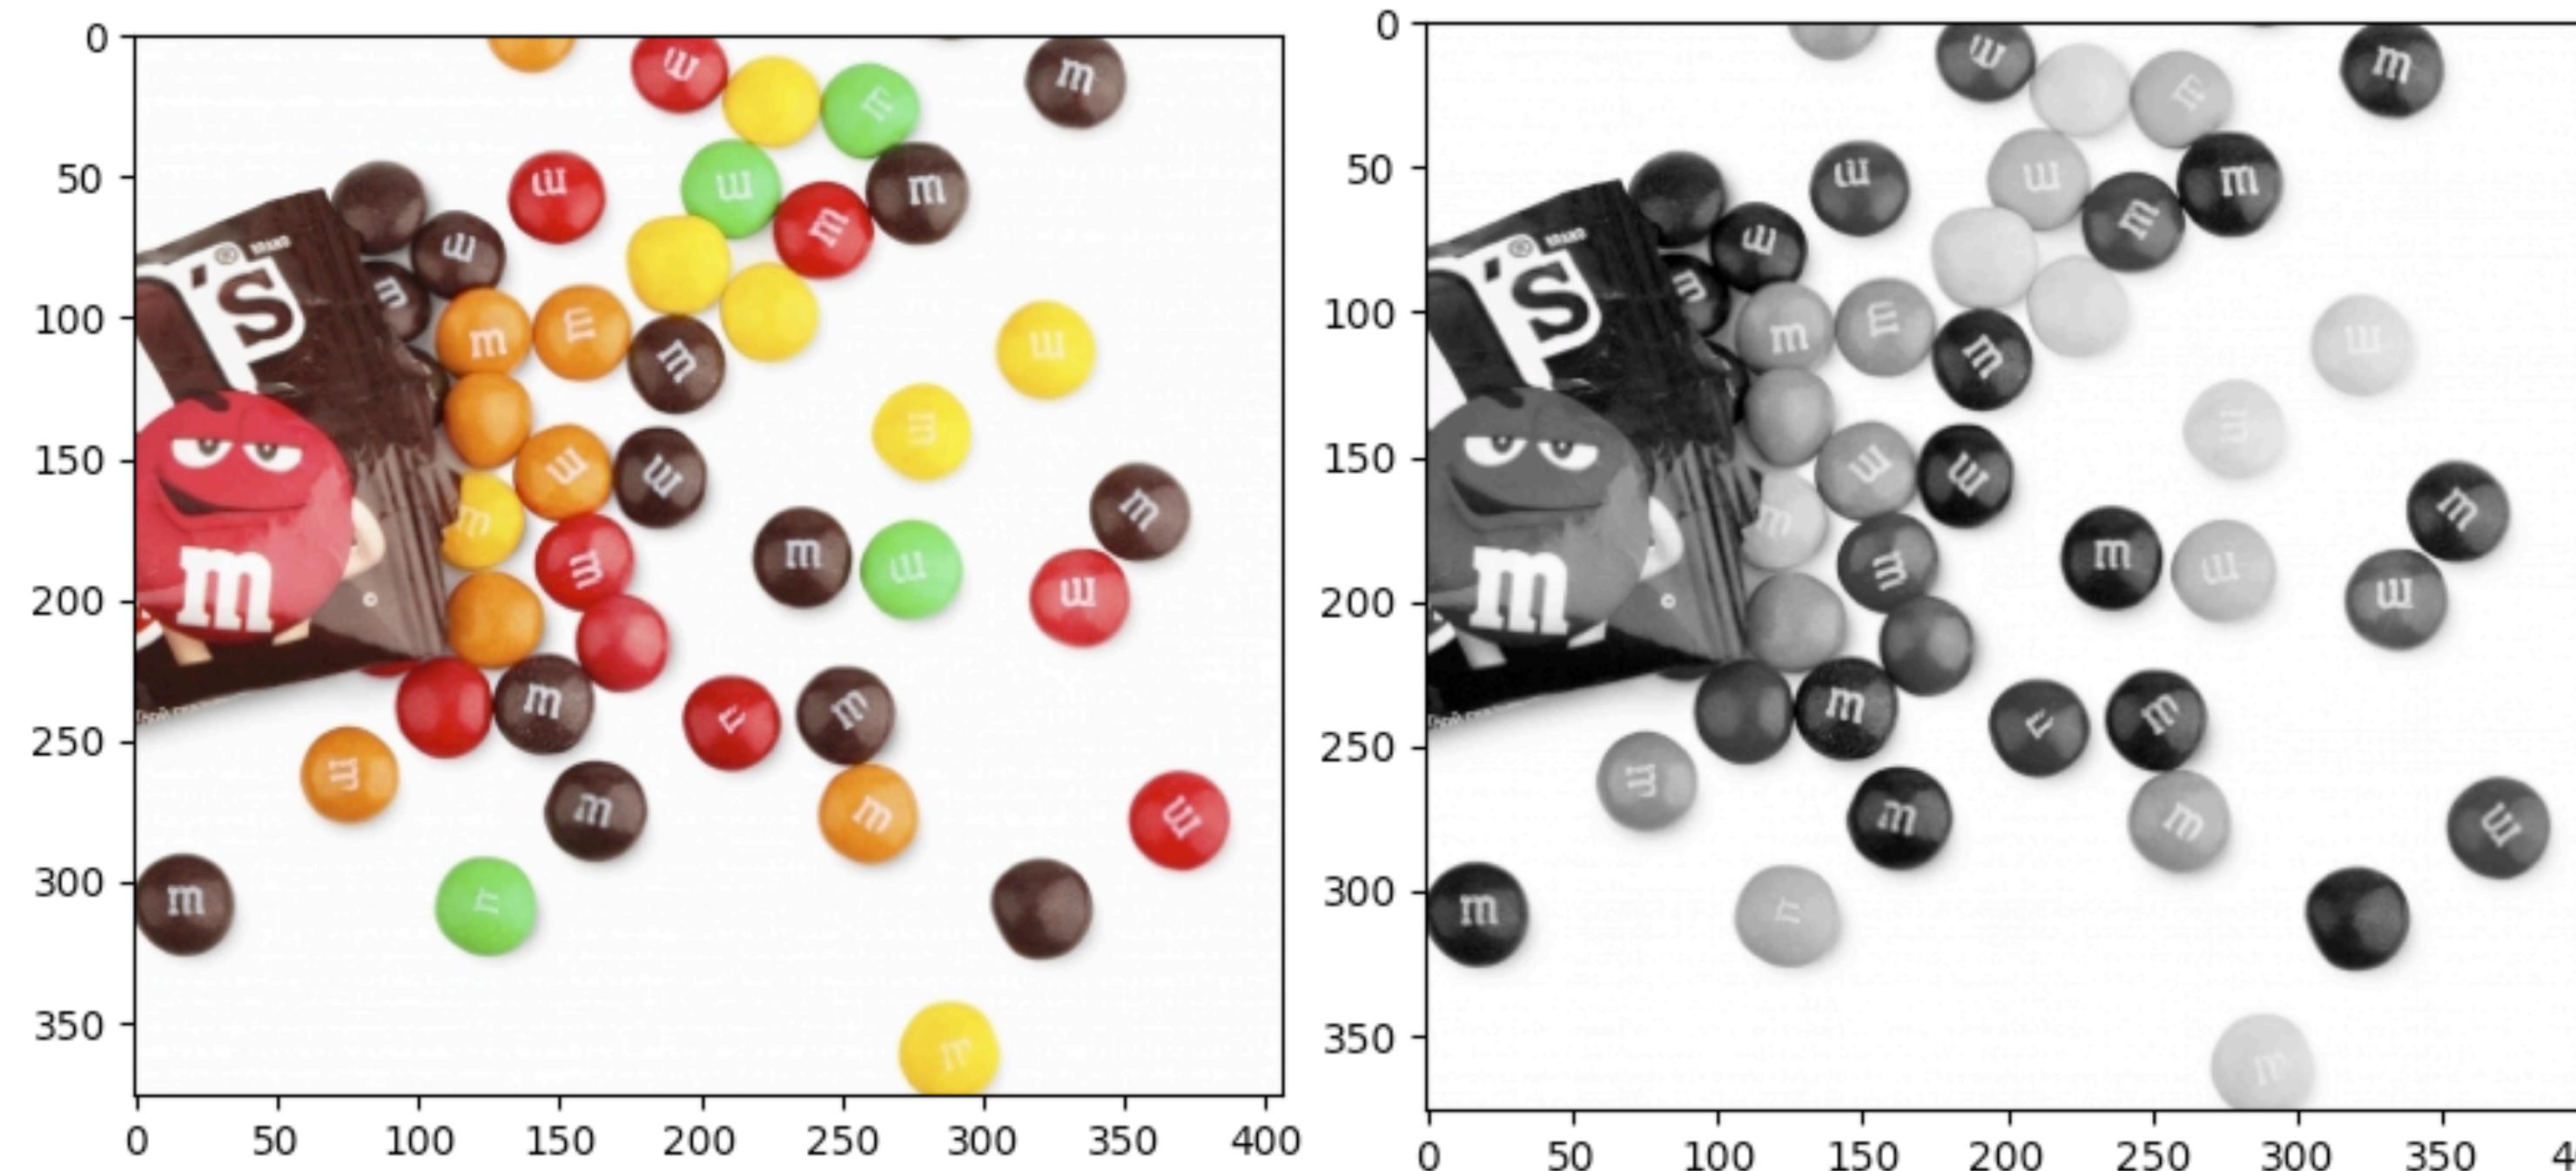
8 bit image: 256 different values.

Simplest way to display: map each number to a gray value:

- ▶ 0 → (0.0, 0.0, 0.0) or (0,0,0)
- ▶ 1 → (0.0039, 0.0039, 0.0039) or (1,1,1)
- ▶ 2 → (0.0078, 0.0078, 0.0078) or (2,2,2)
- ▶ ...
- ▶ 255 → (1.0, 1.0, 1.0) or (255,255,255)

This is called a grayscale mapping.

Color to gray and colormaps



```
File Edit Options Buffers Tools Python Help
from skimage import io
import matplotlib.pyplot as plt

im = io.imread('mnms.jpeg');
plt.figure(1);
plt.imshow(im)

gray = im[:, :, 0]*0.3 + im[:, :, 1]*0.6 + im[:, :, 2]*0.1;
plt.figure(2);
plt.imshow(gray, cmap='gray')
plt.show()
```

R : G : B :: 0.3 : 0.6 : 0.1

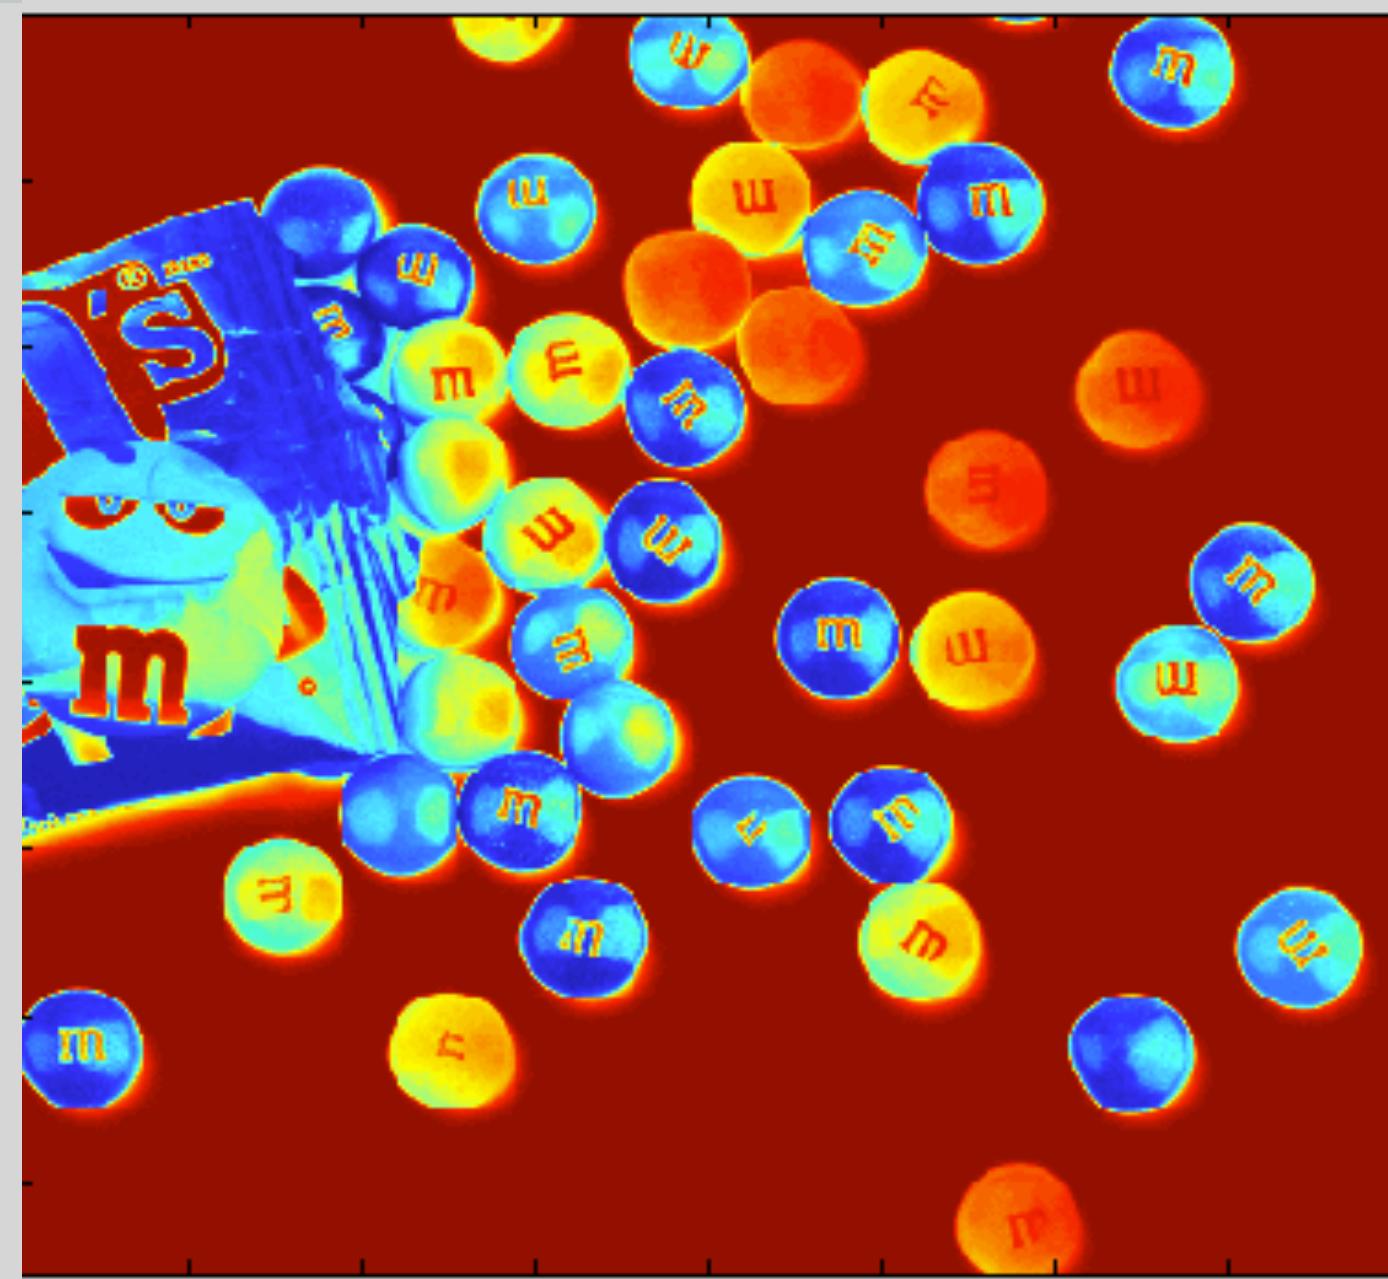
Non-gray lookup tables

We can also use other mappings:

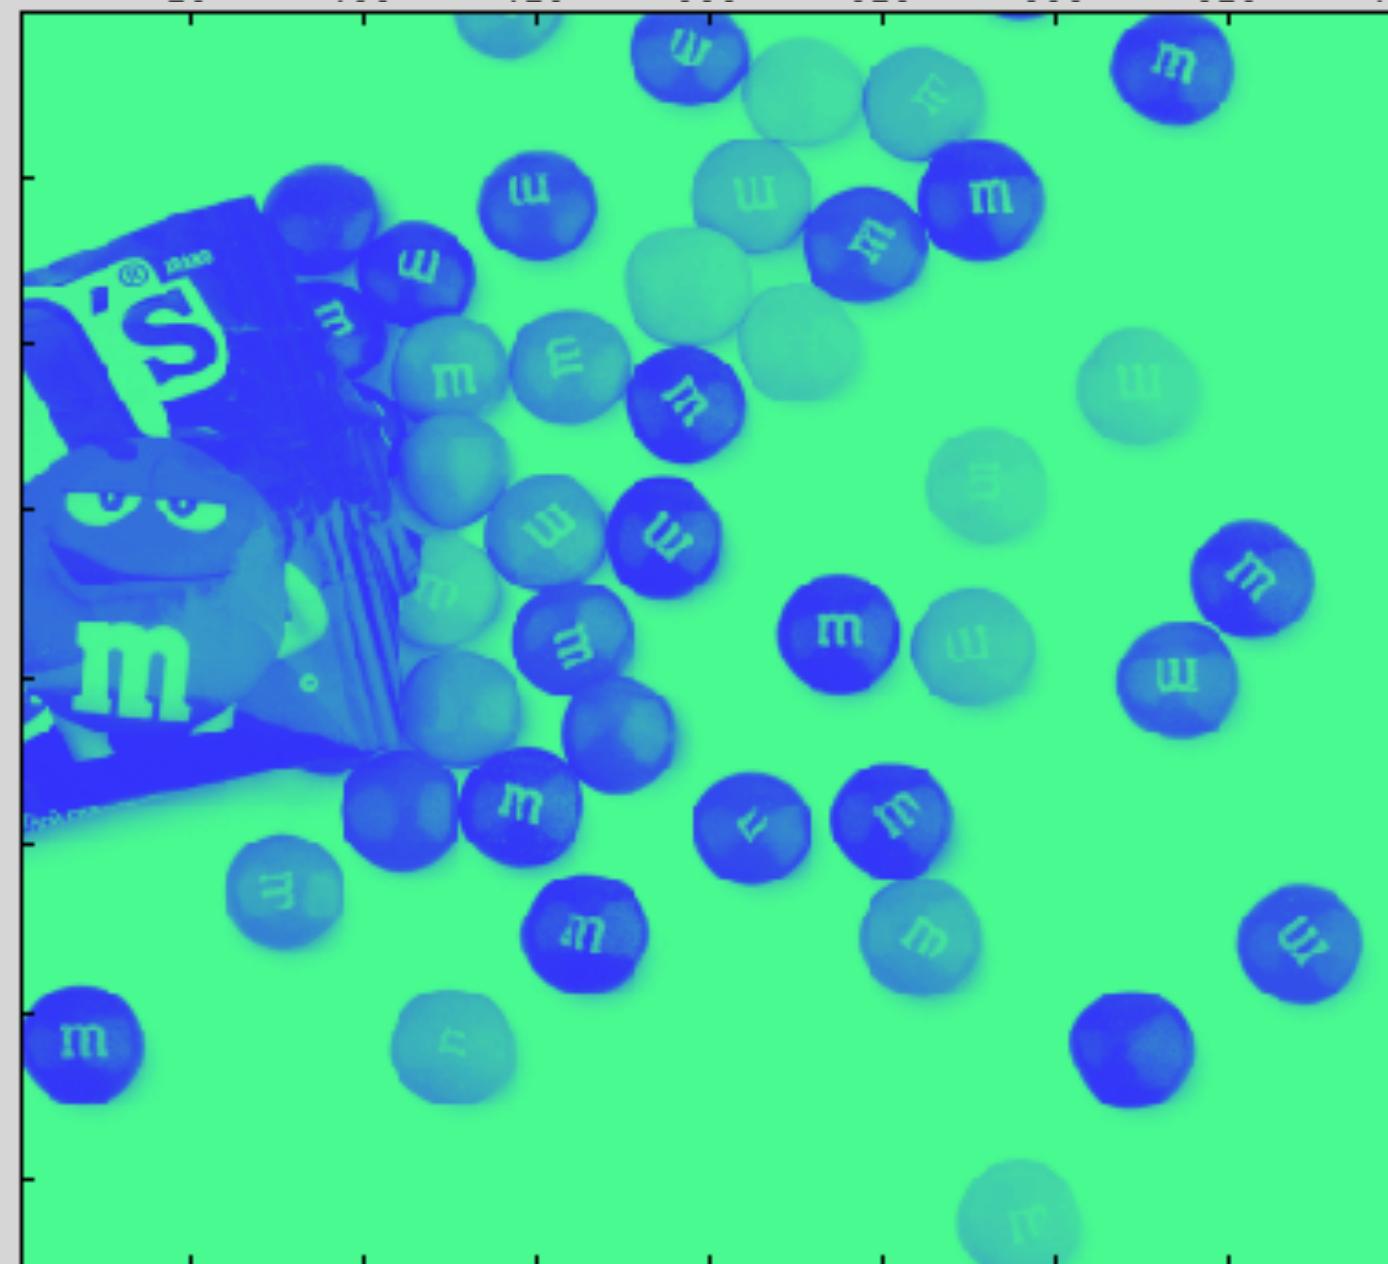
- ▶ $0 \rightarrow (17, 25, 89)$
- ▶ $1 \rightarrow (45, 32, 200)$
- ▶ ...
- ▶ $255 \rightarrow (233, 1, 4)$

These are called **lookup tables**.

More colormaps



jet



winter

Colormap Name	Color Scale
parula	
jet	
hsv	
hot	
cool	
spring	
summer	
autumn	
winter	
gray	
bone	
copper	
pink	
lines	
colorcube	
prism	
flag	
white	

Enhancing images

What can we do to “enhance” an image after it has already been digitized?

- ▶ We can make the information that is there easier to visualize.
- ▶ We can guess at data that is not there, but we cannot be sure, in general.



Increasing the contrast



Removing motion blur

Contrast enhancement

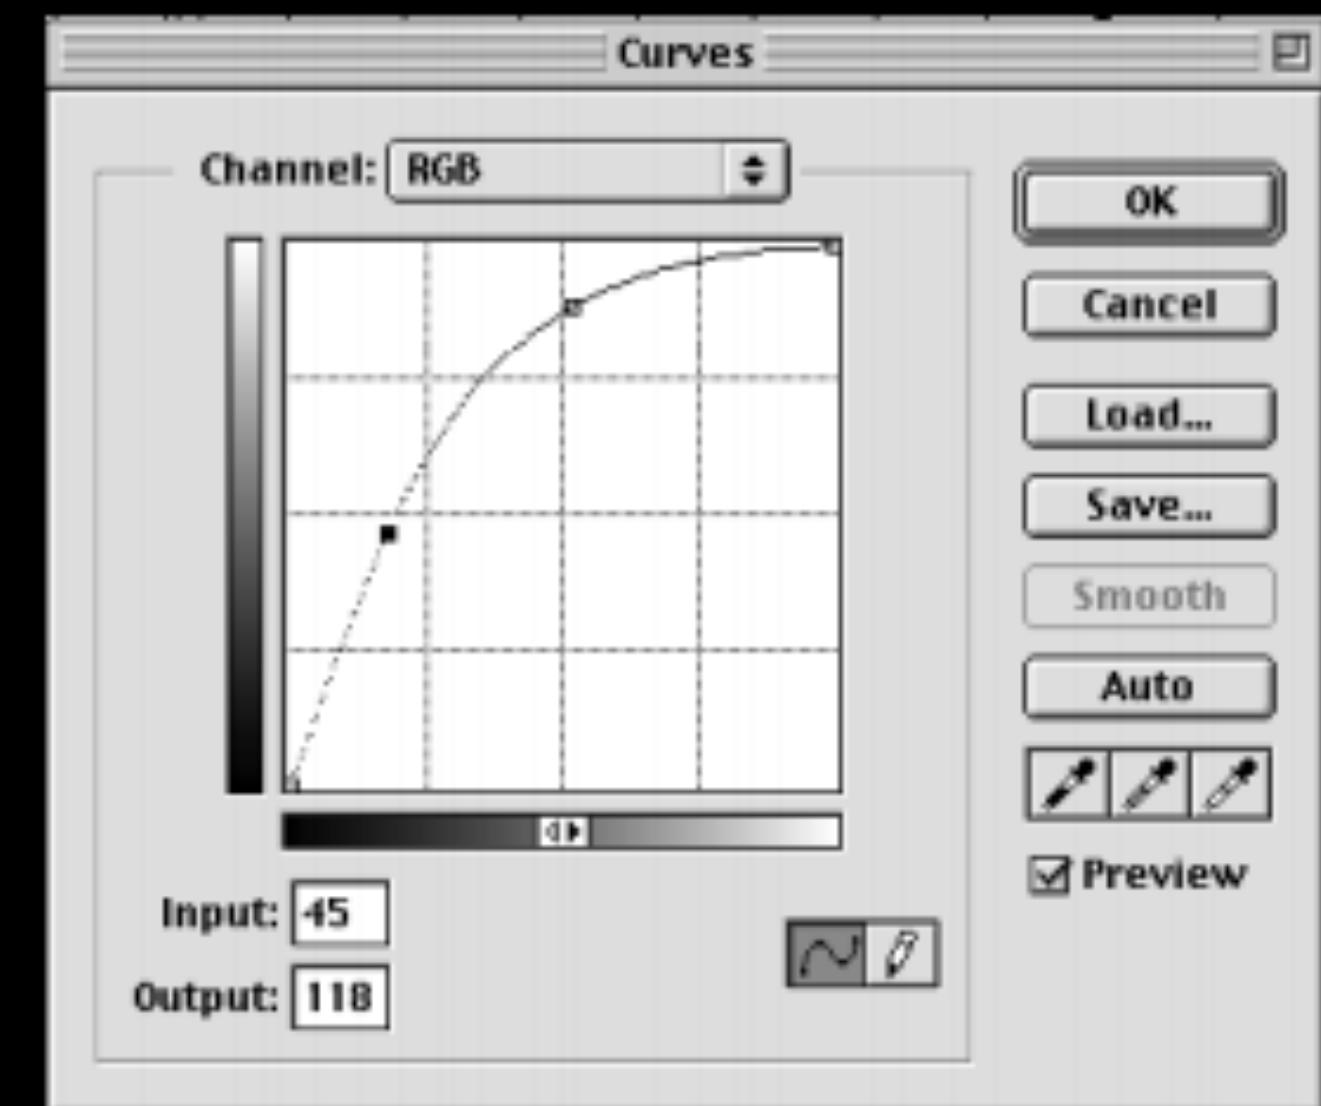
Two methods:

- ▶ Normalize the data (non-linear mapping, contrast stretching)
- ▶ Transform the data (histogram equalization)

Logarithmic quantization



Quantization Curve



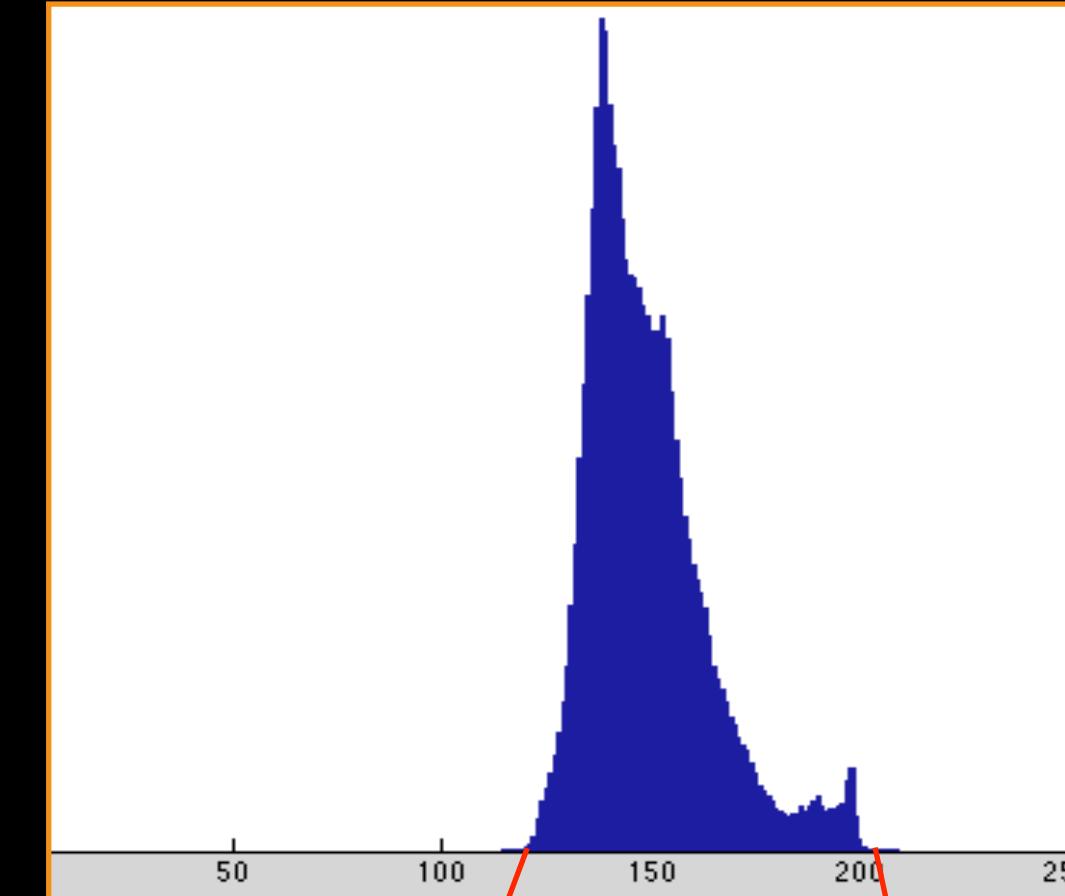
Contrast stretching



Original

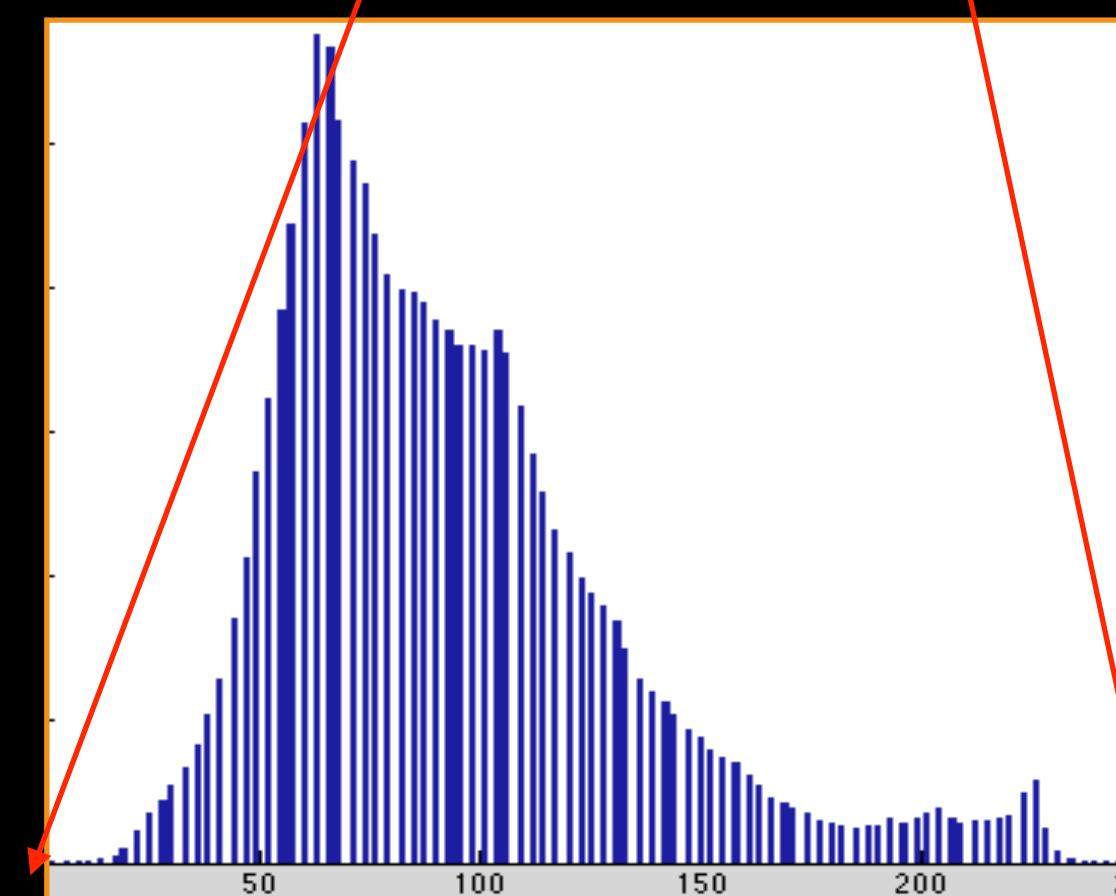


After contrast stretching



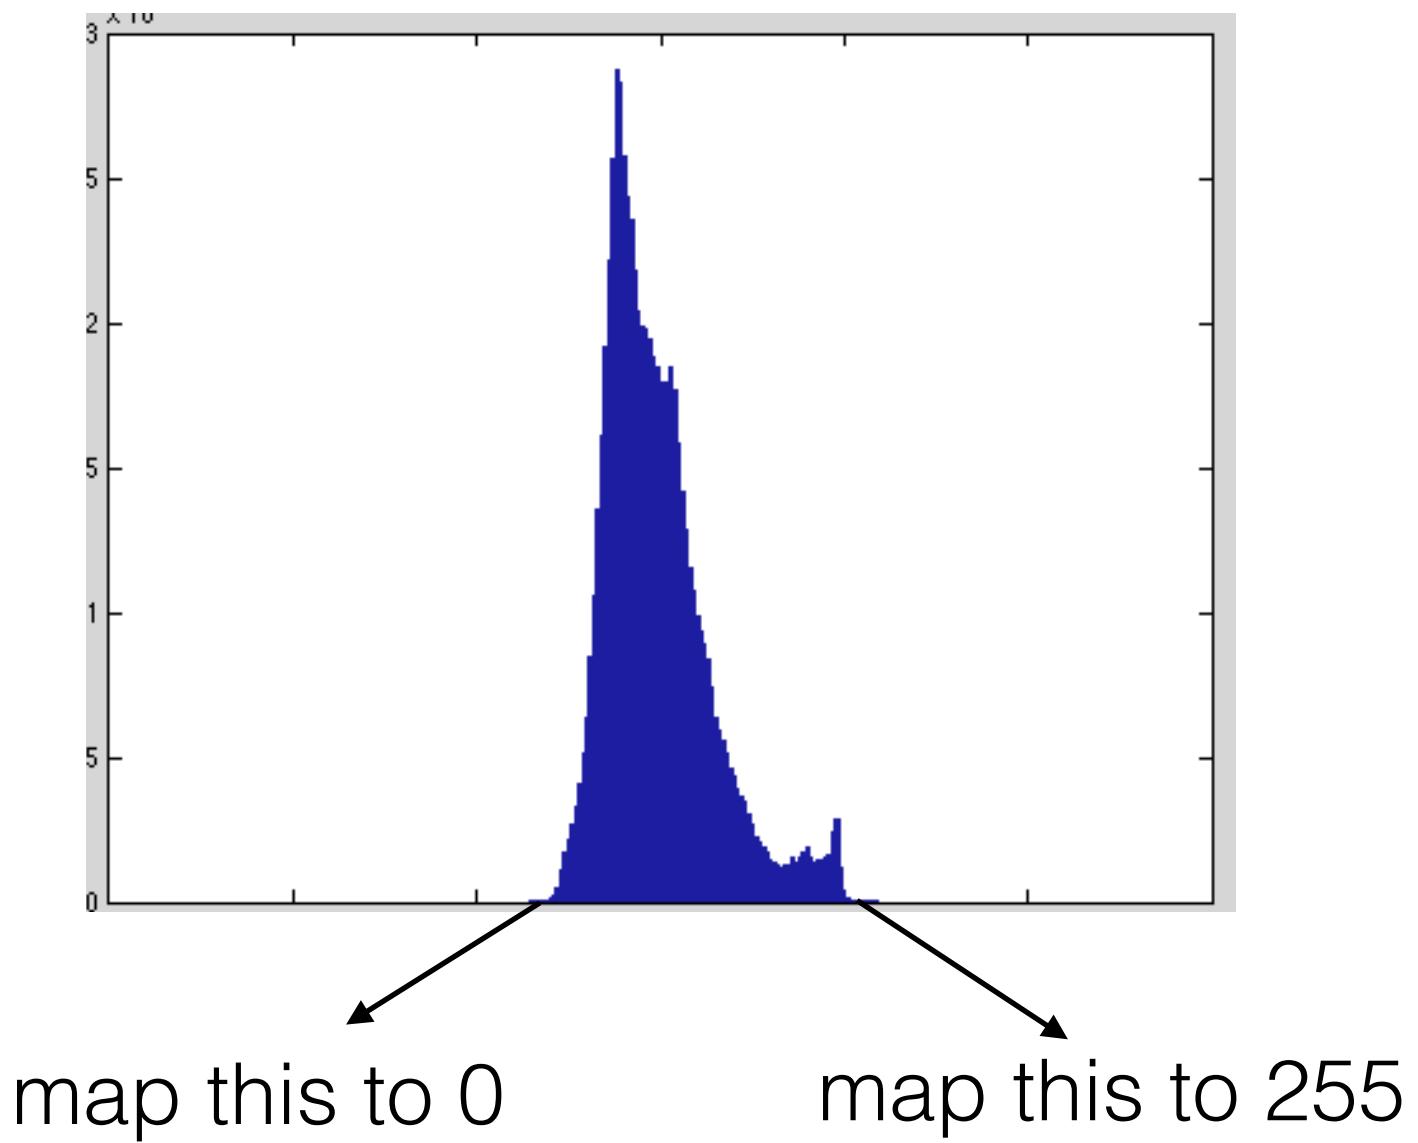
map this to 0

map this to 255



Contrast stretching

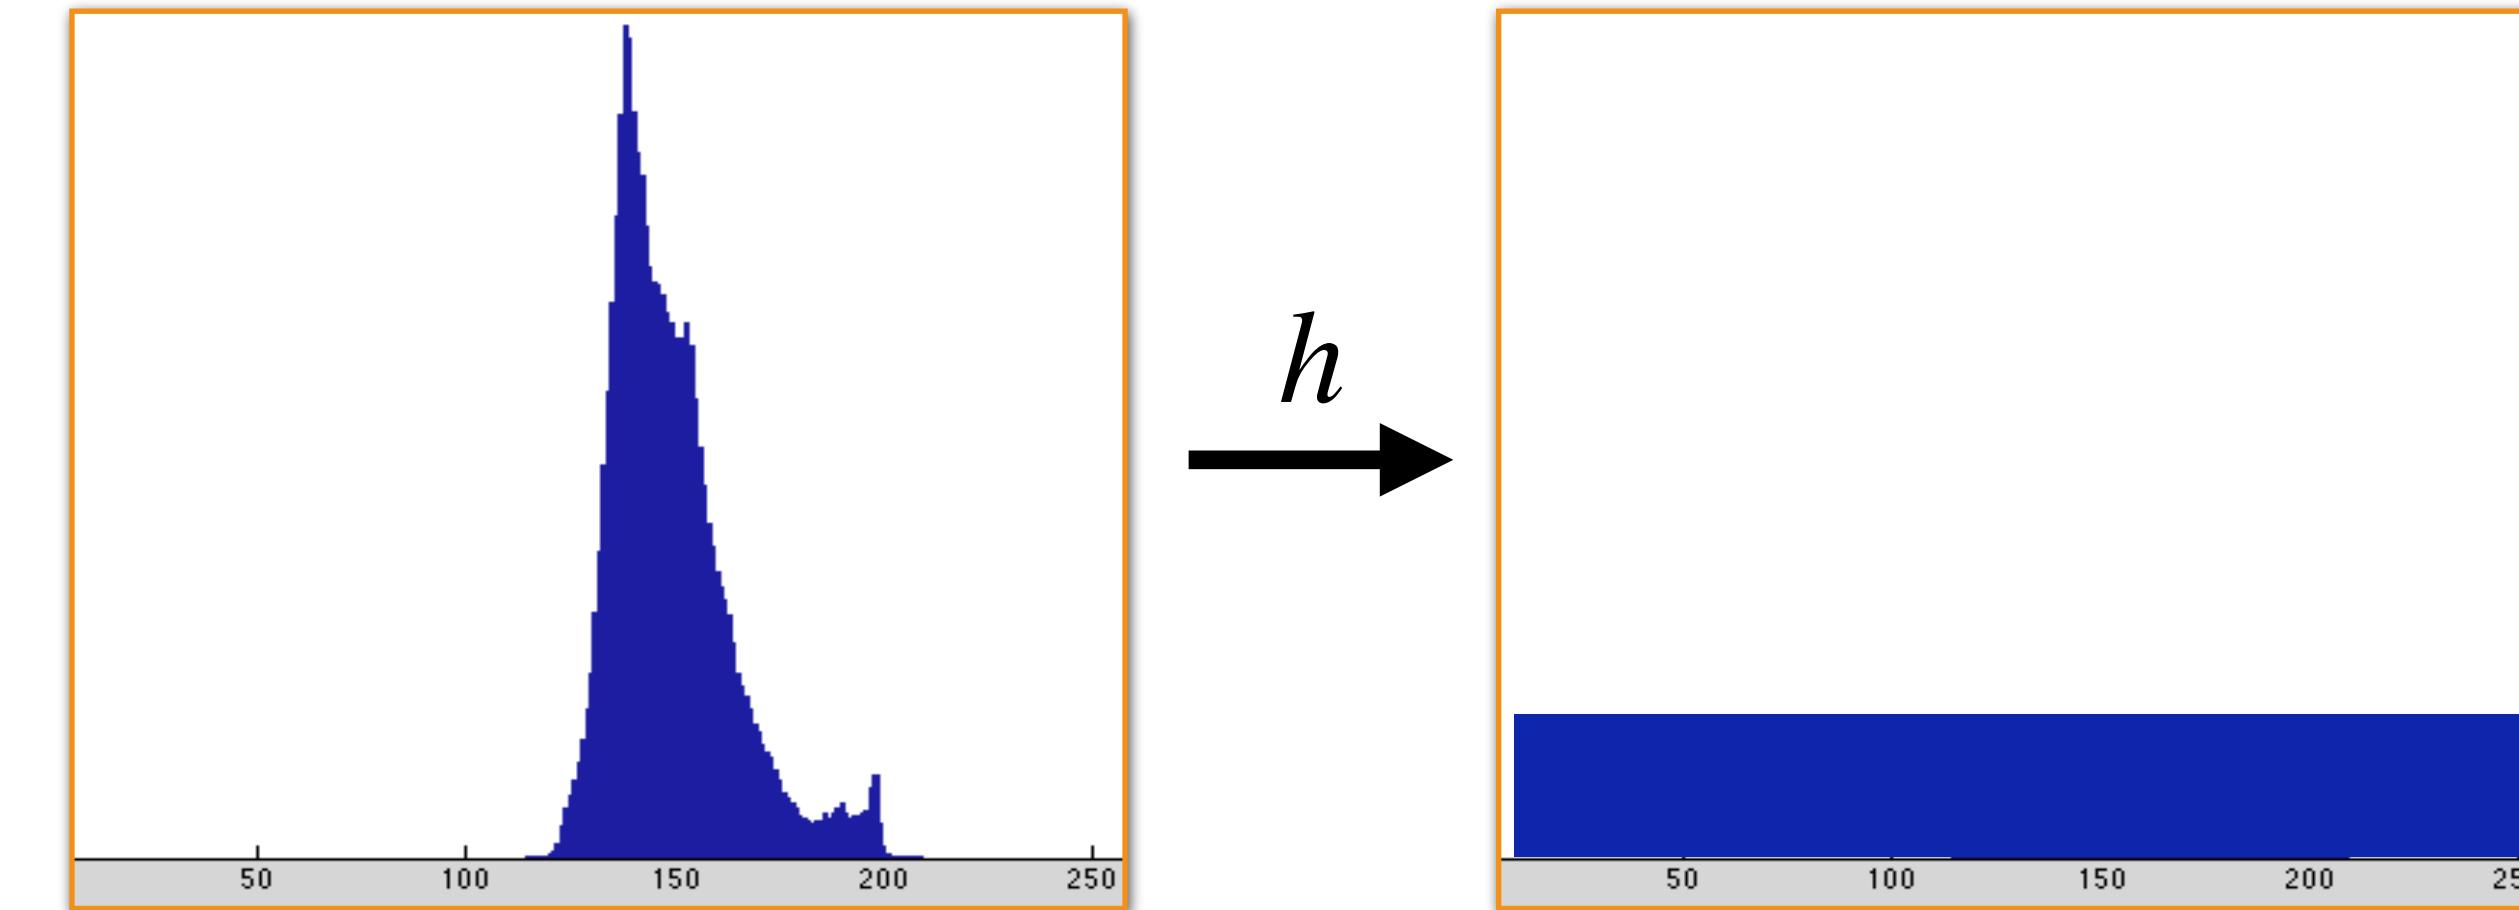
Basic idea: scale the brightness range of the image to occupy the full range of values



$$I \leftarrow \text{floor} \left(\frac{I - \min(I)}{\max(I) - \min(I)} \times 255 \right)$$

Question: When is contrast stretching not effective?

Histogram equalization



Remap data to create a uniform distribution
Why is this good?

https://en.wikipedia.org/wiki/Histogram_equalization

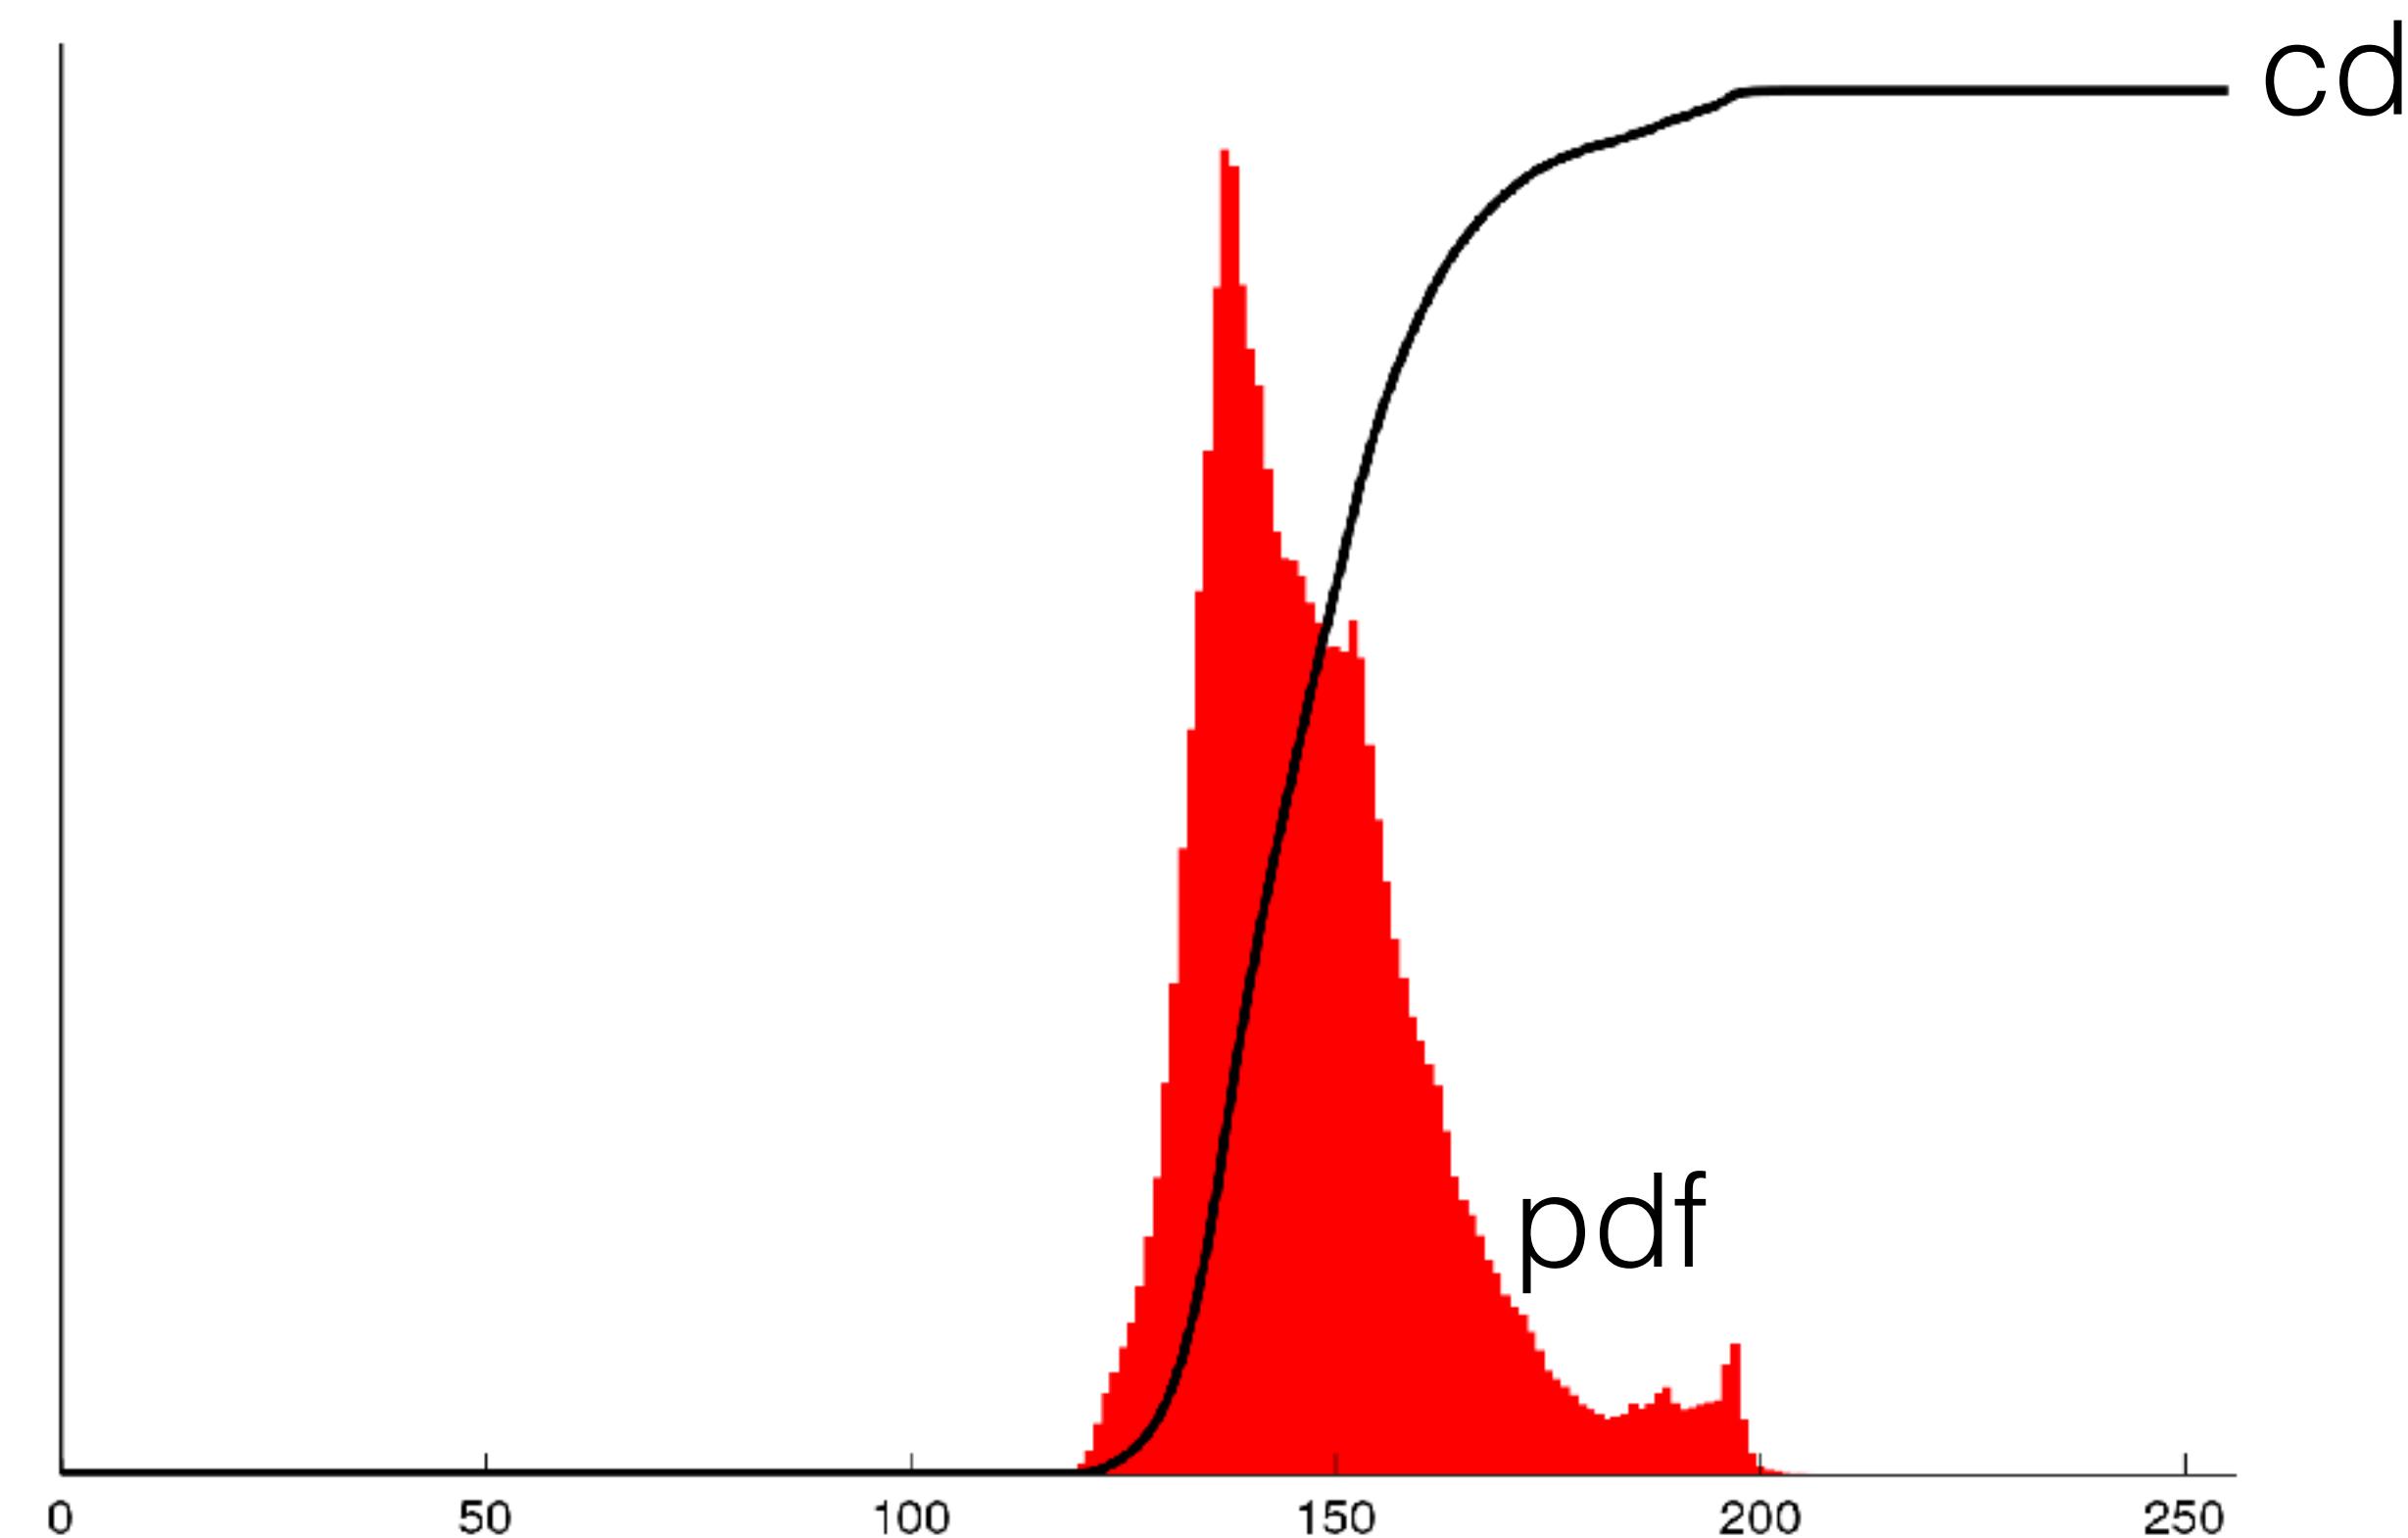
Cumulative distribution function

$\text{pdf}(v) = \#\text{pixels with value } v$

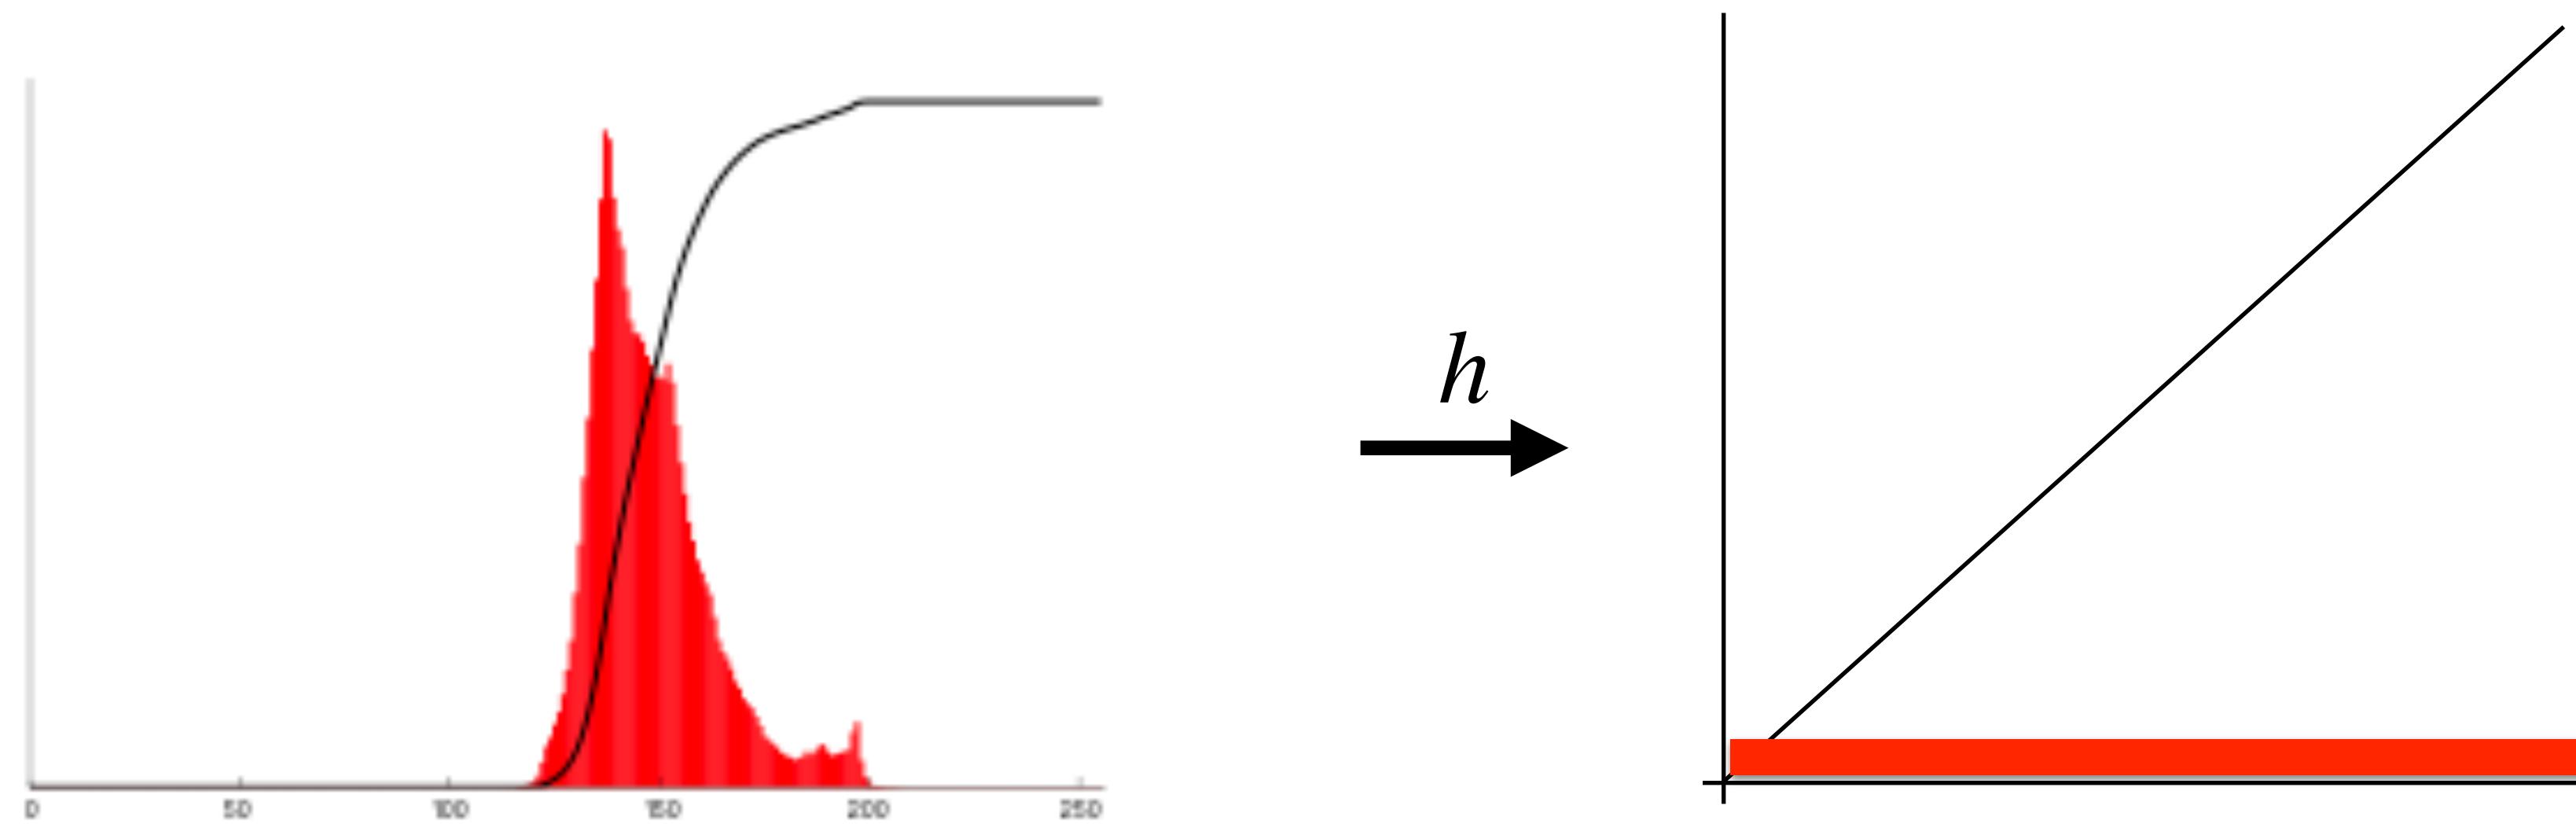
Probability density function (aka histogram)

$\text{cdf}(v) = \#\text{pixels with value } \leq v$

Cumulative distribution function



Histogram equalization ...



What happens to the **cdf** after equalization?
What value should pixels= v be mapped to?

$$h(v) = \text{round} \left(\frac{cdf(v) - cdf_{min}}{(M \times N) - cdf_{min}} \times (L - 1) \right)$$

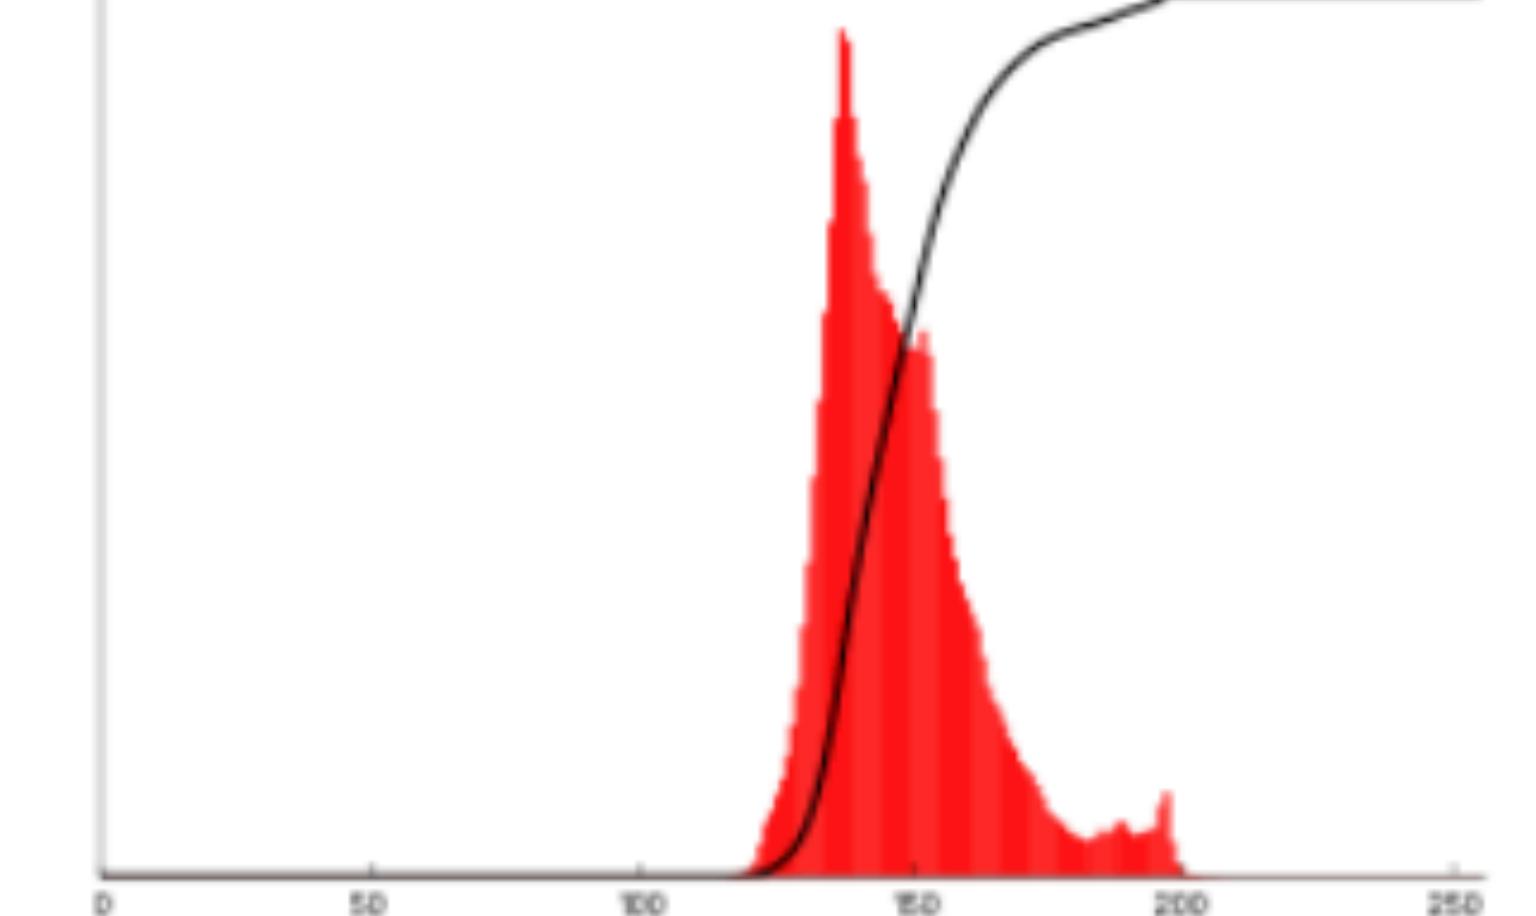
$M \times N \text{ pixels}$
 $L \text{ levels}$



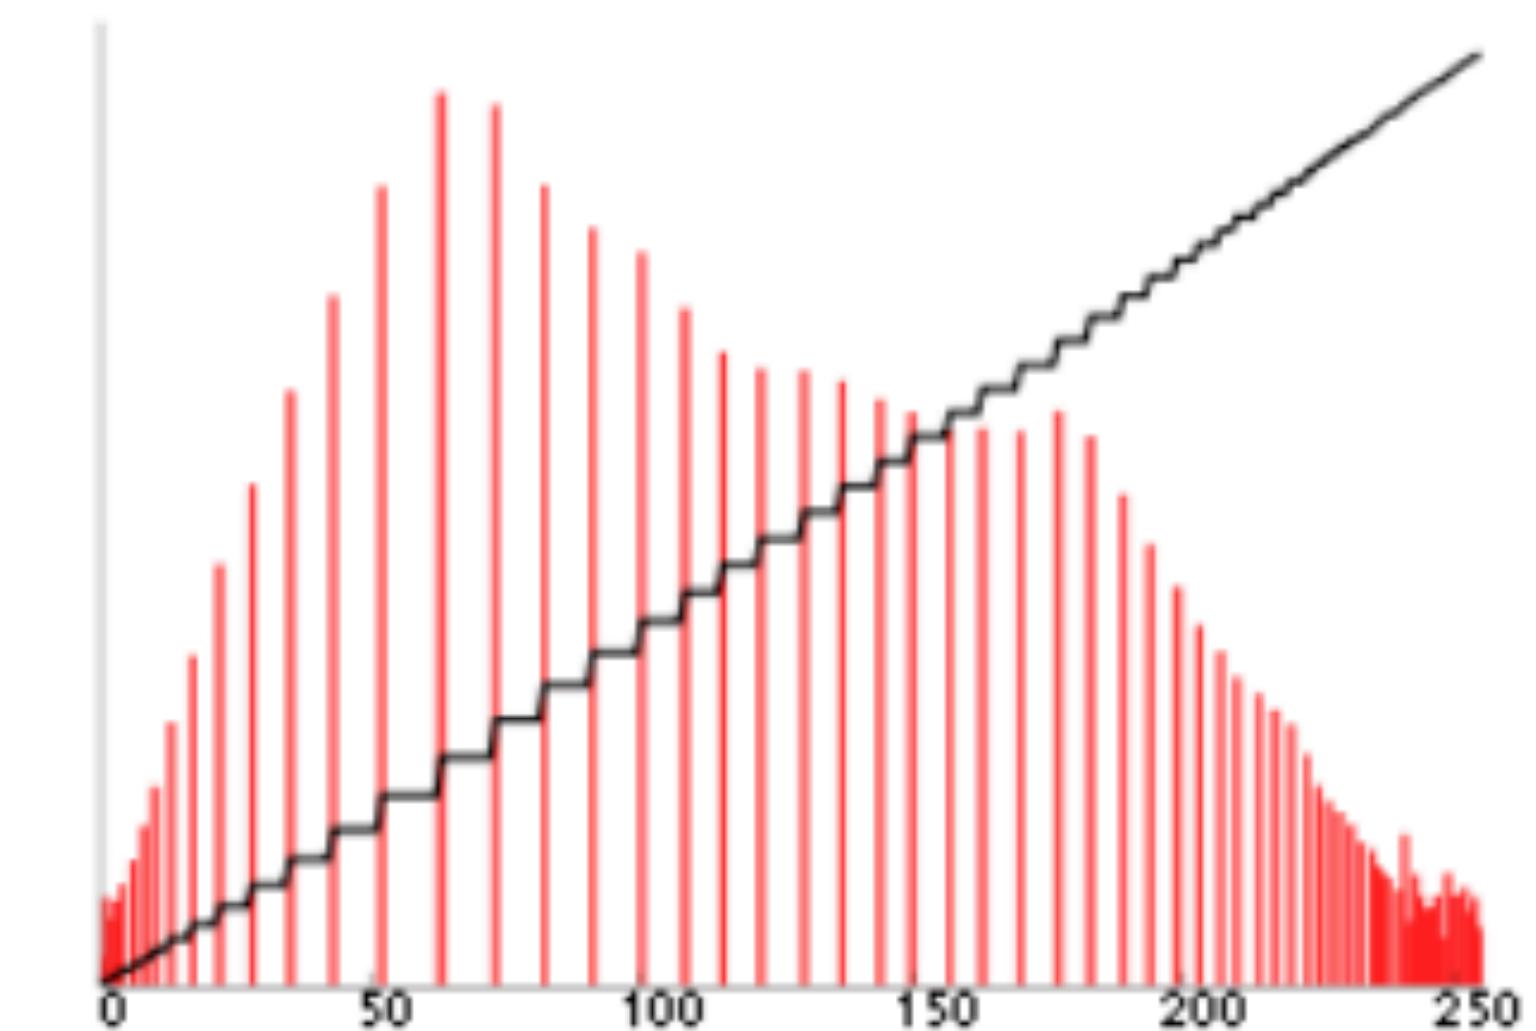
Before Histogram Equalization



After Histogram Equalization



Corresponding histogram (red) and cumulative histogram (black)



Corresponding histogram (red) and cumulative histogram (black)

Denoising

How can we reduce noise in a photograph?



Moving average

Let's replace each pixel with a *weighted* average of its neighborhood

The weights are called the *filter kernel*

Weights for the average of a 3x3 neighborhood

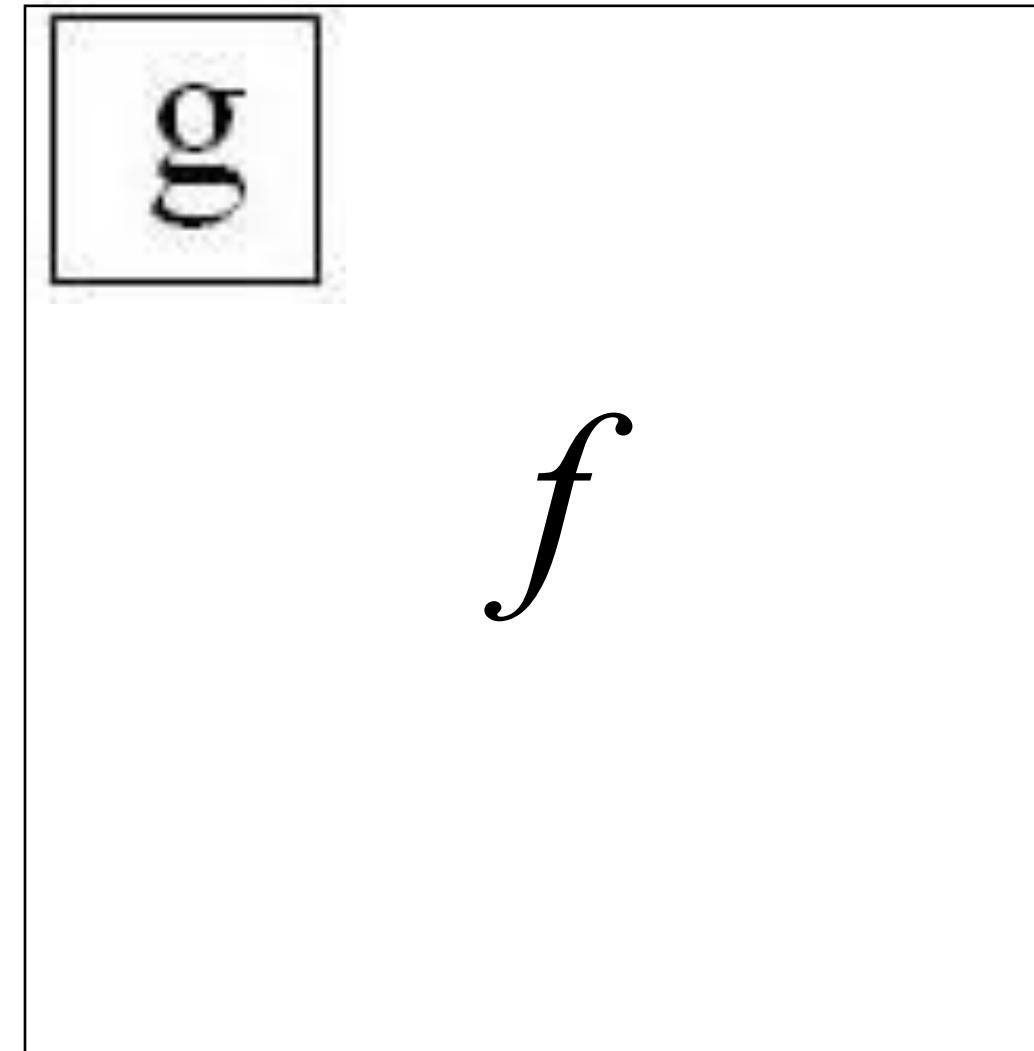
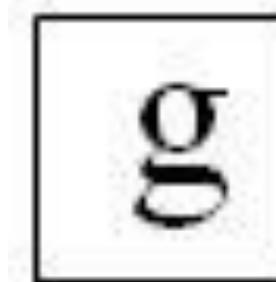
$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

“box filter”

Filtering

Let f be the image and g be the kernel. The output of filtering f with g denoted $f * g$ is given by:

$$(f * g)[m, n] = \sum_{k,l} f[m + k, n + l]g[k, l]$$



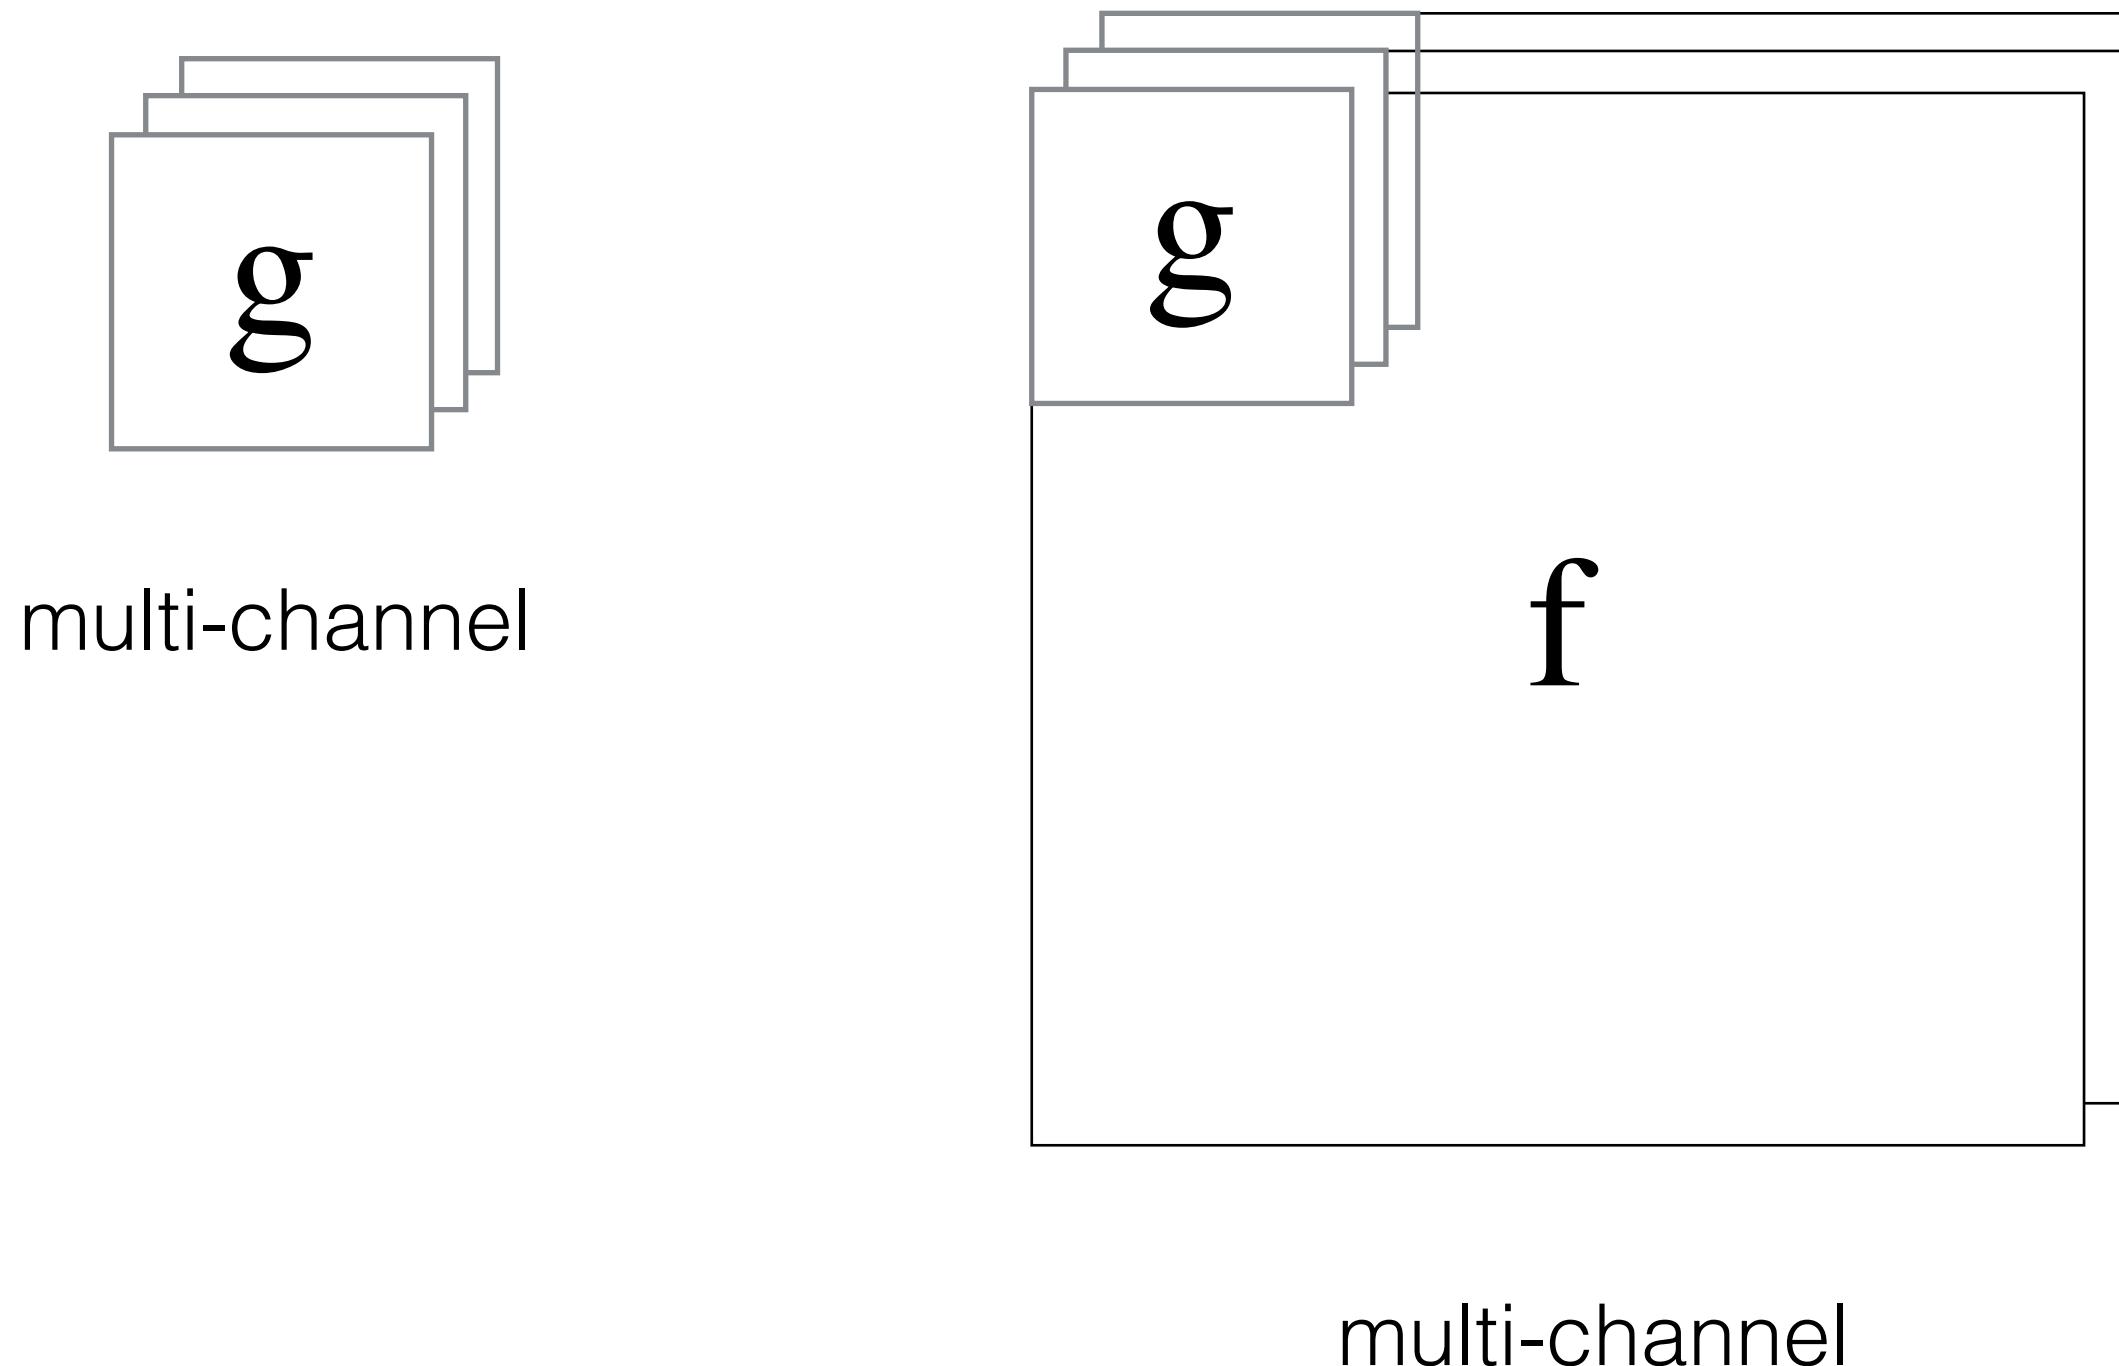
Filtering computes the correlation between the g and f at each location

Convolution is filtering with a flipped g (by notation)

Filtering: multi-channel case

Let f be the image and g be the kernel. The output of filtering f with g denoted $f * g$ is given by:

$$(f * g)[m, n] = \sum_{k, l, c} f[m + k, n + l, c]g[k, l, c]$$



Key properties

Linearity: $\text{filter}(f_1 + f_2) = \text{filter}(f_1) + \text{filter}(f_2)$

Shift invariance: same behavior regardless of pixel location: $\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$

Theoretical result: any linear shift-invariant operator can be represented as a convolution

Properties in more detail

Commutative: $a * b = b * a$

- ◆ Conceptually no difference between filter and signal

Associative: $a * (b * c) = (a * b) * c$

- ◆ Often apply several filters one after another: $((a * b_1) * b_2) * b_3$

- ◆ This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$

Distributes over addition: $a * (b + c) = (a * b) + (a * c)$

Scalars factor out: $ka * b = a * kb = k(a * b)$

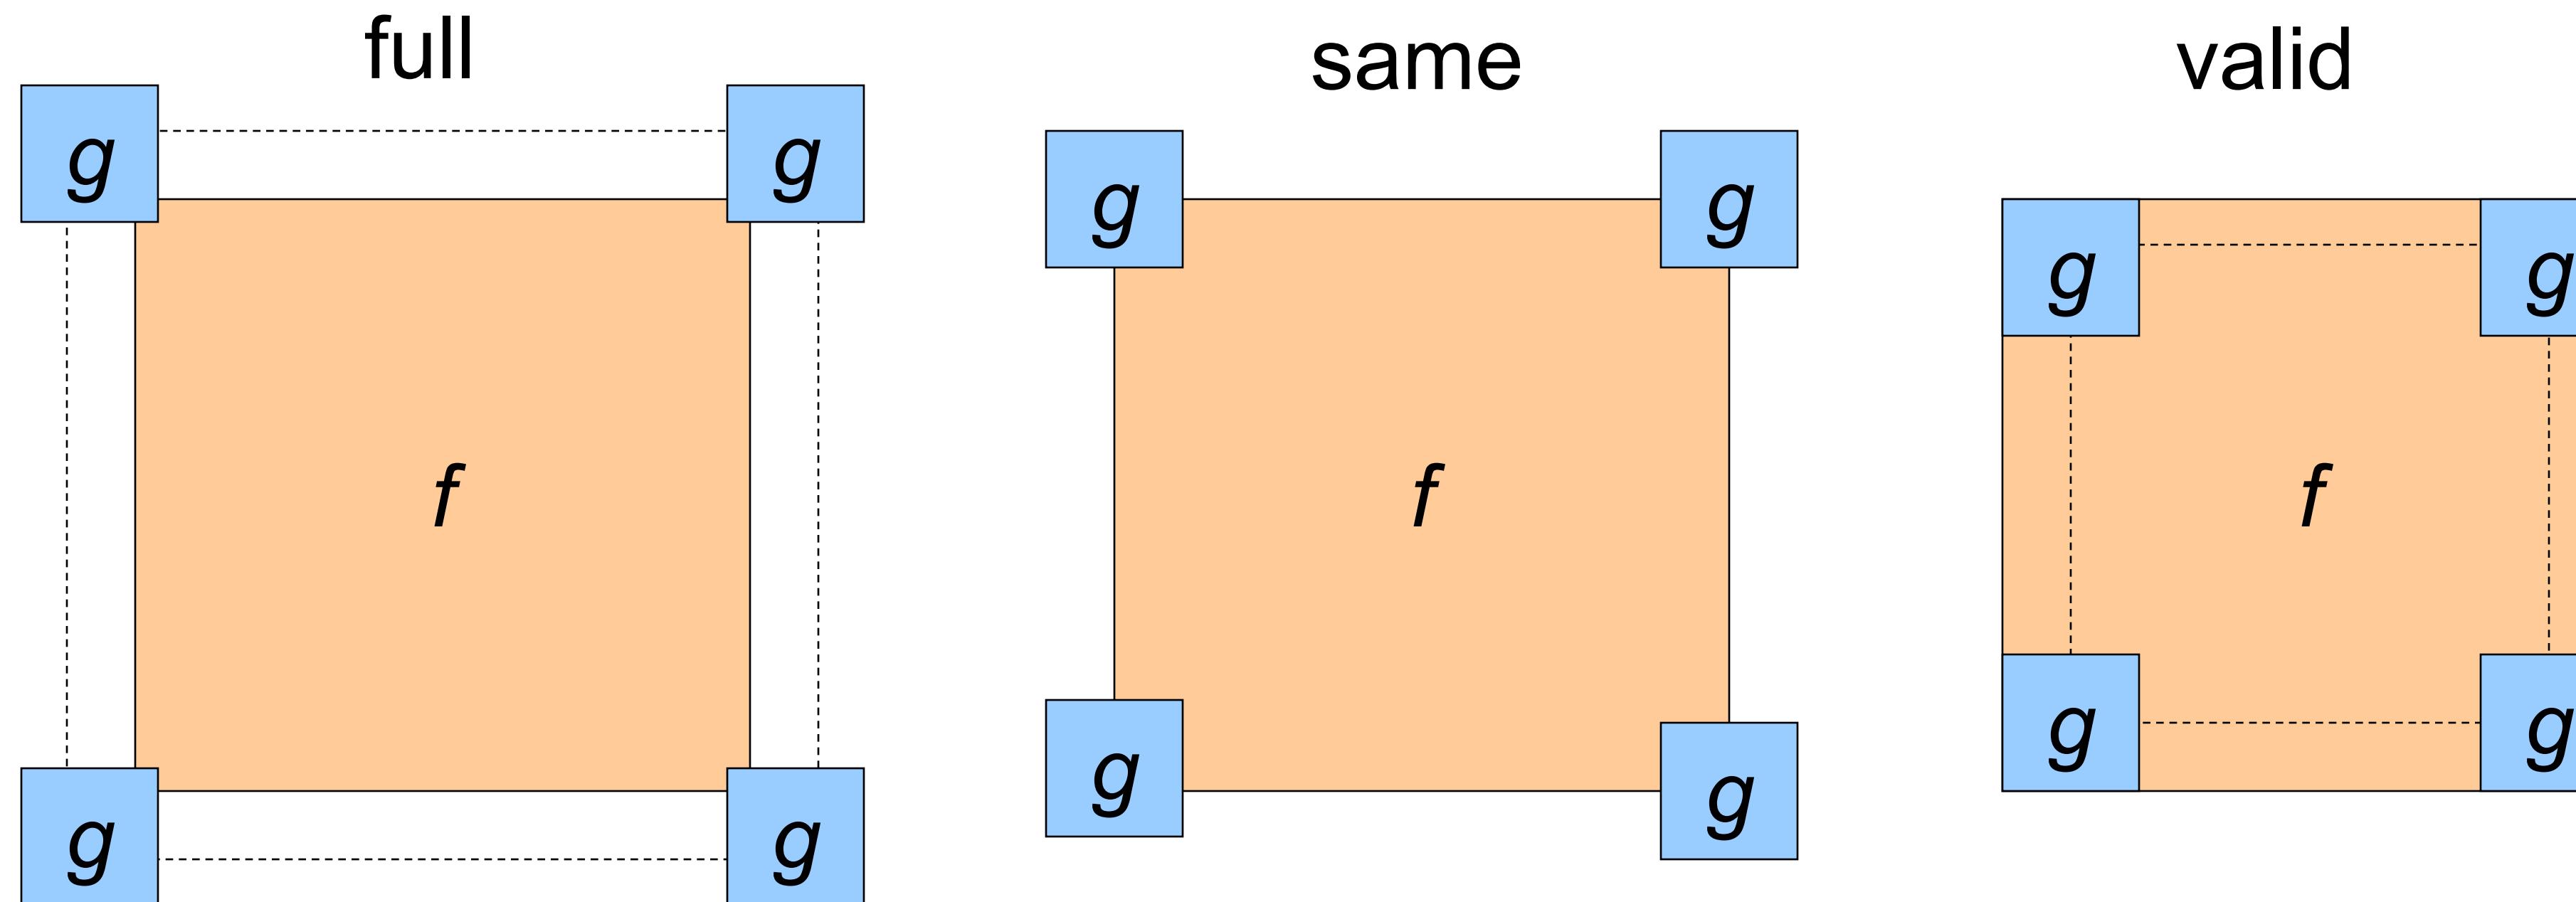
Identity: unit impulse $e = [..., 0, 0, 1, 0, 0, ...]$, $a * e = a$

Annoying details

What is the size of the output?

Python: `scipy.ndimage.correlate / convolve`

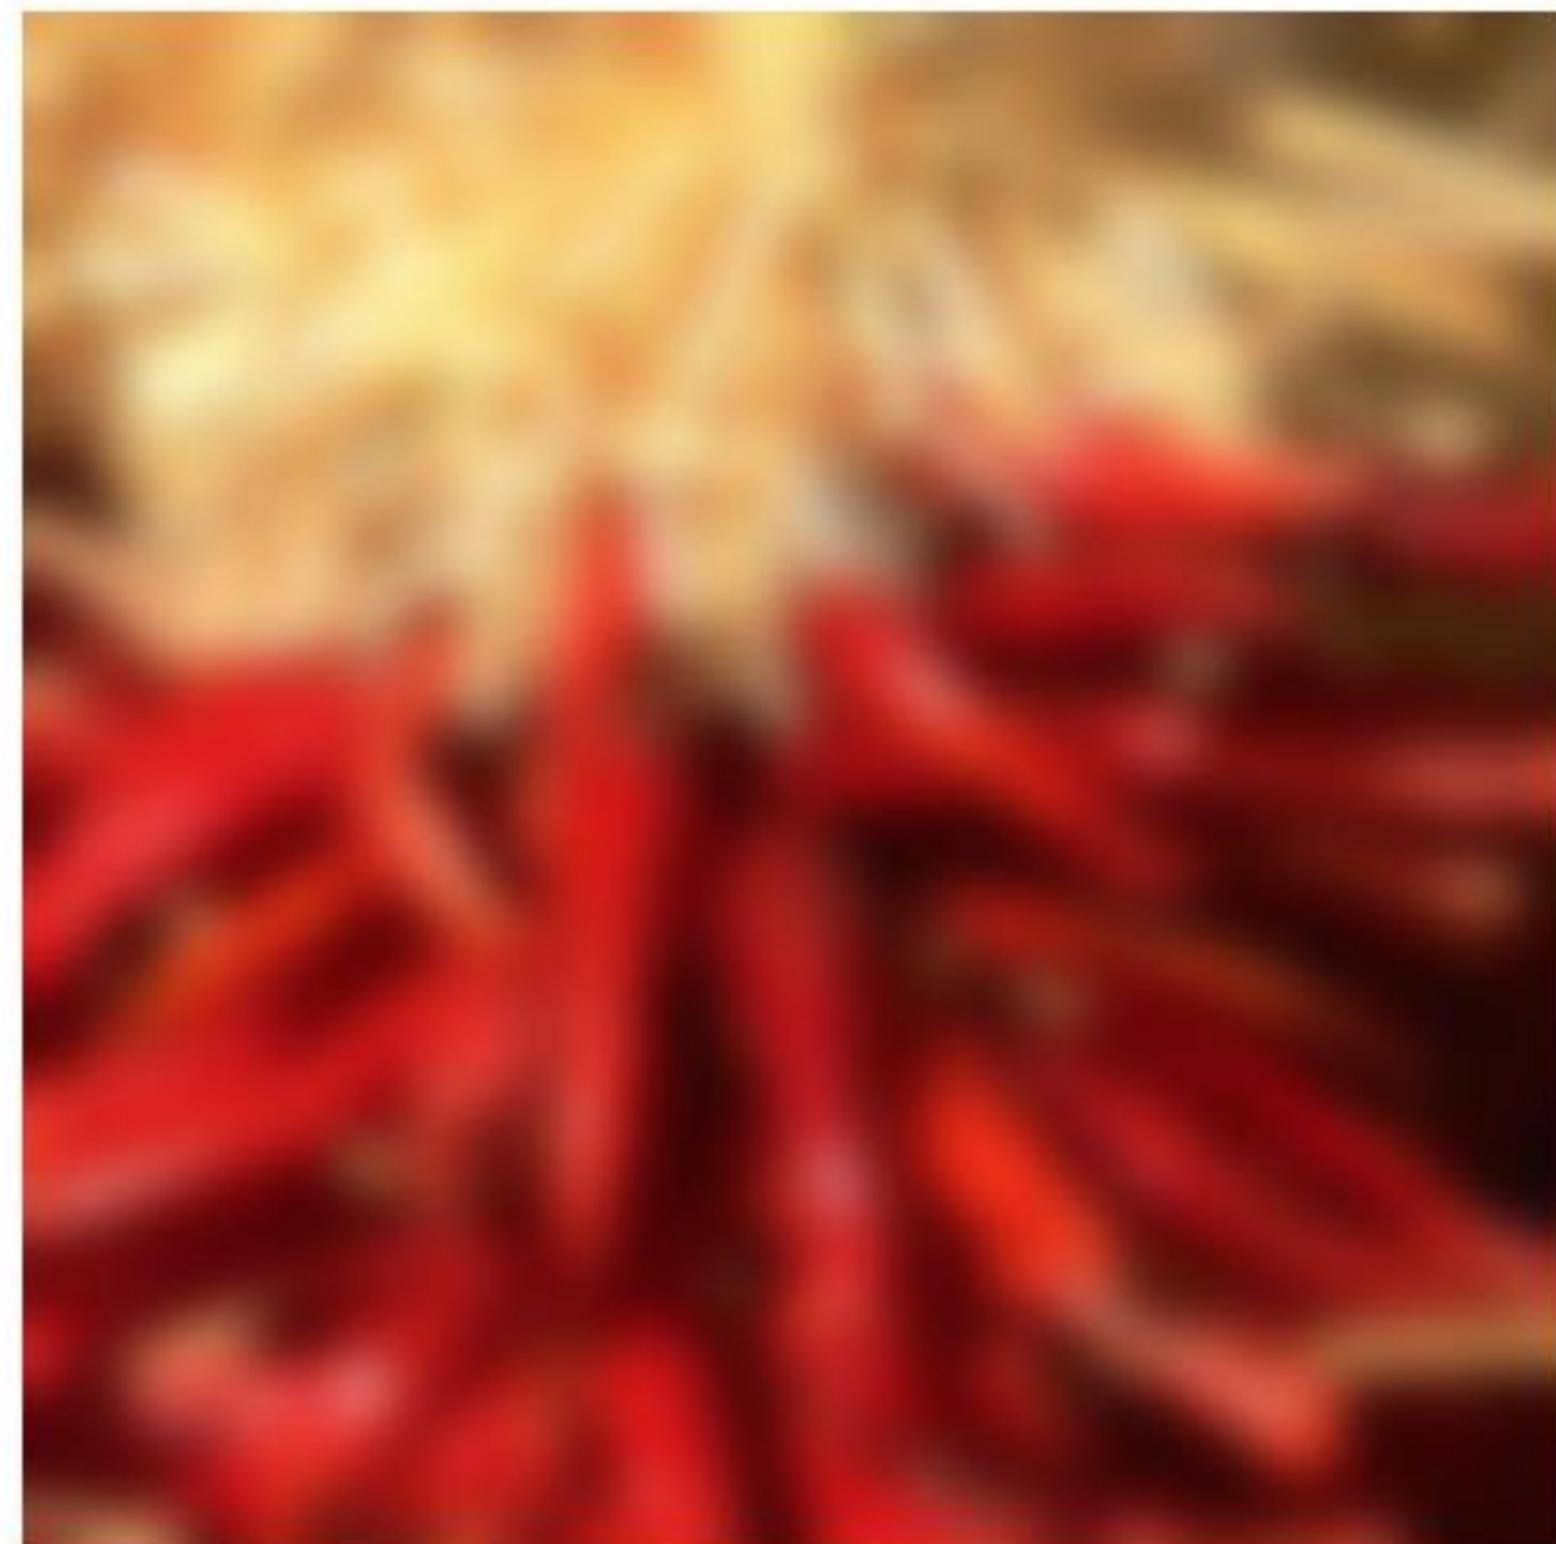
- `shape = 'full'`: output size is sum of sizes of f and g
- `shape = 'same'`: output size is same as f
- `shape = 'valid'`: output size is difference of sizes of f and g



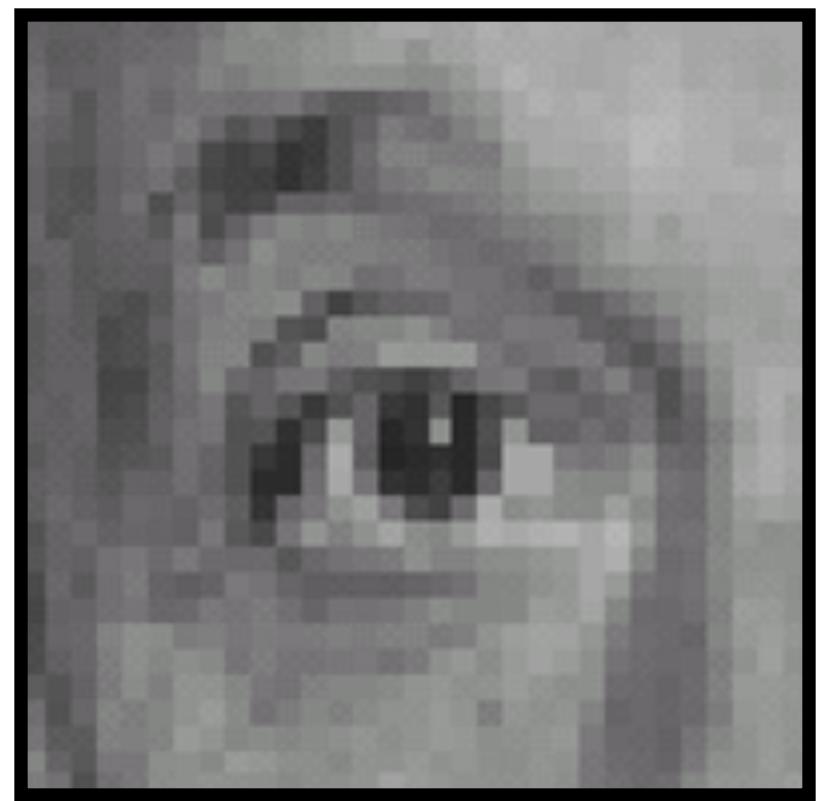
Annoying details

What about near the edge?

- the filter window falls off the edge of the image
- need to extrapolate
- methods:
 - clip filter (black) — `correlate(f, g, mode='constant', cval=0.0)`
 - wrap around — `correlate(f, g, mode='wrap')`
 - copy edge — `correlate(f, g, mode='nearest')`
 - reflect across edge — `correlate (f, g, mode='reflect')`



Practice with linear filters

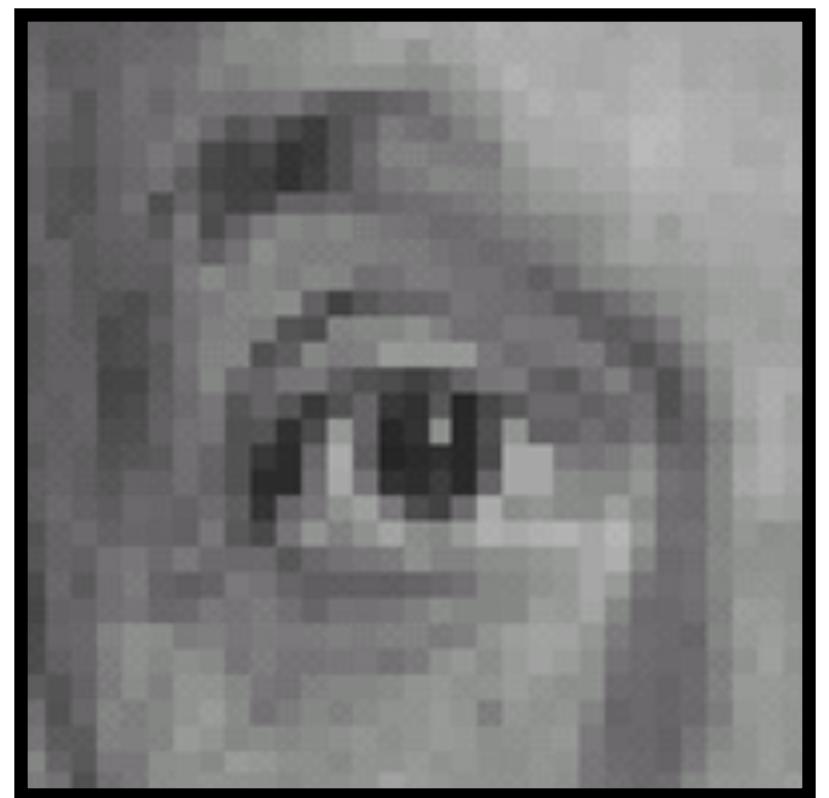


Original

0	0	0
0	1	0
0	0	0

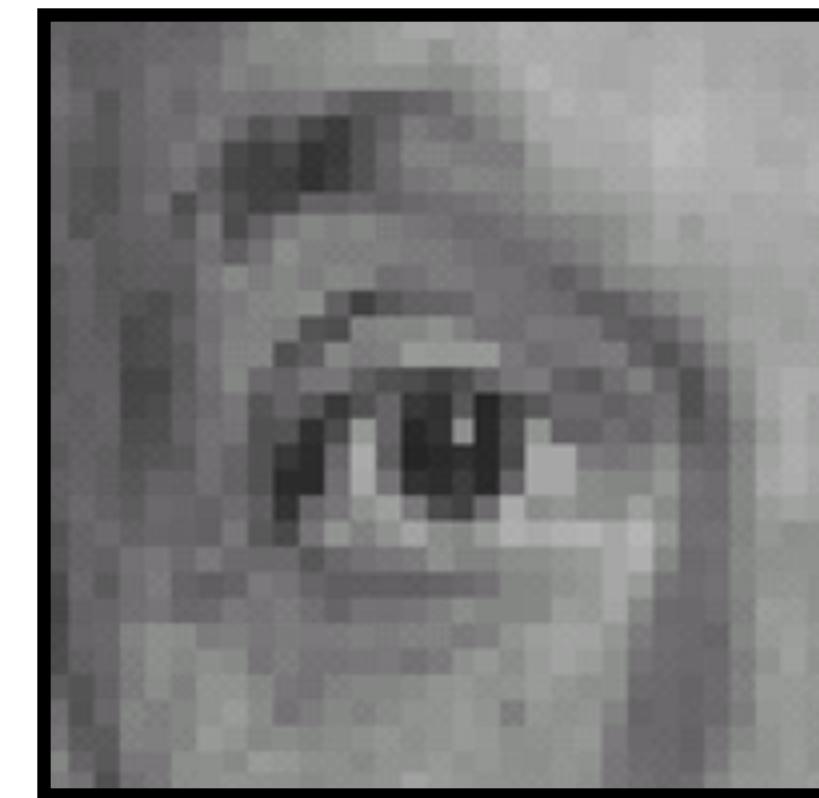
?

Practice with linear filters



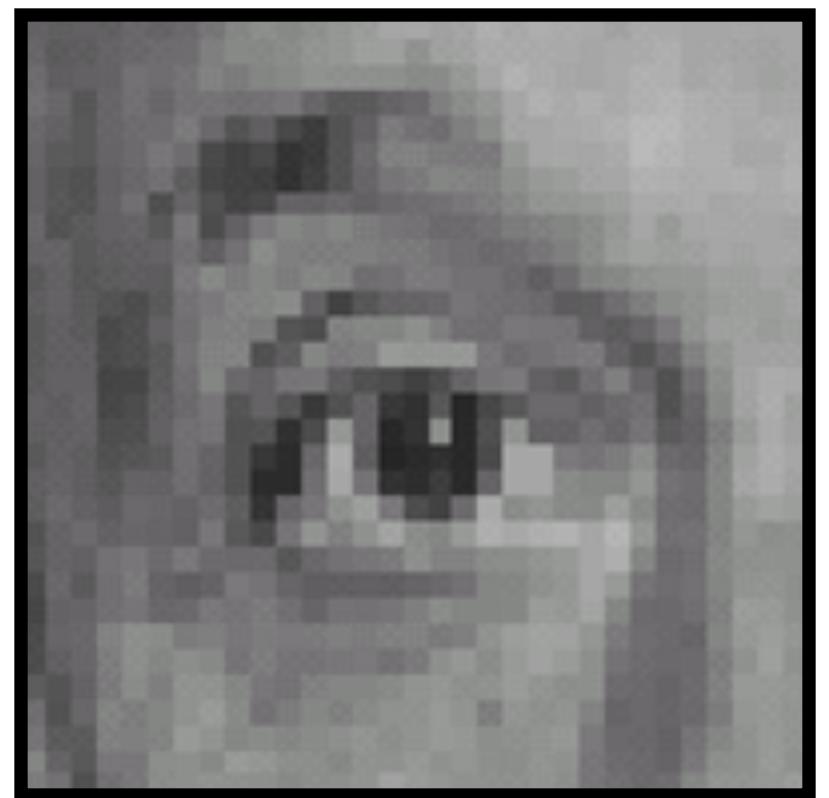
Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Practice with linear filters

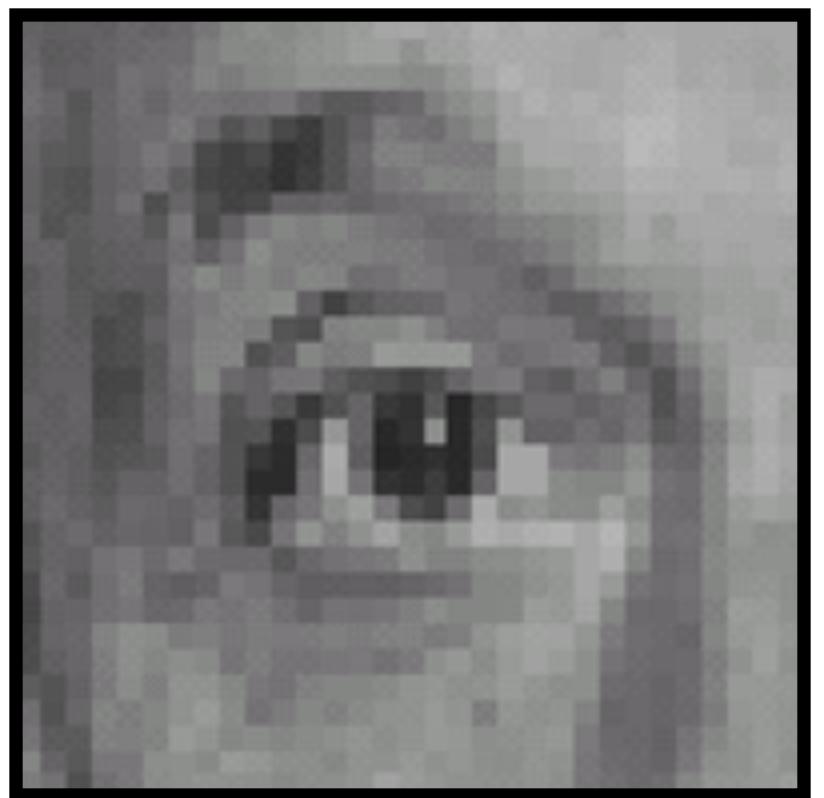


Original

0	0	0
0	0	1
0	0	0

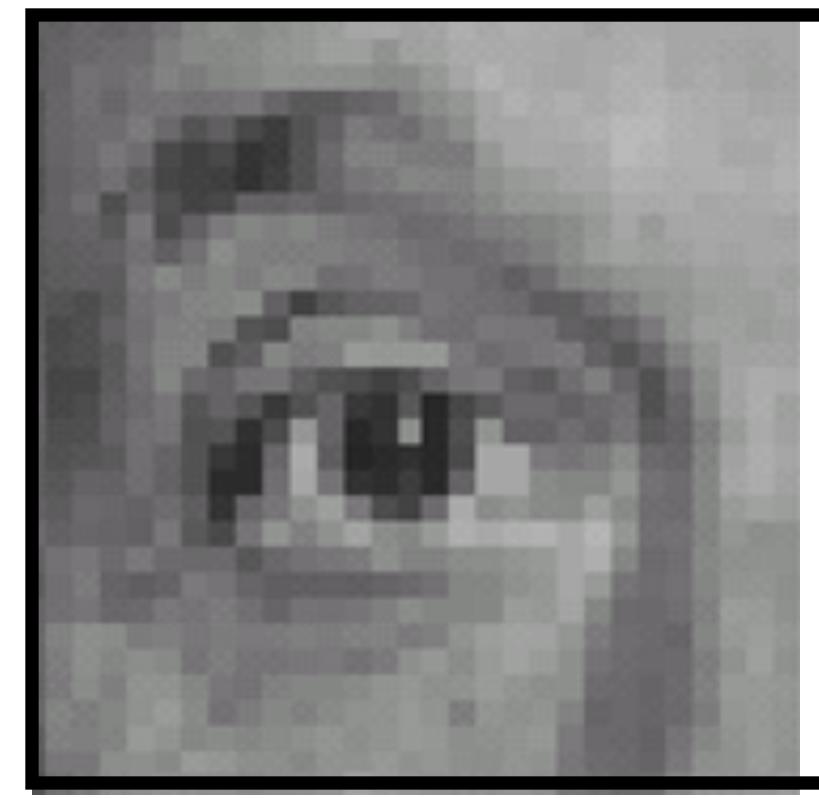
?

Practice with linear filters



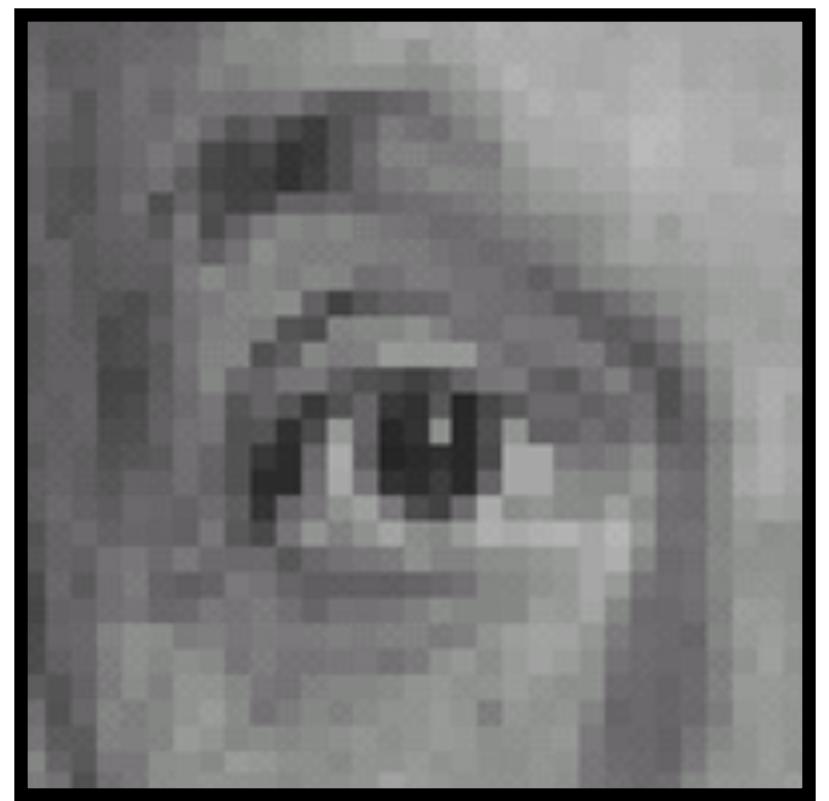
Original

0	0	0
0	0	1
0	0	0



Shifted *left*
By 1 pixel

Practice with linear filters

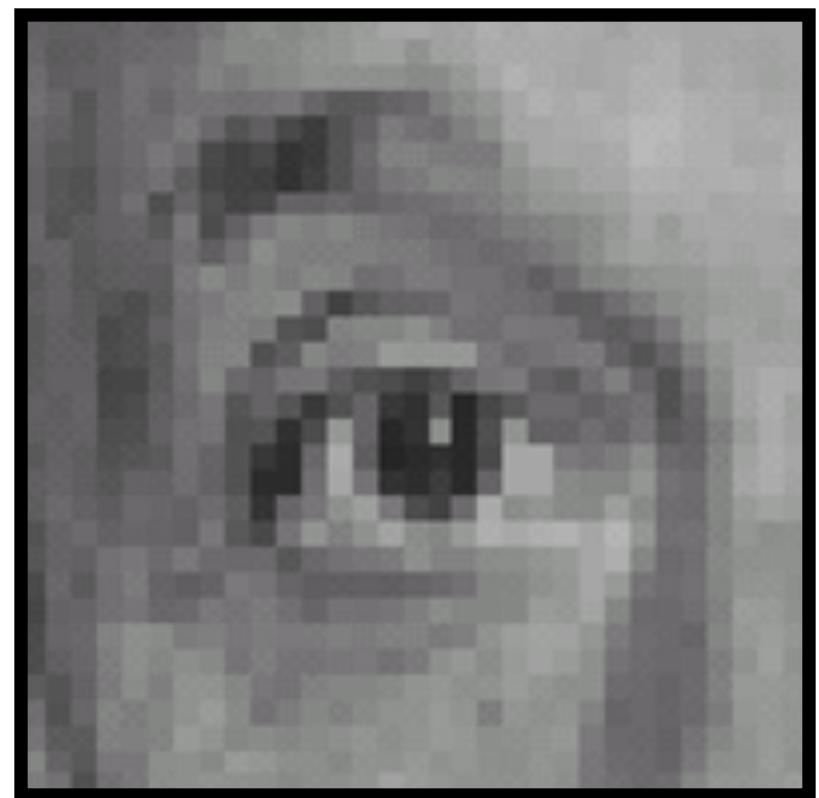


$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

?

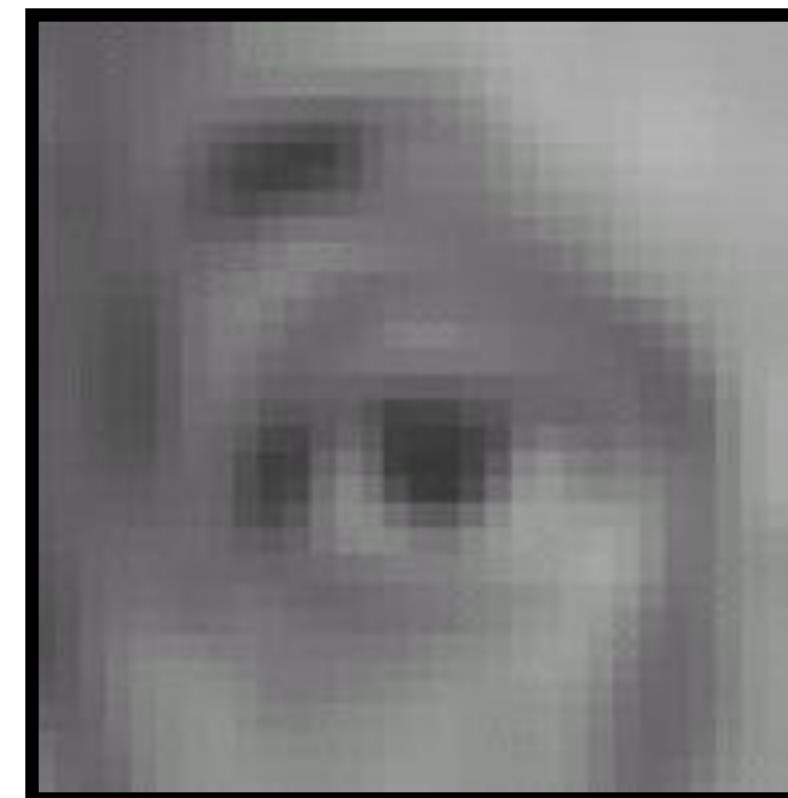
Original

Practice with linear filters



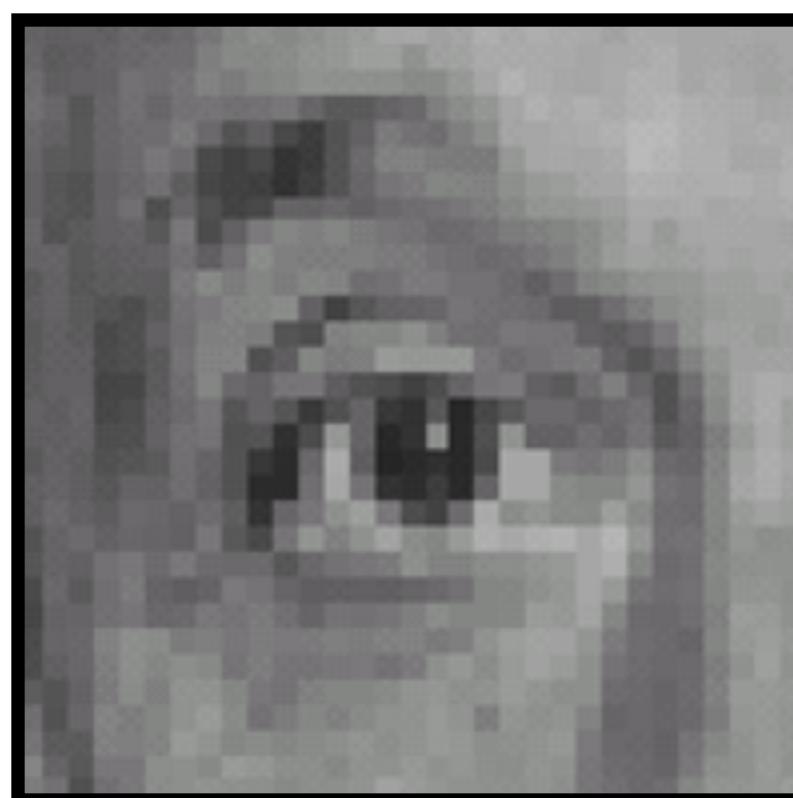
Original

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$



Blur (with a
box filter)

Practice with linear filters



Original

$$\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix}$$

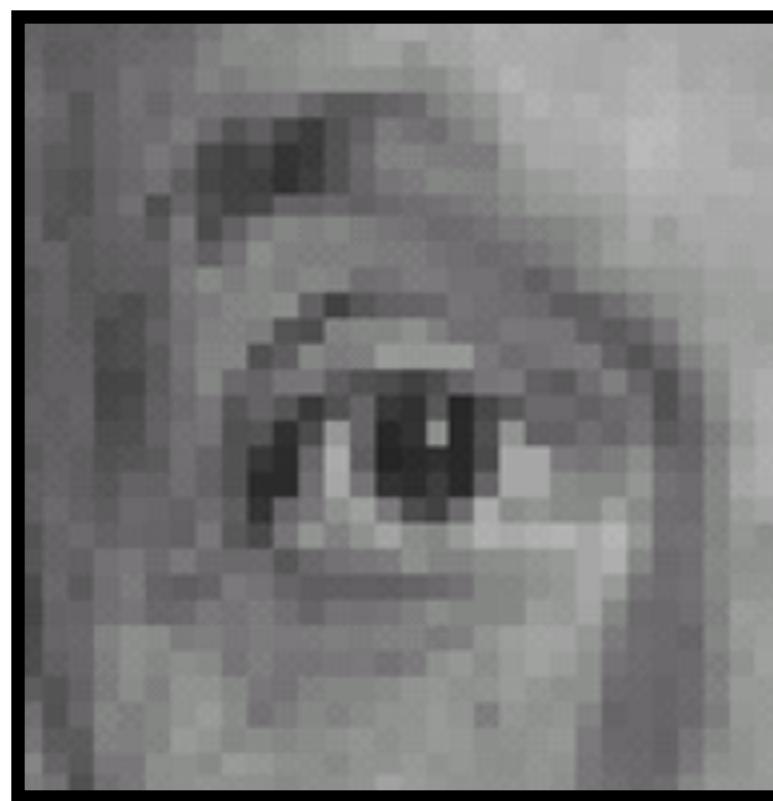
-

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

?

(Note that filter sums to 1)

Practice with linear filters



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix}$$

-

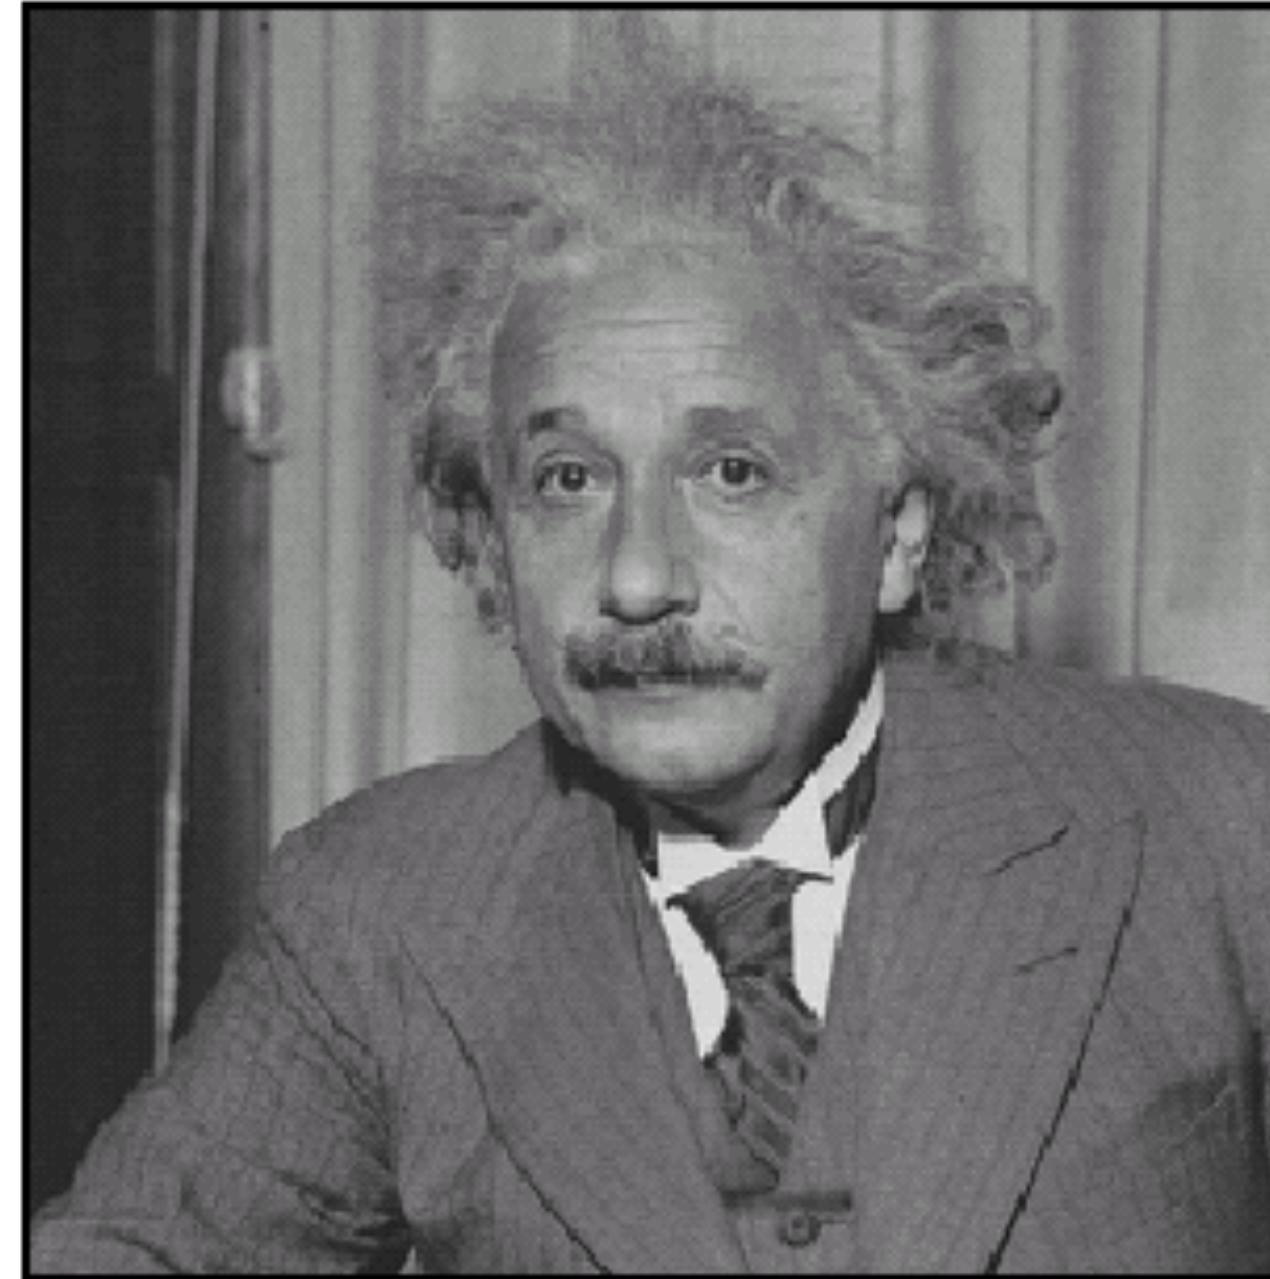
$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



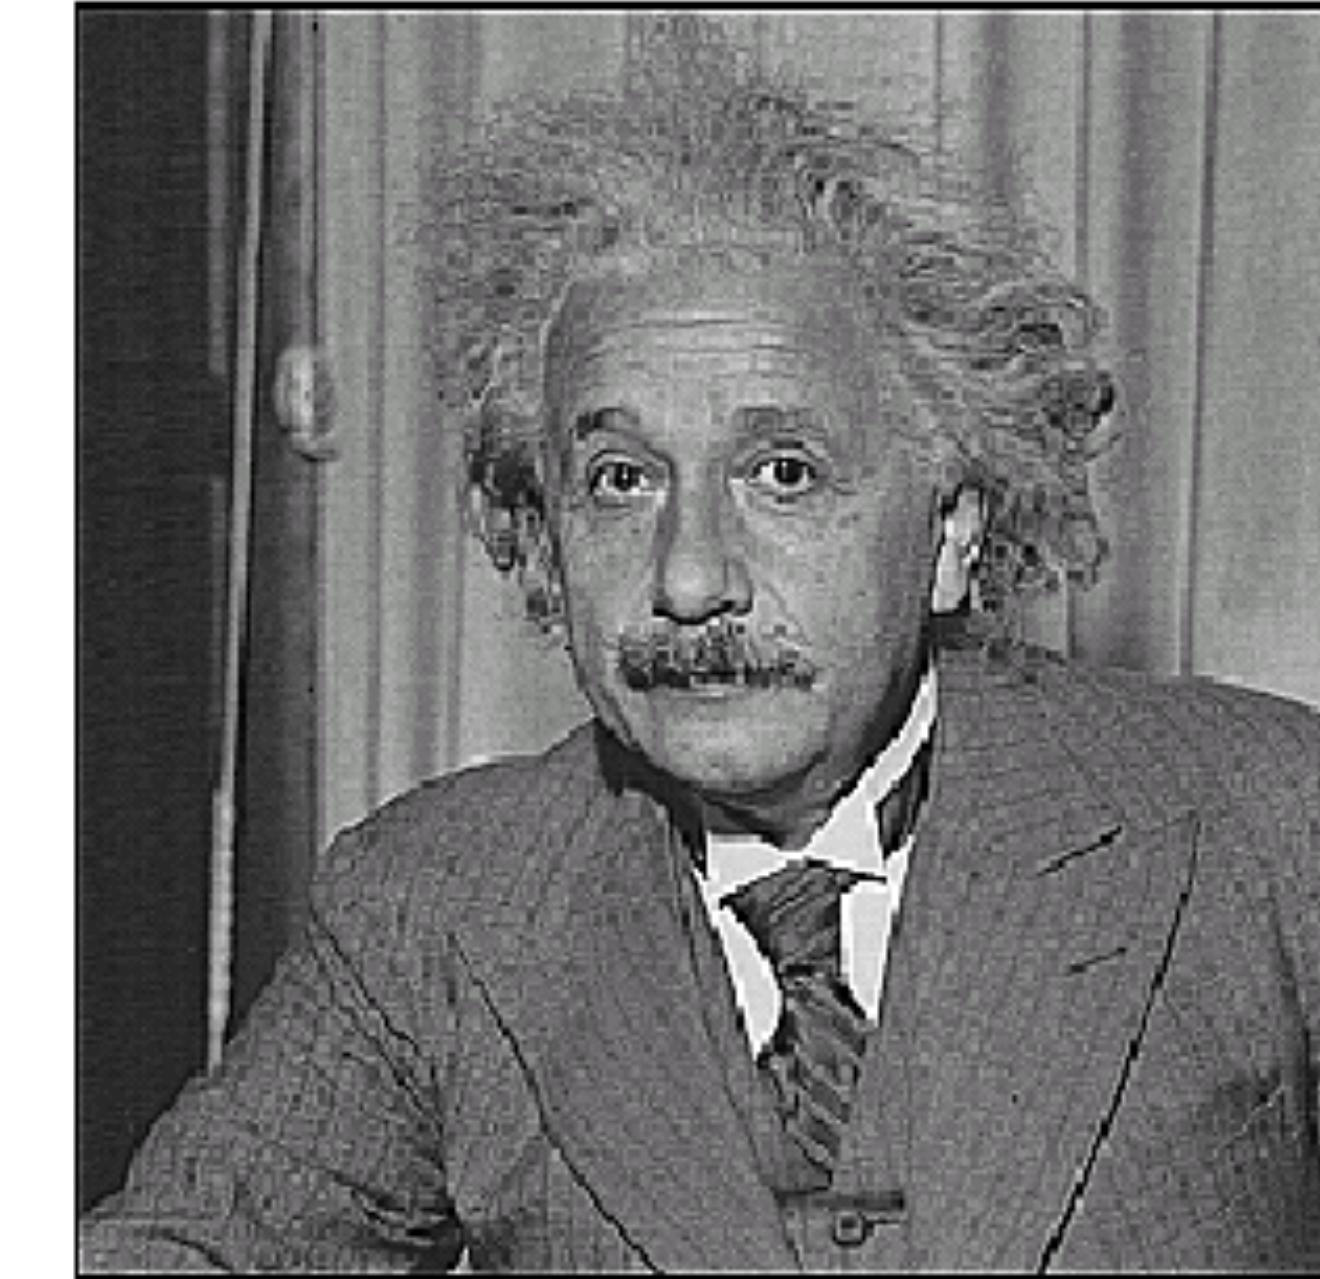
Original

Sharpening filter: accentuates differences with local average

Sharpening



before

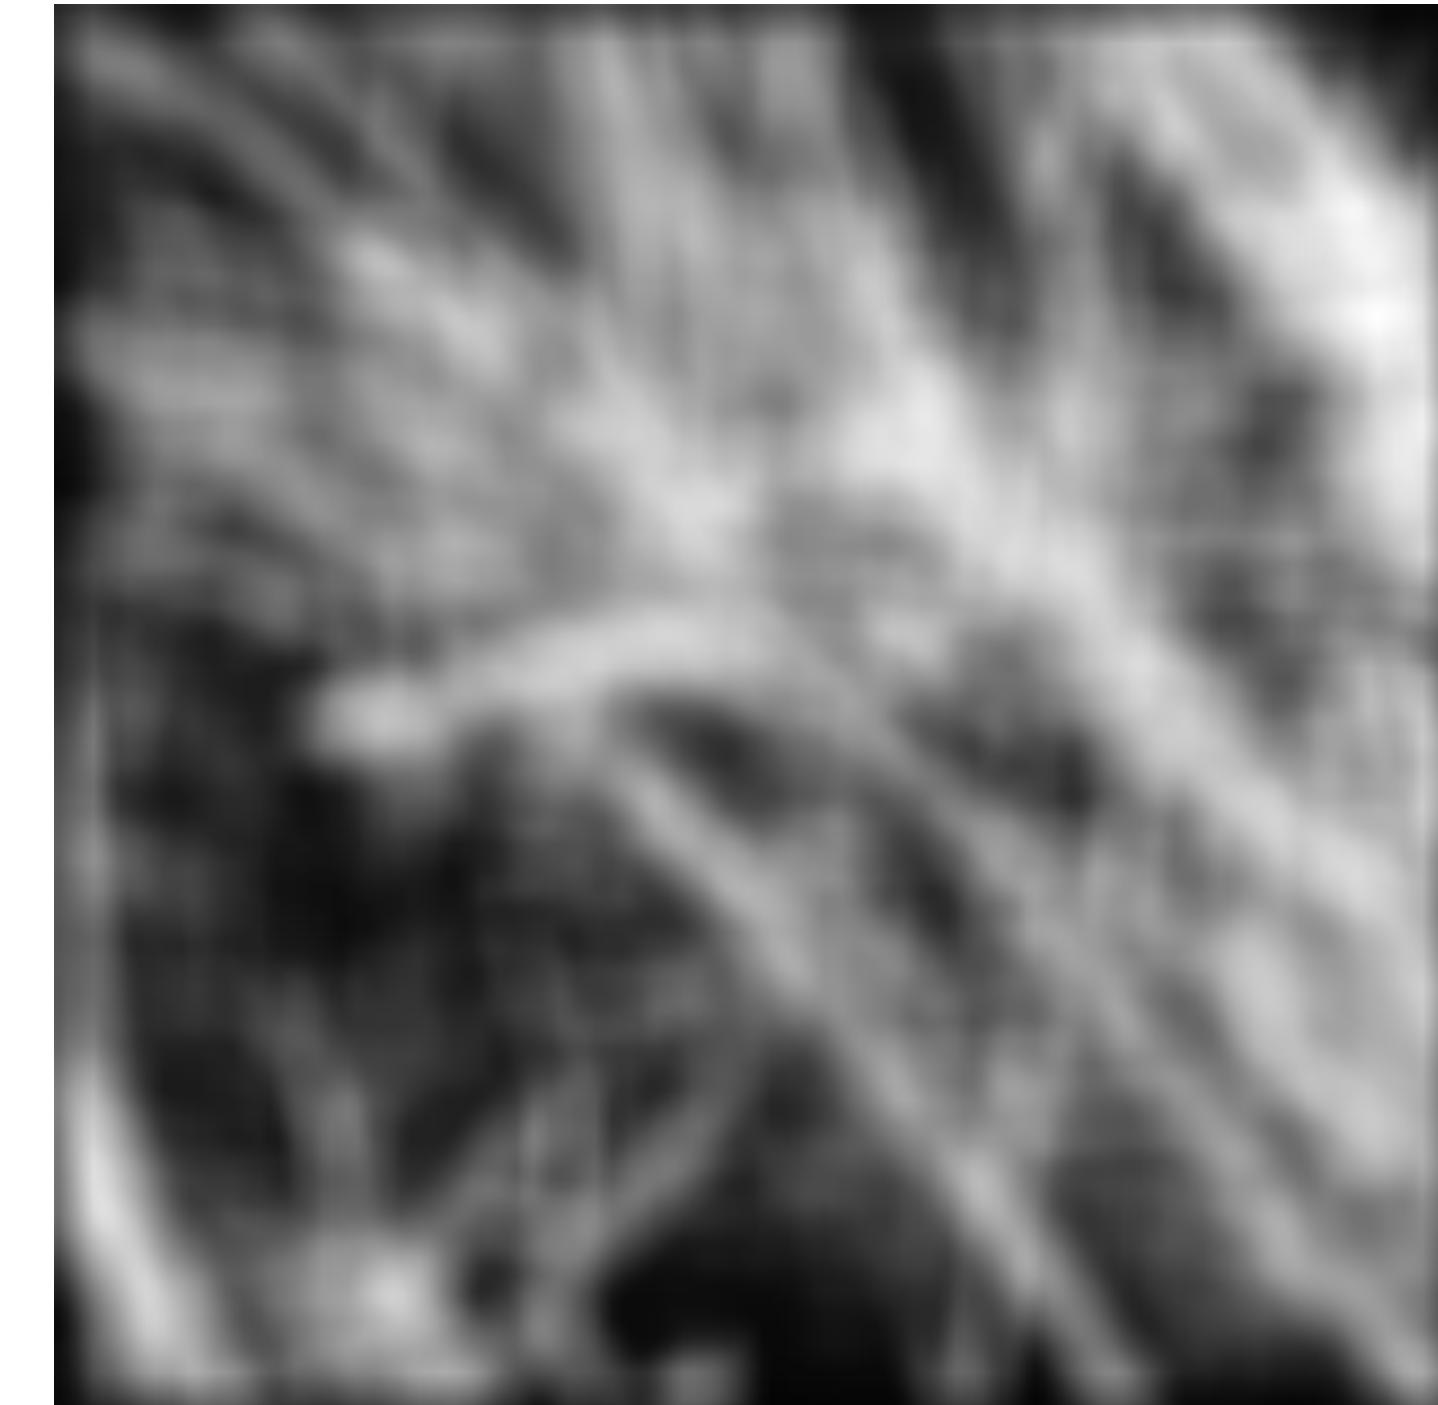


after

Smoothing with box filter revisited

What's wrong with this picture?

What's the solution?

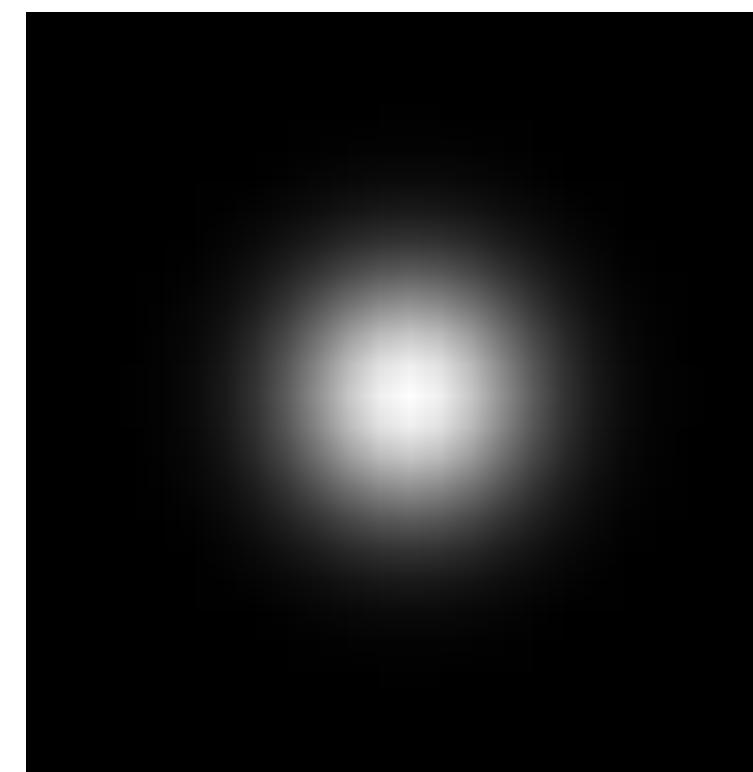


Smoothing with box filter revisited

What's wrong with this picture?

What's the solution?

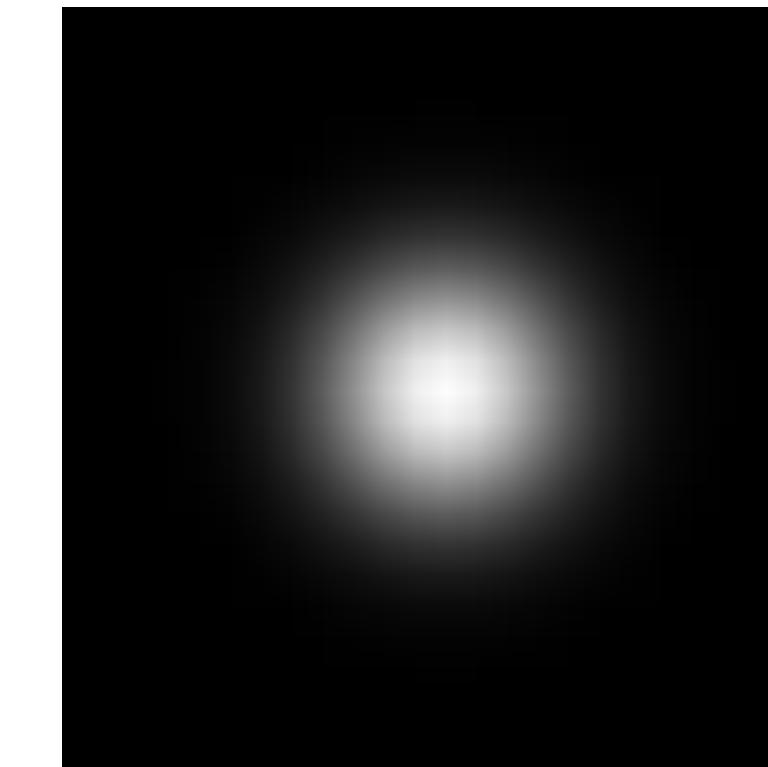
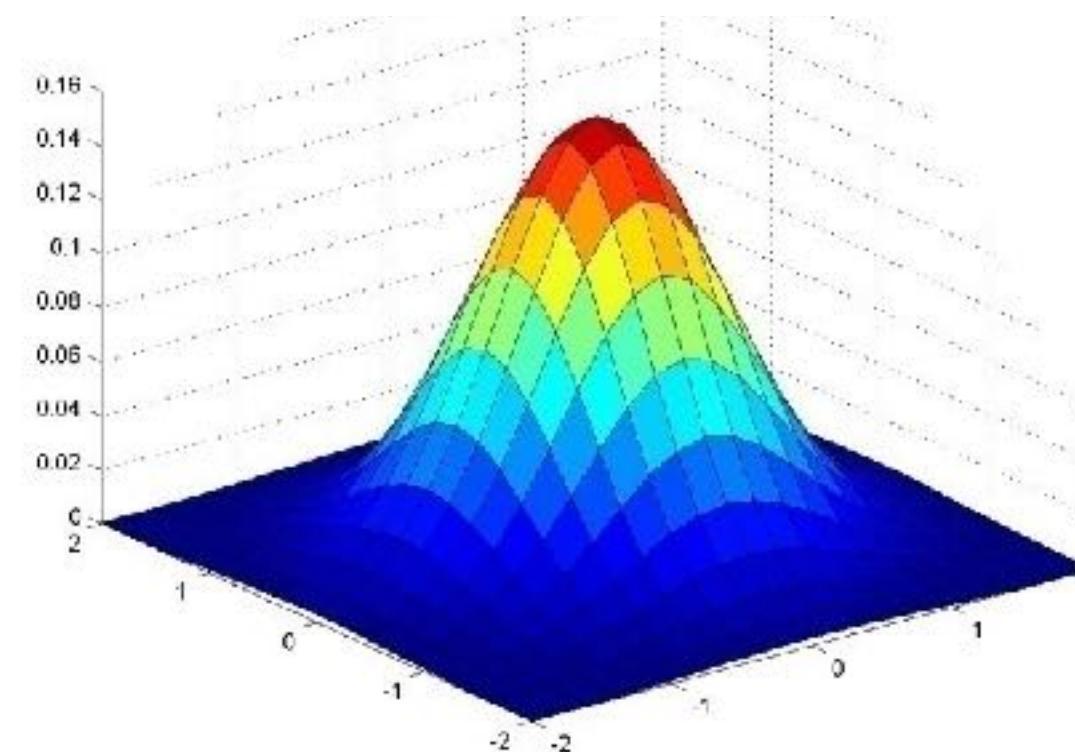
- To eliminate edge effects, weight contribution of neighborhood pixels according to their closeness to the center



“fuzzy blob”

Gaussian kernel

Constant factor at front makes volume sum to 1 (can be ignored when computing the filter values, as we should *renormalize* weights to sum to 1 in any case)



0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

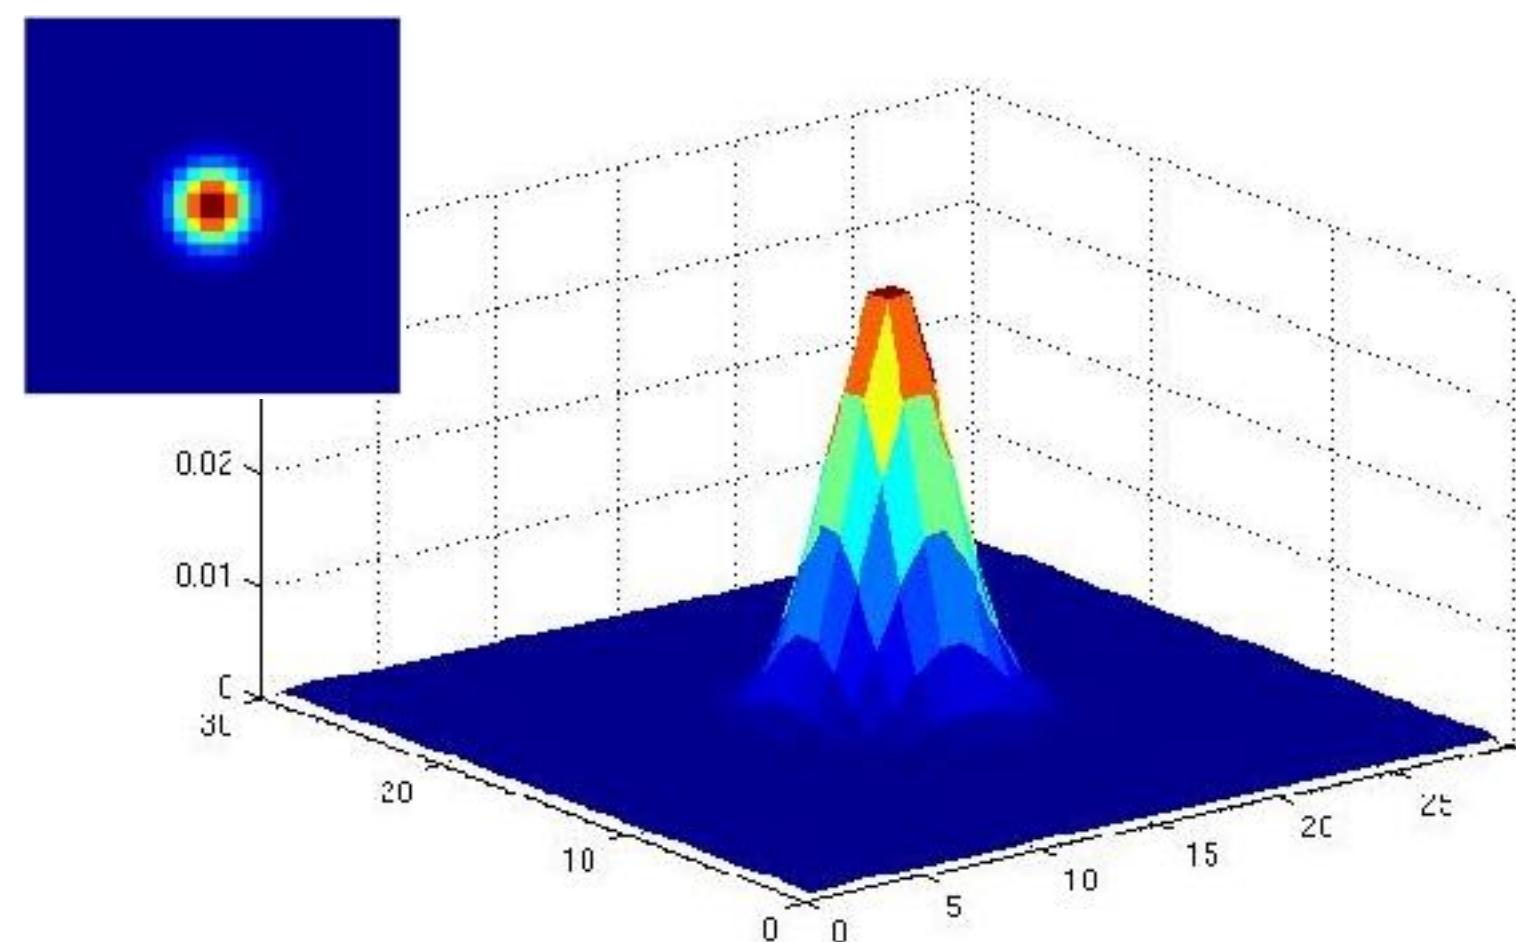
5 × 5, $\sigma = 1$

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

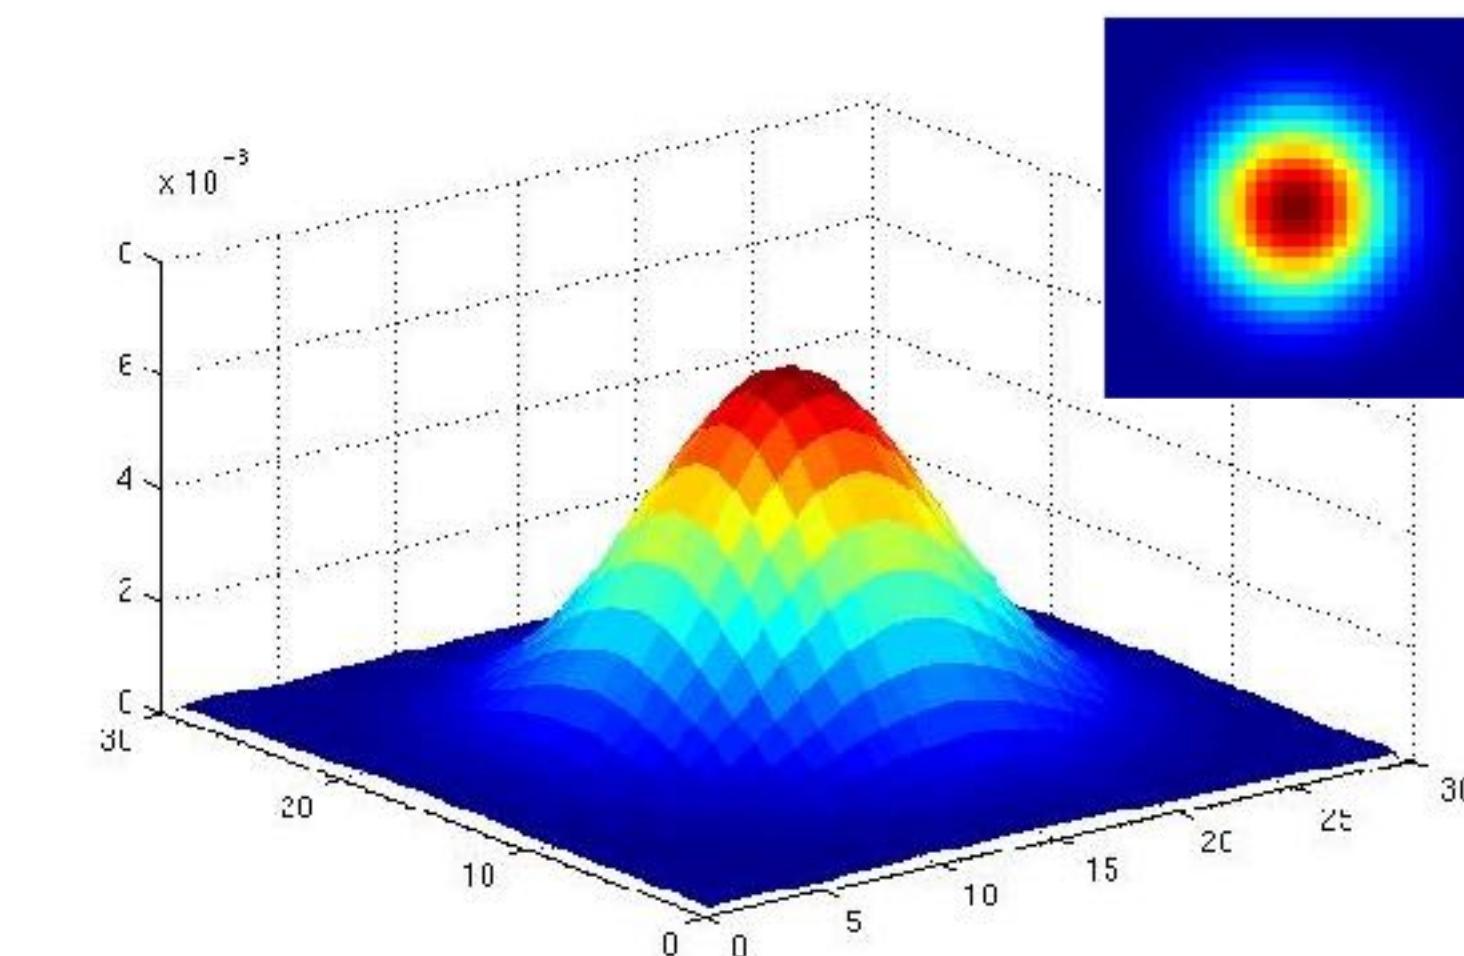
Source: C. Rasmussen

Gaussian kernel

Standard deviation σ : determines extent of smoothing



$\sigma = 2$ with 30×30 kernel

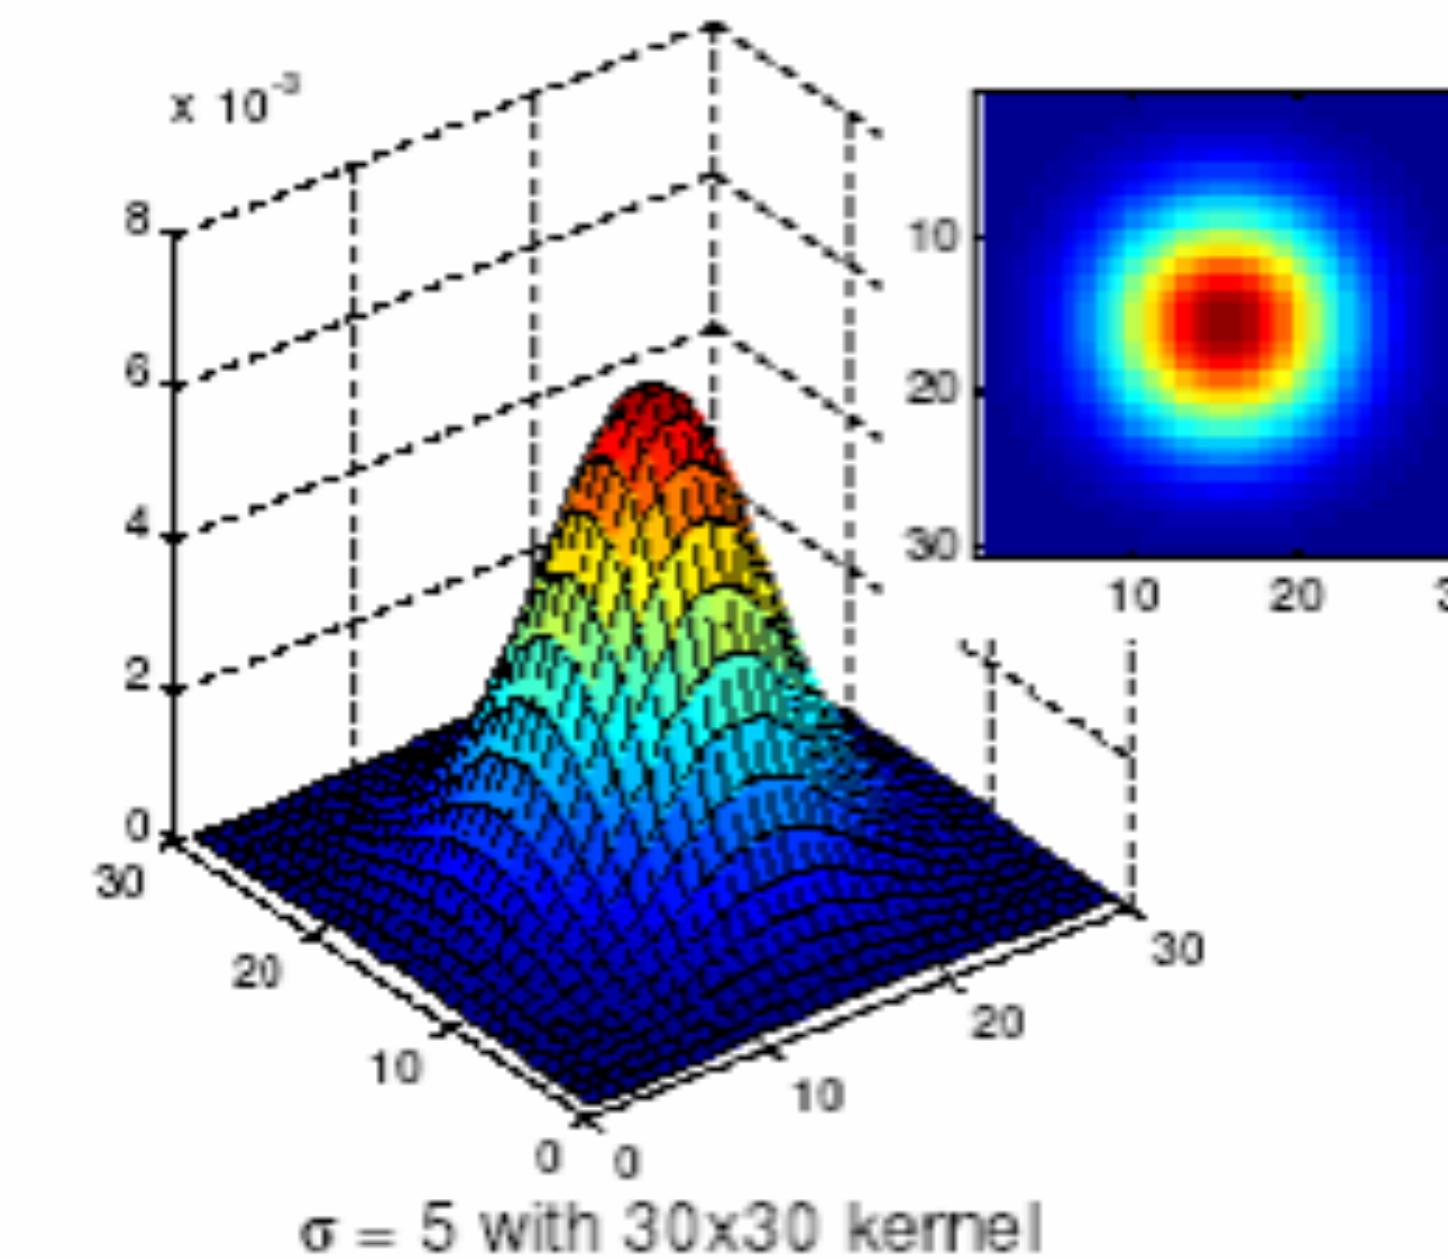
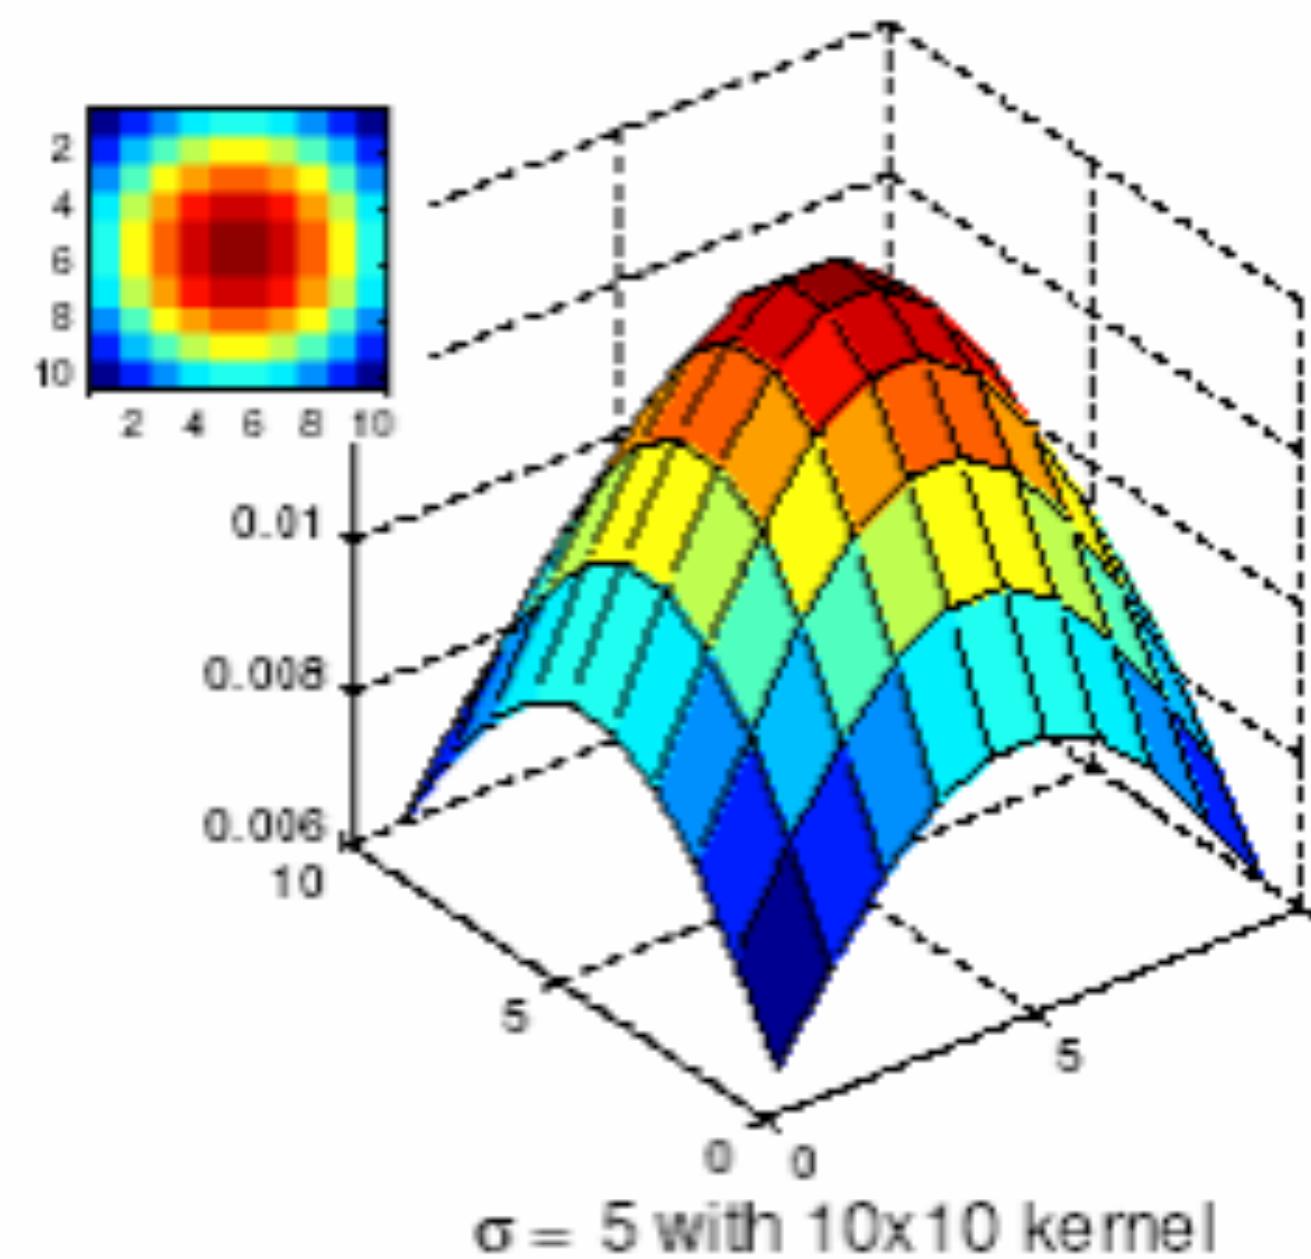


$\sigma = 5$ with 30×30 kernel

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

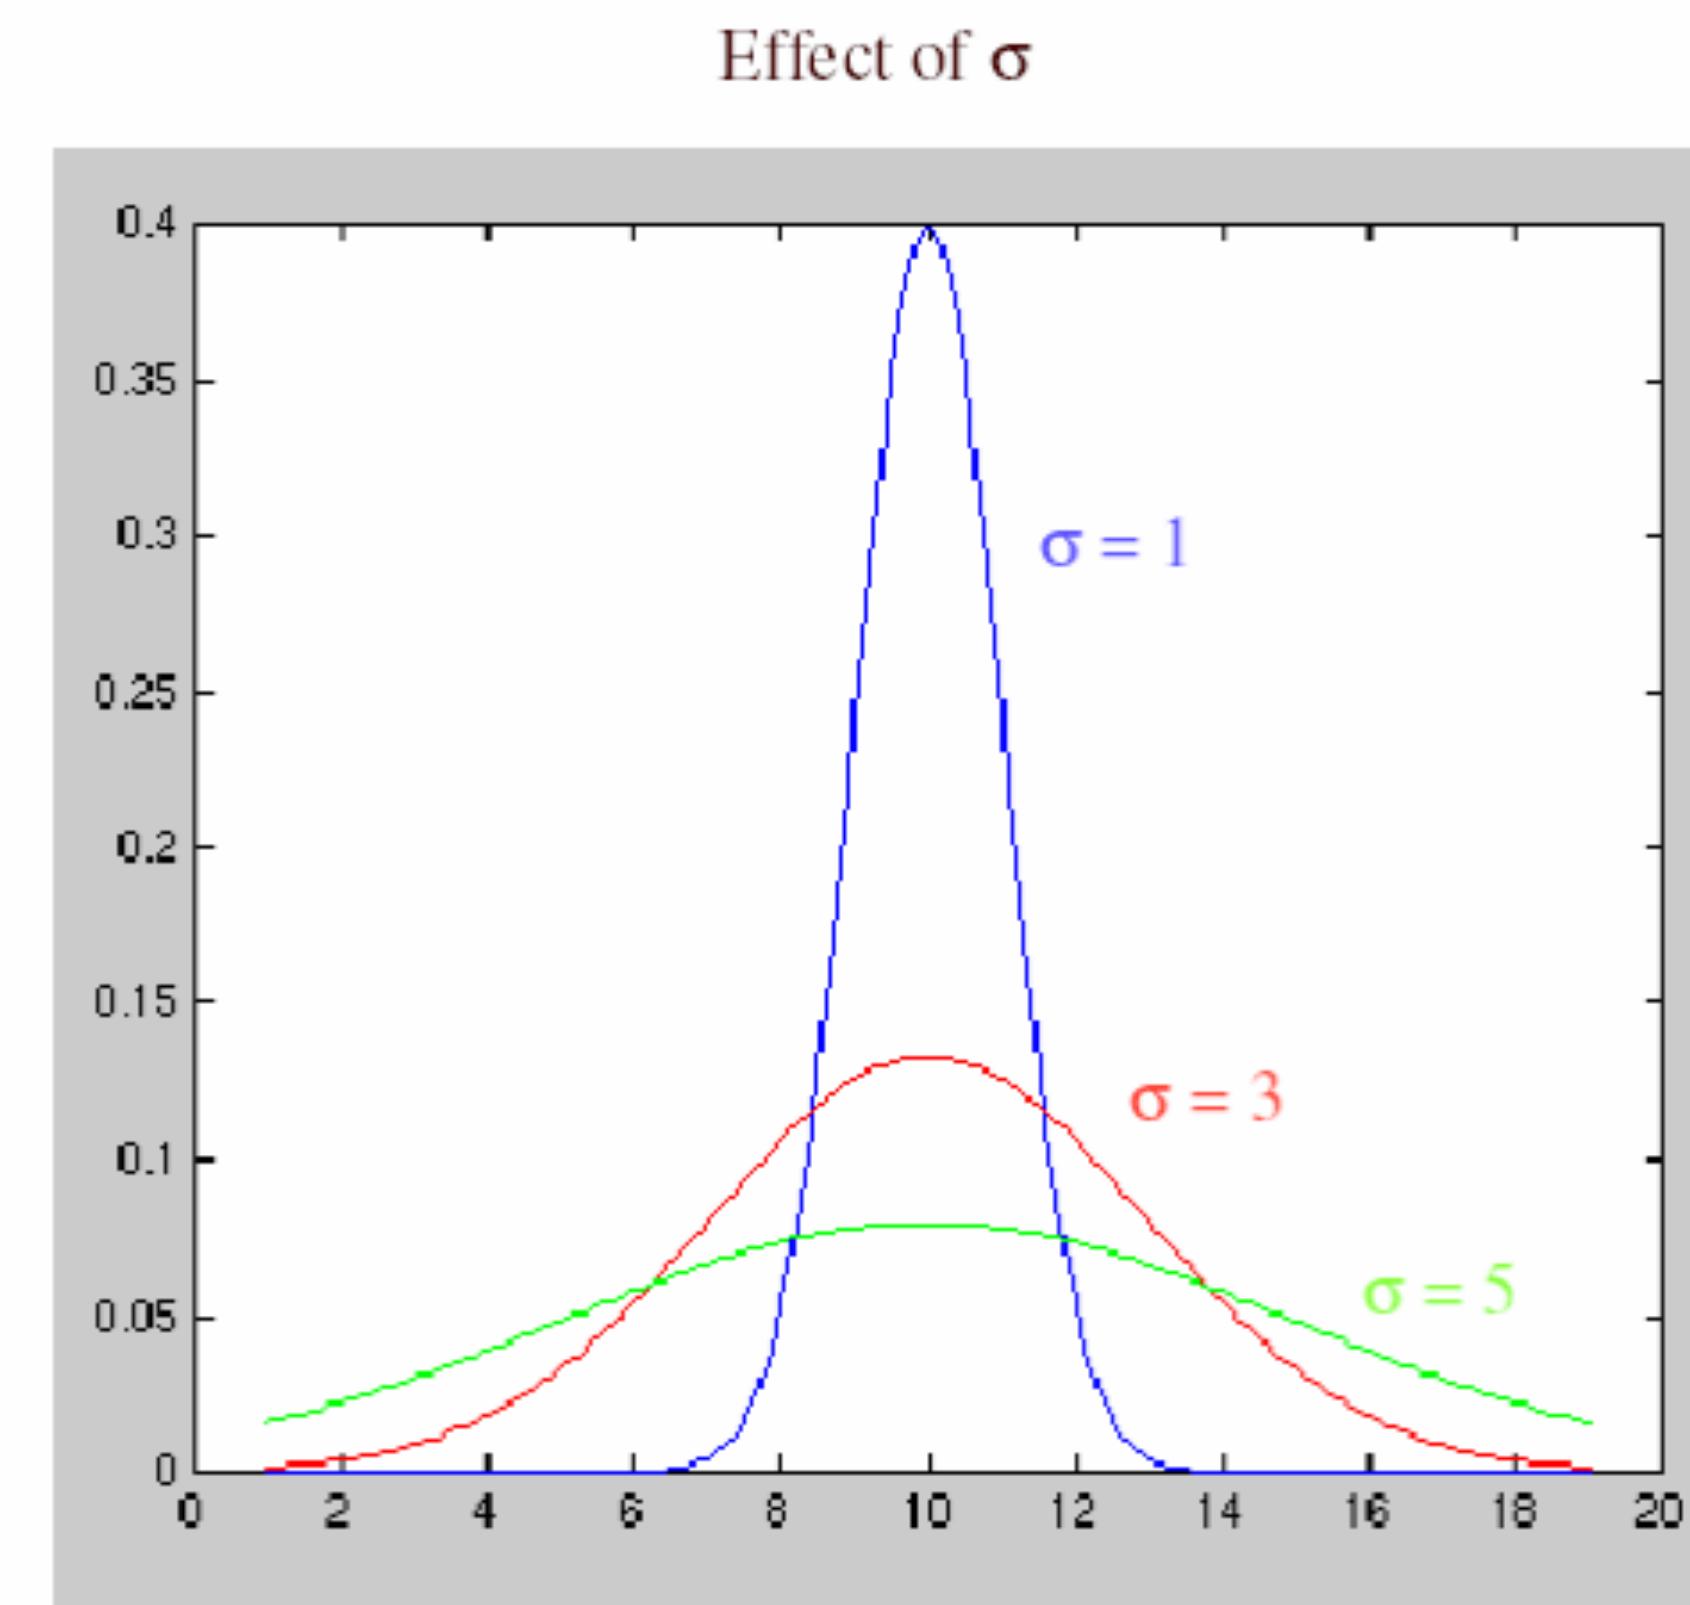
Choosing kernel width

The Gaussian function has infinite support, but discrete filters use finite kernels



Choosing kernel width

Rule of thumb: set filter half-width to about 3σ



Gaussian filters

Remove high-frequency components from the image (*low-pass filter*)

Convolution with self is another Gaussian

- So can smooth with small- σ kernel, repeat, and get same result as larger- σ kernel would have
- Convoluting two times with Gaussian kernel with std. dev. σ is same as convoluting once with kernel with std. dev. $\sigma\sqrt{2}$

Separable kernel

- Factors into product of two 1D Gaussians
- Discrete example:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

`scipy.ndimage.gaussian_filter(input, sigma, order=0, output=None, mode='reflect', cval=0.0, truncate=4.0)`

Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

Source: D. Lowe

Why is separability useful?

Separability means that a 2D convolution can be reduced to two 1D convolutions (one among rows and one among columns)

What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?

- $O(n^2 m^2)$

What if the kernel is separable?

- $O(n^2 m)$

Question: Is the box filter separable?

Types of noise



Original



Salt and pepper noise



Impulse noise



Gaussian noise

Salt and pepper noise: contains random occurrences of black and white pixels

Impulse noise: contains random occurrences of white pixels

Gaussian noise: variations in intensity drawn from a Gaussian normal distribution

Source: S. Seitz

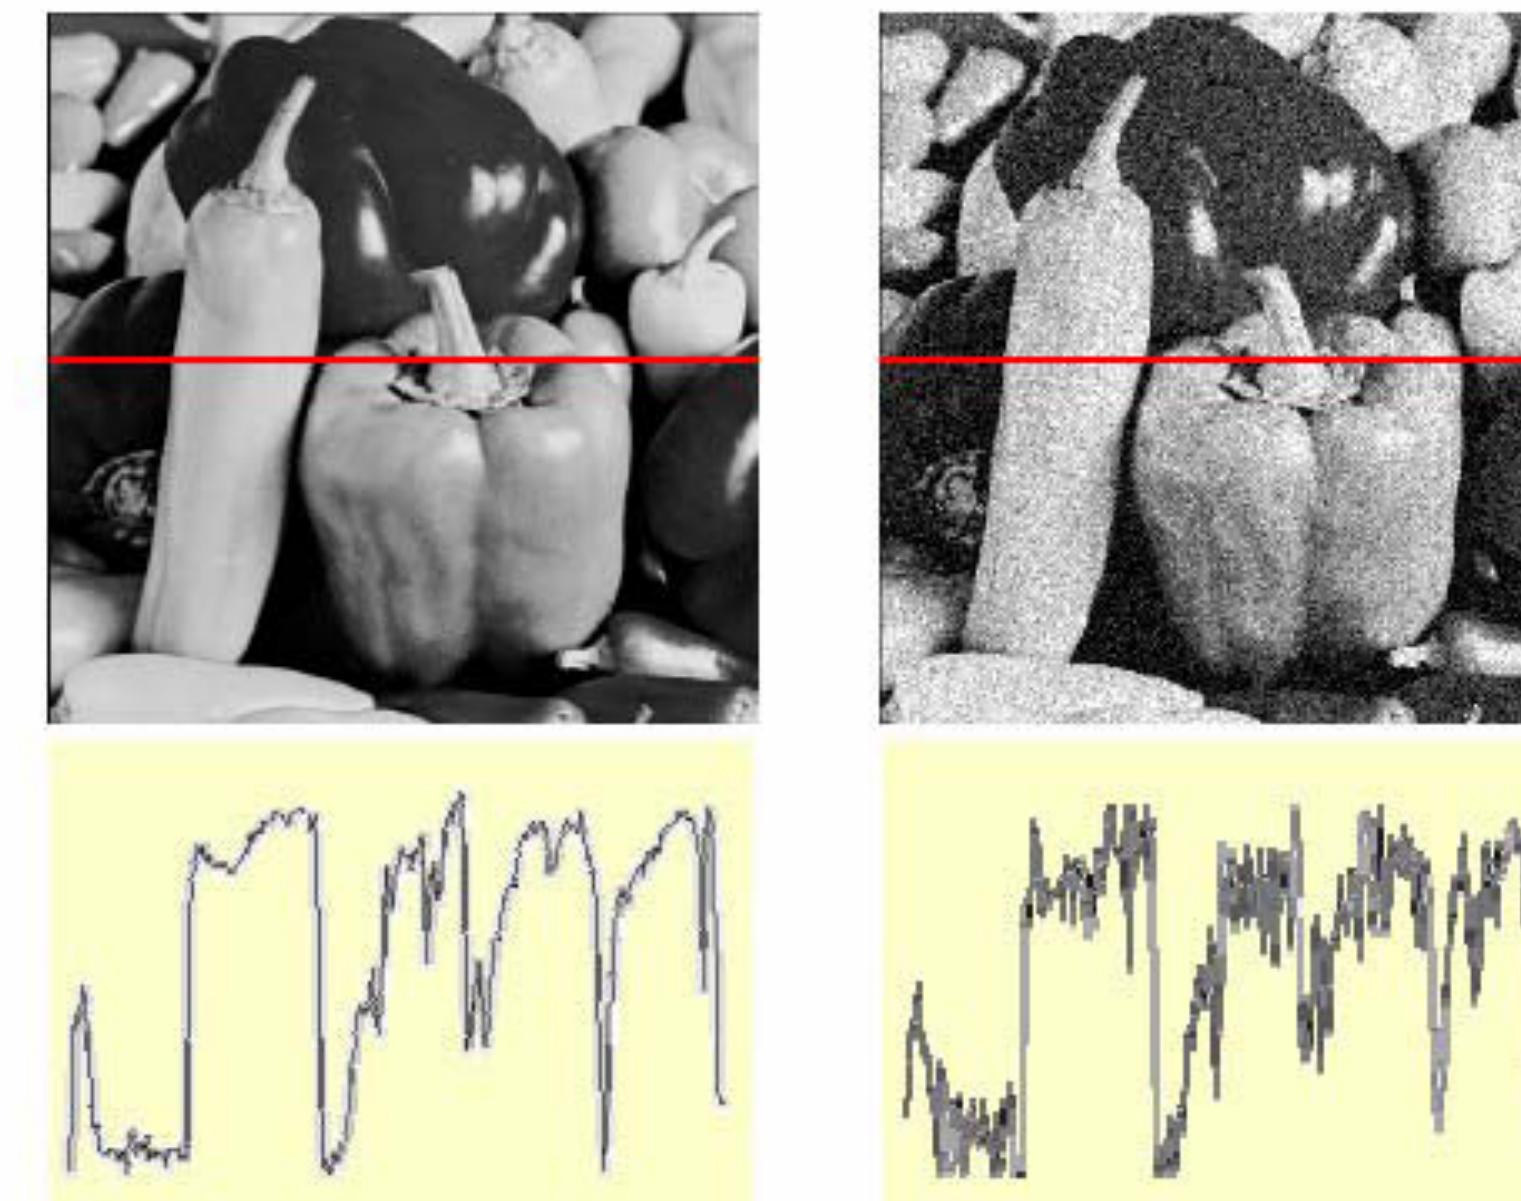
Gaussian noise

Mathematical model: sum of many independent factors

Good for small standard deviations

Assumption: independent, zero-mean noise

Image
Noise

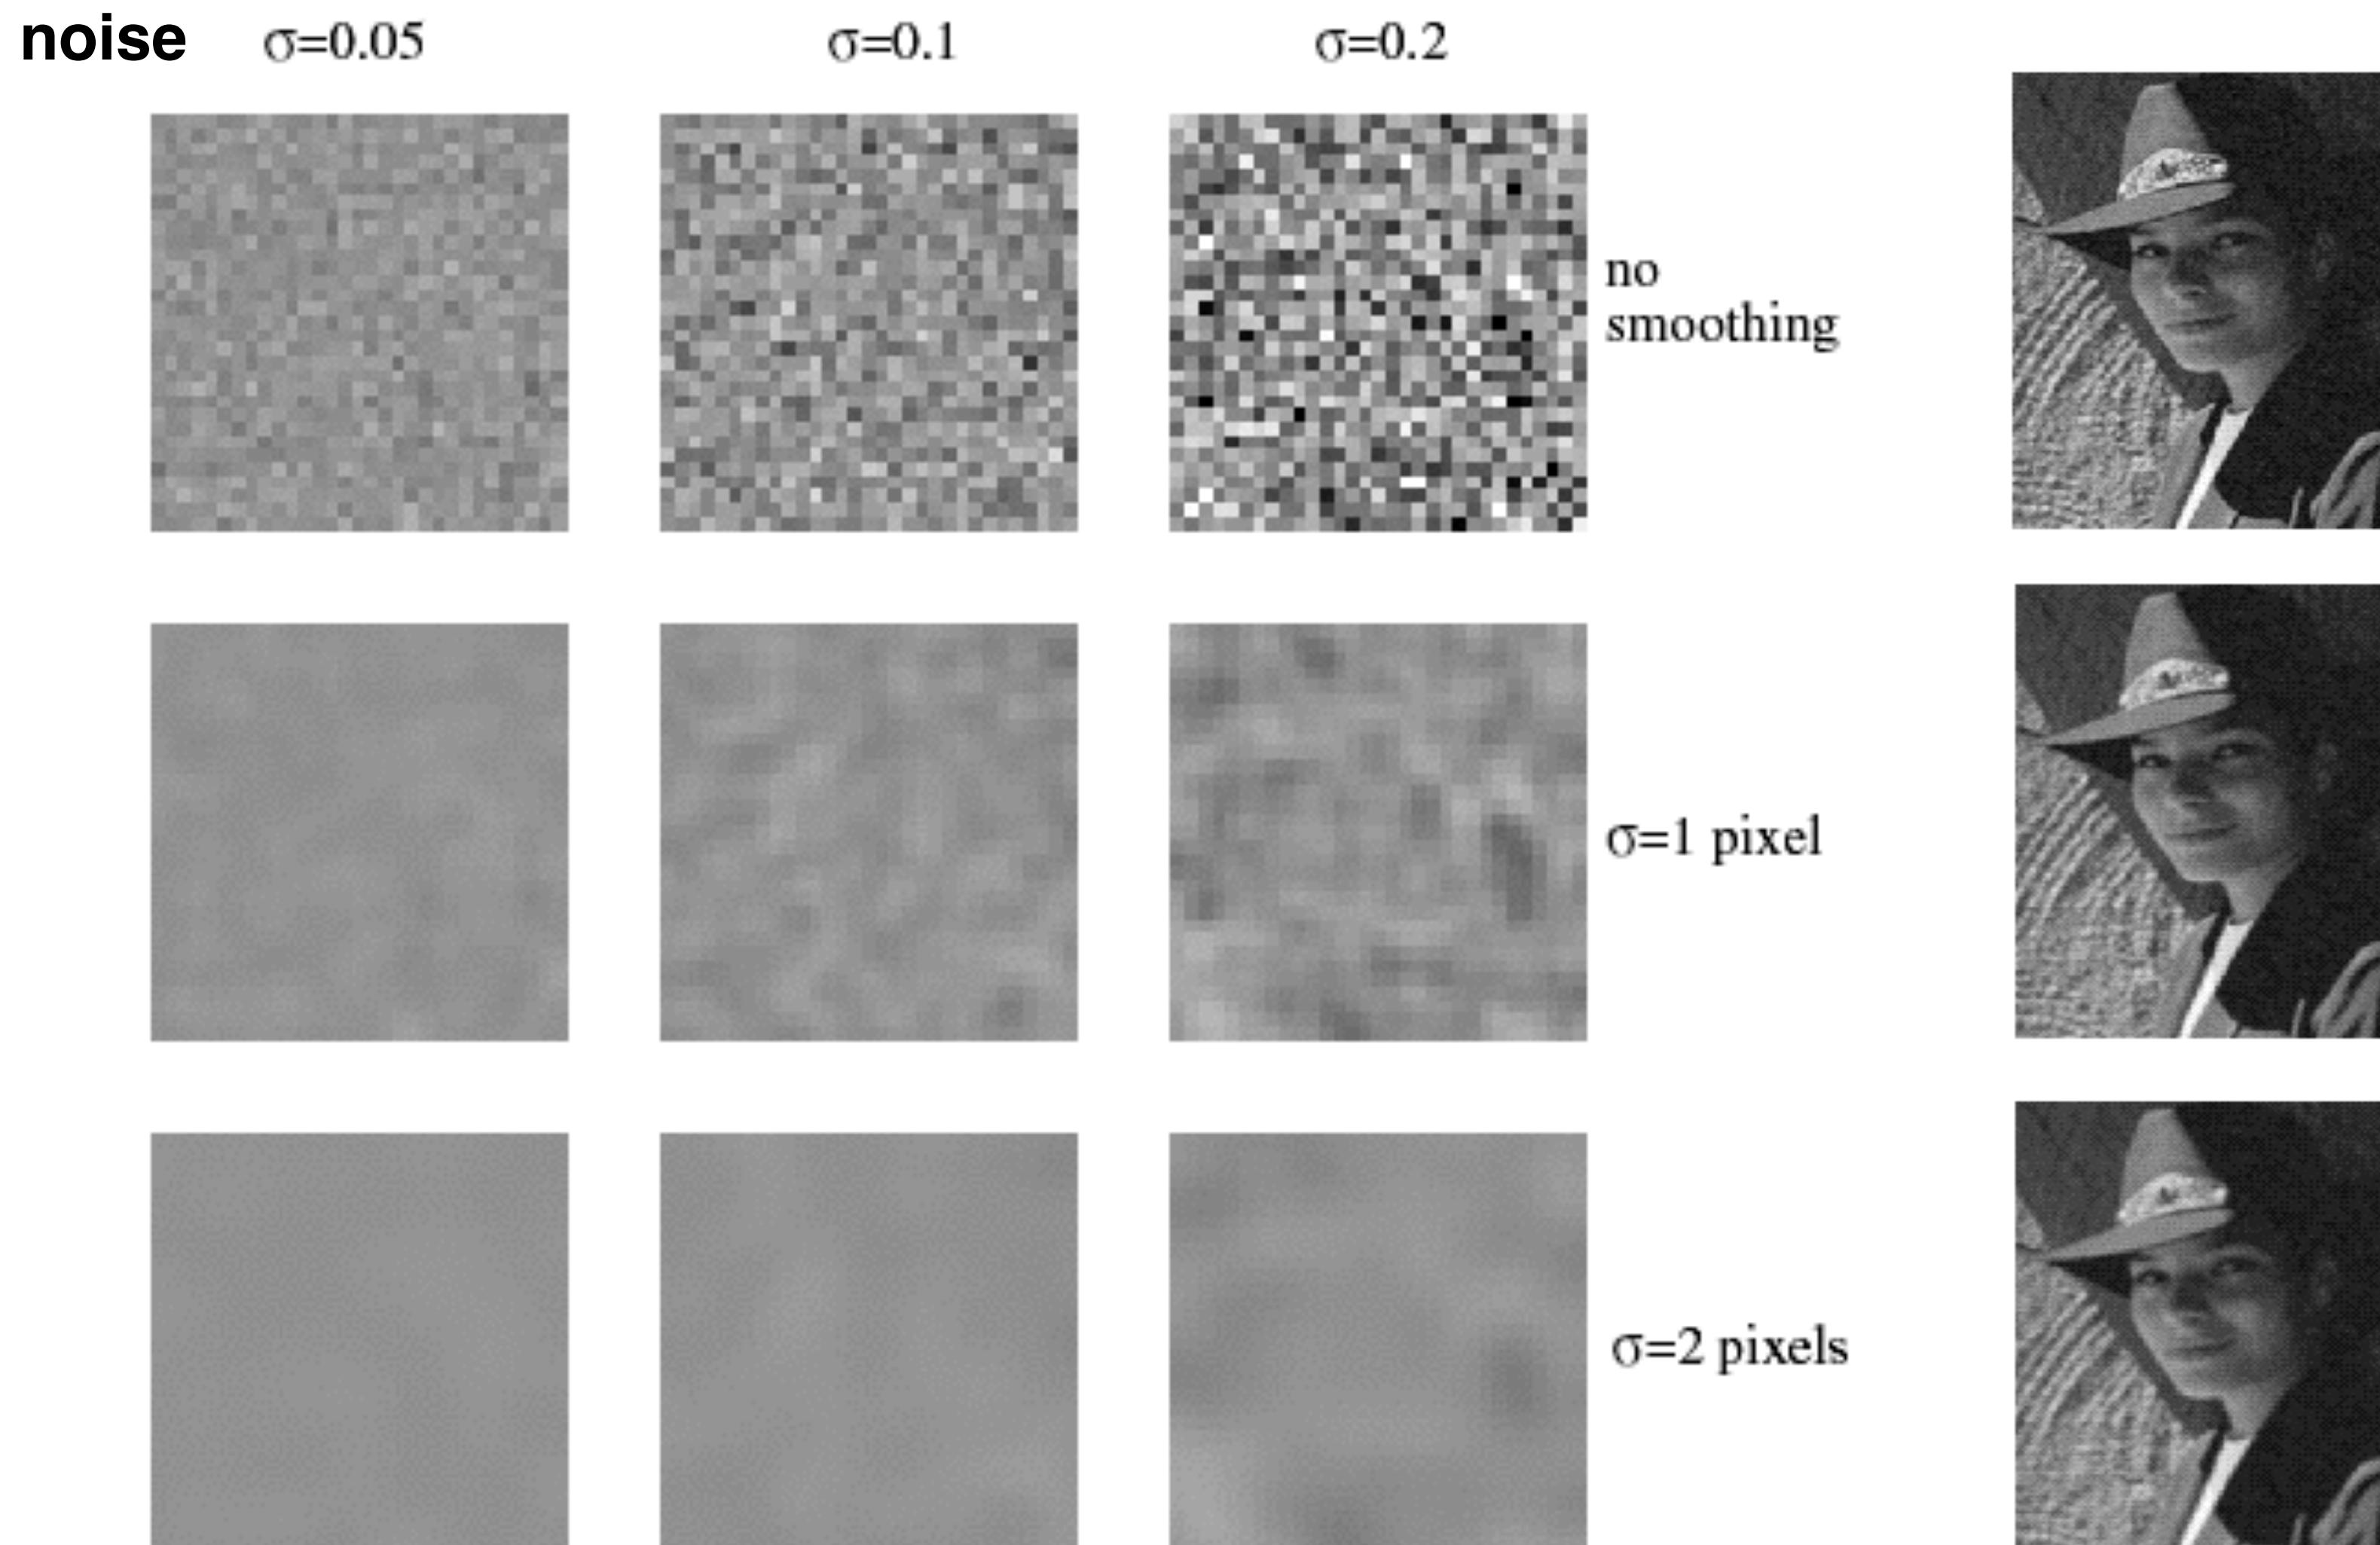


$$f(x, y) = \widehat{f(x, y)} + \widehat{\eta(x, y)}$$

Gaussian i.i.d. ("white") noise:
 $\eta(x, y) \sim \mathcal{N}(\mu, \sigma)$

Source: M. Hebert

Reducing Gaussian noise



Smoothing with larger standard deviations suppresses noise, but also blurs the image

Reducing salt-and-pepper noise

What's wrong with the results?

3x3



5x5



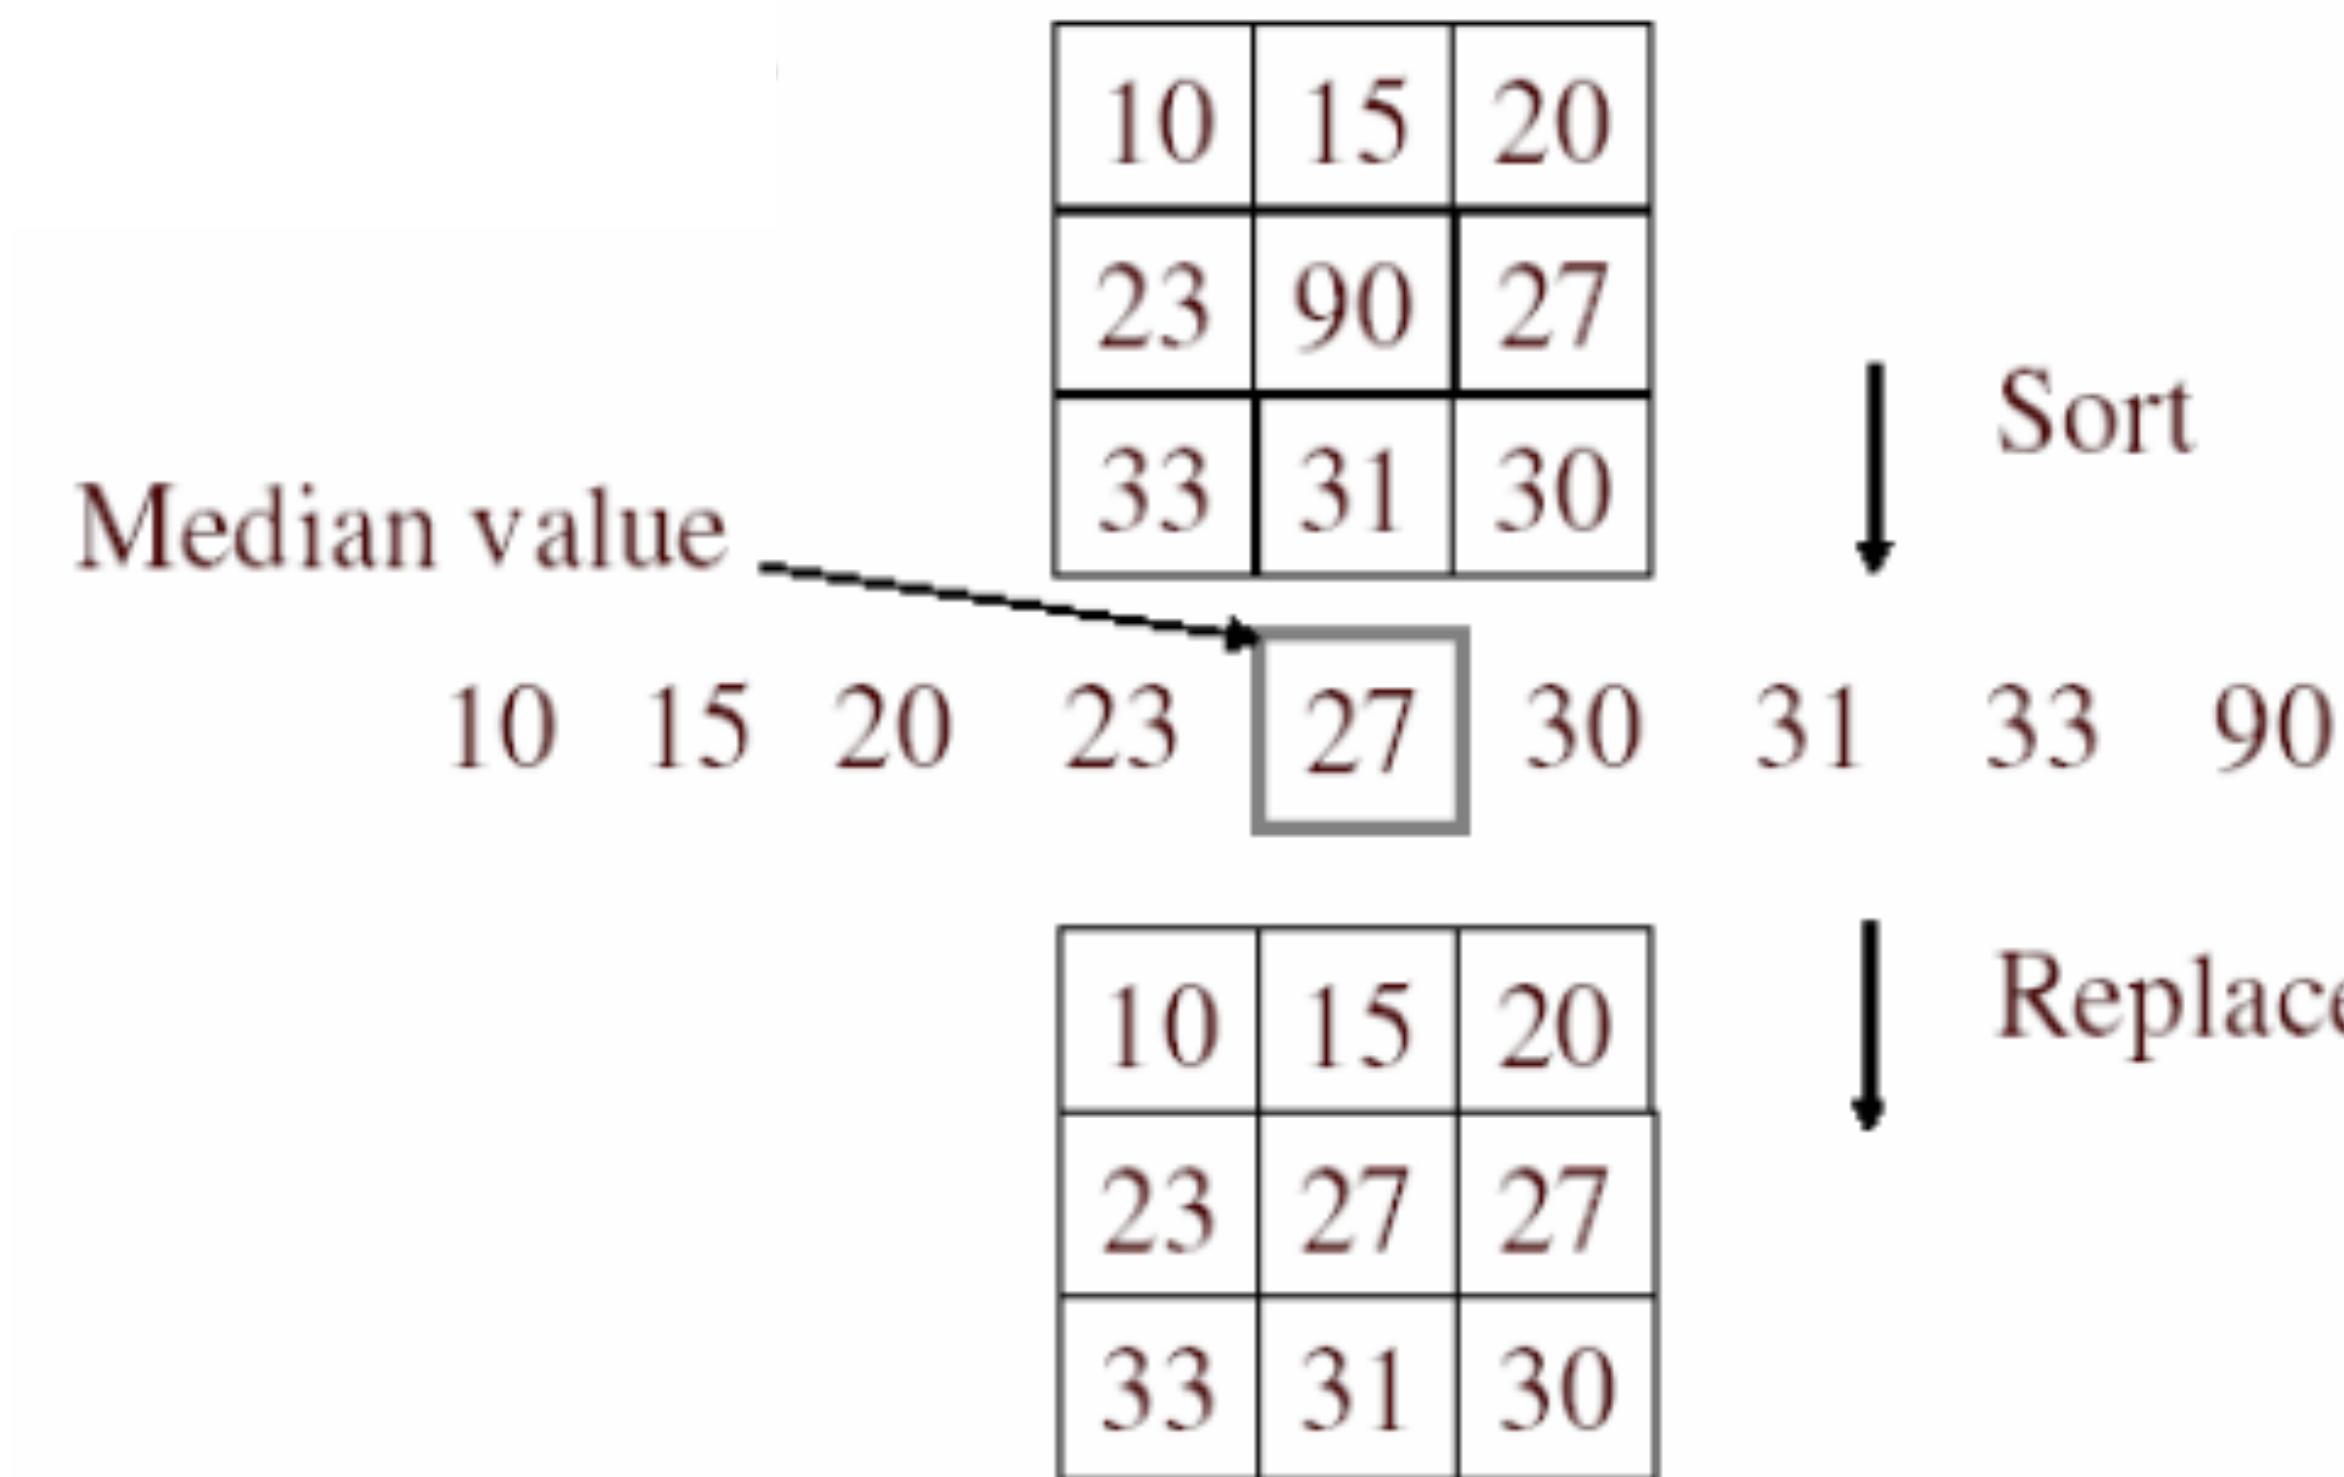
7x7



Gaussian smoothing with increasing standard deviation

Alternative idea: Median filtering

A **median filter** operates over a window by selecting the median intensity in the window



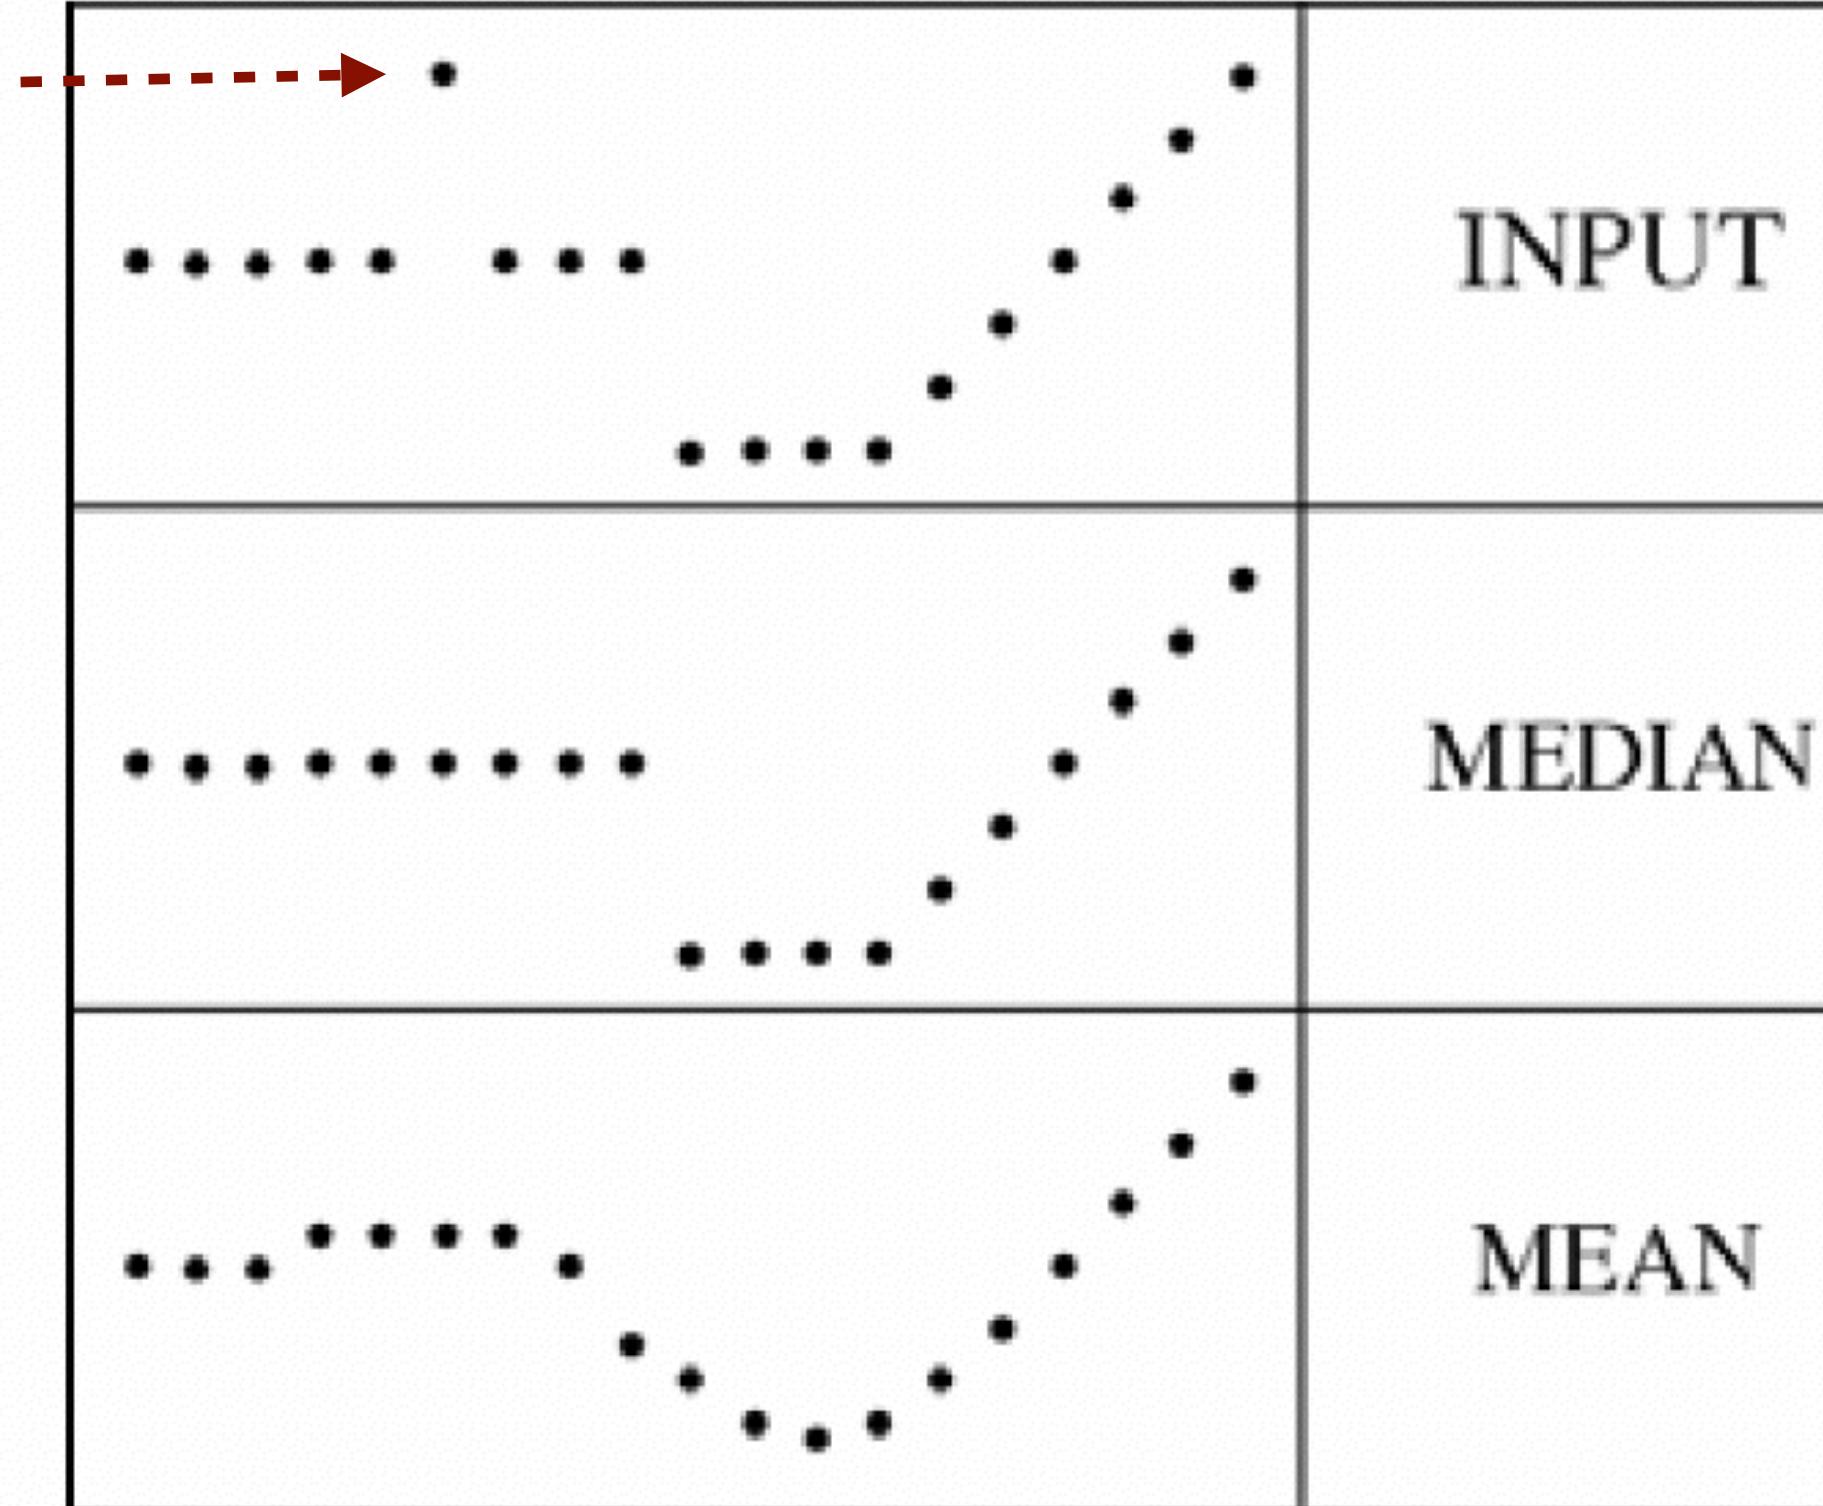
Question: is median filtering linear?

Median filter

What advantage does median filtering have over Gaussian filtering?

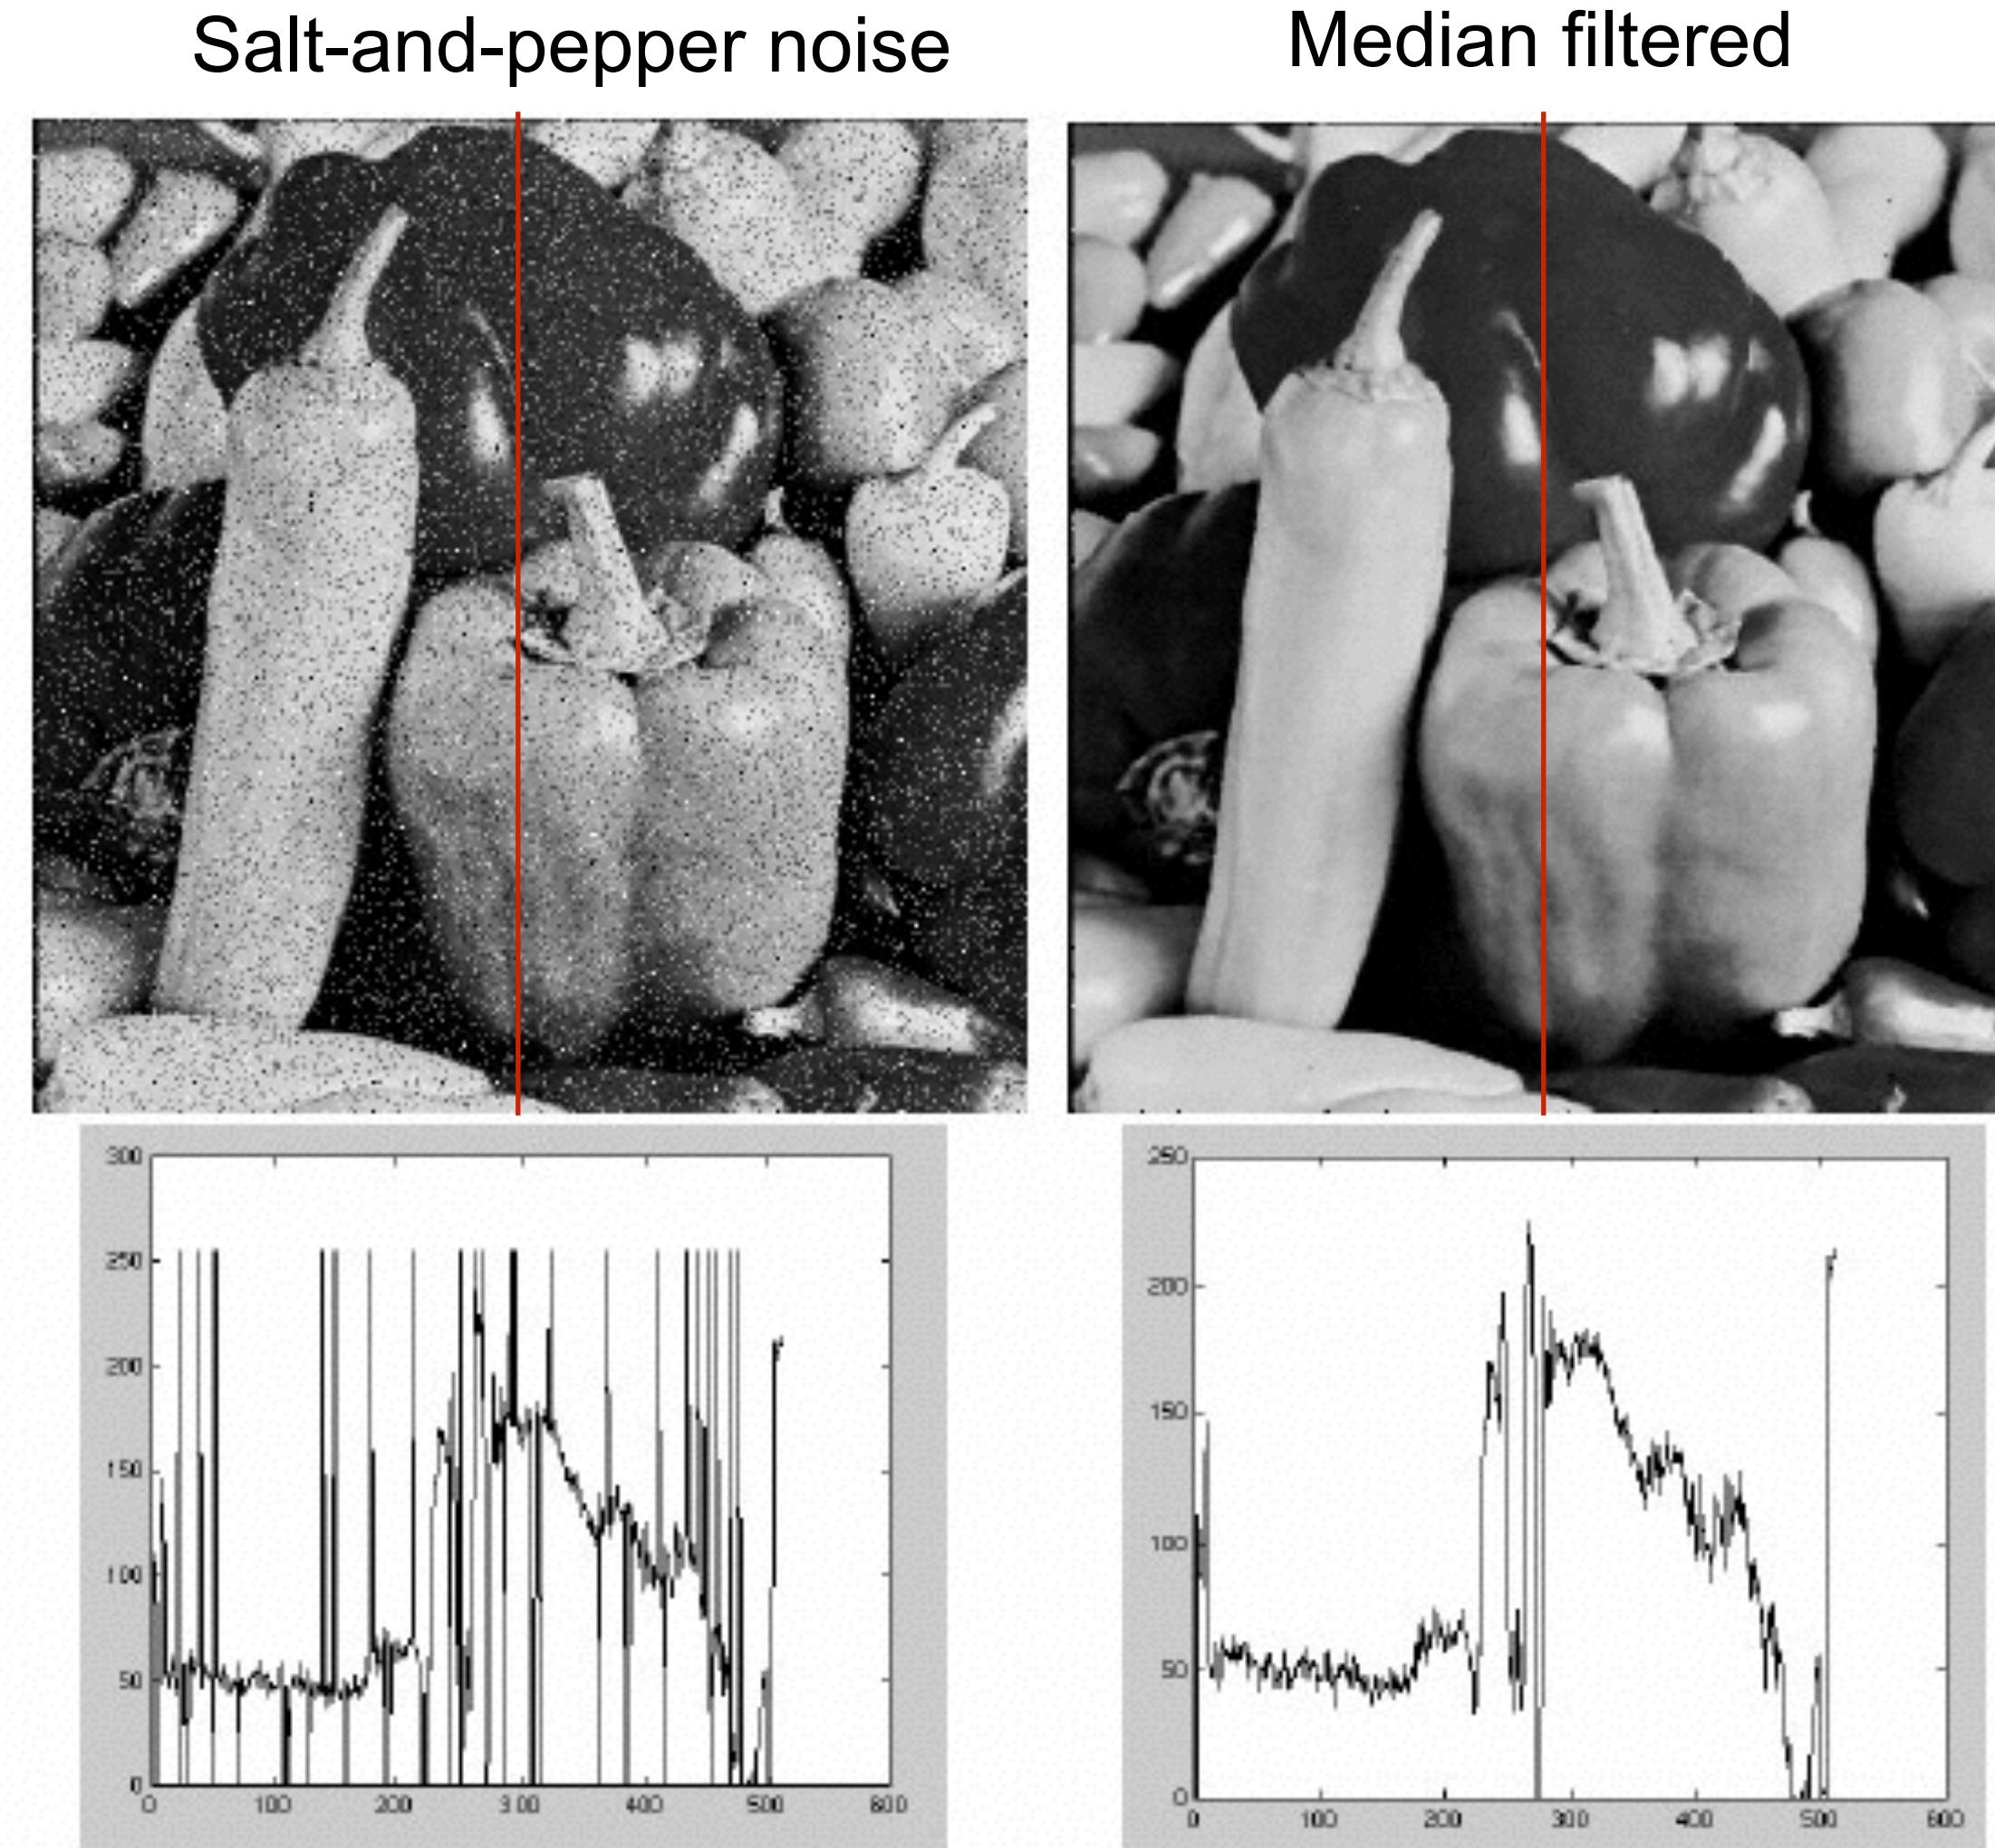
filters have width 5 :

Robustness to outliers



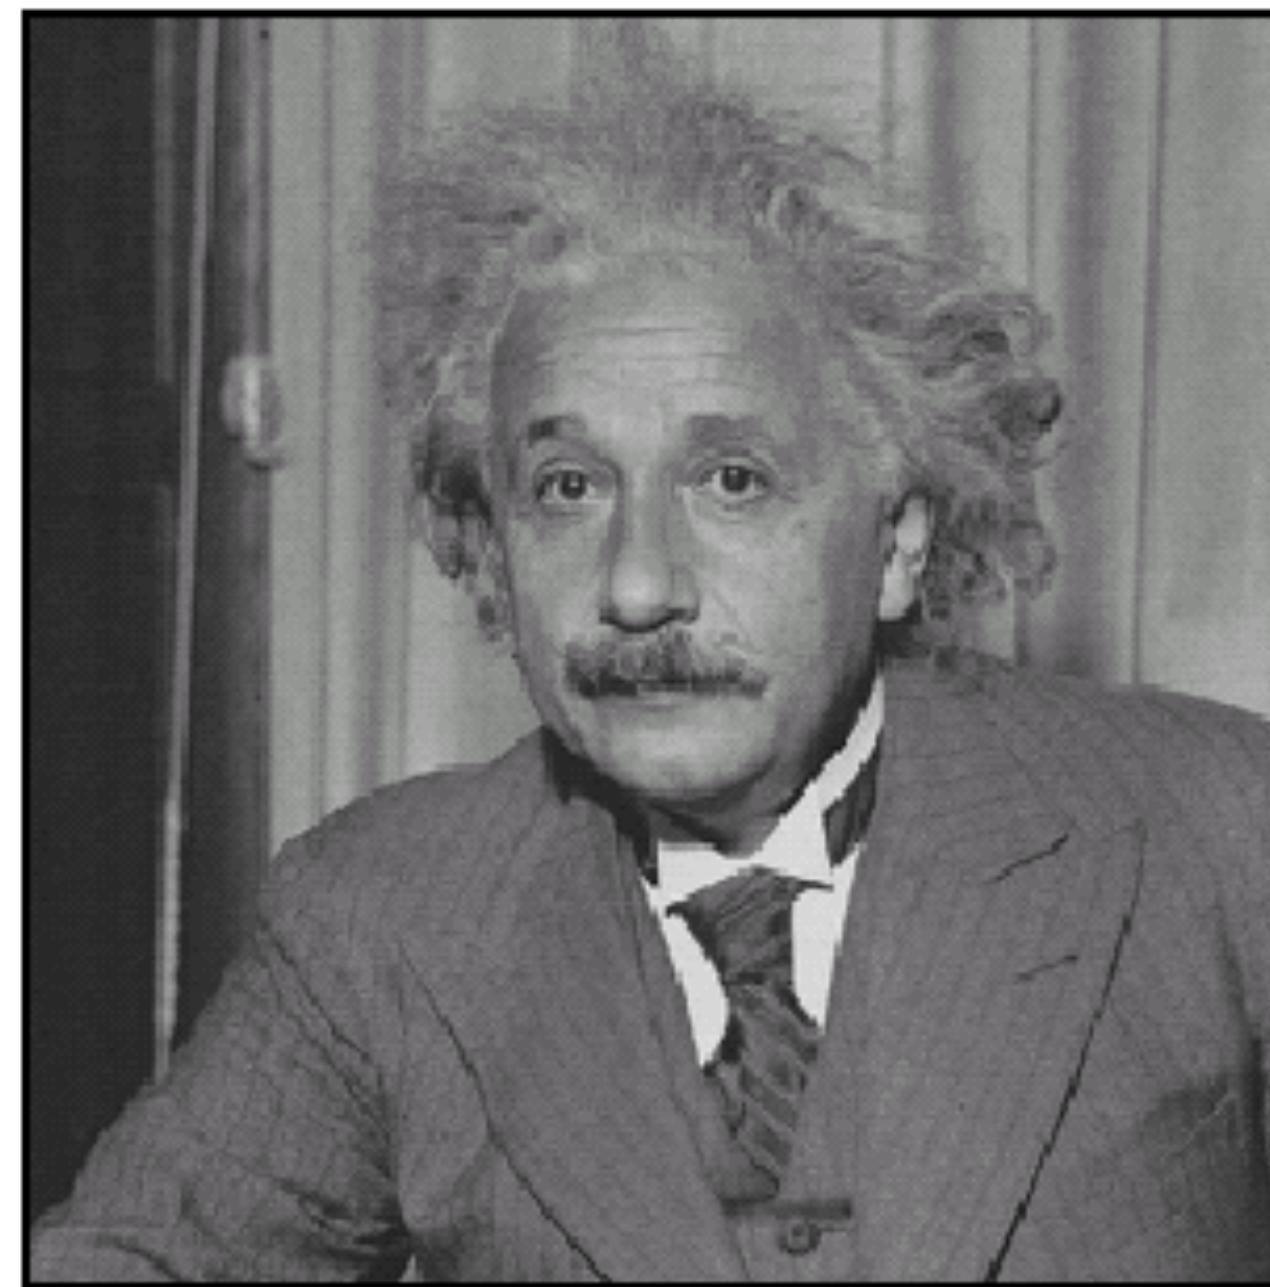
Source: K. Grauman

Median filter

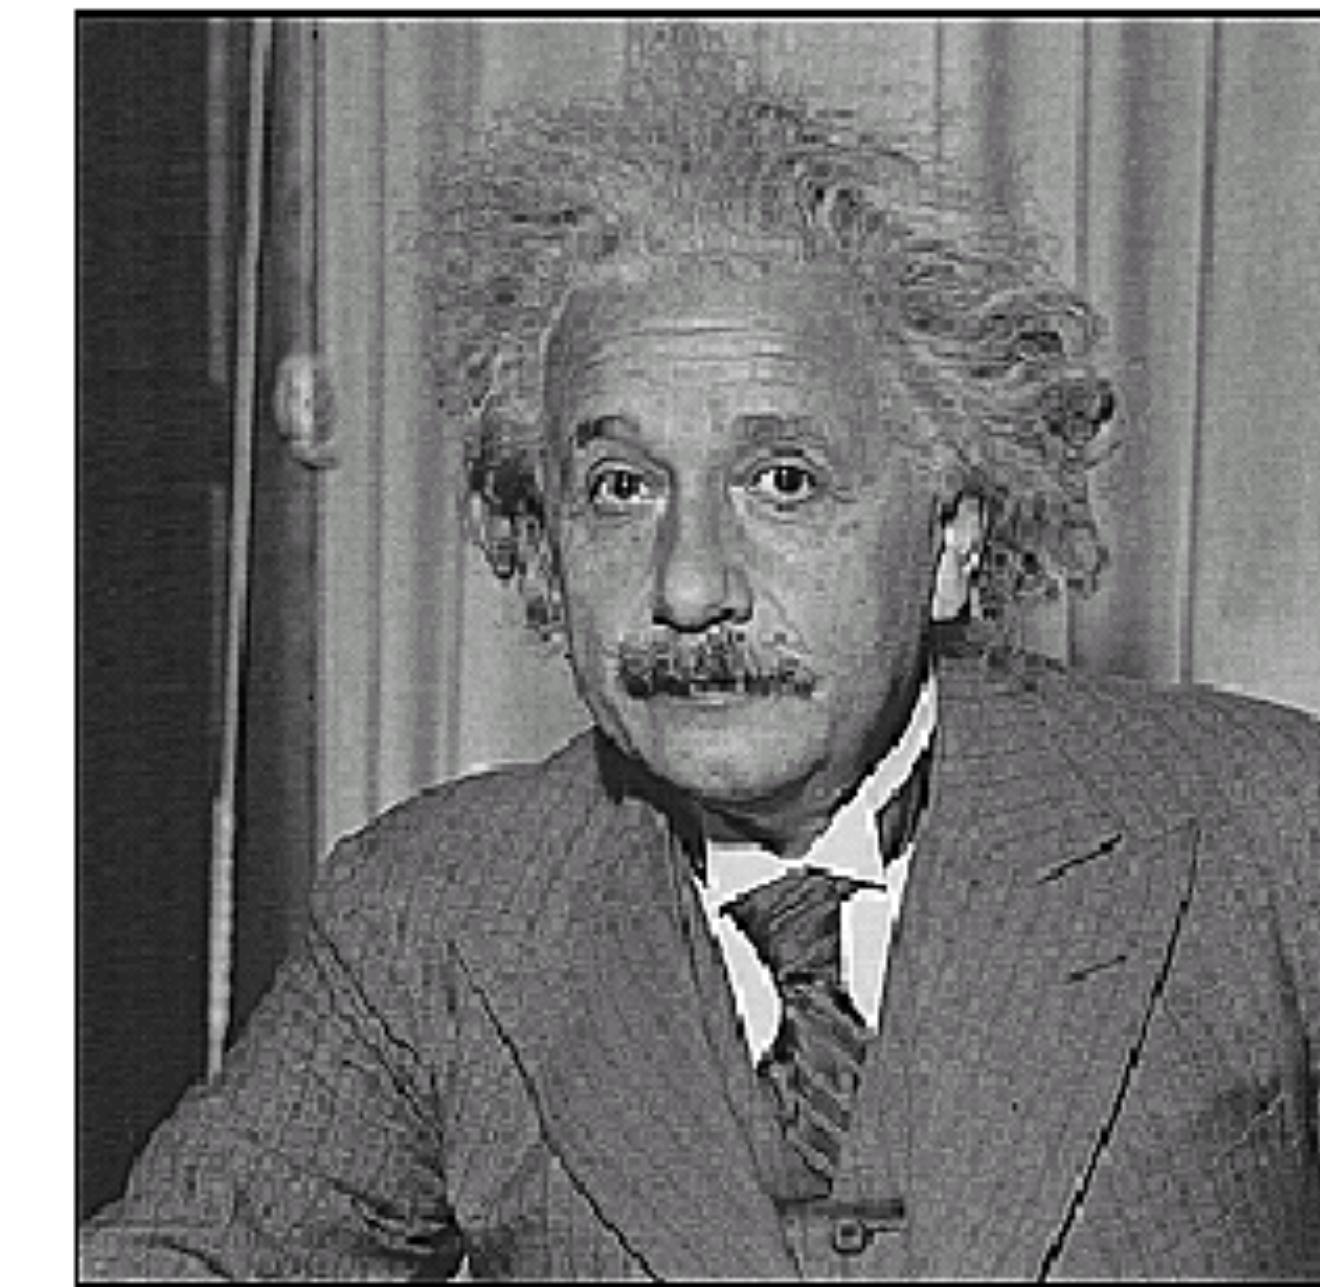


```
scipy.ndimage.median_filter(input, size=None, footprint=None, output=None, mode='reflect', cval=0.0, origin=0)
```

Sharpening



before



after

Source: D. Lowe

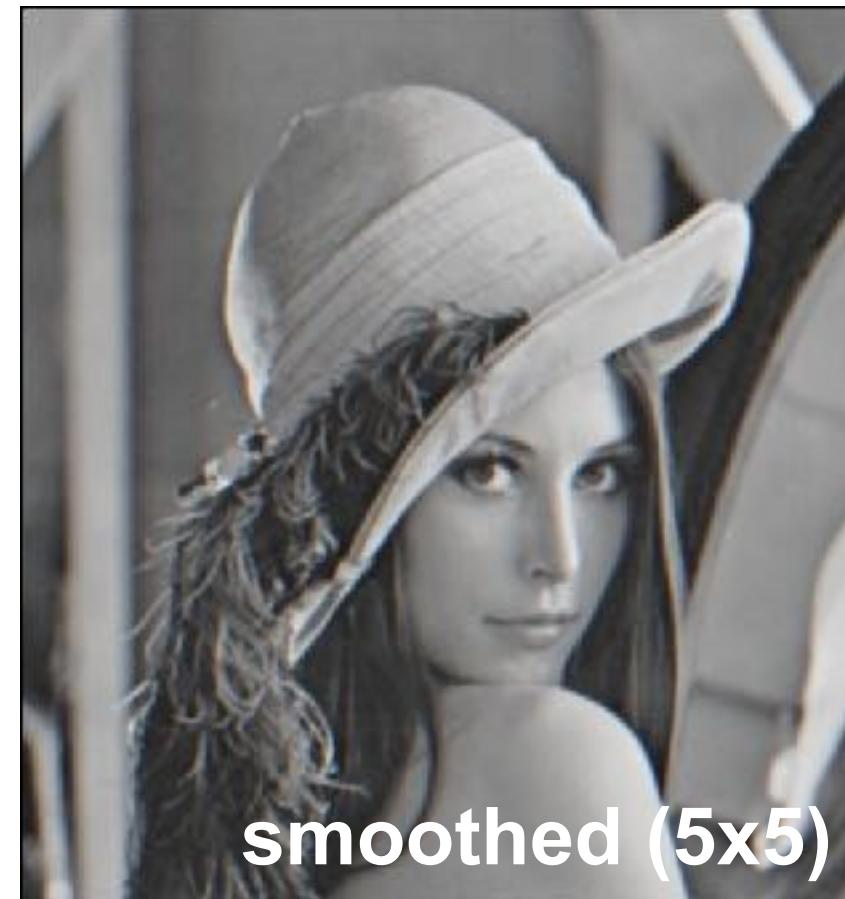
Sharpening

What does blurring take away?



original

-



smoothed (5x5)

=



detail

Let's add it back:



original

+ α



detail

=



sharpened

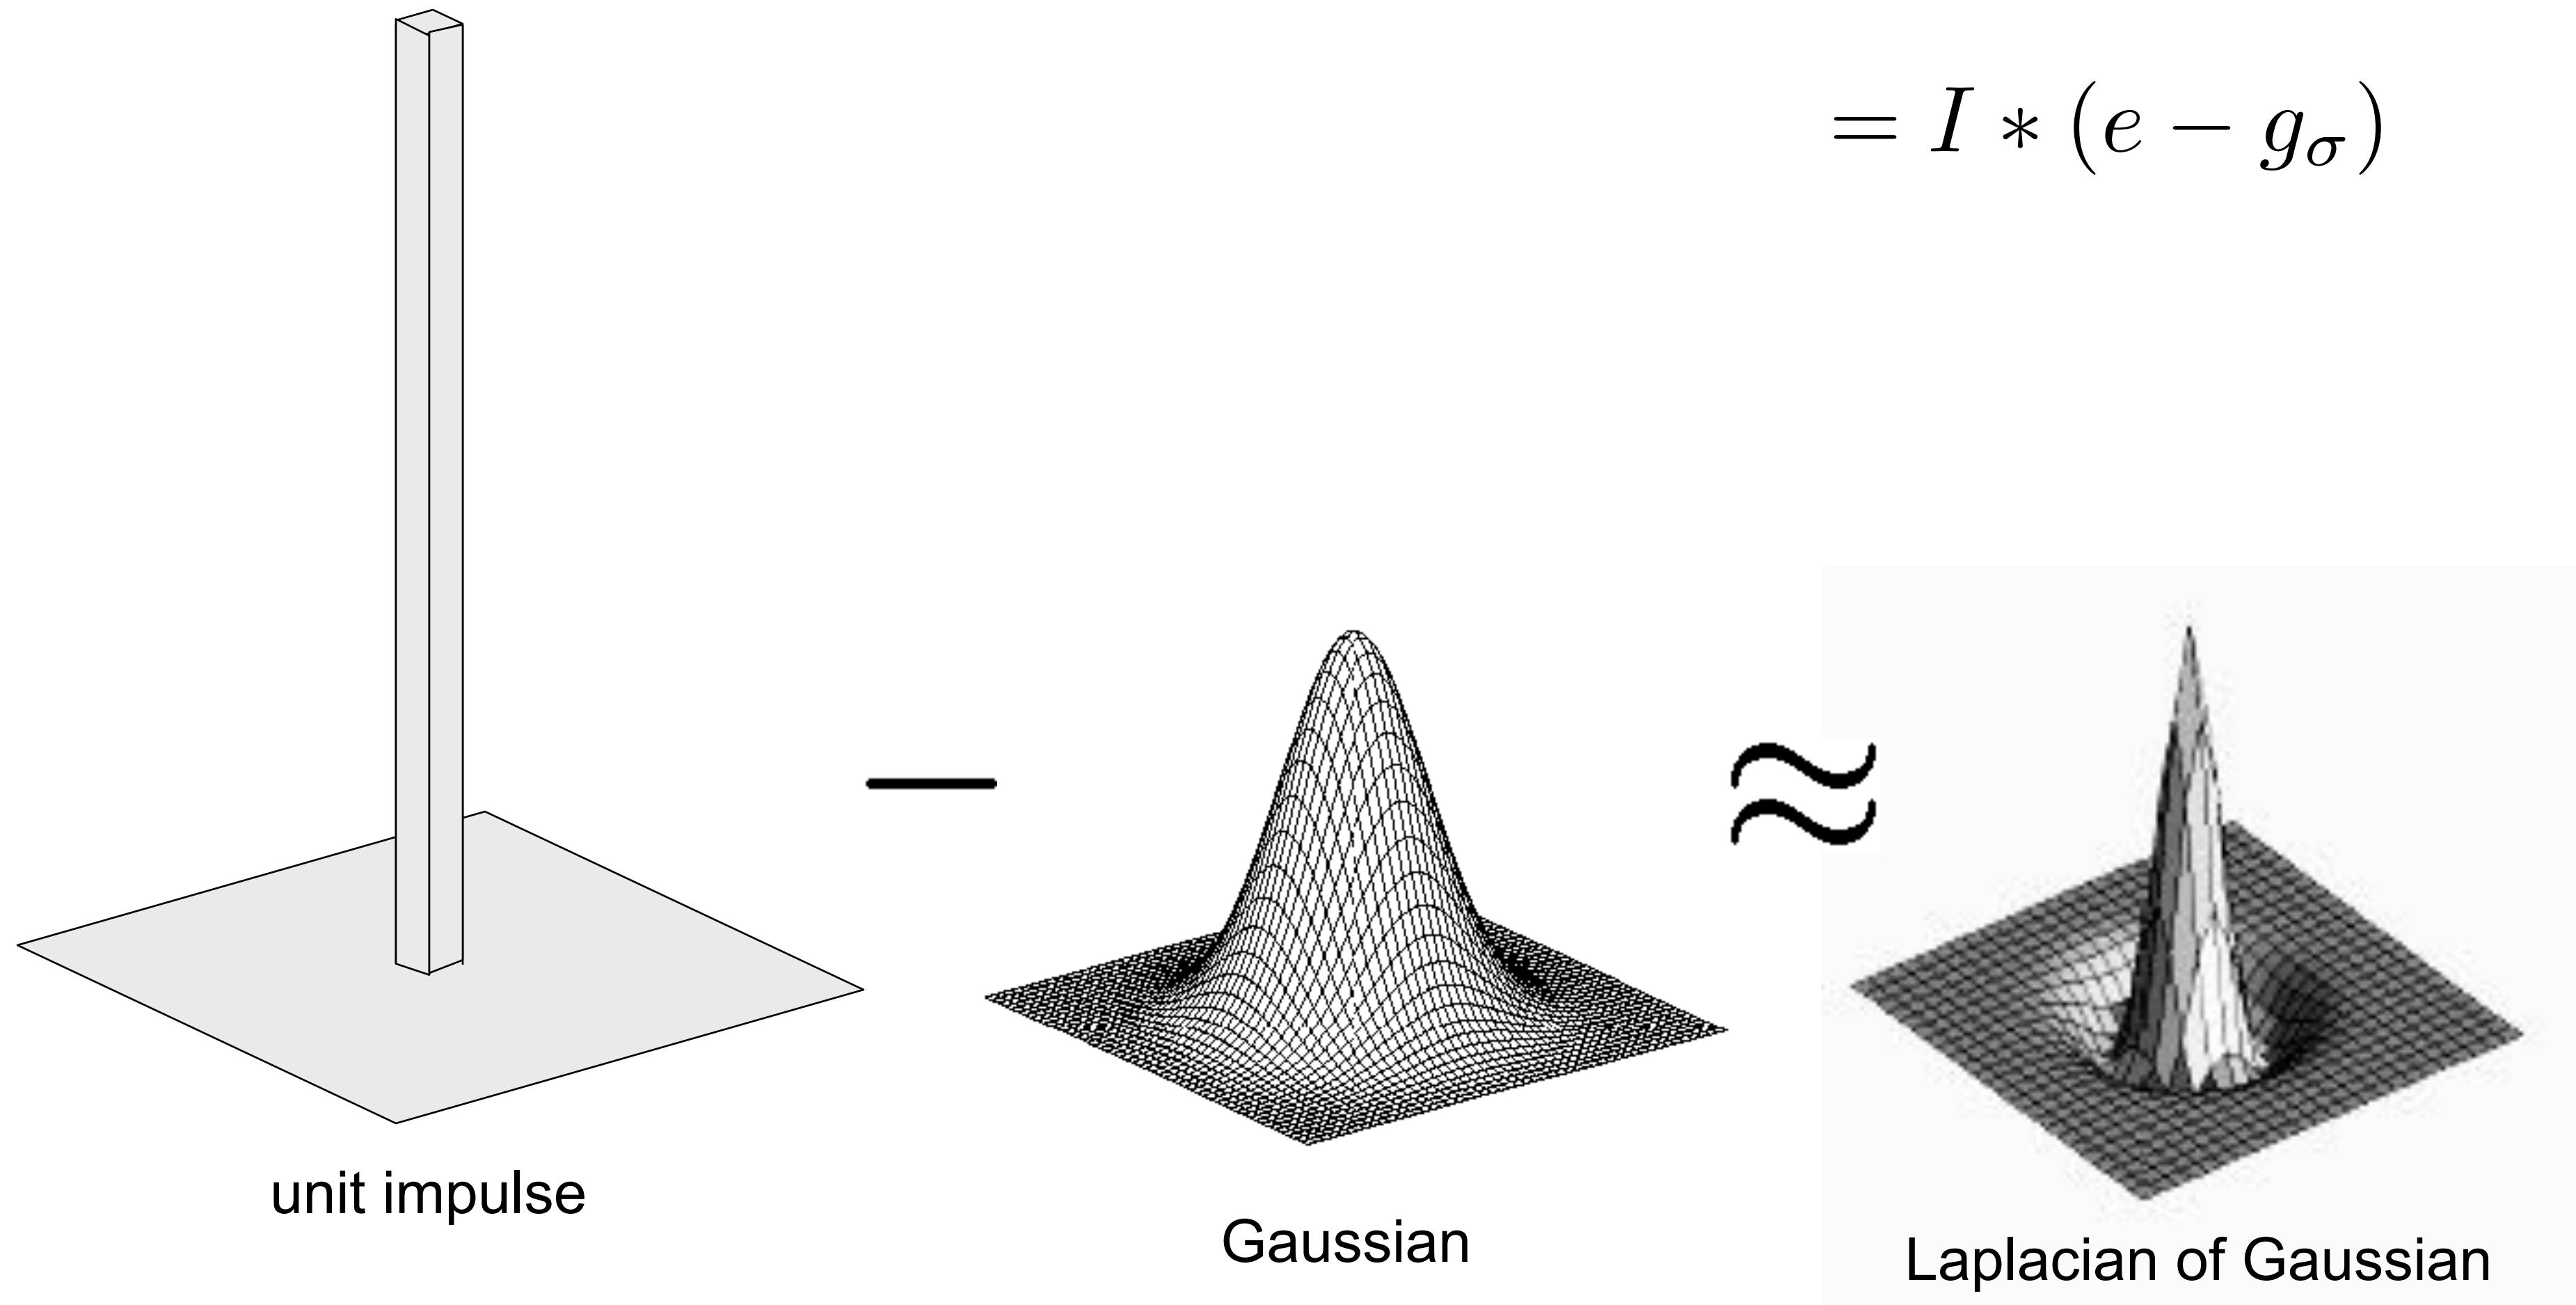
Sharpening filter

$$I = \text{blurry}(I) + \text{sharp}(I)$$

$$\text{sharp}(I) = I - \text{blurry}(I)$$

$$= I * e - I * g_\sigma$$

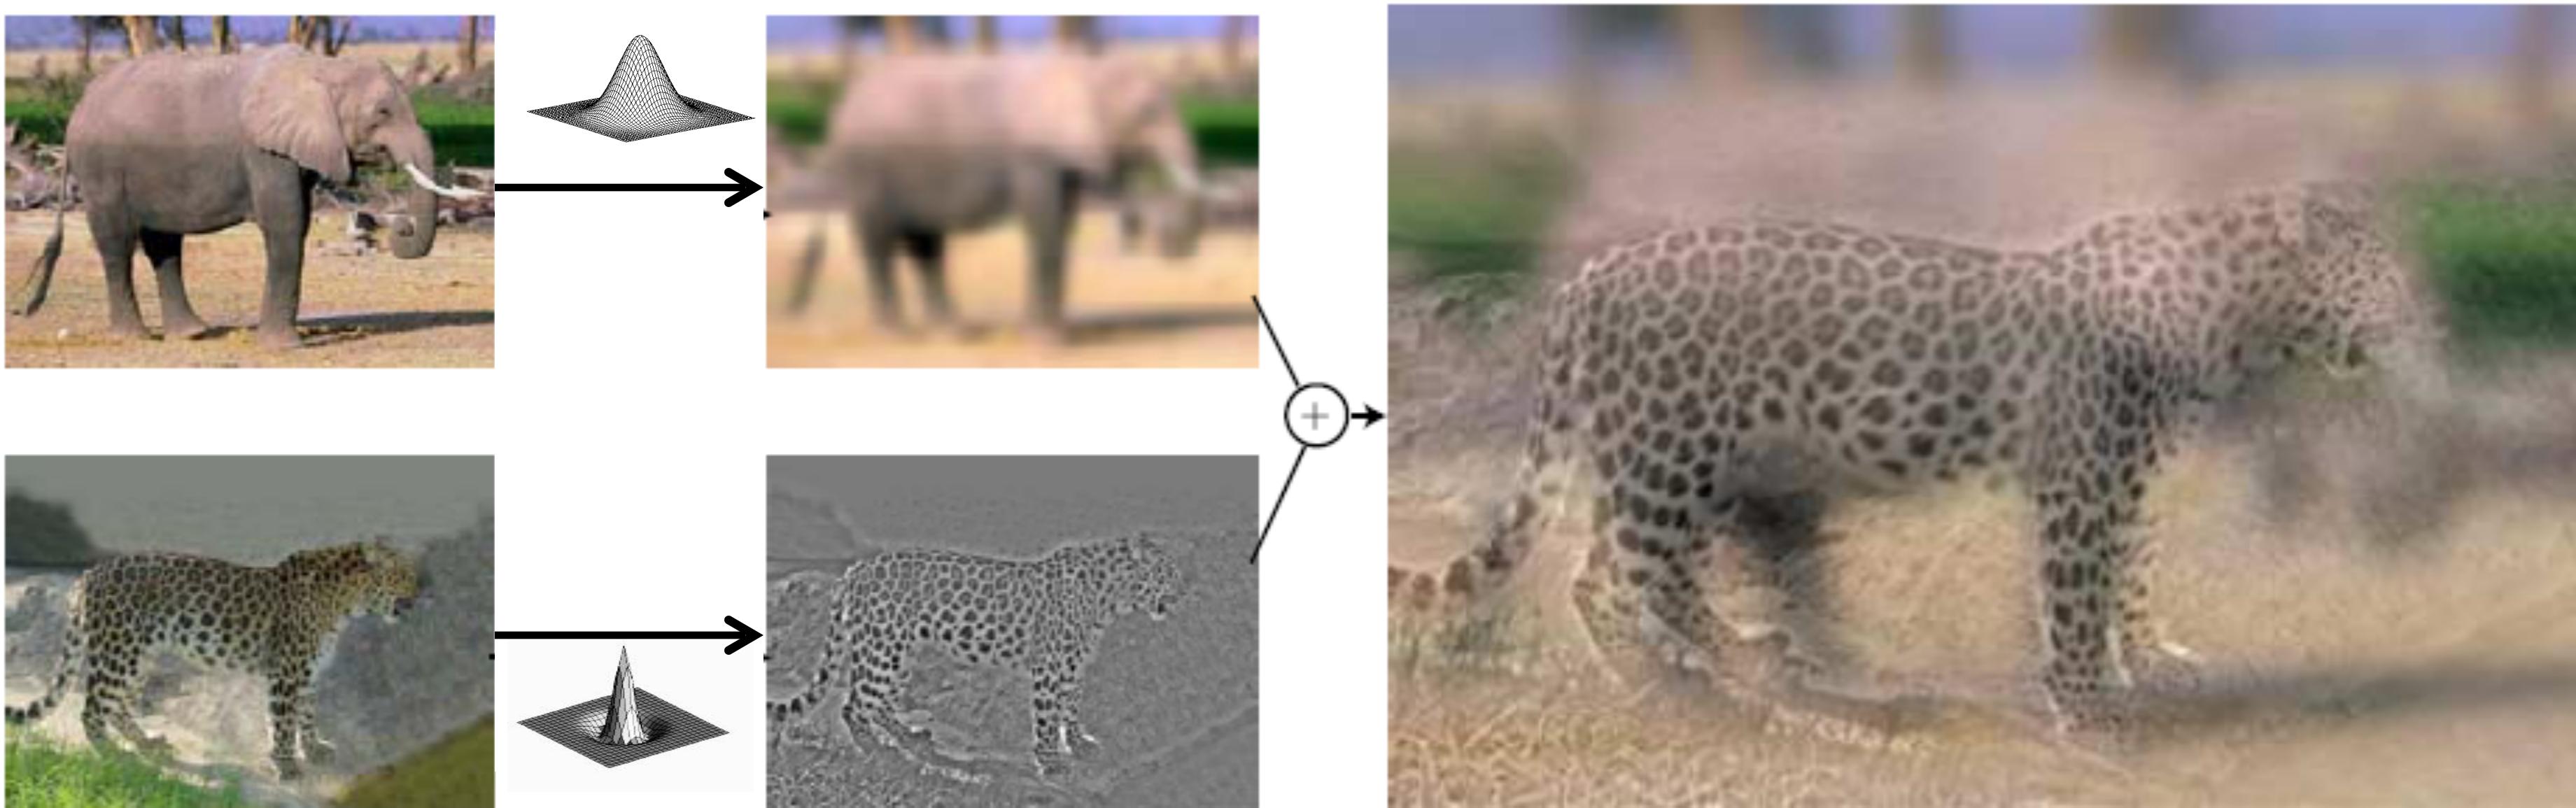
$$= I * (e - g_\sigma)$$



Hybrid Images

A. Oliva, A. Torralba, P.G. Schyns, “[Hybrid Images](#),” SIGGRAPH 2006

Gaussian Filter



Laplacian Filter



Changing expression



SIGGRAPH2006

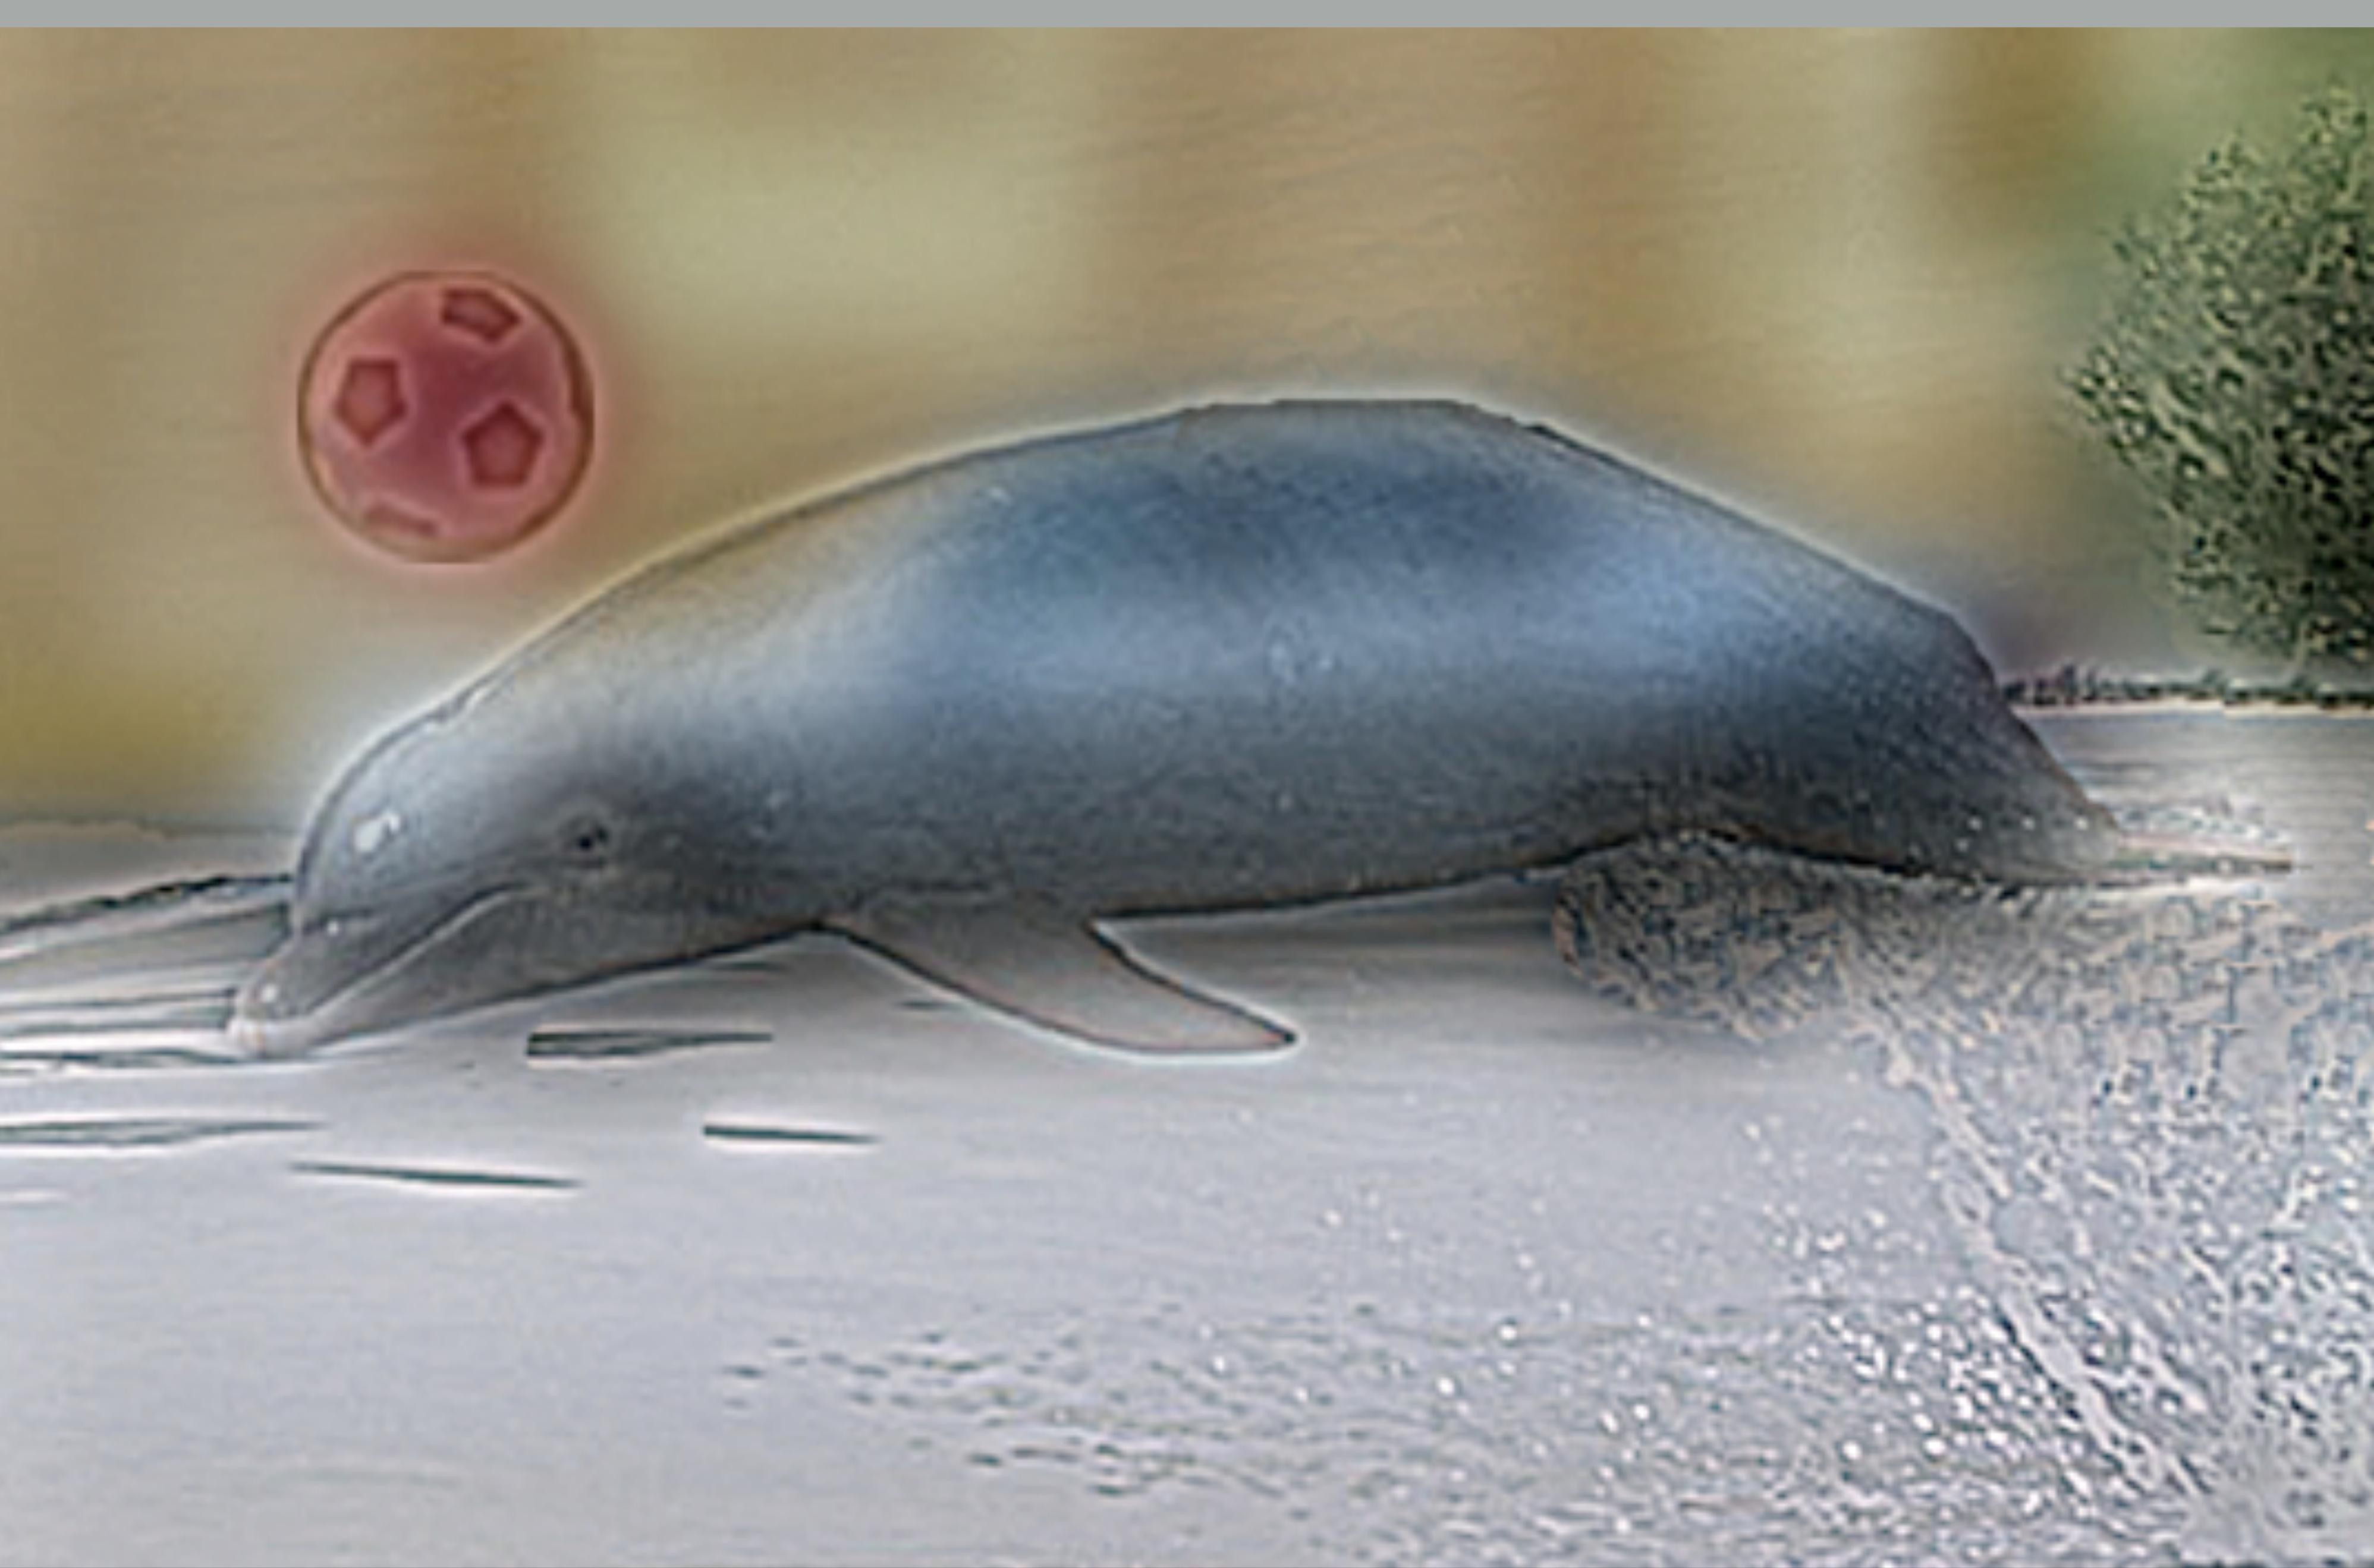
Sad



Surprised

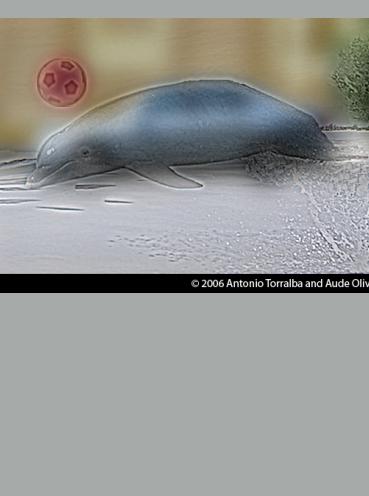






dolphin and car

© 2006 Antonio Torralba and Aude Oliva



© 2006 Antonio Torralba and Aude Oliva

© 2006 Antonio Torralba and Aude Oliva

Edge detection



[Winter in Kraków photographed by Marcin Ryczek](#)

Edge detection

Goal: Identify sudden changes (discontinuities) in an image

- Intuitively, most semantic and shape information from the image can be encoded in the edges
- More compact than pixels

Ideal: artist's line drawing (but artist is also using object-level knowledge)

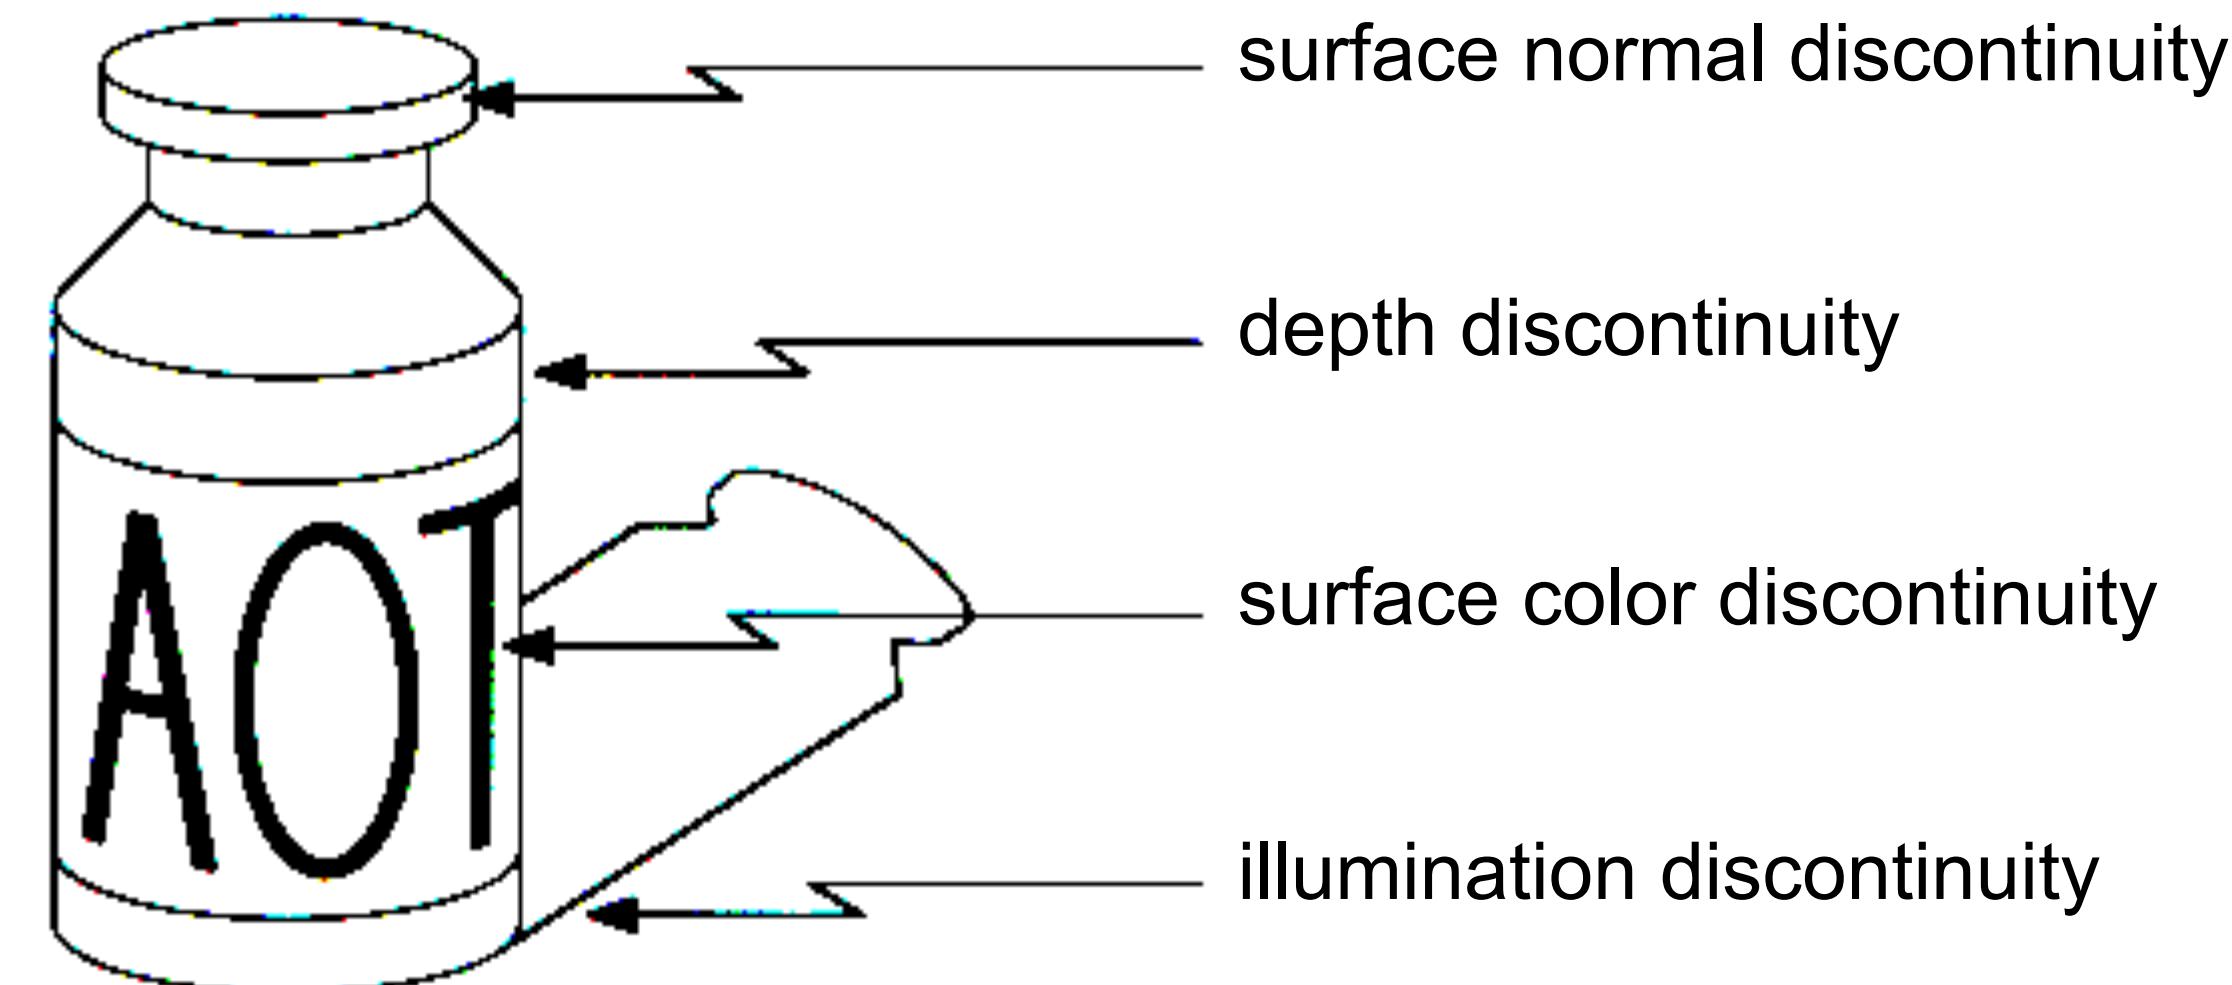


Attneave's Cat (1954)

Source: D. Lowe

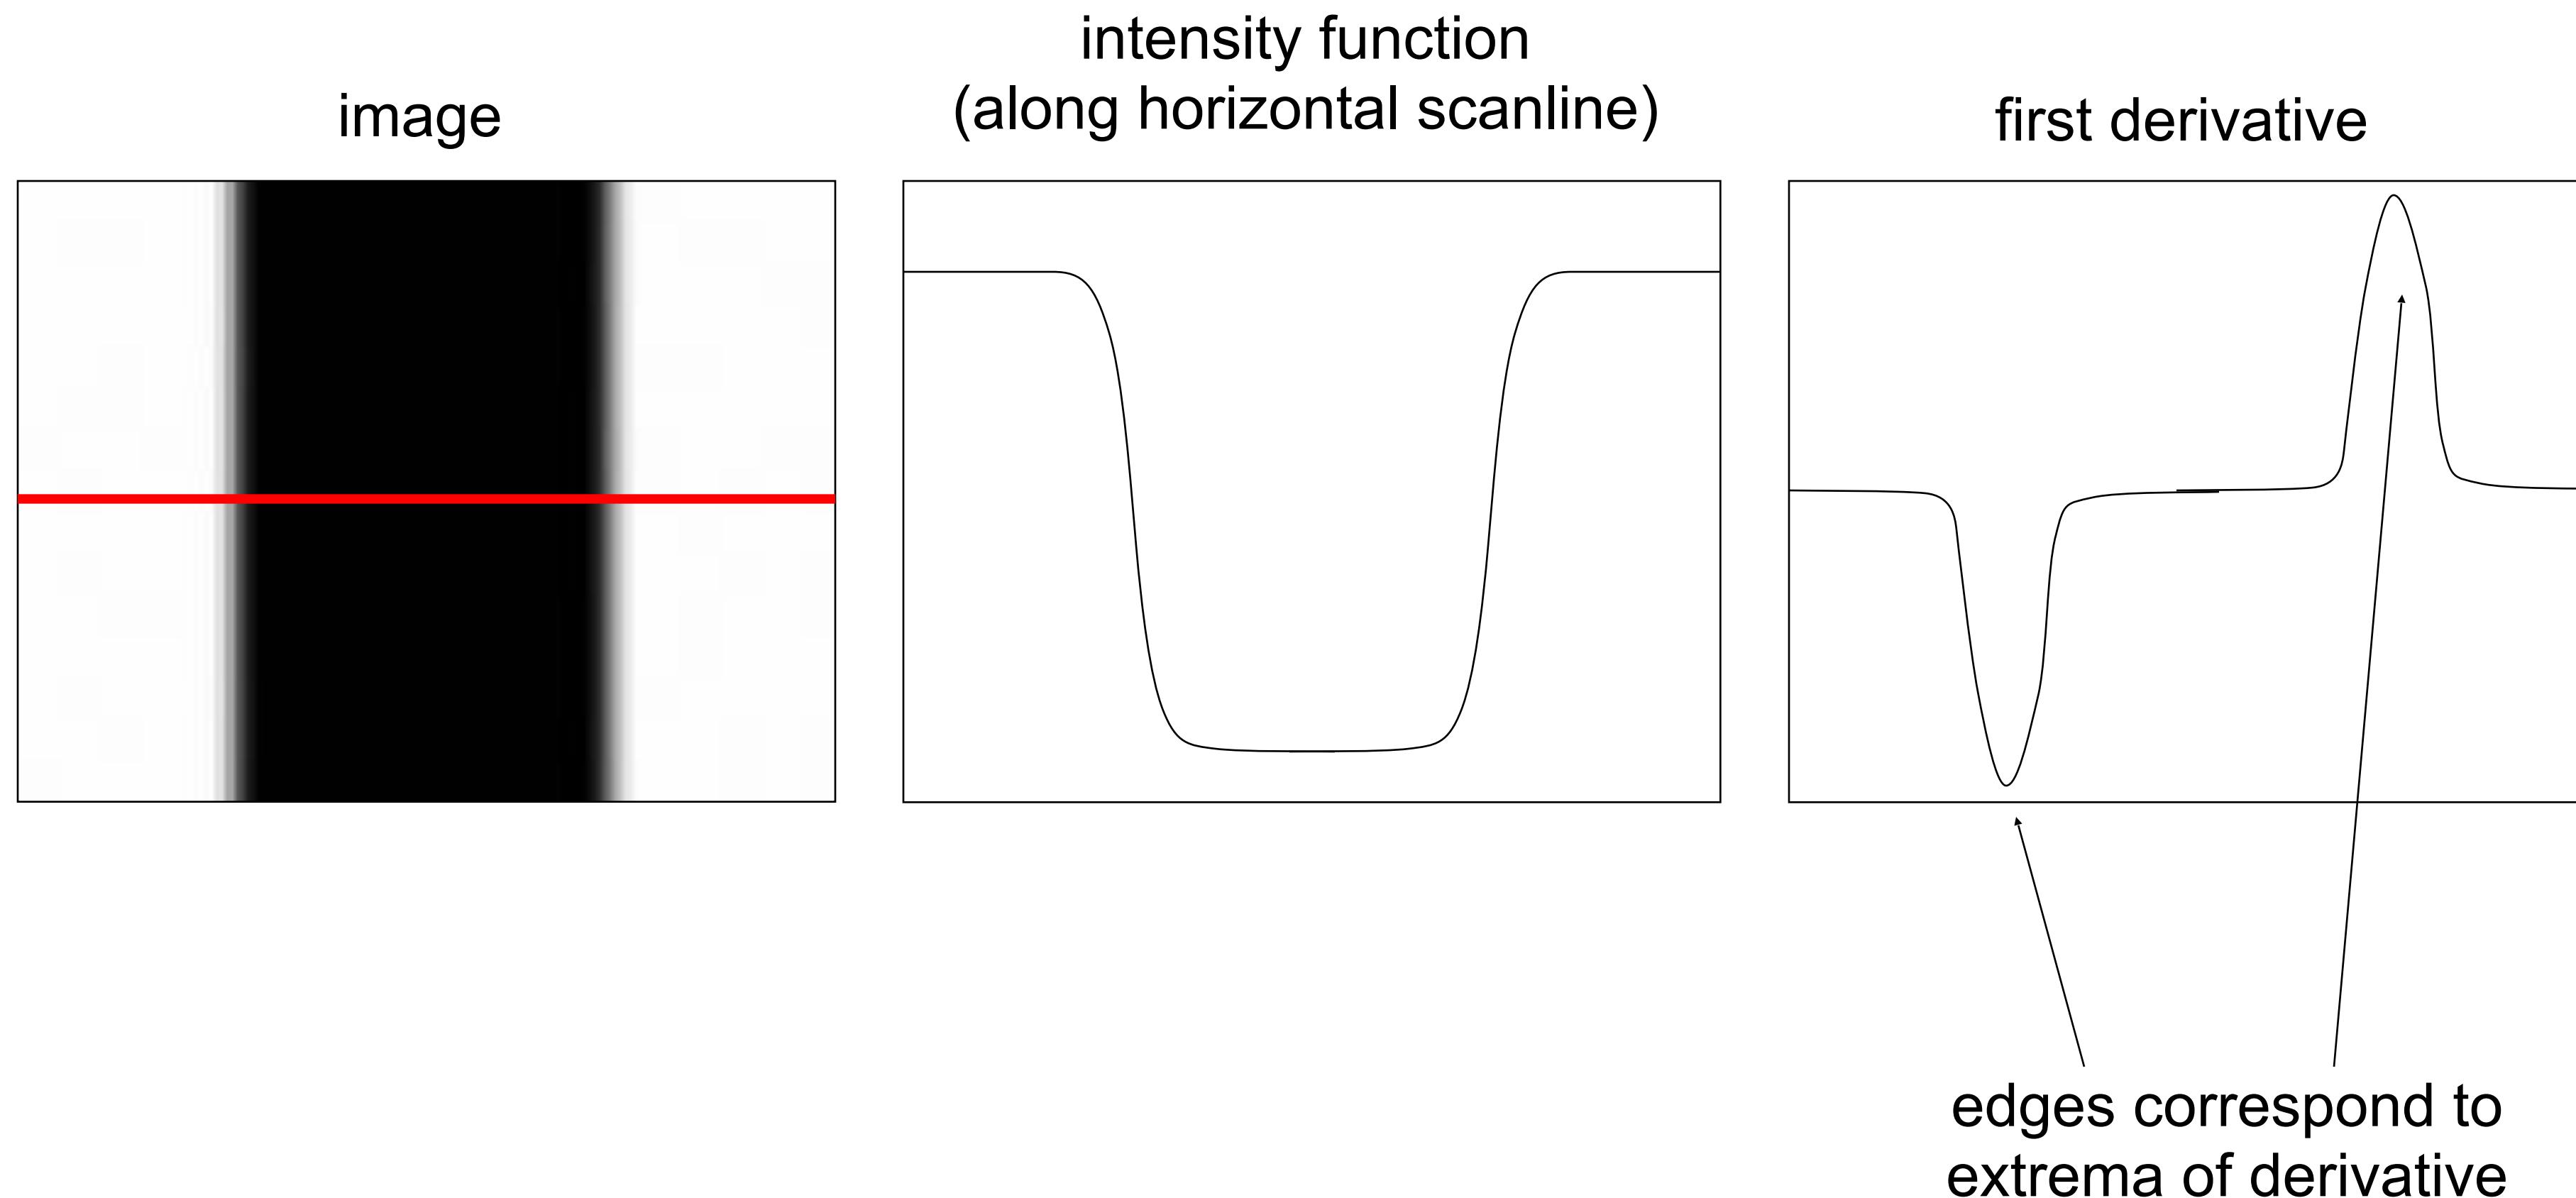
Origin of edges

Edges are caused by a variety of factors:



Edge detection

An edge is a place of rapid change in the image intensity function



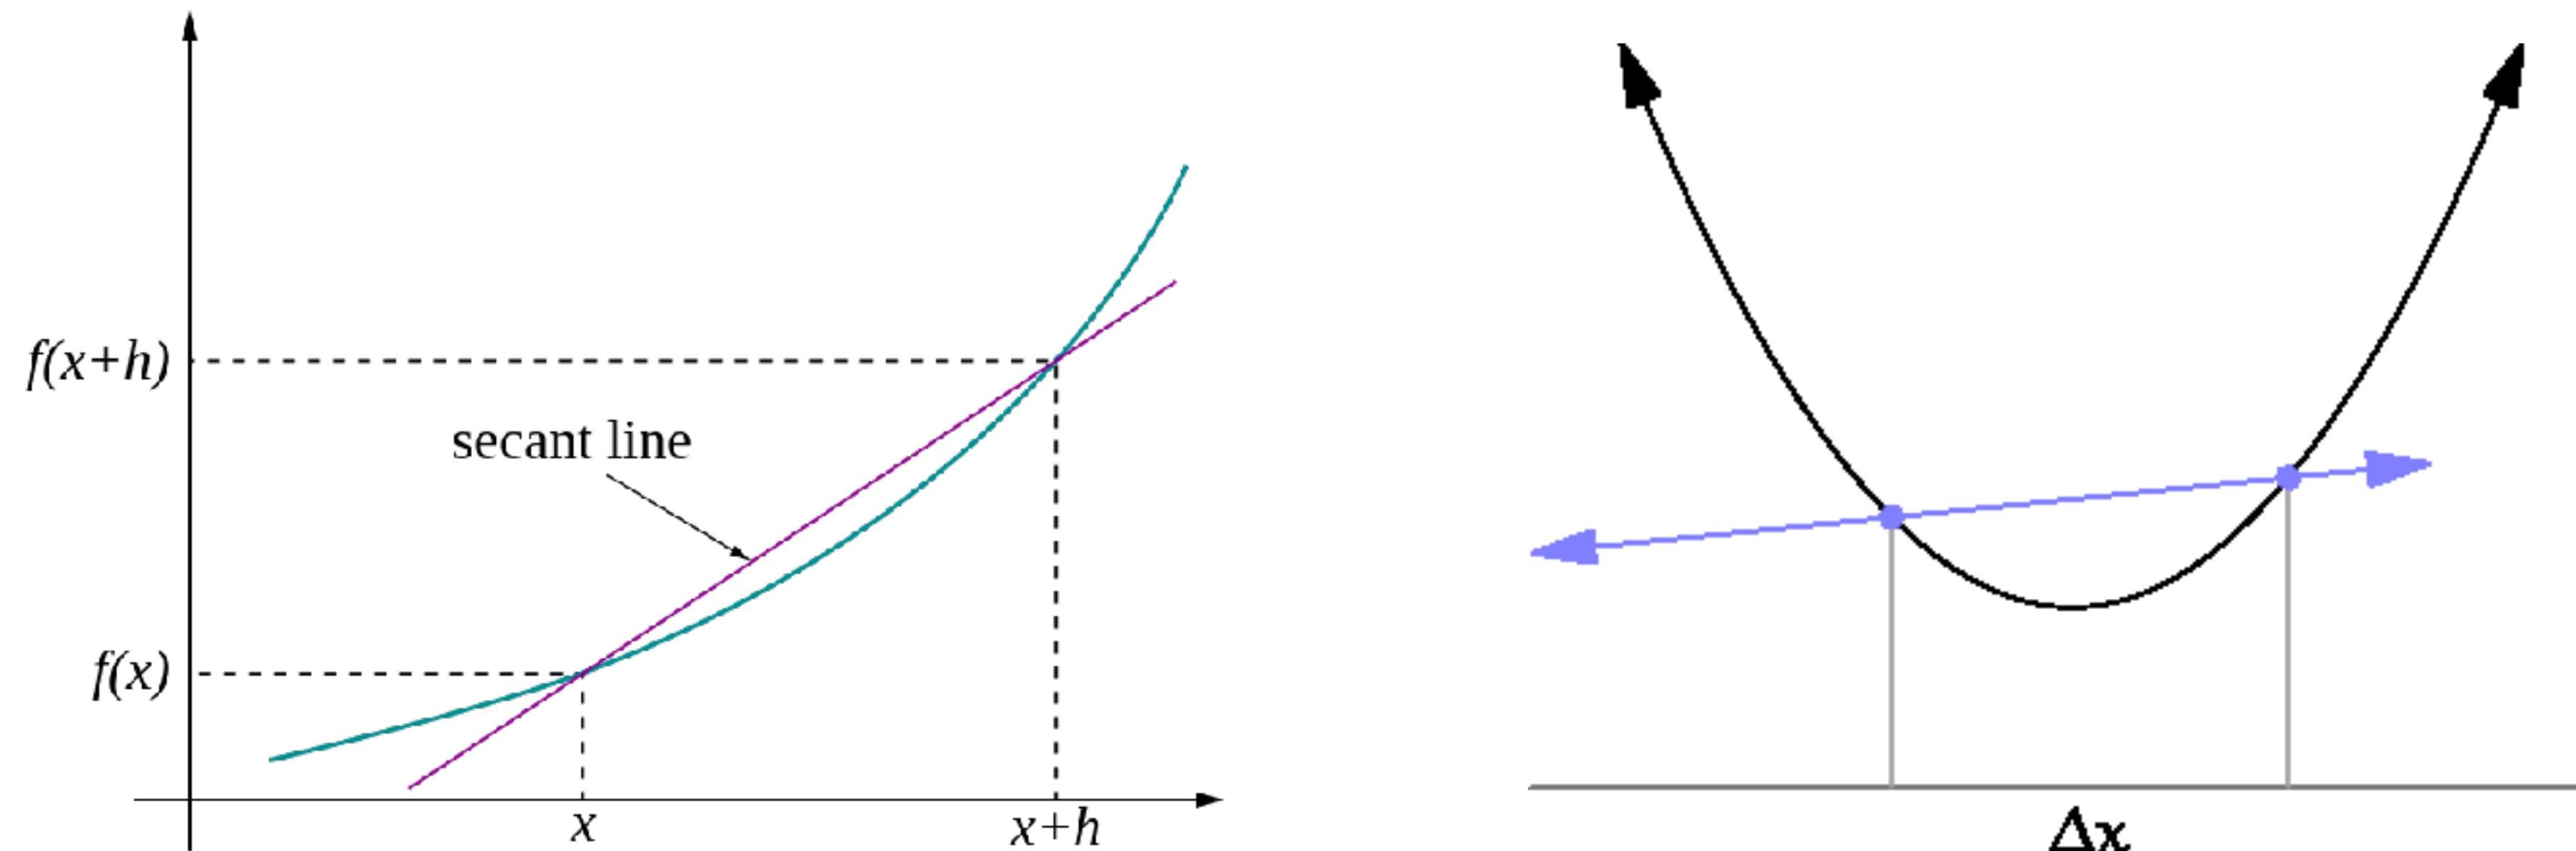
One dimensional derivatives

From Calc101

$$y = f(x)$$

Gradient $m = \frac{\Delta y}{\Delta x}$

$$m = \frac{f(x + h) - f(x)}{(x + h) - x} = \frac{f(x + h) - f(x)}{h}$$

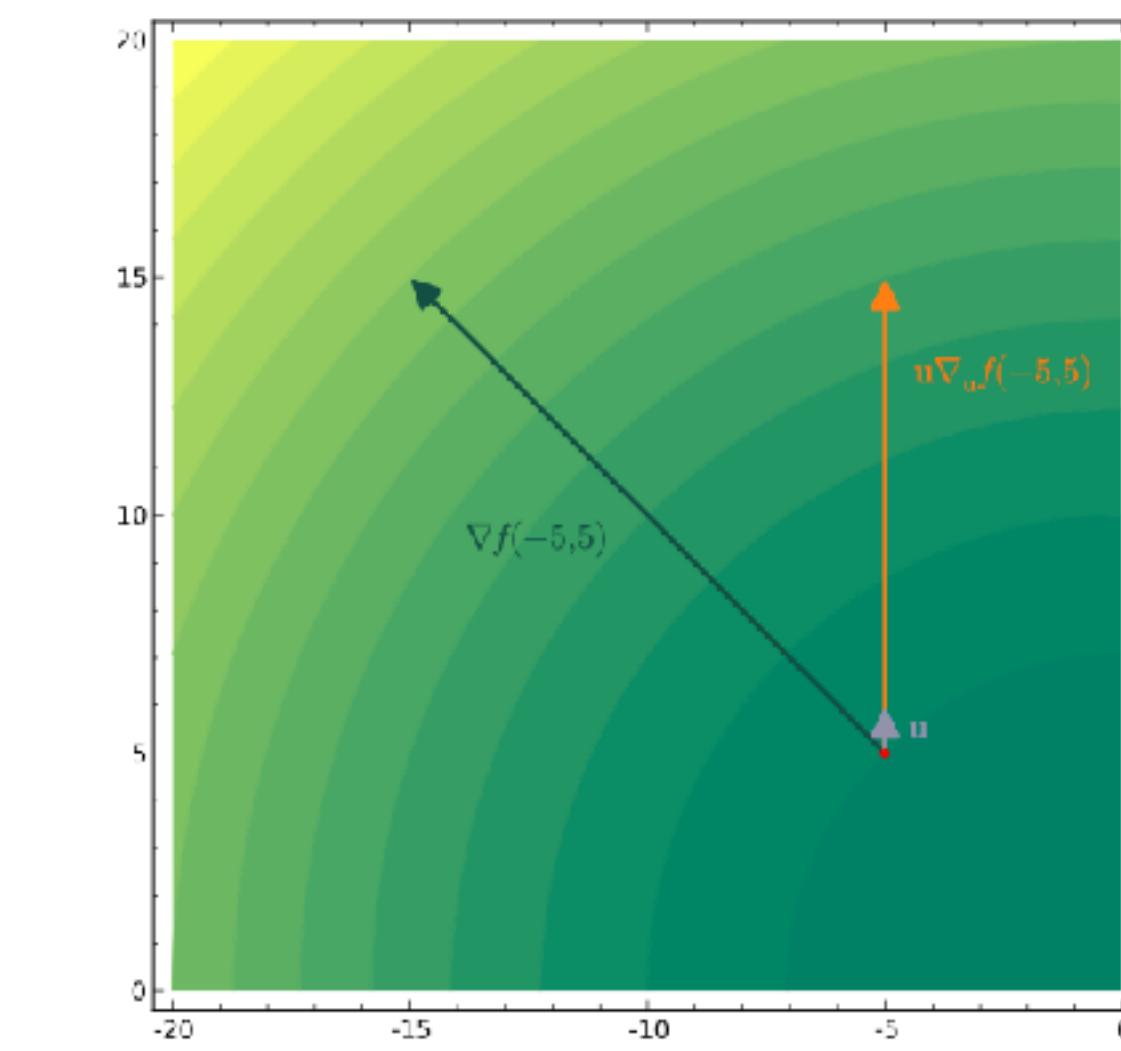


<https://en.wikipedia.org/wiki/Derivative>

Two dimensional derivatives

For 2D function $f(\mathbf{x})$, one can compute a derivative for each direction \mathbf{v}

$$\nabla_{\mathbf{v}} f(\mathbf{x}) = \lim_{h \rightarrow 0} \frac{f(\mathbf{x} + h\mathbf{v}) - f(\mathbf{x})}{h}.$$



Directional derivatives of the function along the axes are called partial derivatives.
For example the partial derivative with respect to x is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

Partial derivatives with convolutions

For 2D function $f(x,y)$, the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

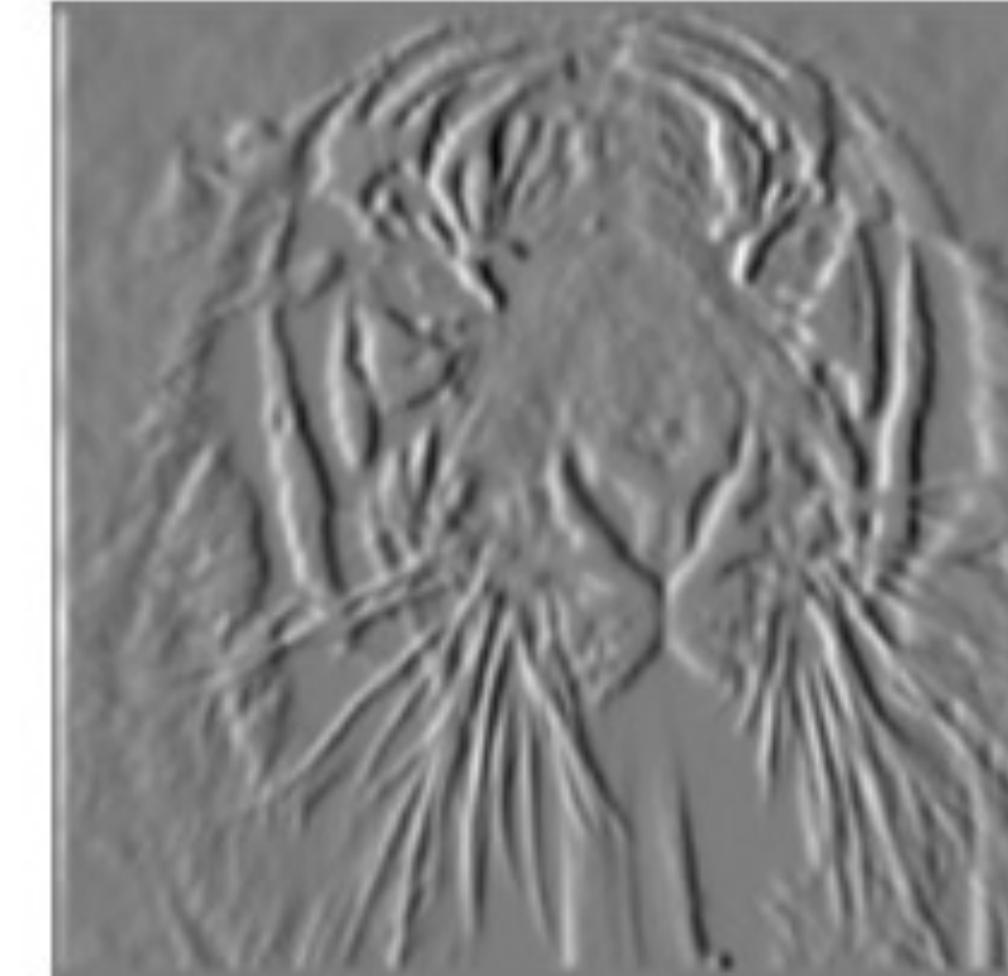
Question: To implement the above as correlation, what would be the associated filter?

Partial derivatives of an image



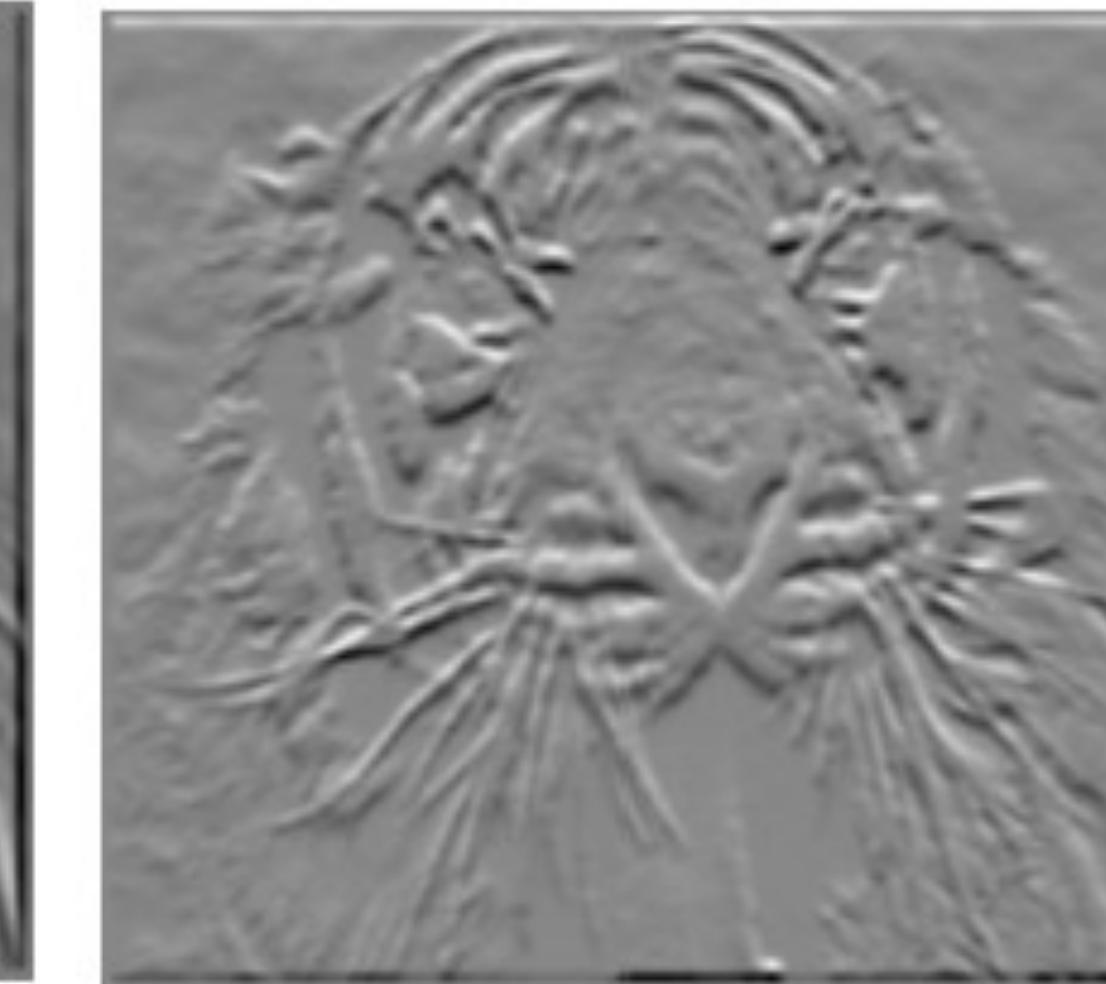
$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

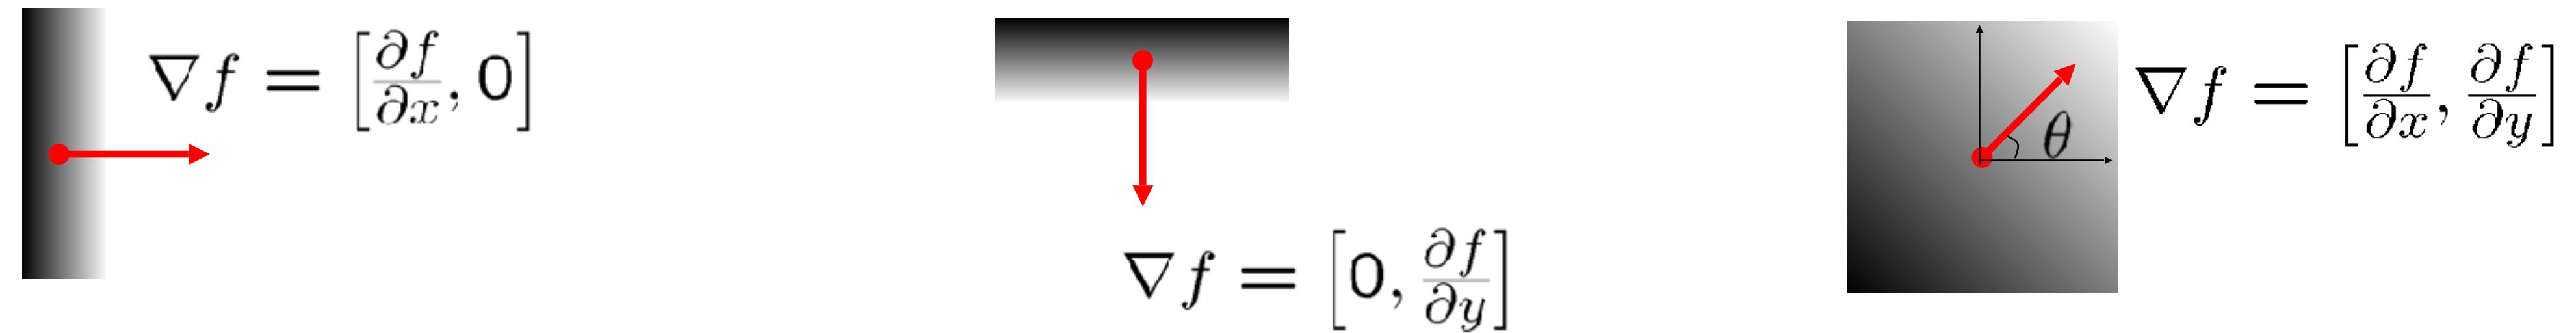
-1	1
1	-1



Which one shows changes with respect to x?

Image gradient

The gradient of an image: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$



The gradient points in the direction of most rapid increase in intensity

How does this direction relate to the direction of the edge? — they are orthogonal

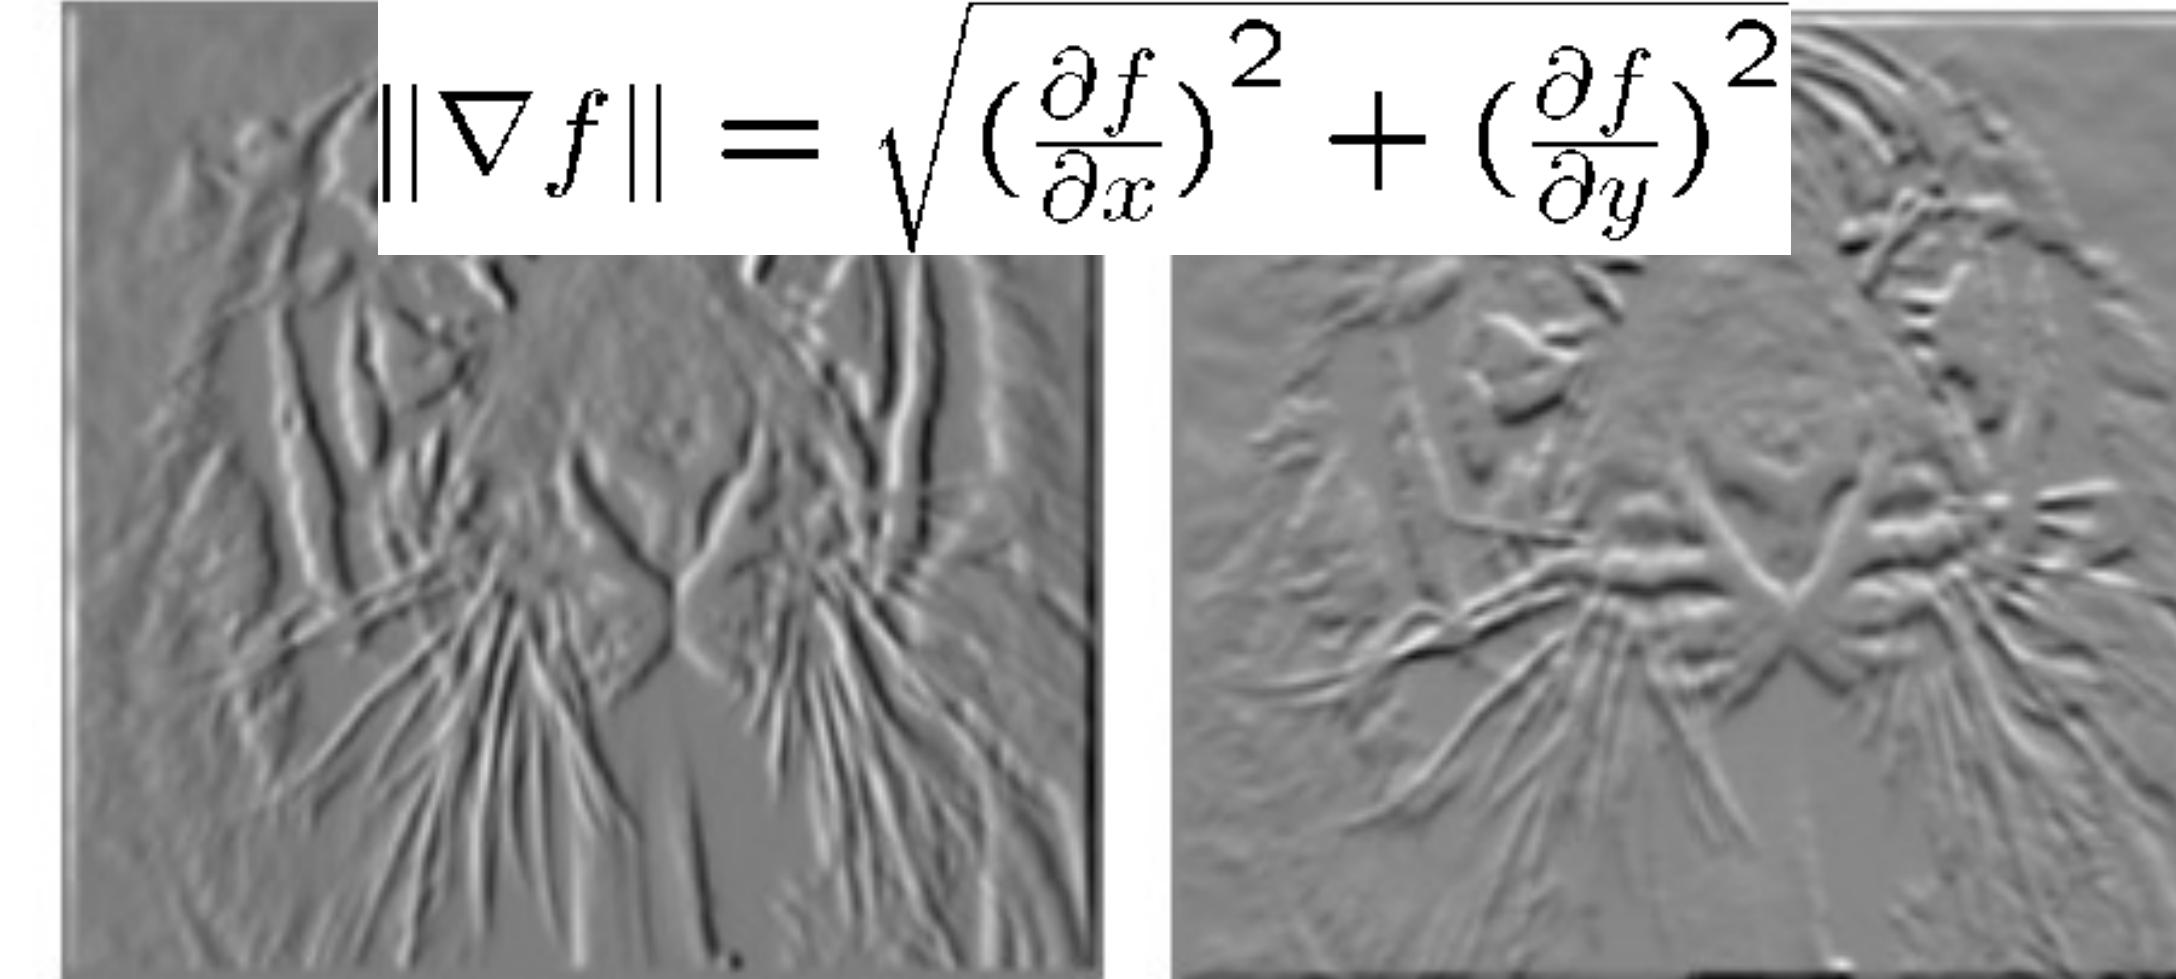
The gradient direction is given by $\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$

The gradient strength is given by the magnitude $\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$

Partial derivatives of an image

$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

-1	1
1	-1

Which one shows changes with respect to x?

Edge detection in Python

```
~ — IPython: Users/smaji — ipython
File Edit Options Buffers Tools Python Help
import numpy as np
import scipy.ndimage as ndi
import matplotlib.pyplot as plt
from skimage import data

im = data.checkerboard()

# Filters along x and y
fx = np.array([[-1, 0, 1], [-1, 0, 1], [-1, 0, 1]])
fy = fx.transpose()

# Apply filters and compute magnitude
gx = ndi.correlate(im, fx)
gy = ndi.correlate(im, fy)
mag = np.sqrt(np.maximum(gx**2 + gy**2, 0))

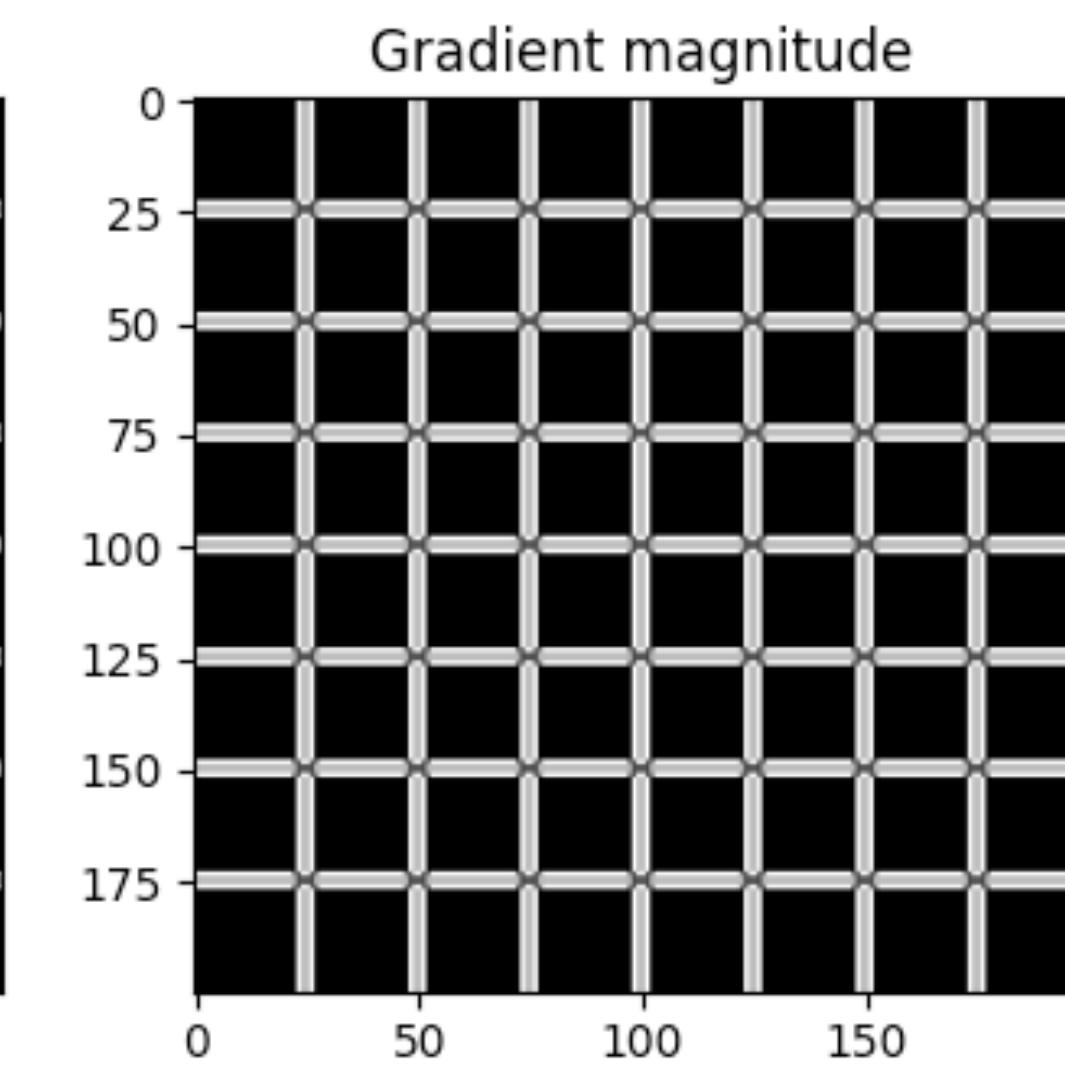
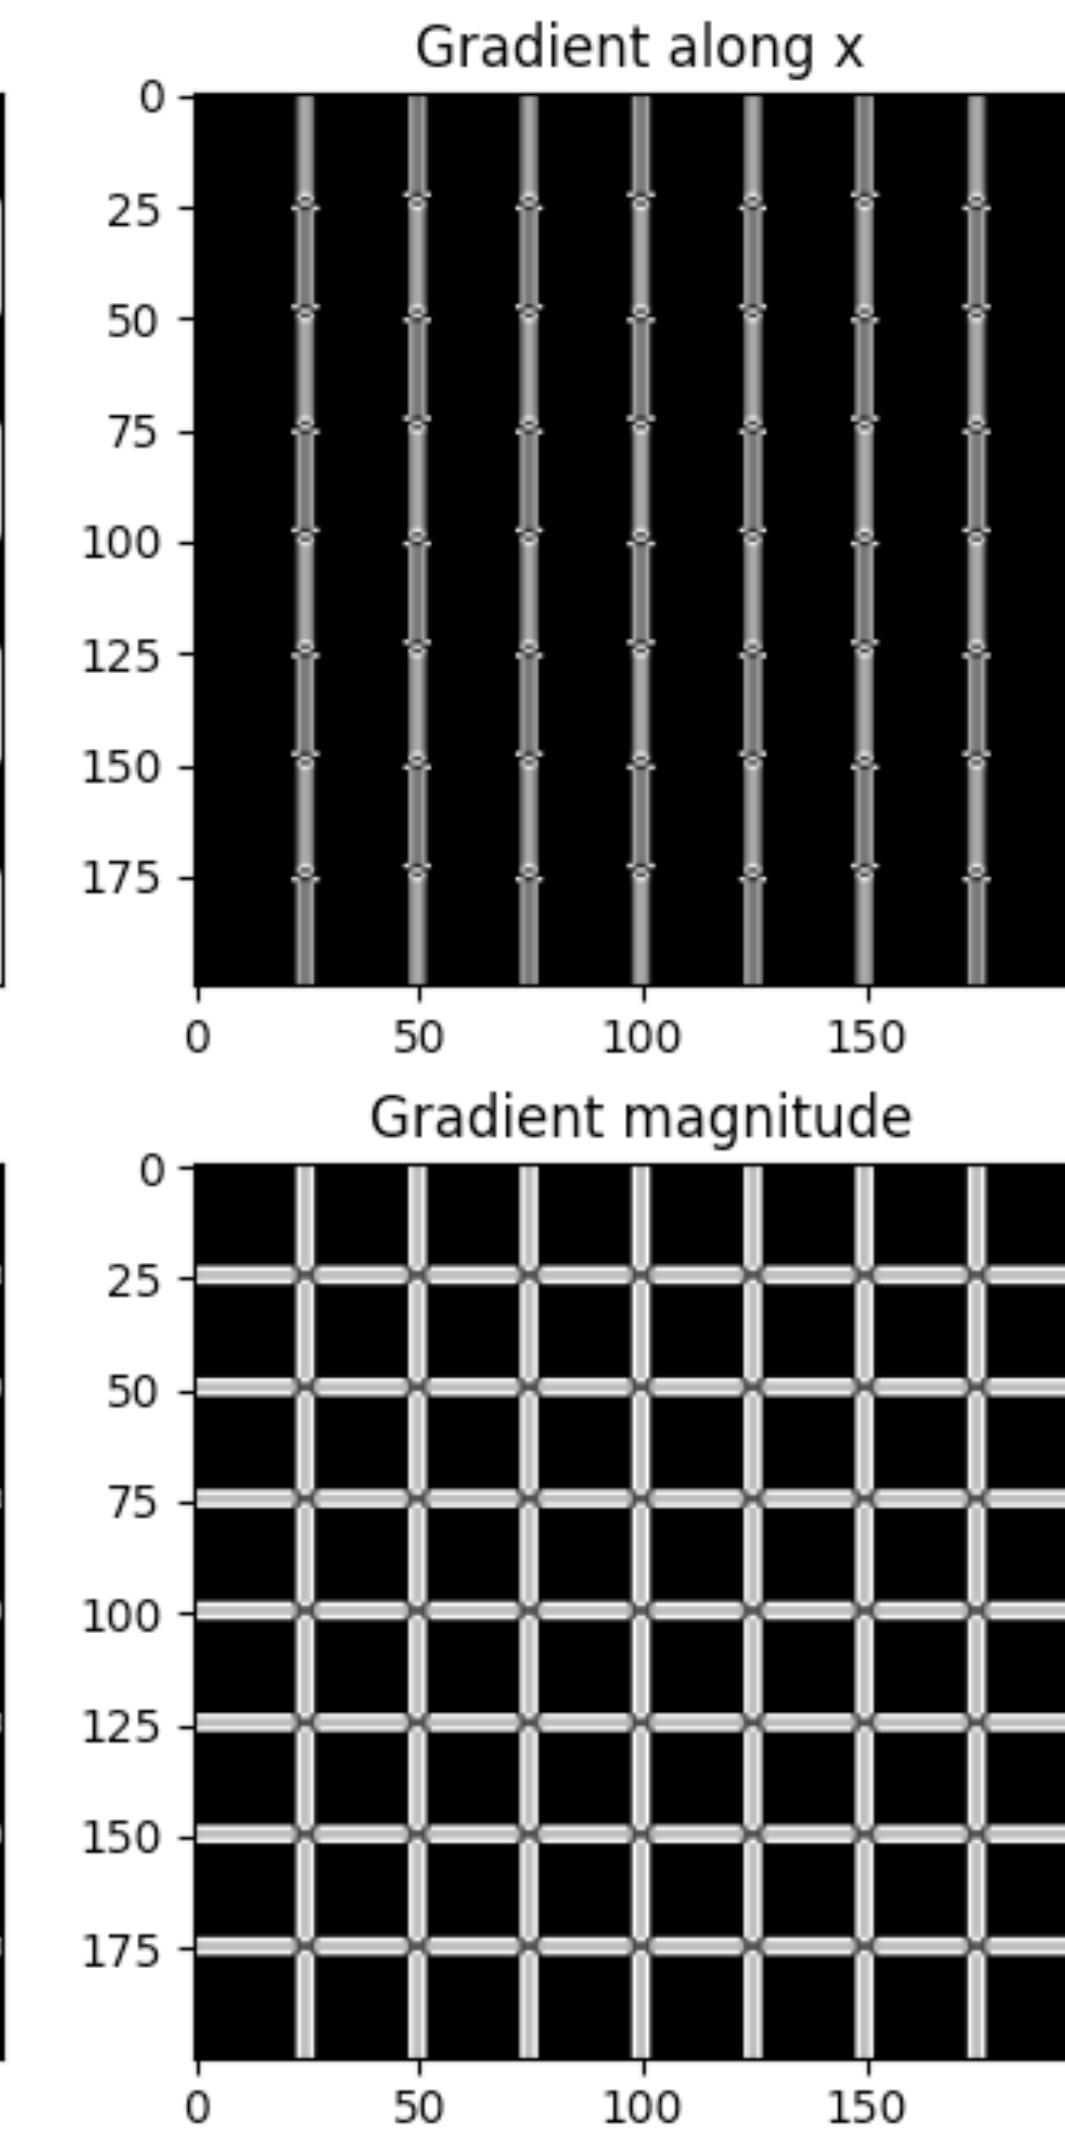
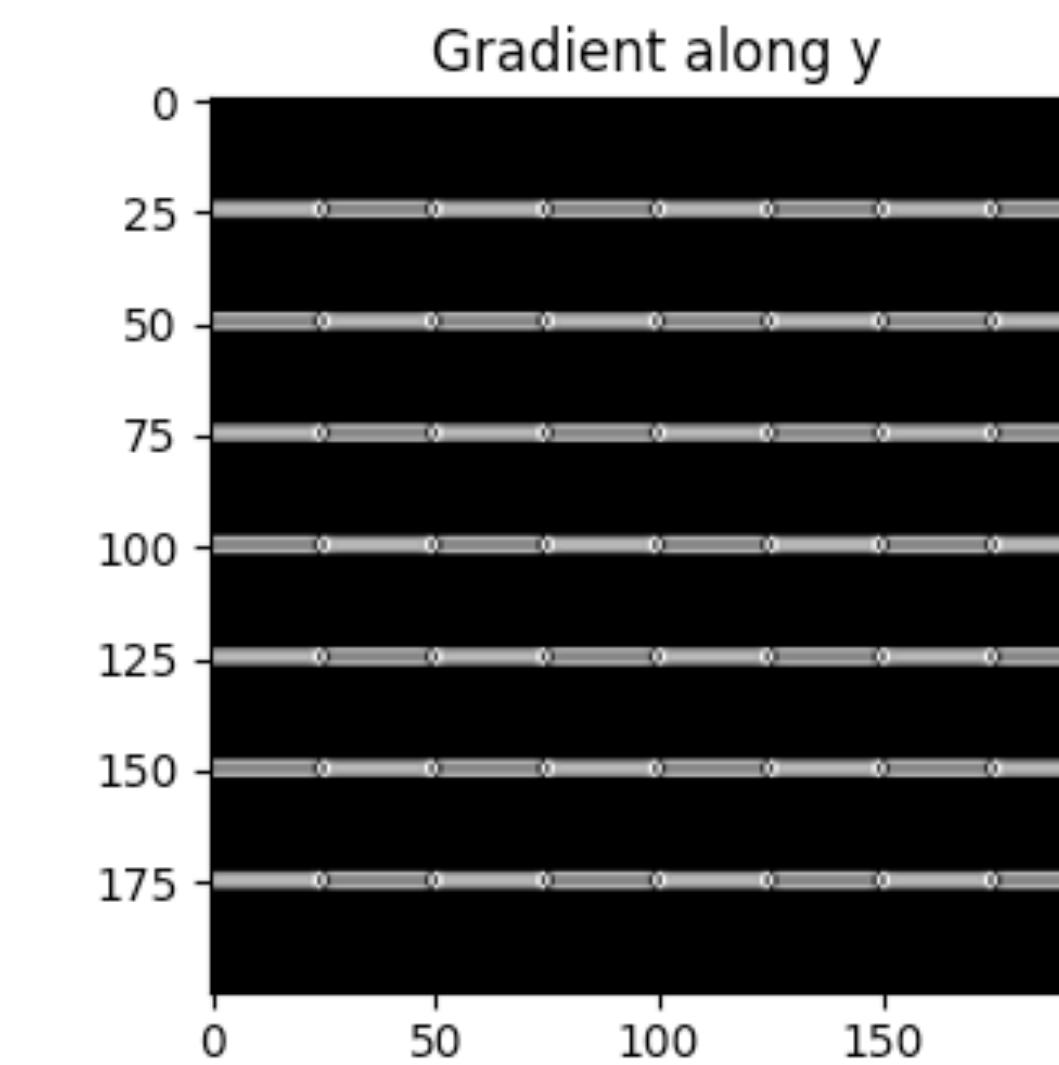
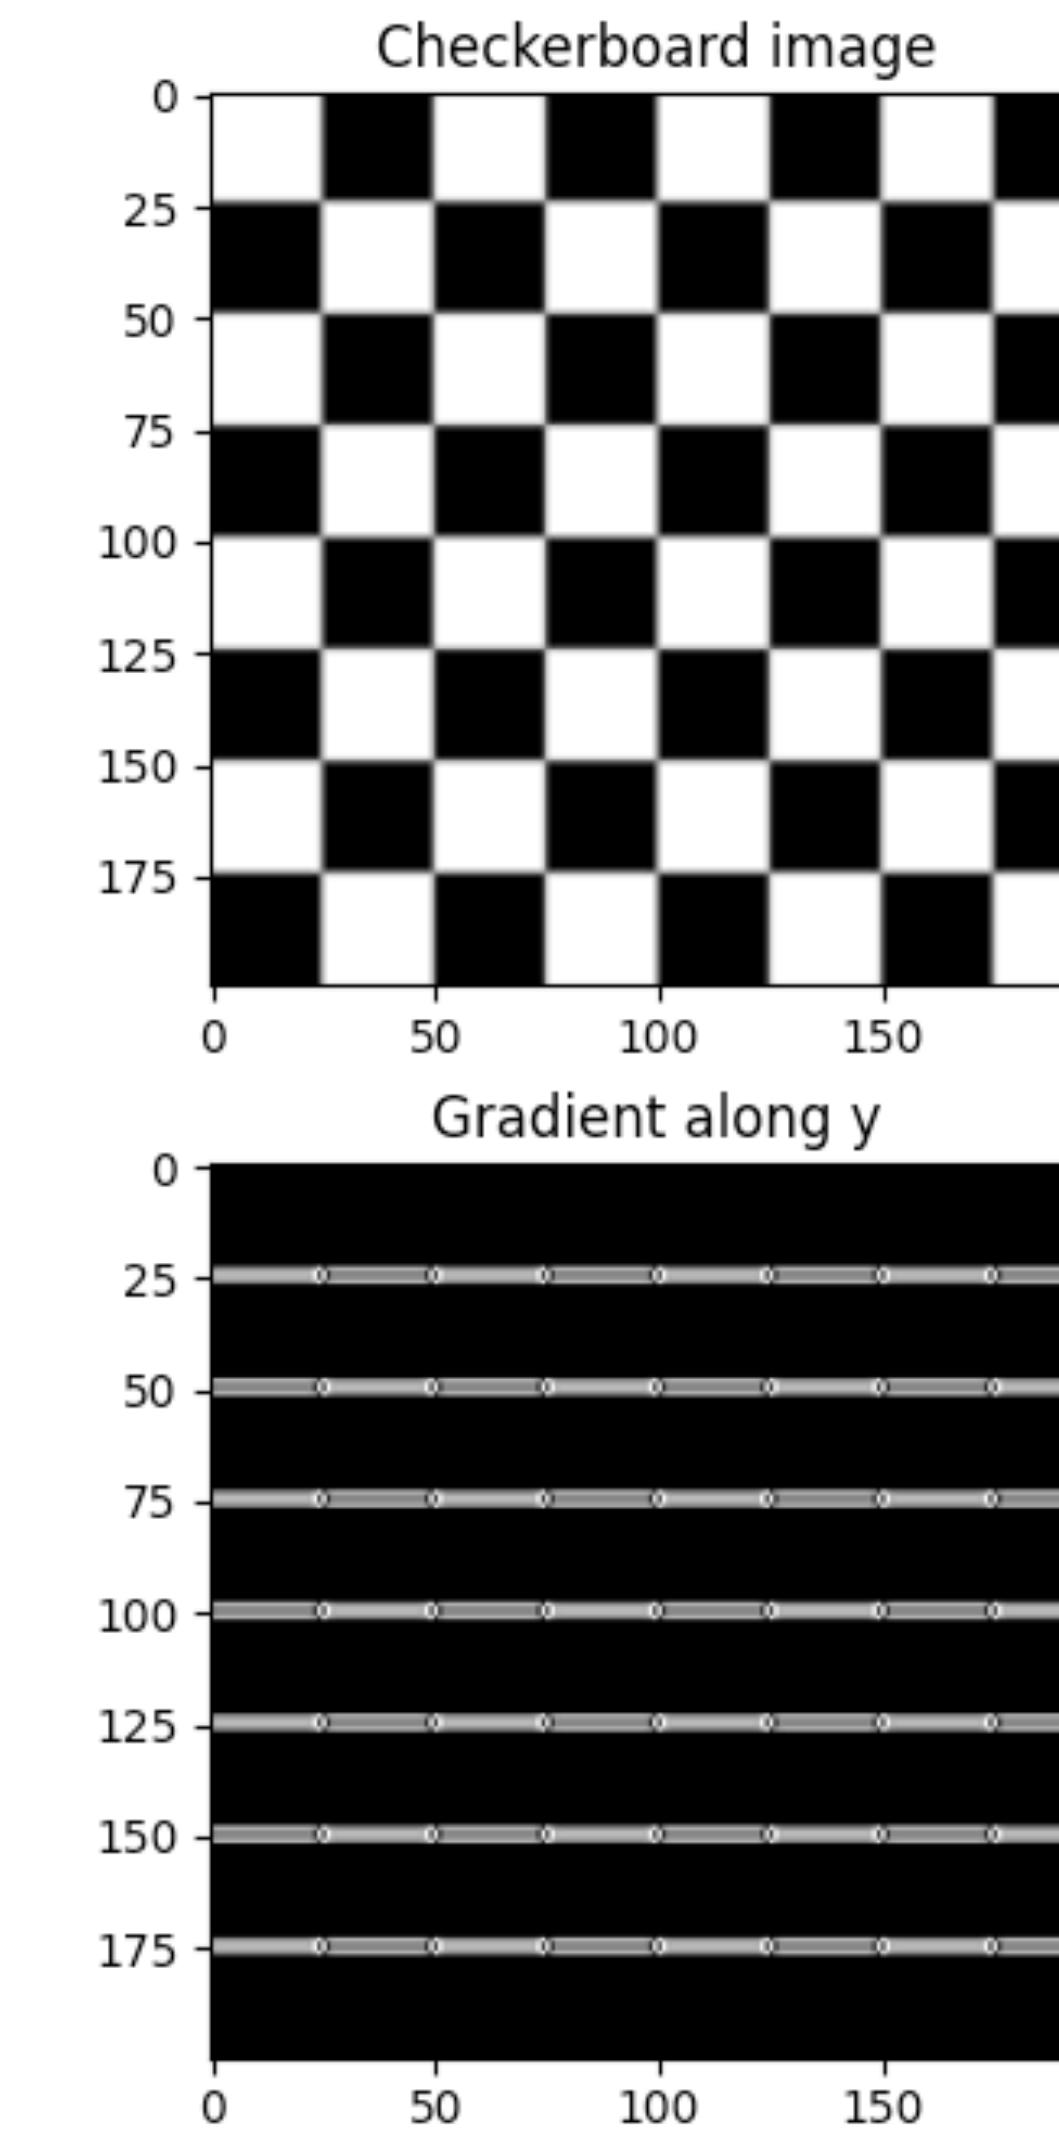
# Optionally convert this to a 0-255 image for display
mag = np.uint8(mag/mag.max()*255)

# Visualize outputs
plt.subplot(2,2,1)
plt.imshow(im, cmap='gray')
plt.title('Checkerboard image')

plt.subplot(2,2,2)
plt.imshow(gx, cmap='gray')
plt.title('Gradient along x')

plt.subplot(2,2,3)
plt.imshow(gy, cmap='gray')
plt.title('Gradient along y')

plt.subplot(2,2,4)
plt.imshow(mag, cmap='gray')
plt.title('Gradient magnitude')
plt.show()
```



Edge detection example



image



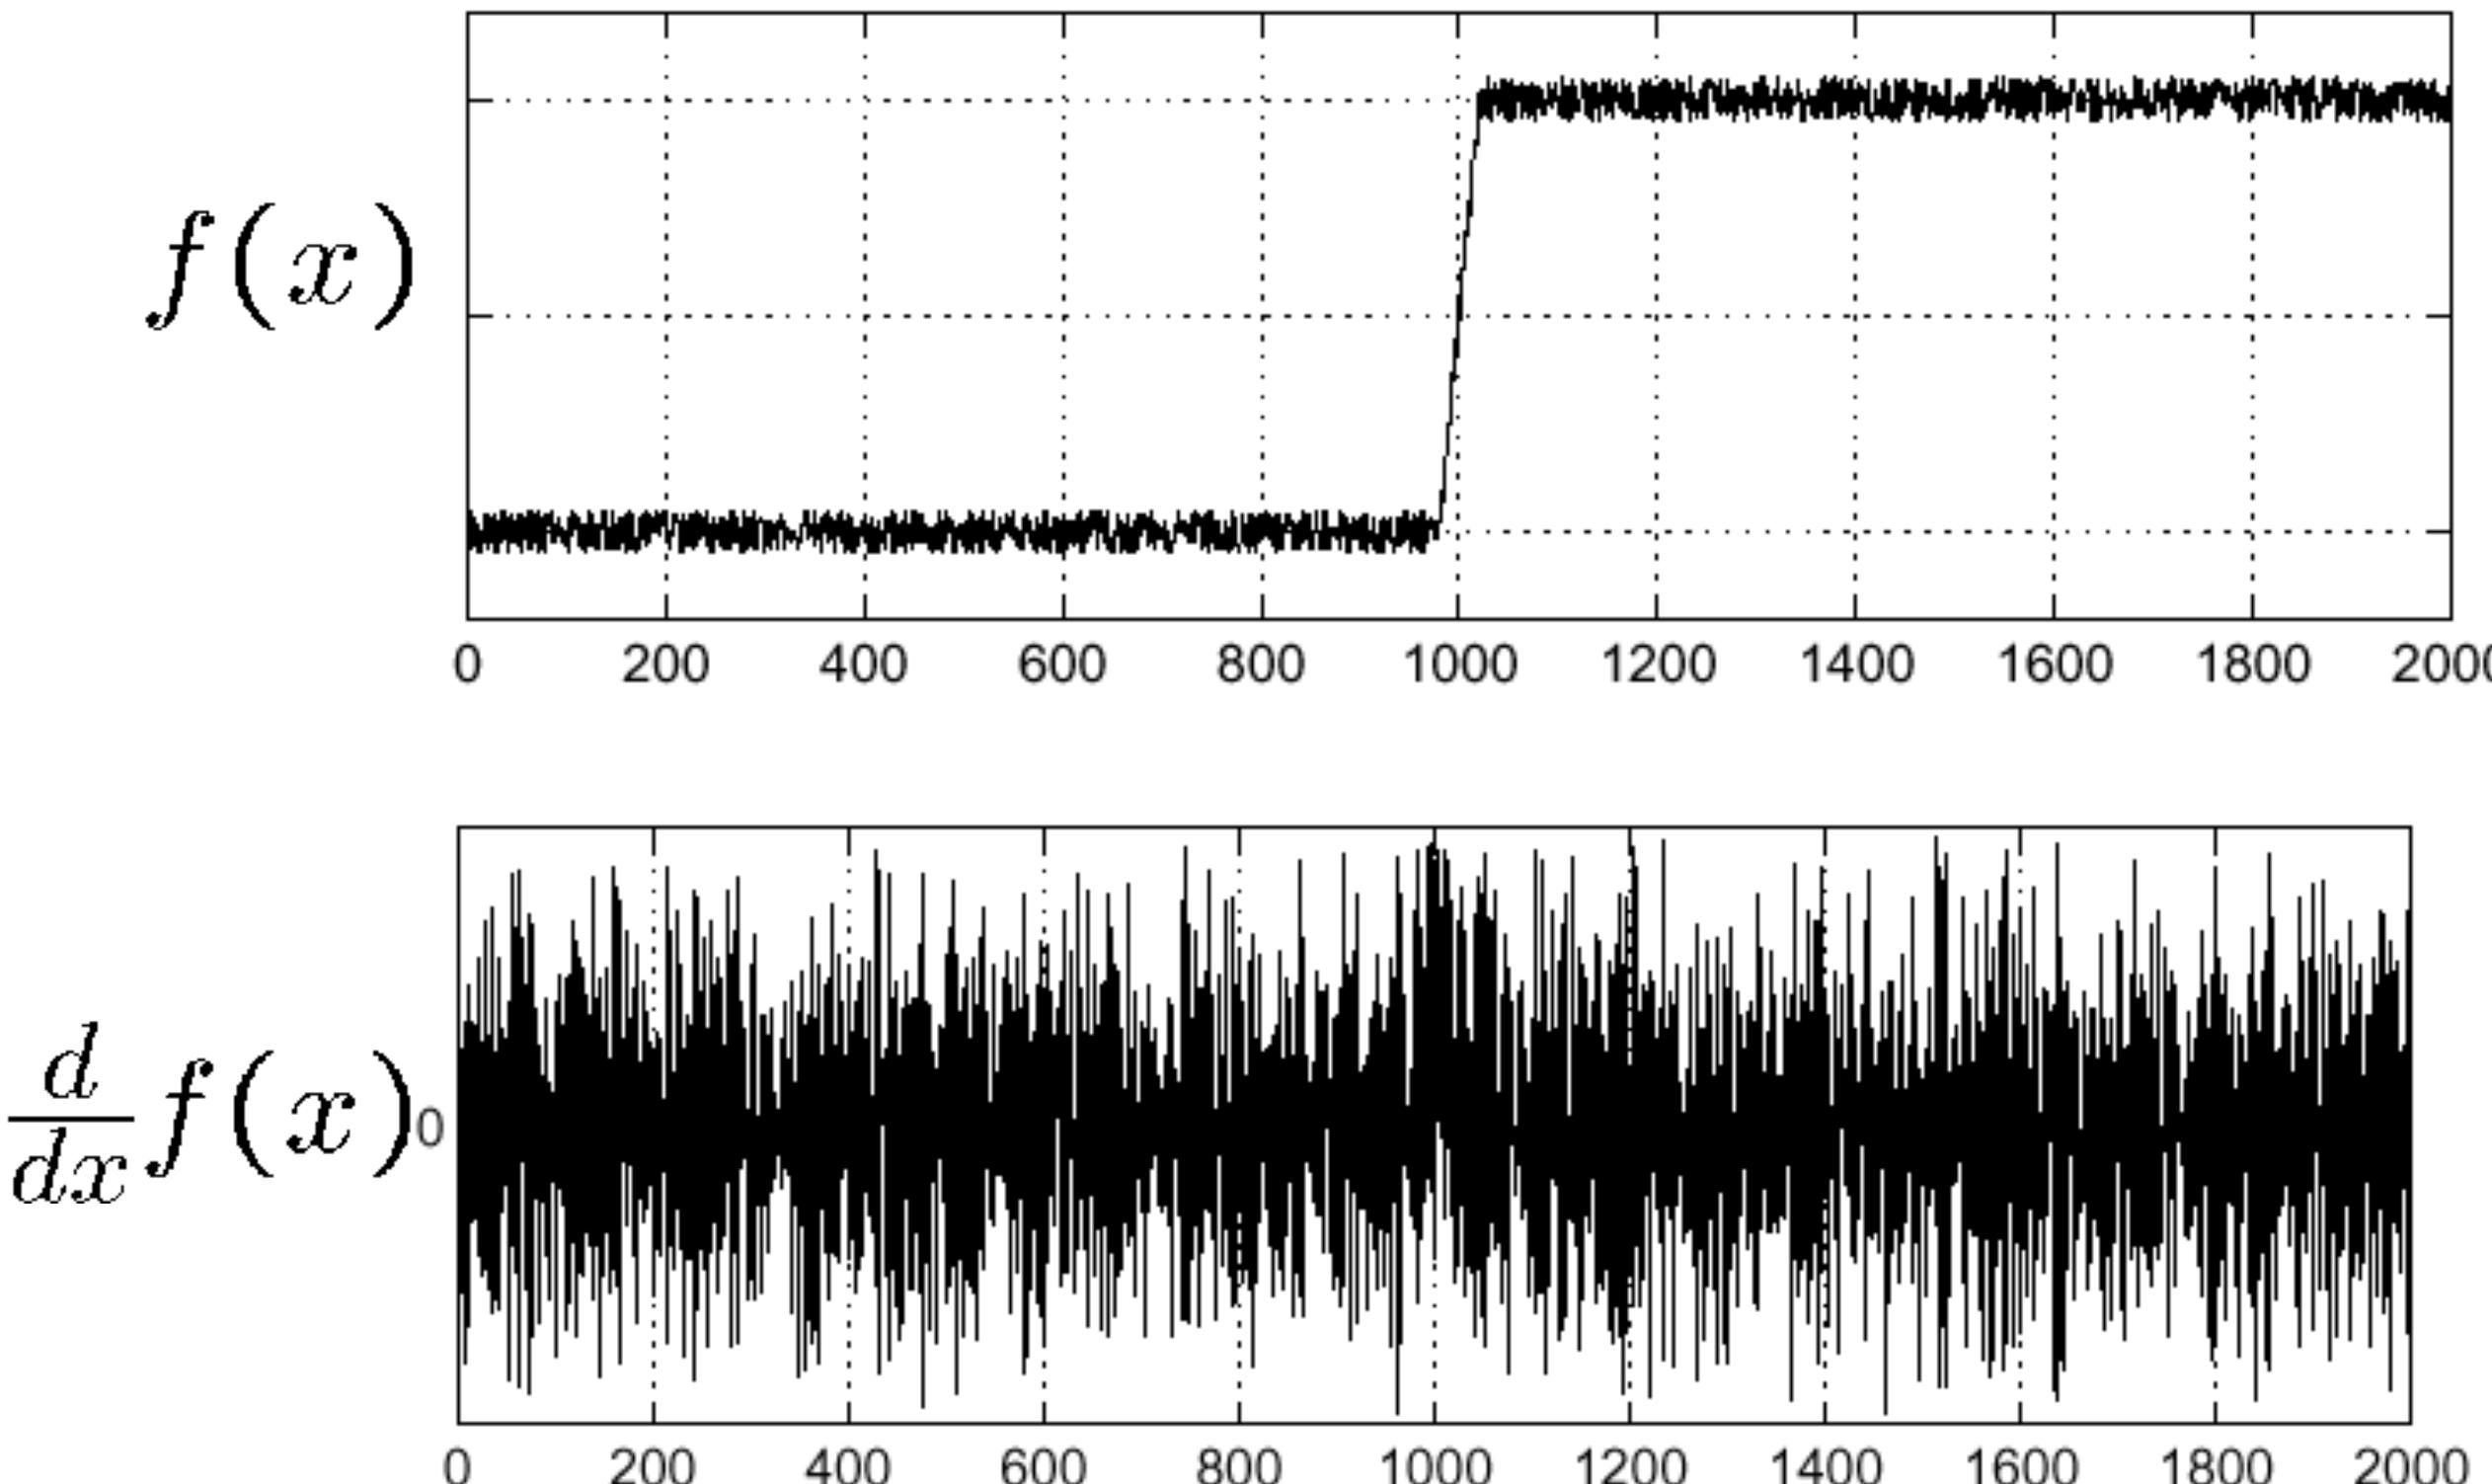
edge magnitude

$$\mathbf{G}_x = \begin{bmatrix} -1 & 0 & +1 \\ -1 & 0 & +1 \\ -1 & 0 & +1 \end{bmatrix} * \mathbf{A} \quad \text{and} \quad \mathbf{G}_y = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ +1 & +1 & +1 \end{bmatrix} * \mathbf{A}$$

https://en.wikipedia.org/wiki/Prewitt_operator

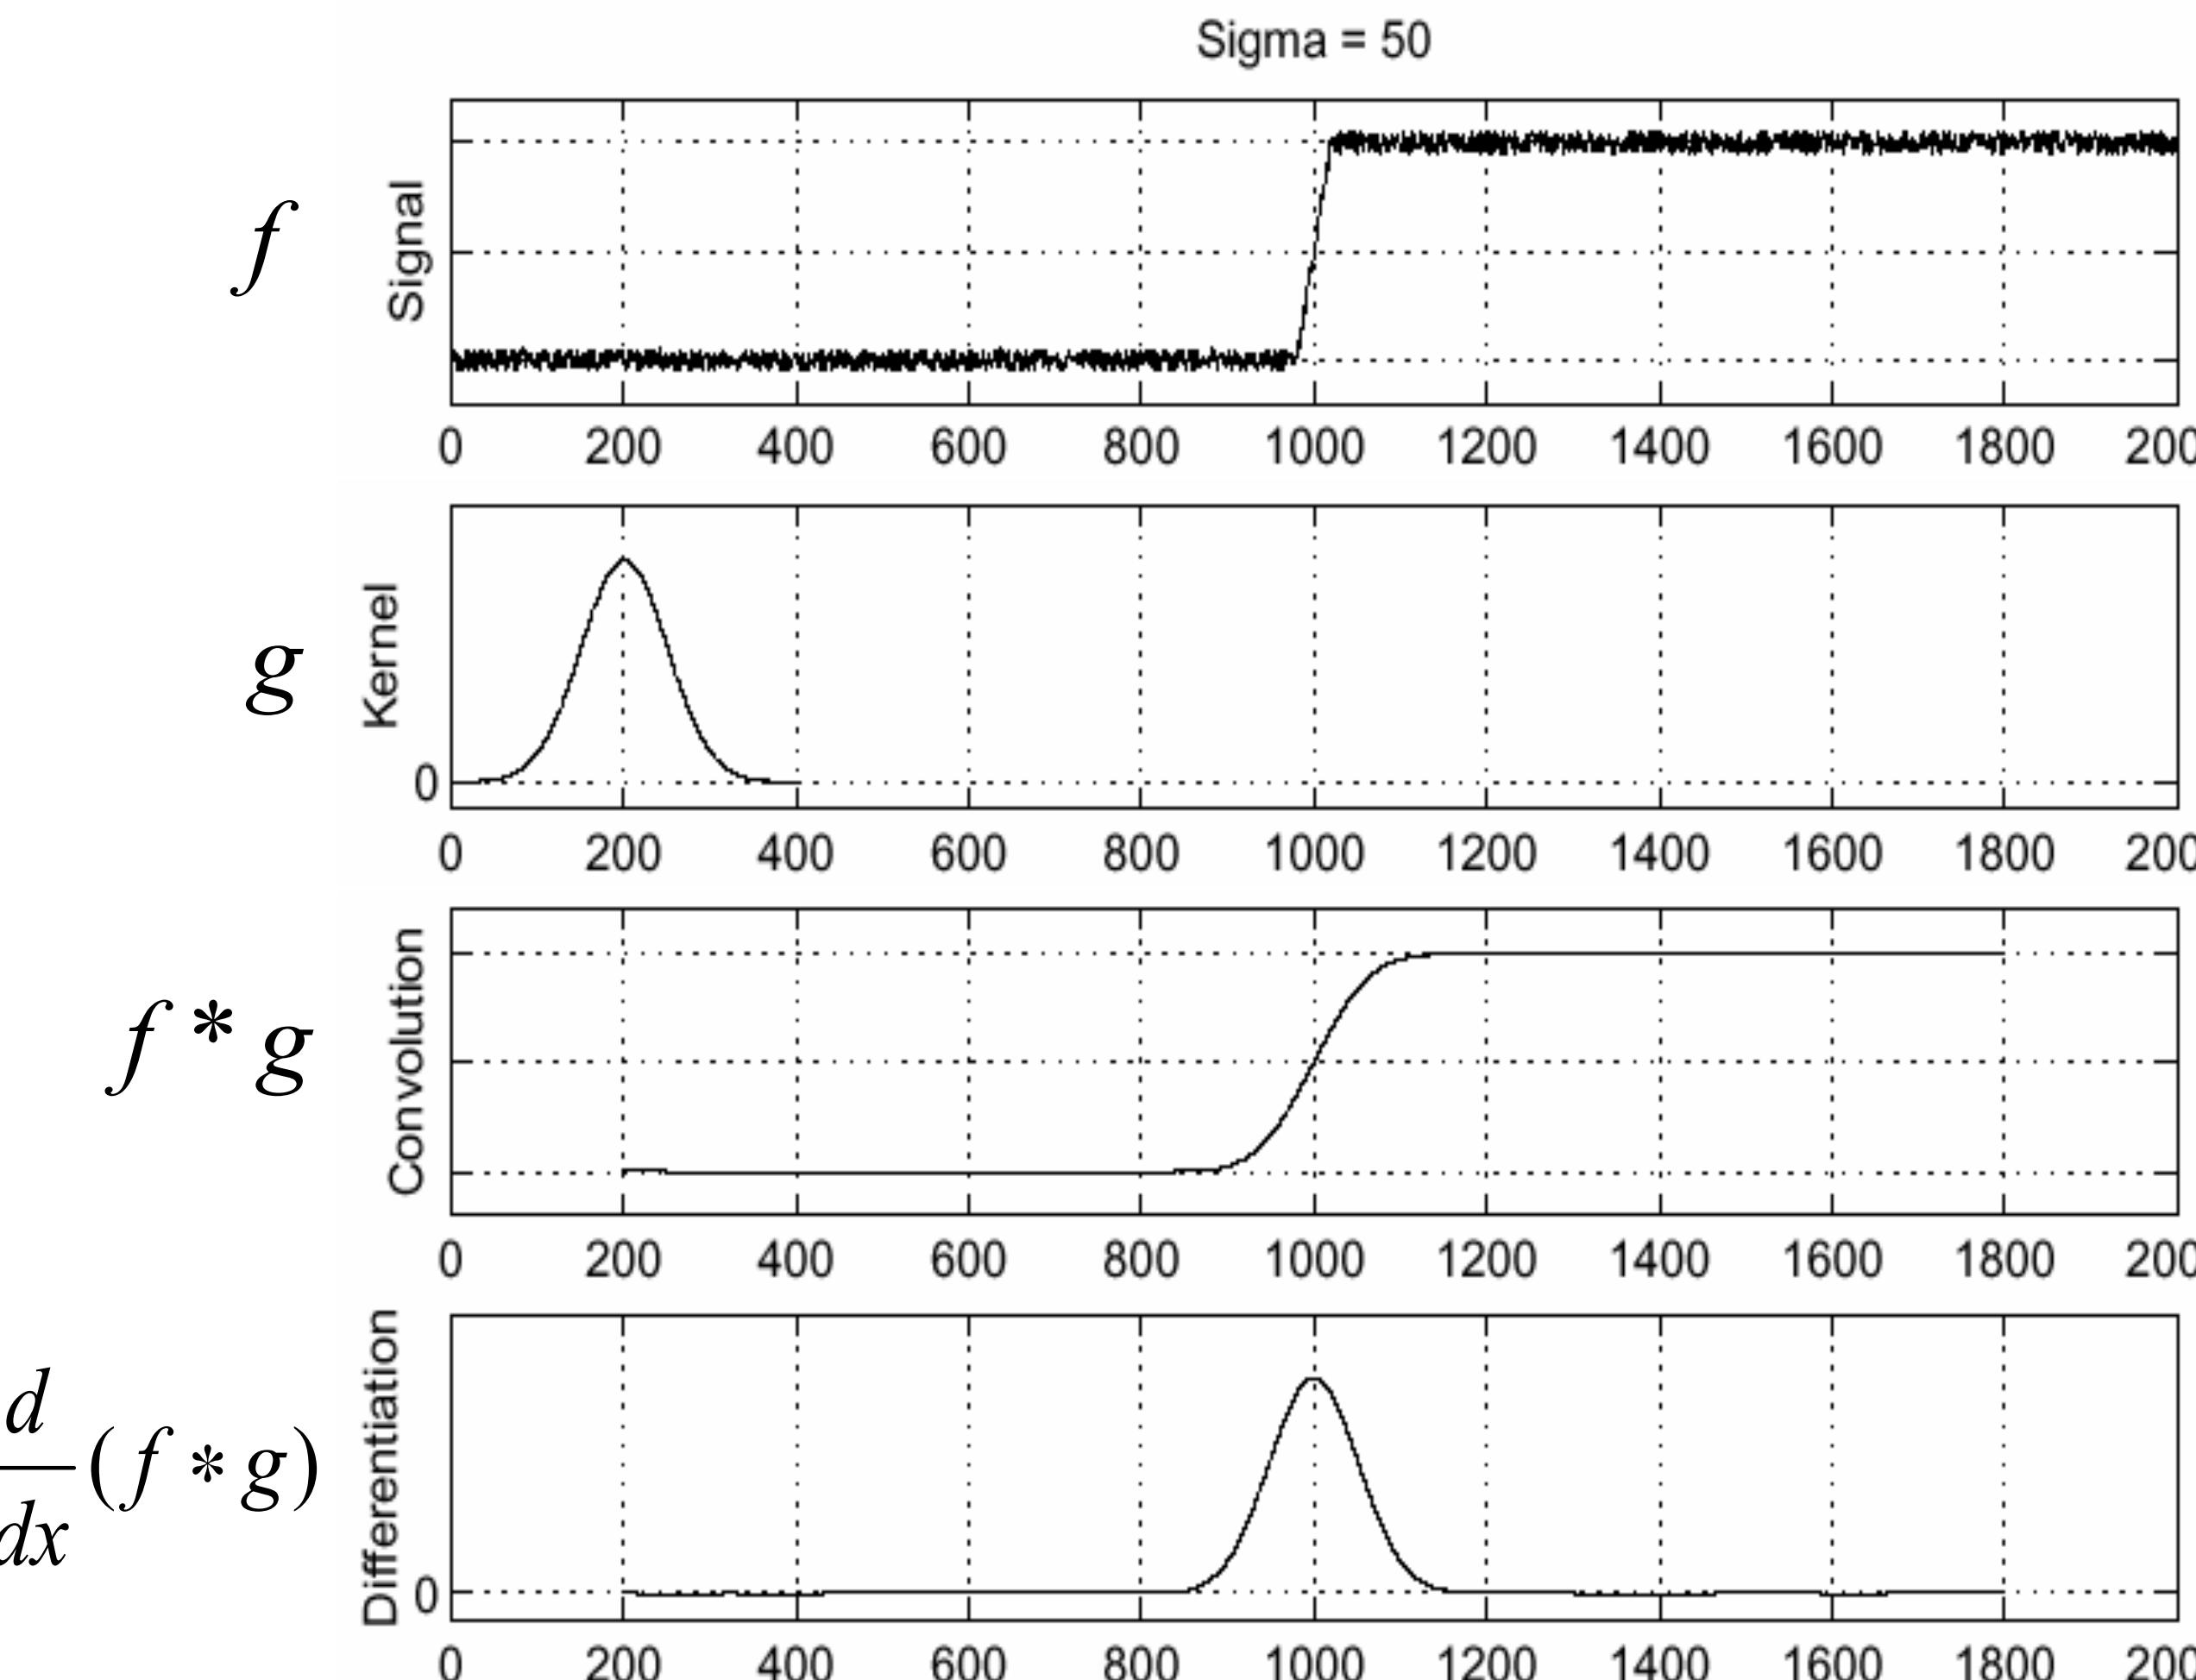
Effects of noise

Consider a single row or column of the image



Where is the edge?

Solution: smooth first



- To find edges, look for peaks in

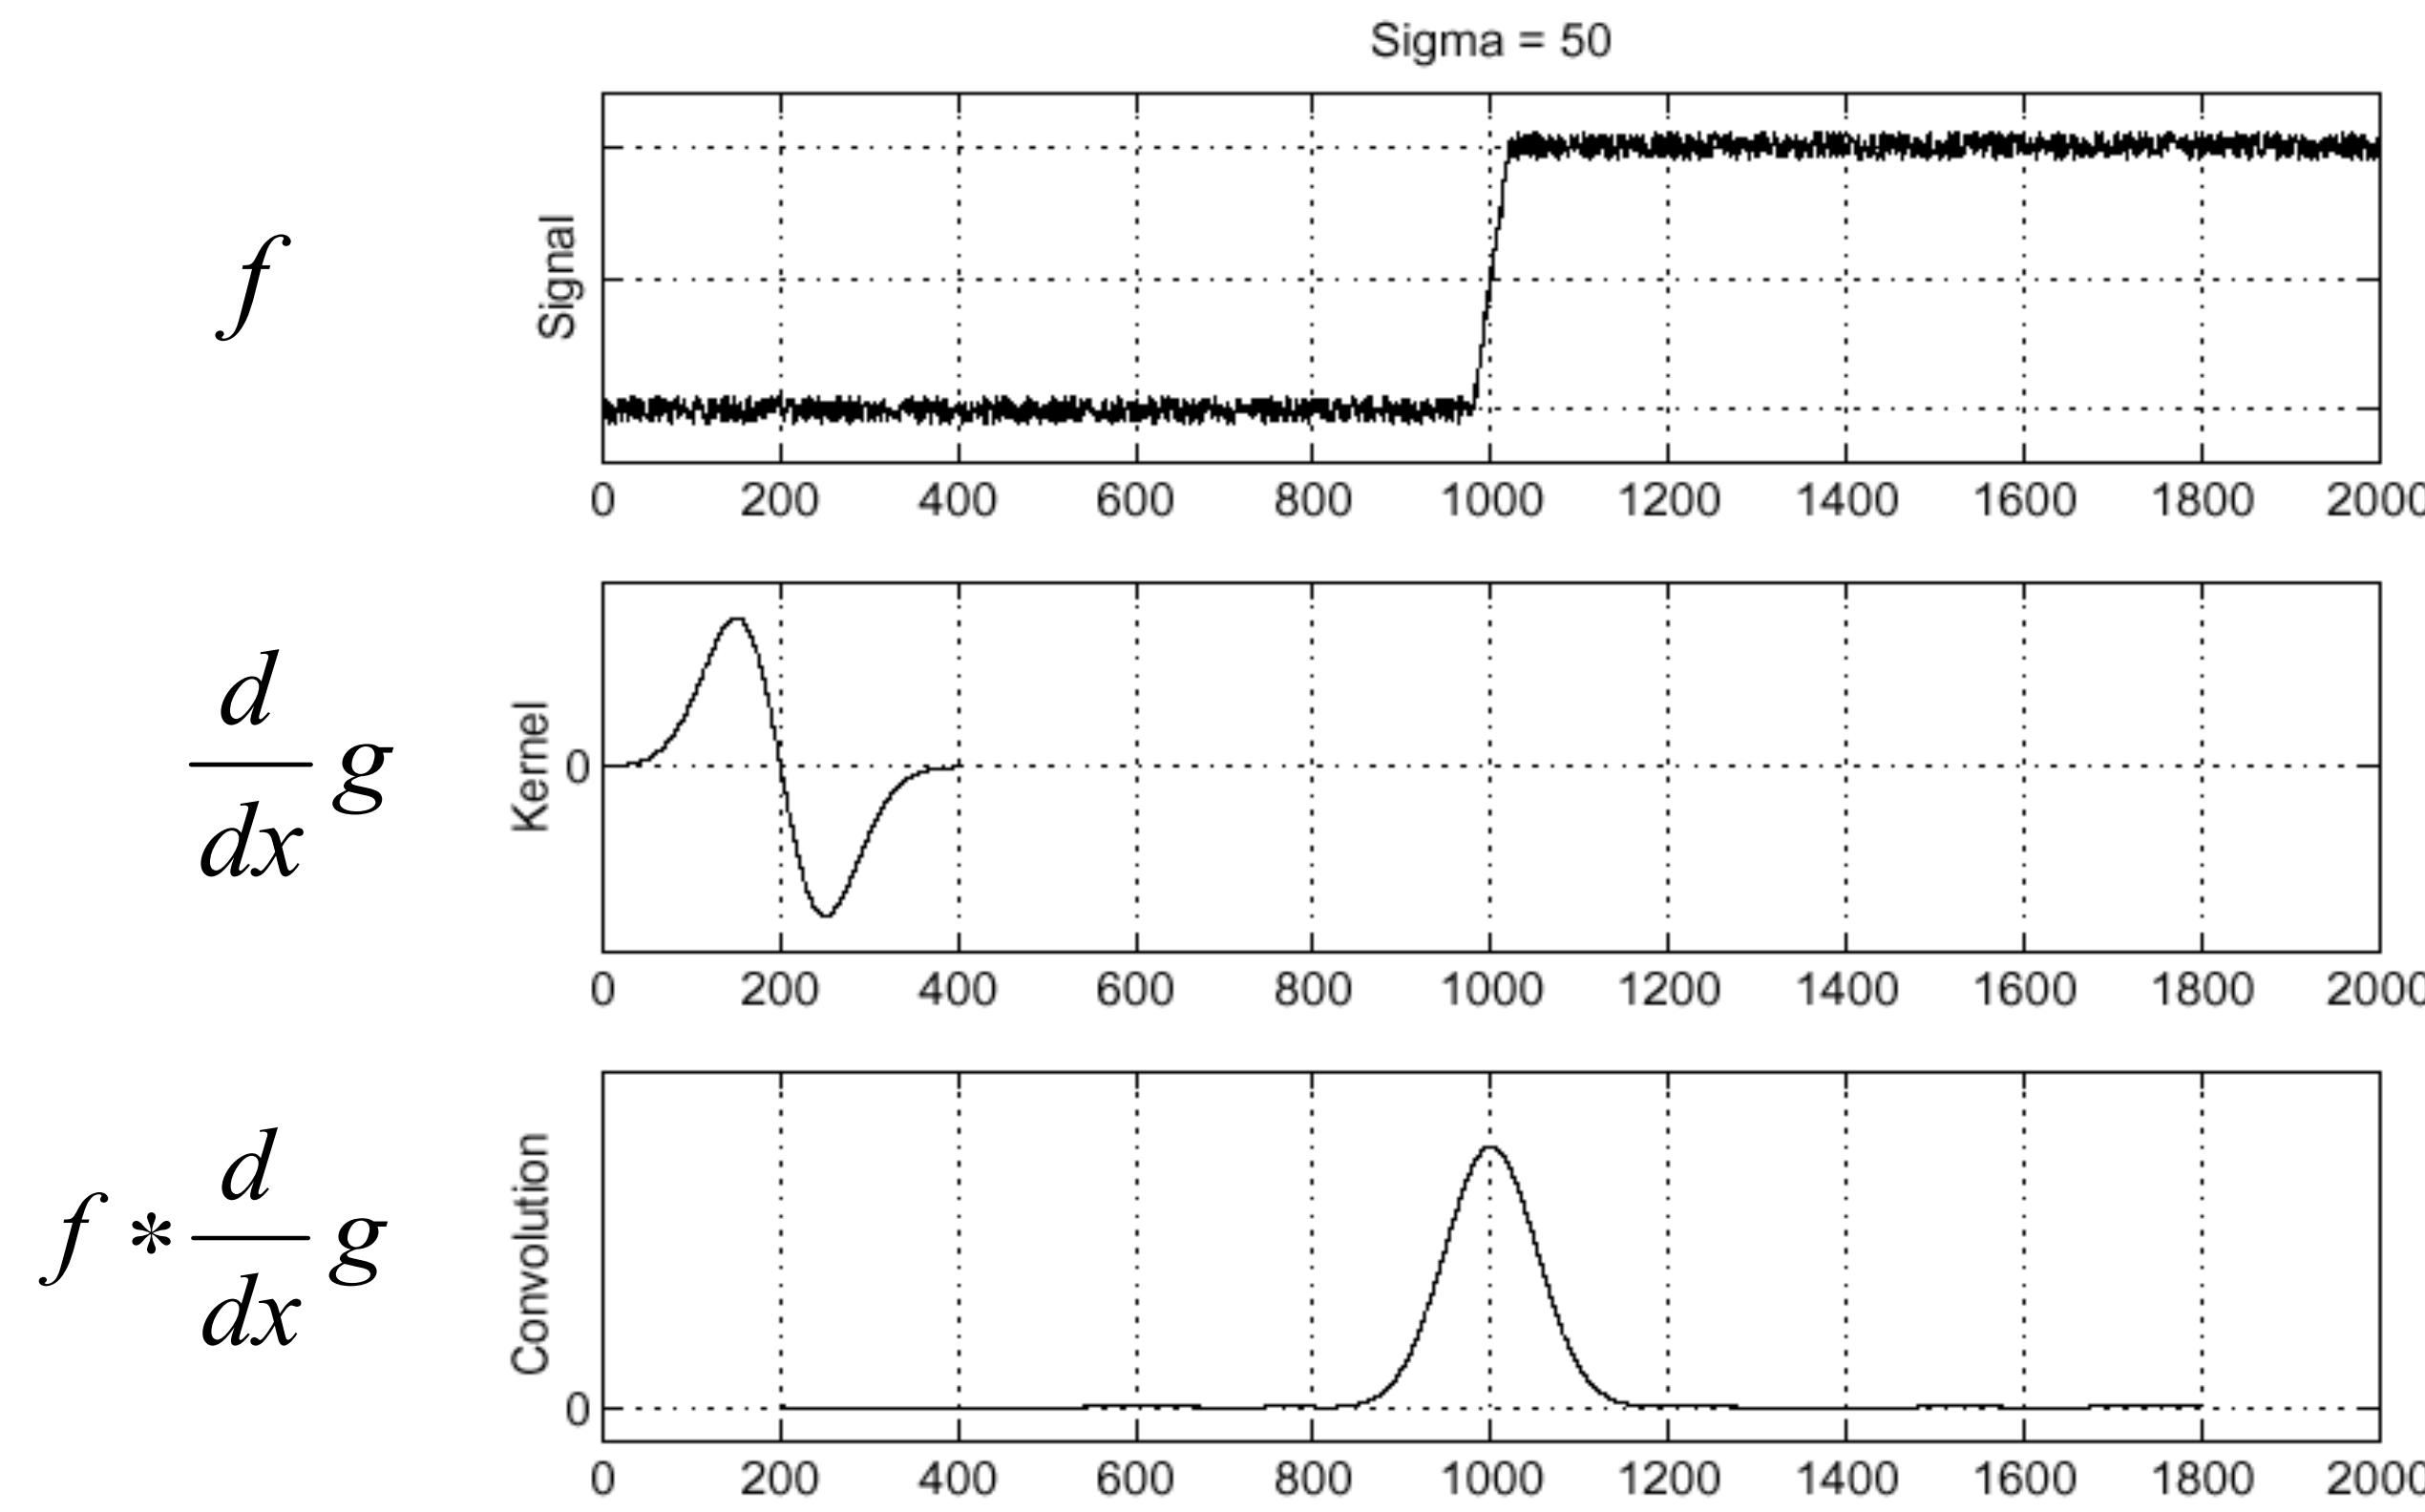
$$\frac{d}{dx}(f * g)$$

Smooth derivative filters

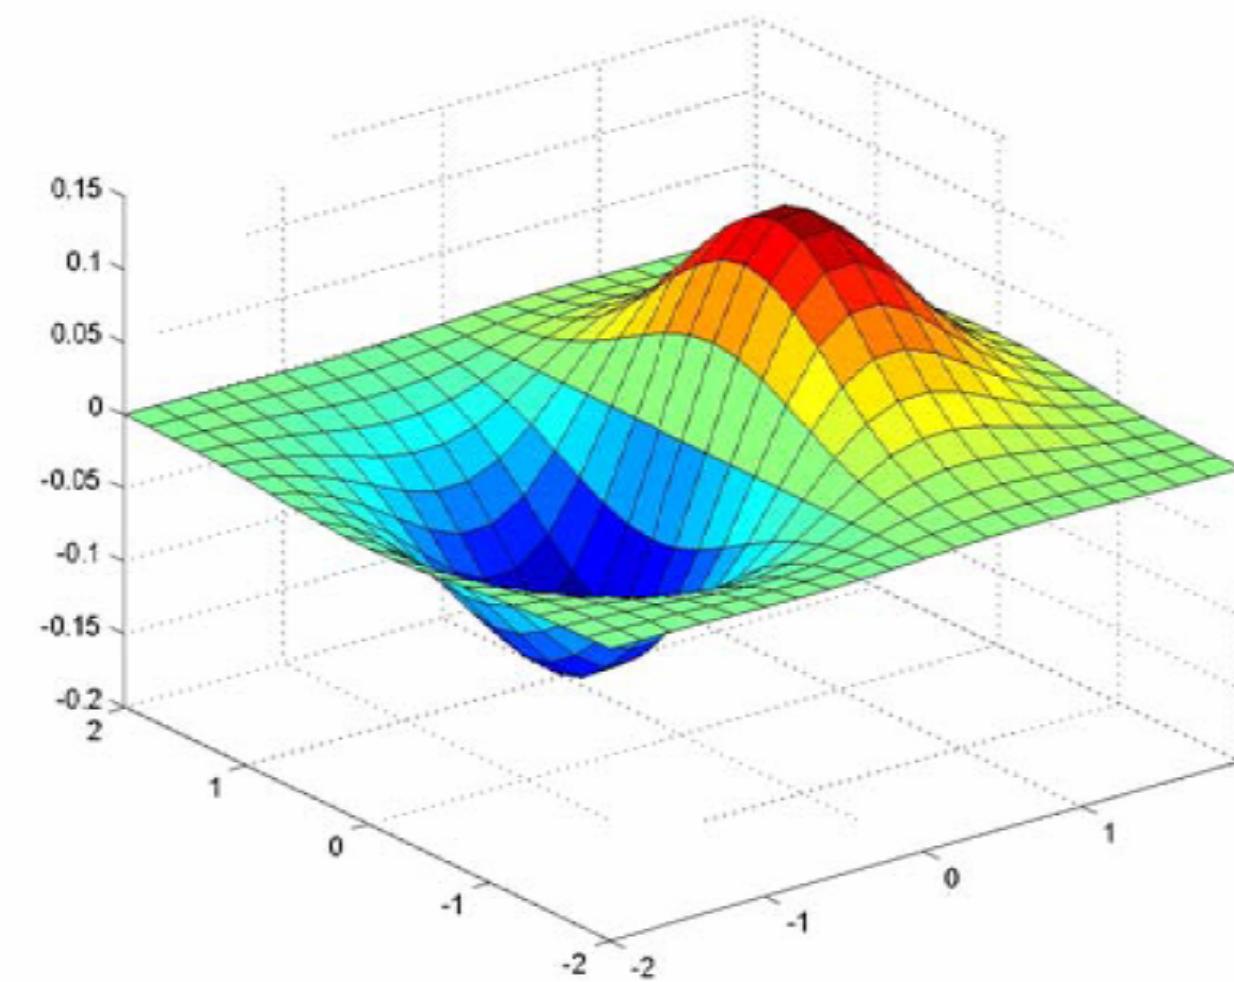
Differentiation is convolution, and convolution is associative:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

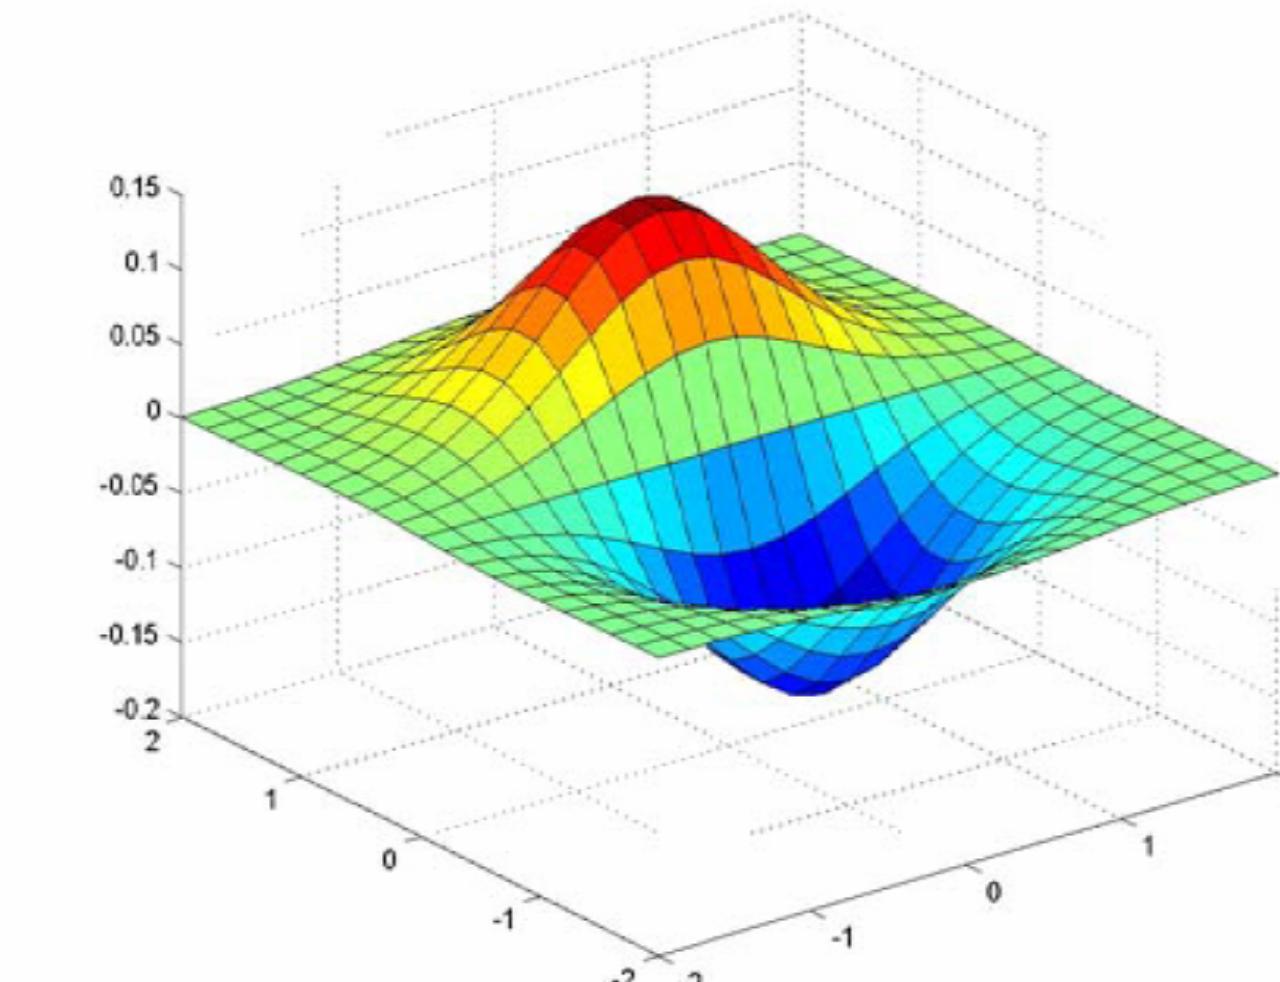
This saves us one operation:



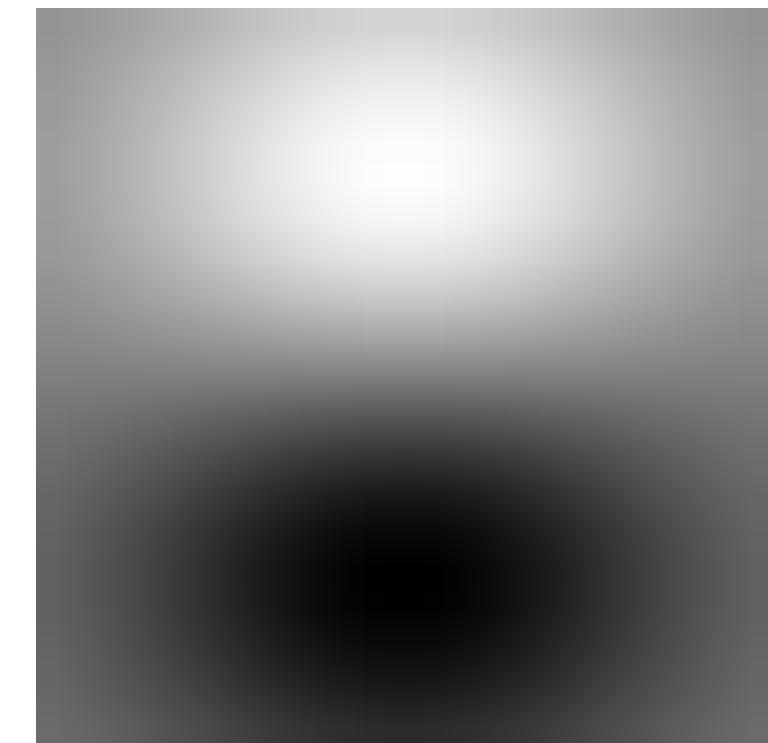
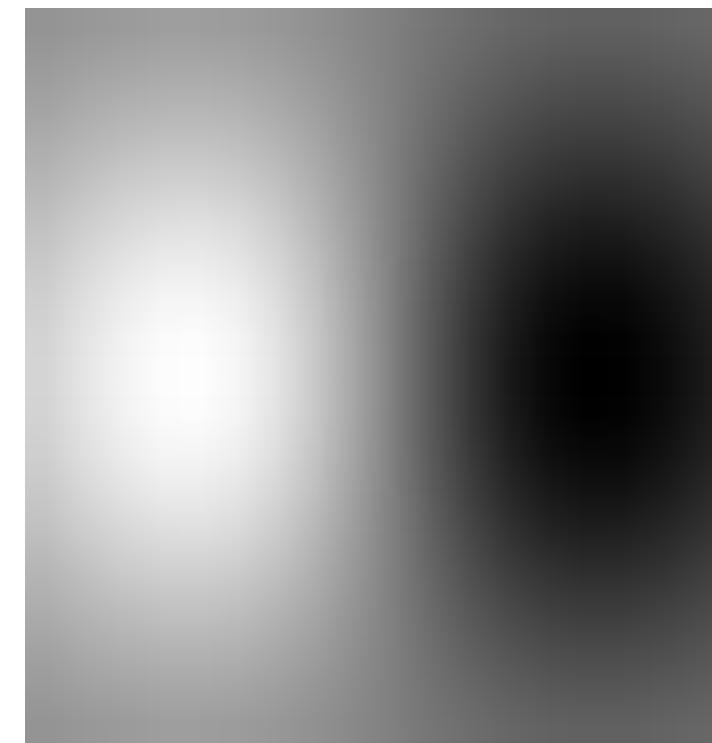
Derivative of Gaussian filters



x-direction

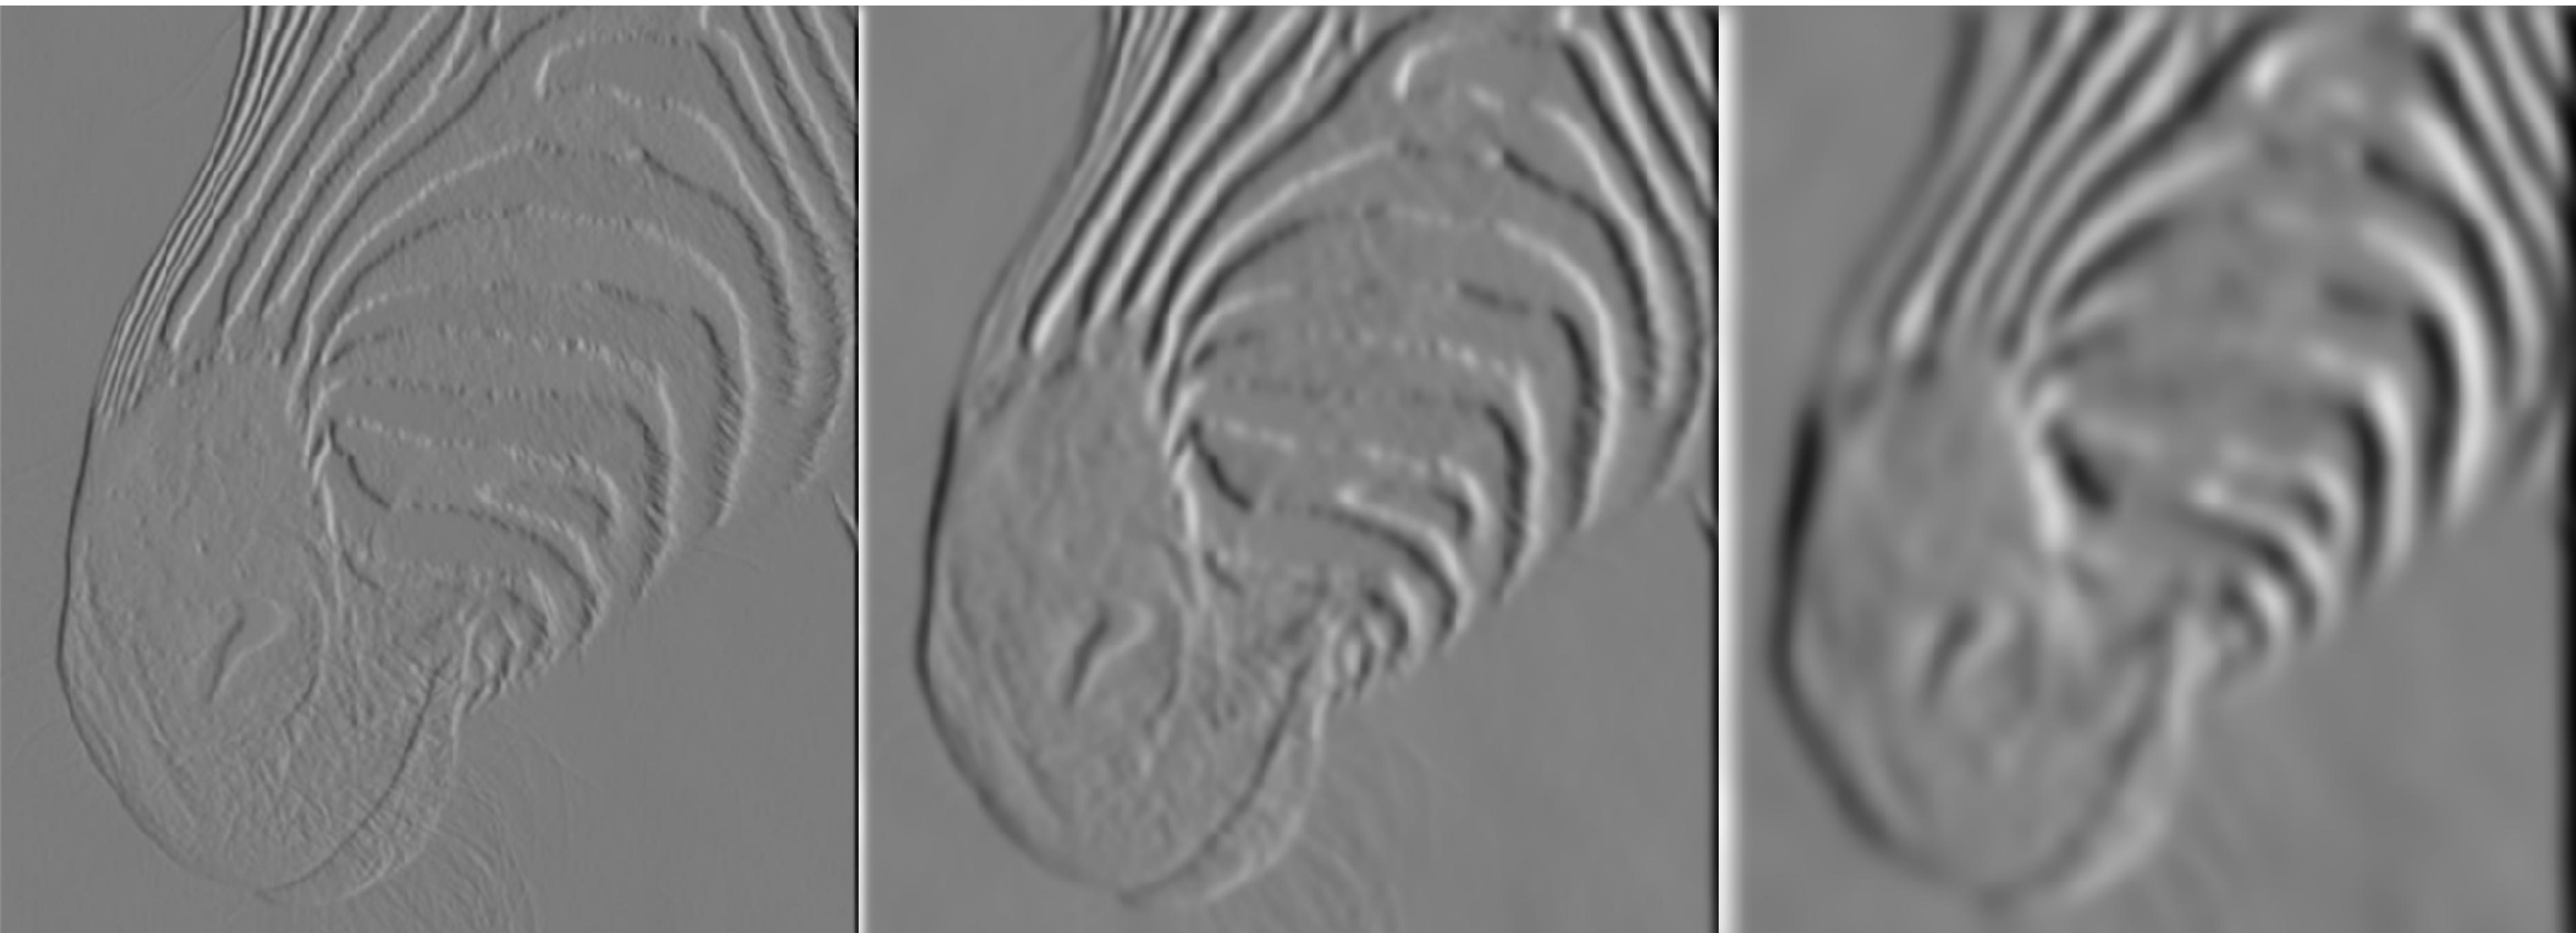


y-direction



- 1) Which one finds horizontal edges?
- 2) Are these filters separable?

Scale of Gaussian derivative filter



1 pixel

3 pixels

7 pixels

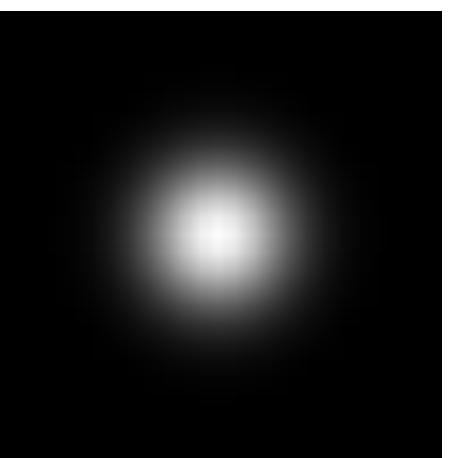
Smoothed derivative removes noise, but blurs edge. Also finds edges at different “scales”

Source: D. Forsyth

Smoothing and derivative filters

Smoothing filters

- Gaussian: remove “high-frequency” components; “low-pass” filter
- Can the values of a smoothing filter be negative?
- What should the values sum to?
 - **One:** constant regions are not affected by the filter



Derivative filters

- Derivatives of Gaussian
- Can the values of a derivative filter be negative?
- What should the values sum to?
 - **Zero:** no response in constant regions
- High absolute value at points of high contrast

