

Local features

370: Intro to Computer Vision

Subhransu Maji

March 6 & 11, 2025

College of
INFORMATION AND
COMPUTER SCIENCES



UMASS
AMHERST

Topics

Why extract features?

Corner detector

Scale-invariant feature detector (or blob detector)

Why extract features?

Motivation: panorama stitching

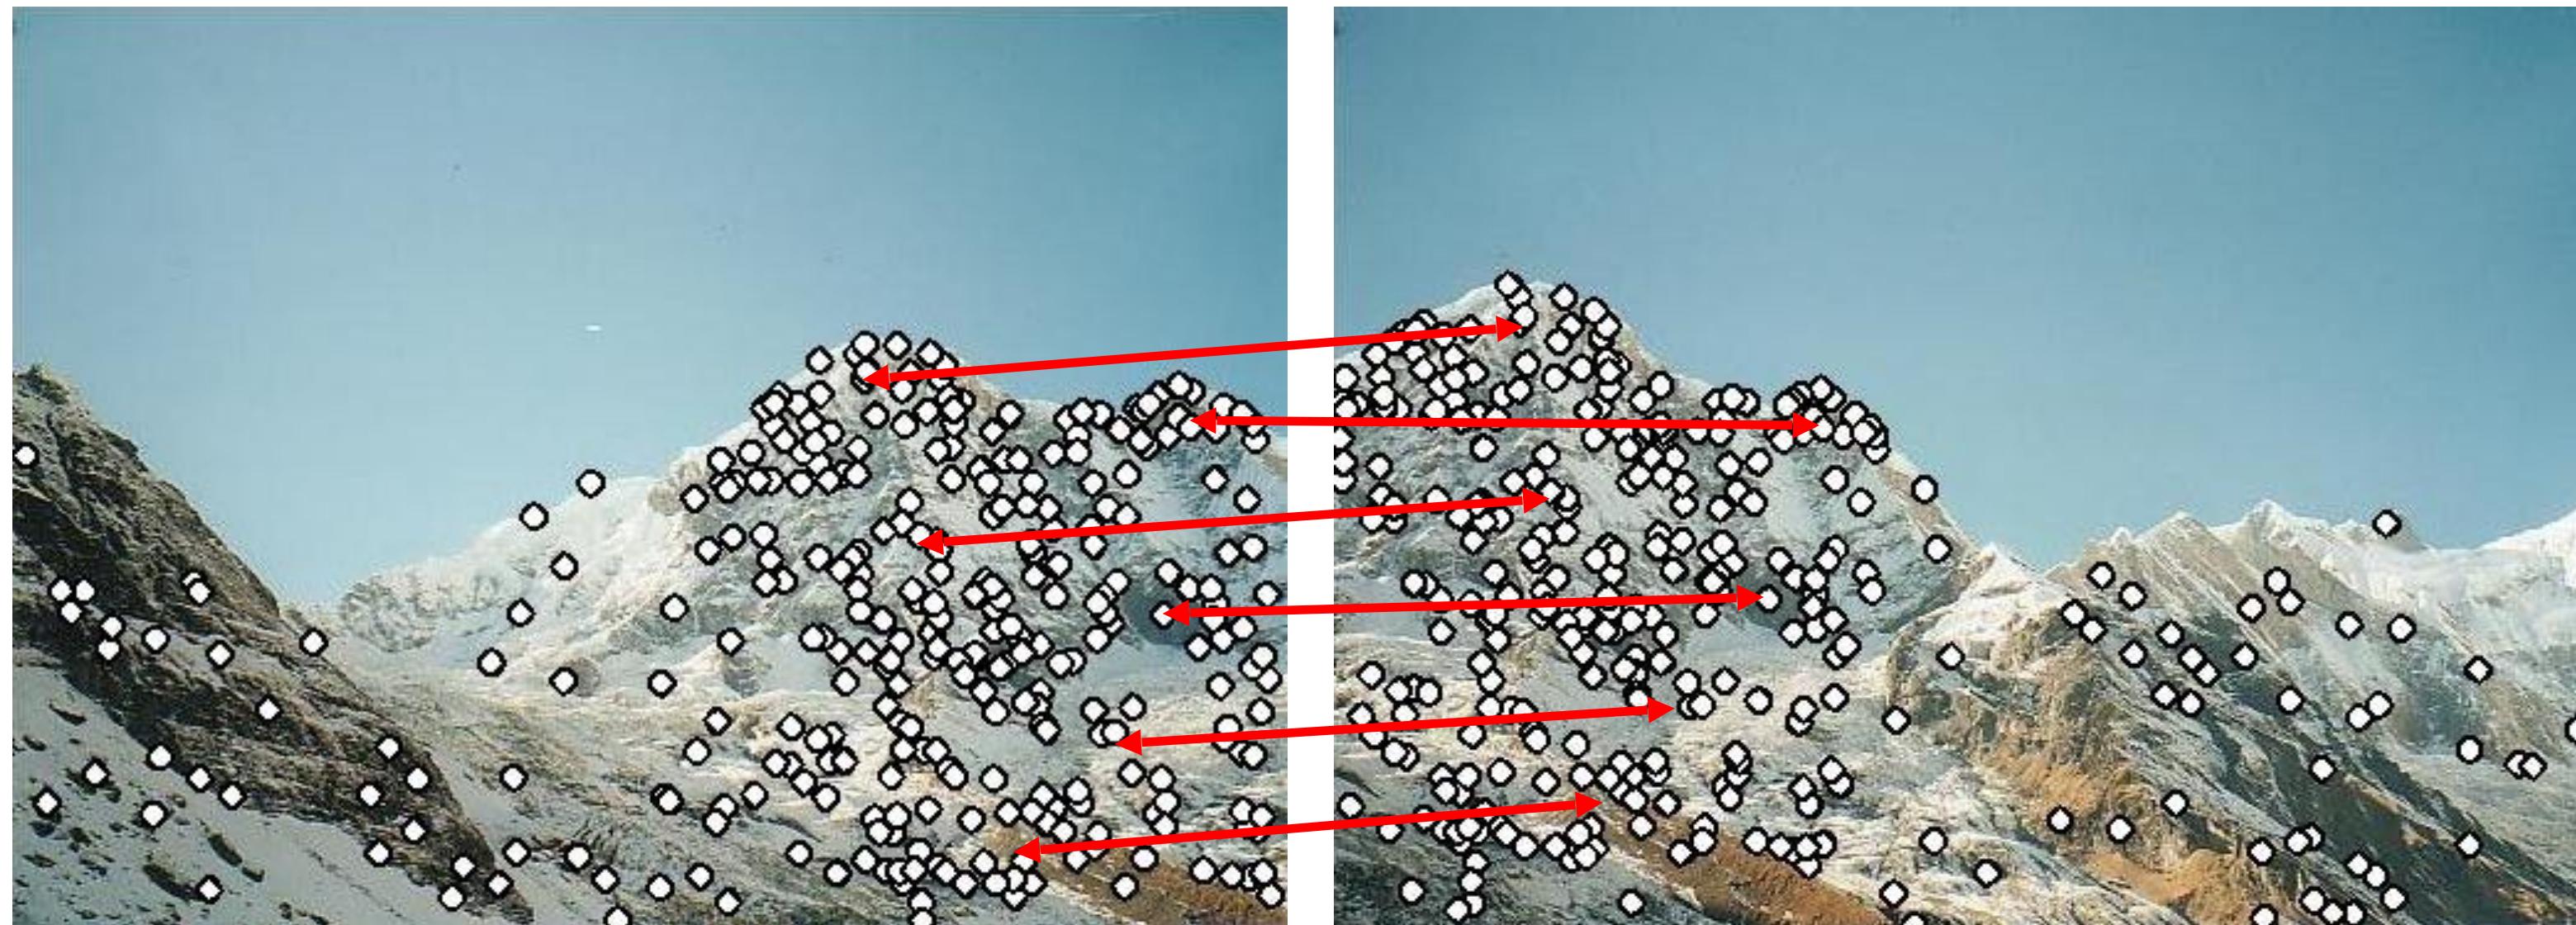
- We have two images – how do we combine them?



Why extract features?

Motivation: panorama stitching

- We have two images – how do we combine them?



Step 1: extract features

Step 2: match features

Why extract features?

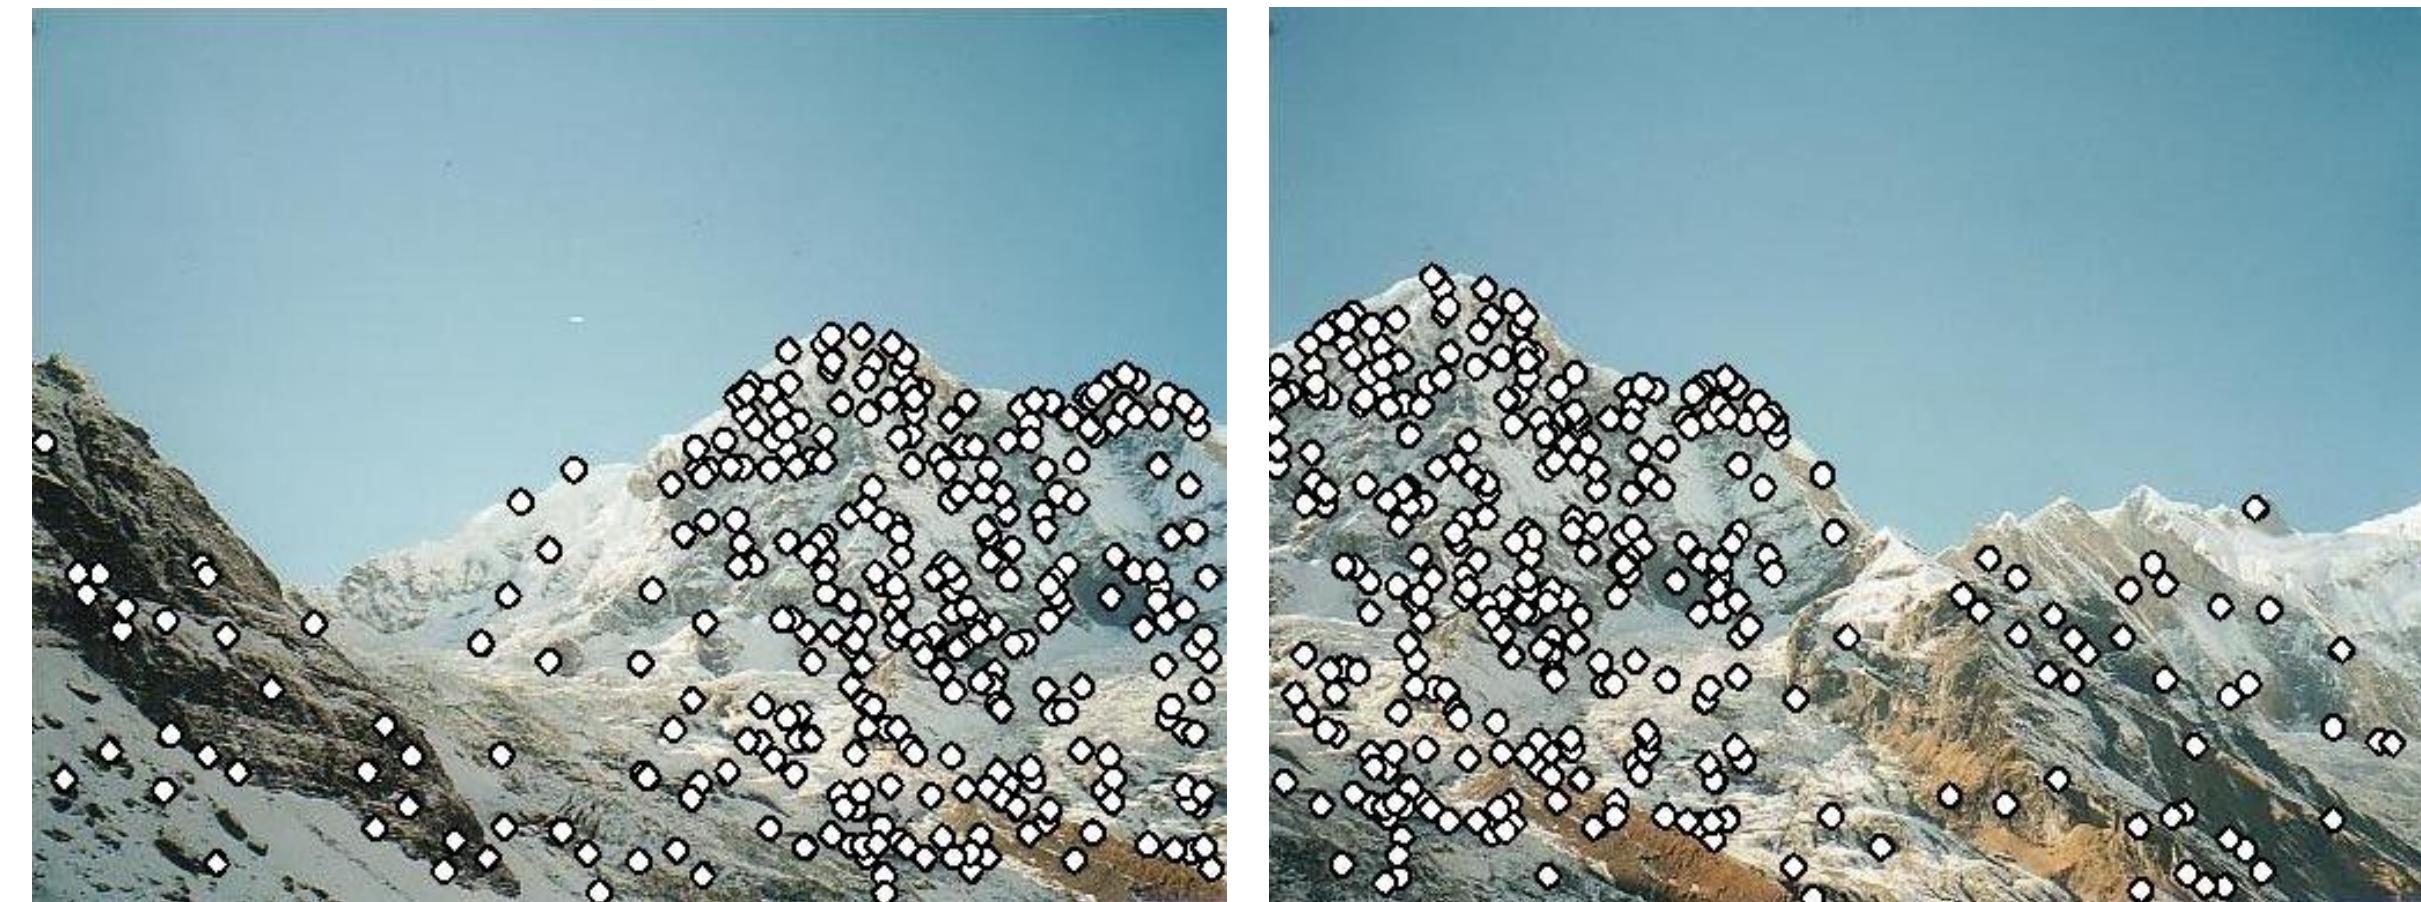
Motivation: panorama stitching

- We have two images – how do we combine them?



- Step 1: extract features
- Step 2: match features
- Step 3: align images

Characteristics of good features



Repeatability

- The same feature can be found in several images despite geometric and photometric transformations

Saliency

- Each feature is distinctive

Compactness and efficiency

- Many fewer features than image pixels

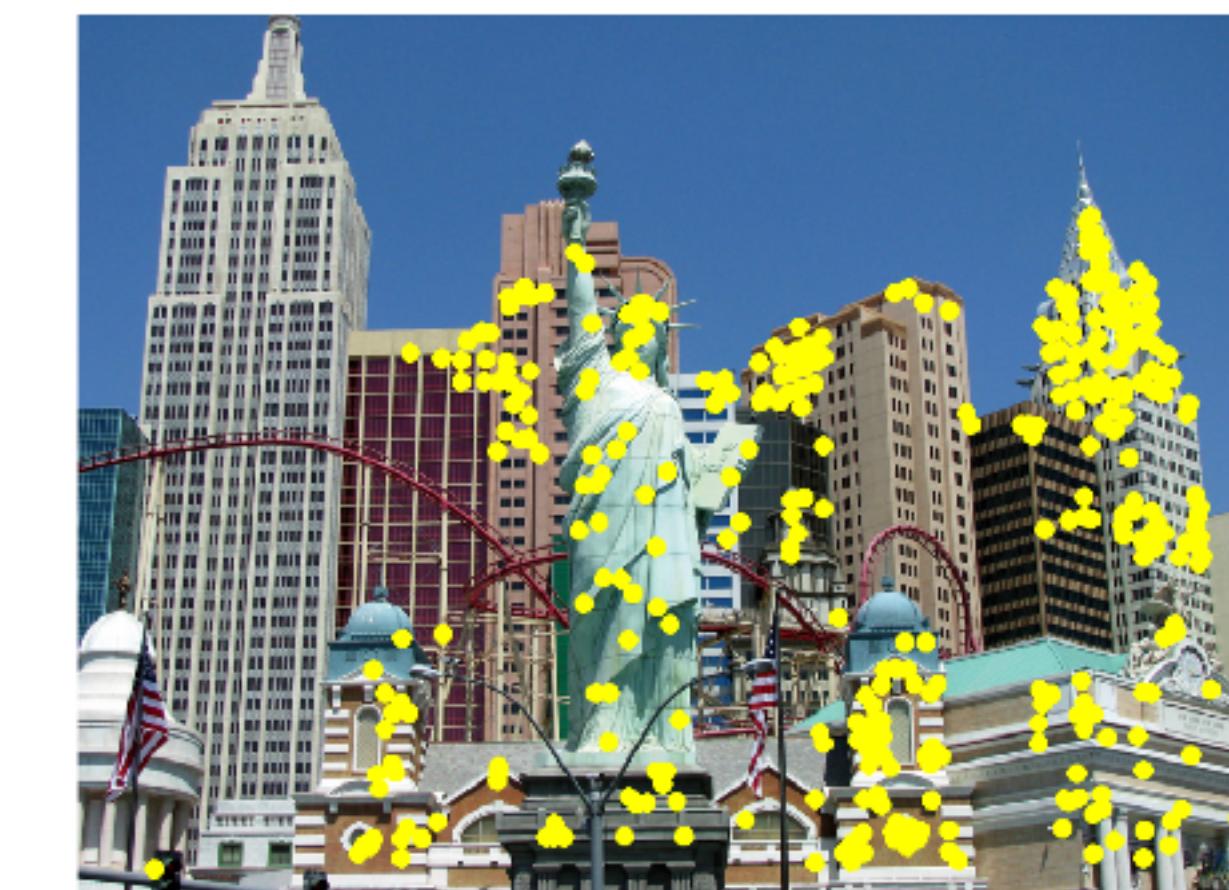
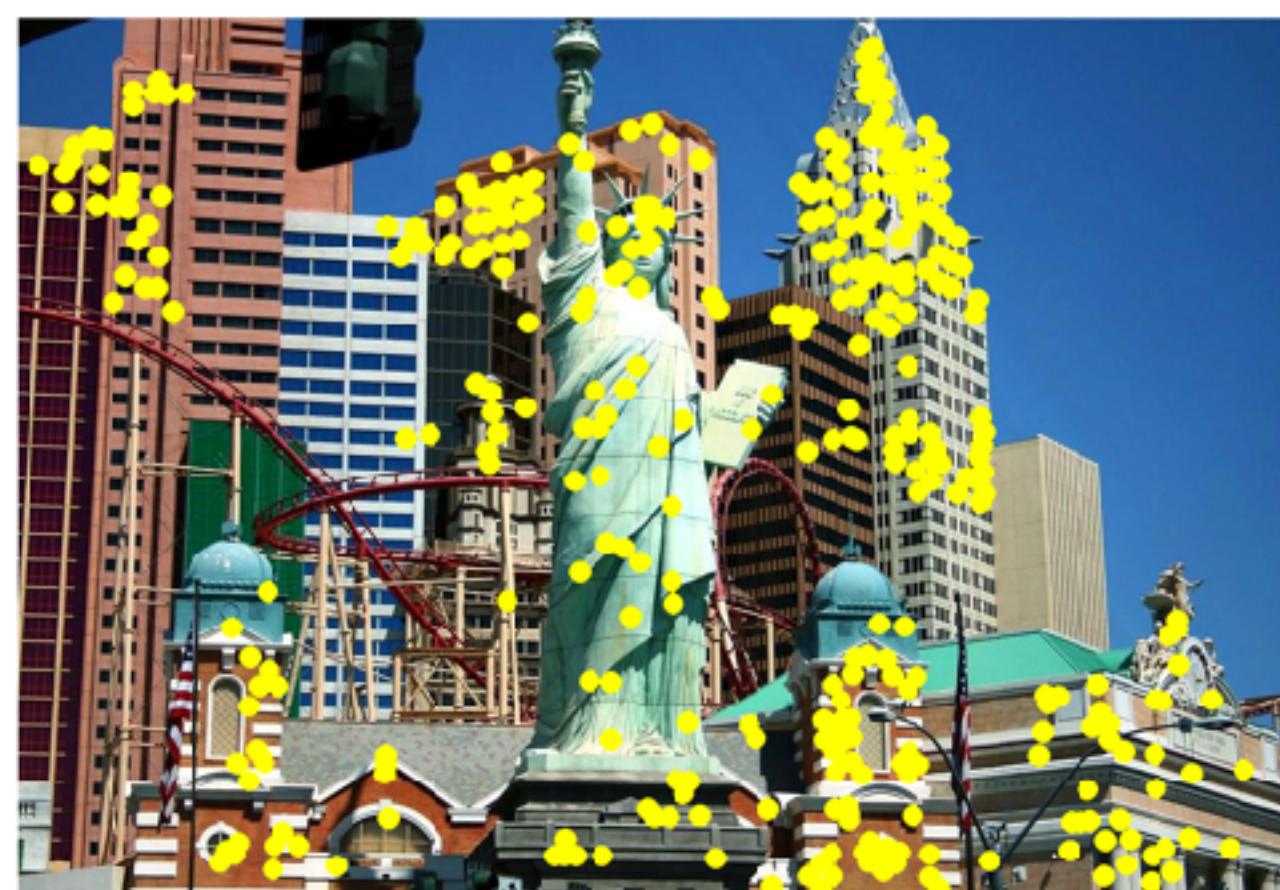
Locality

- A feature occupies a relatively small area of the image; robust to clutter and occlusion

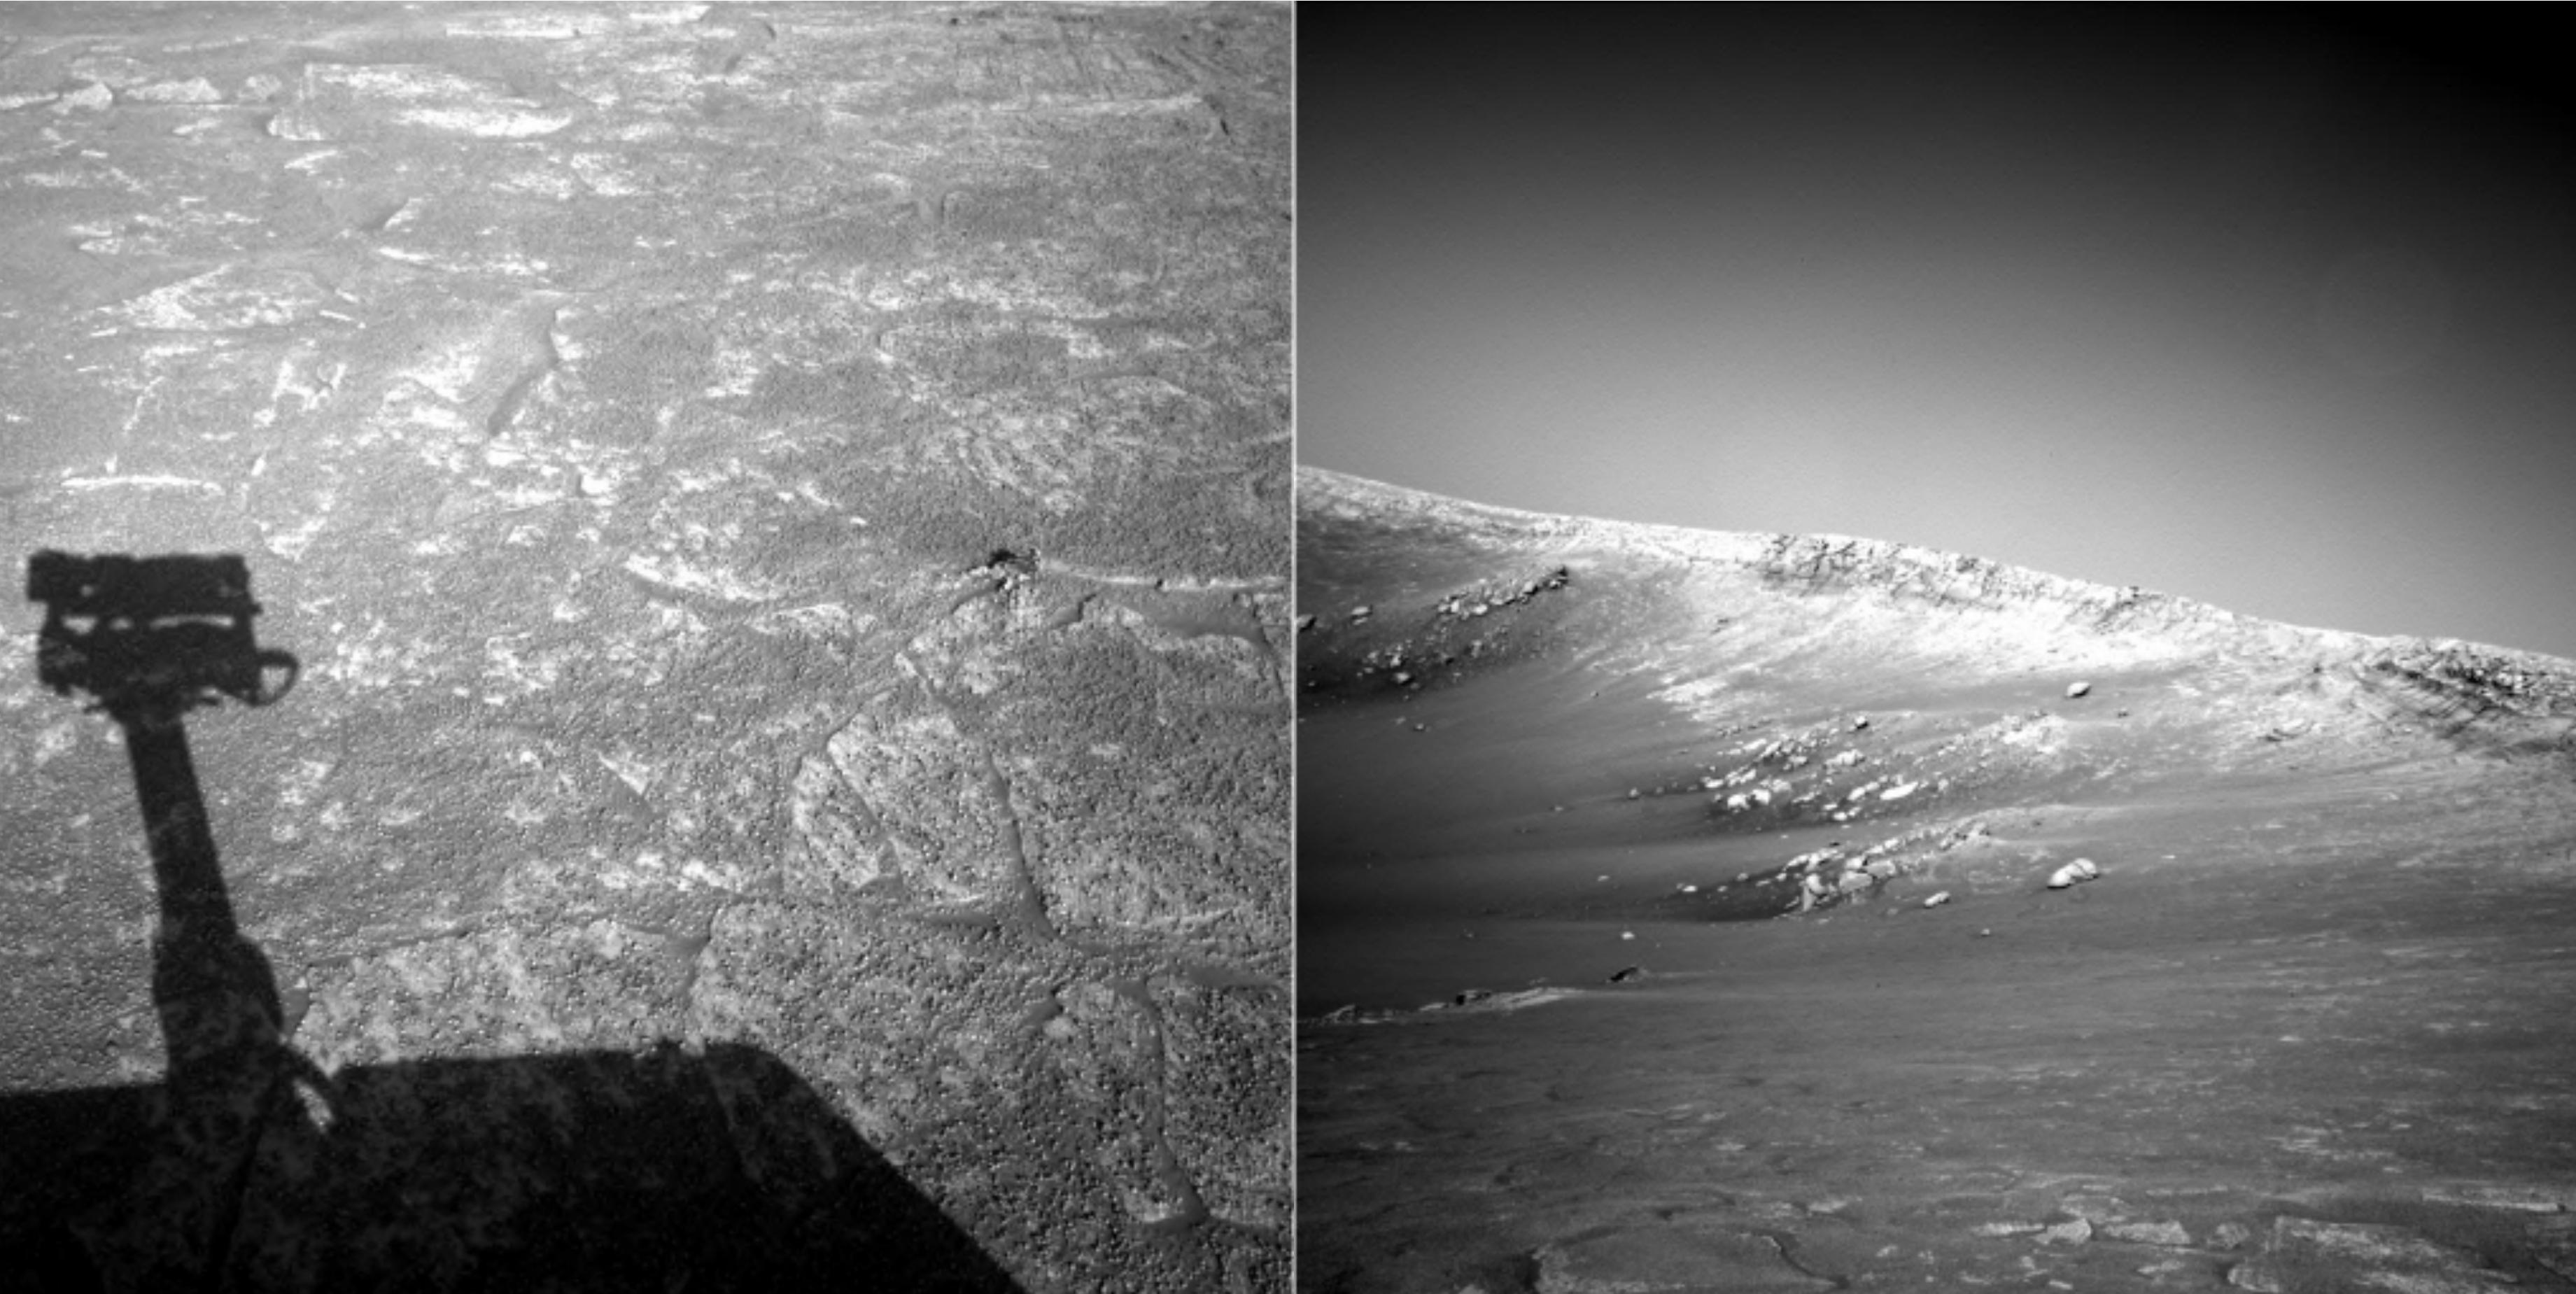
Applications

Feature points are used for:

- Image alignment
- 3D reconstruction
- Motion tracking
- Robot navigation
- Indexing and database retrieval
- Object recognition

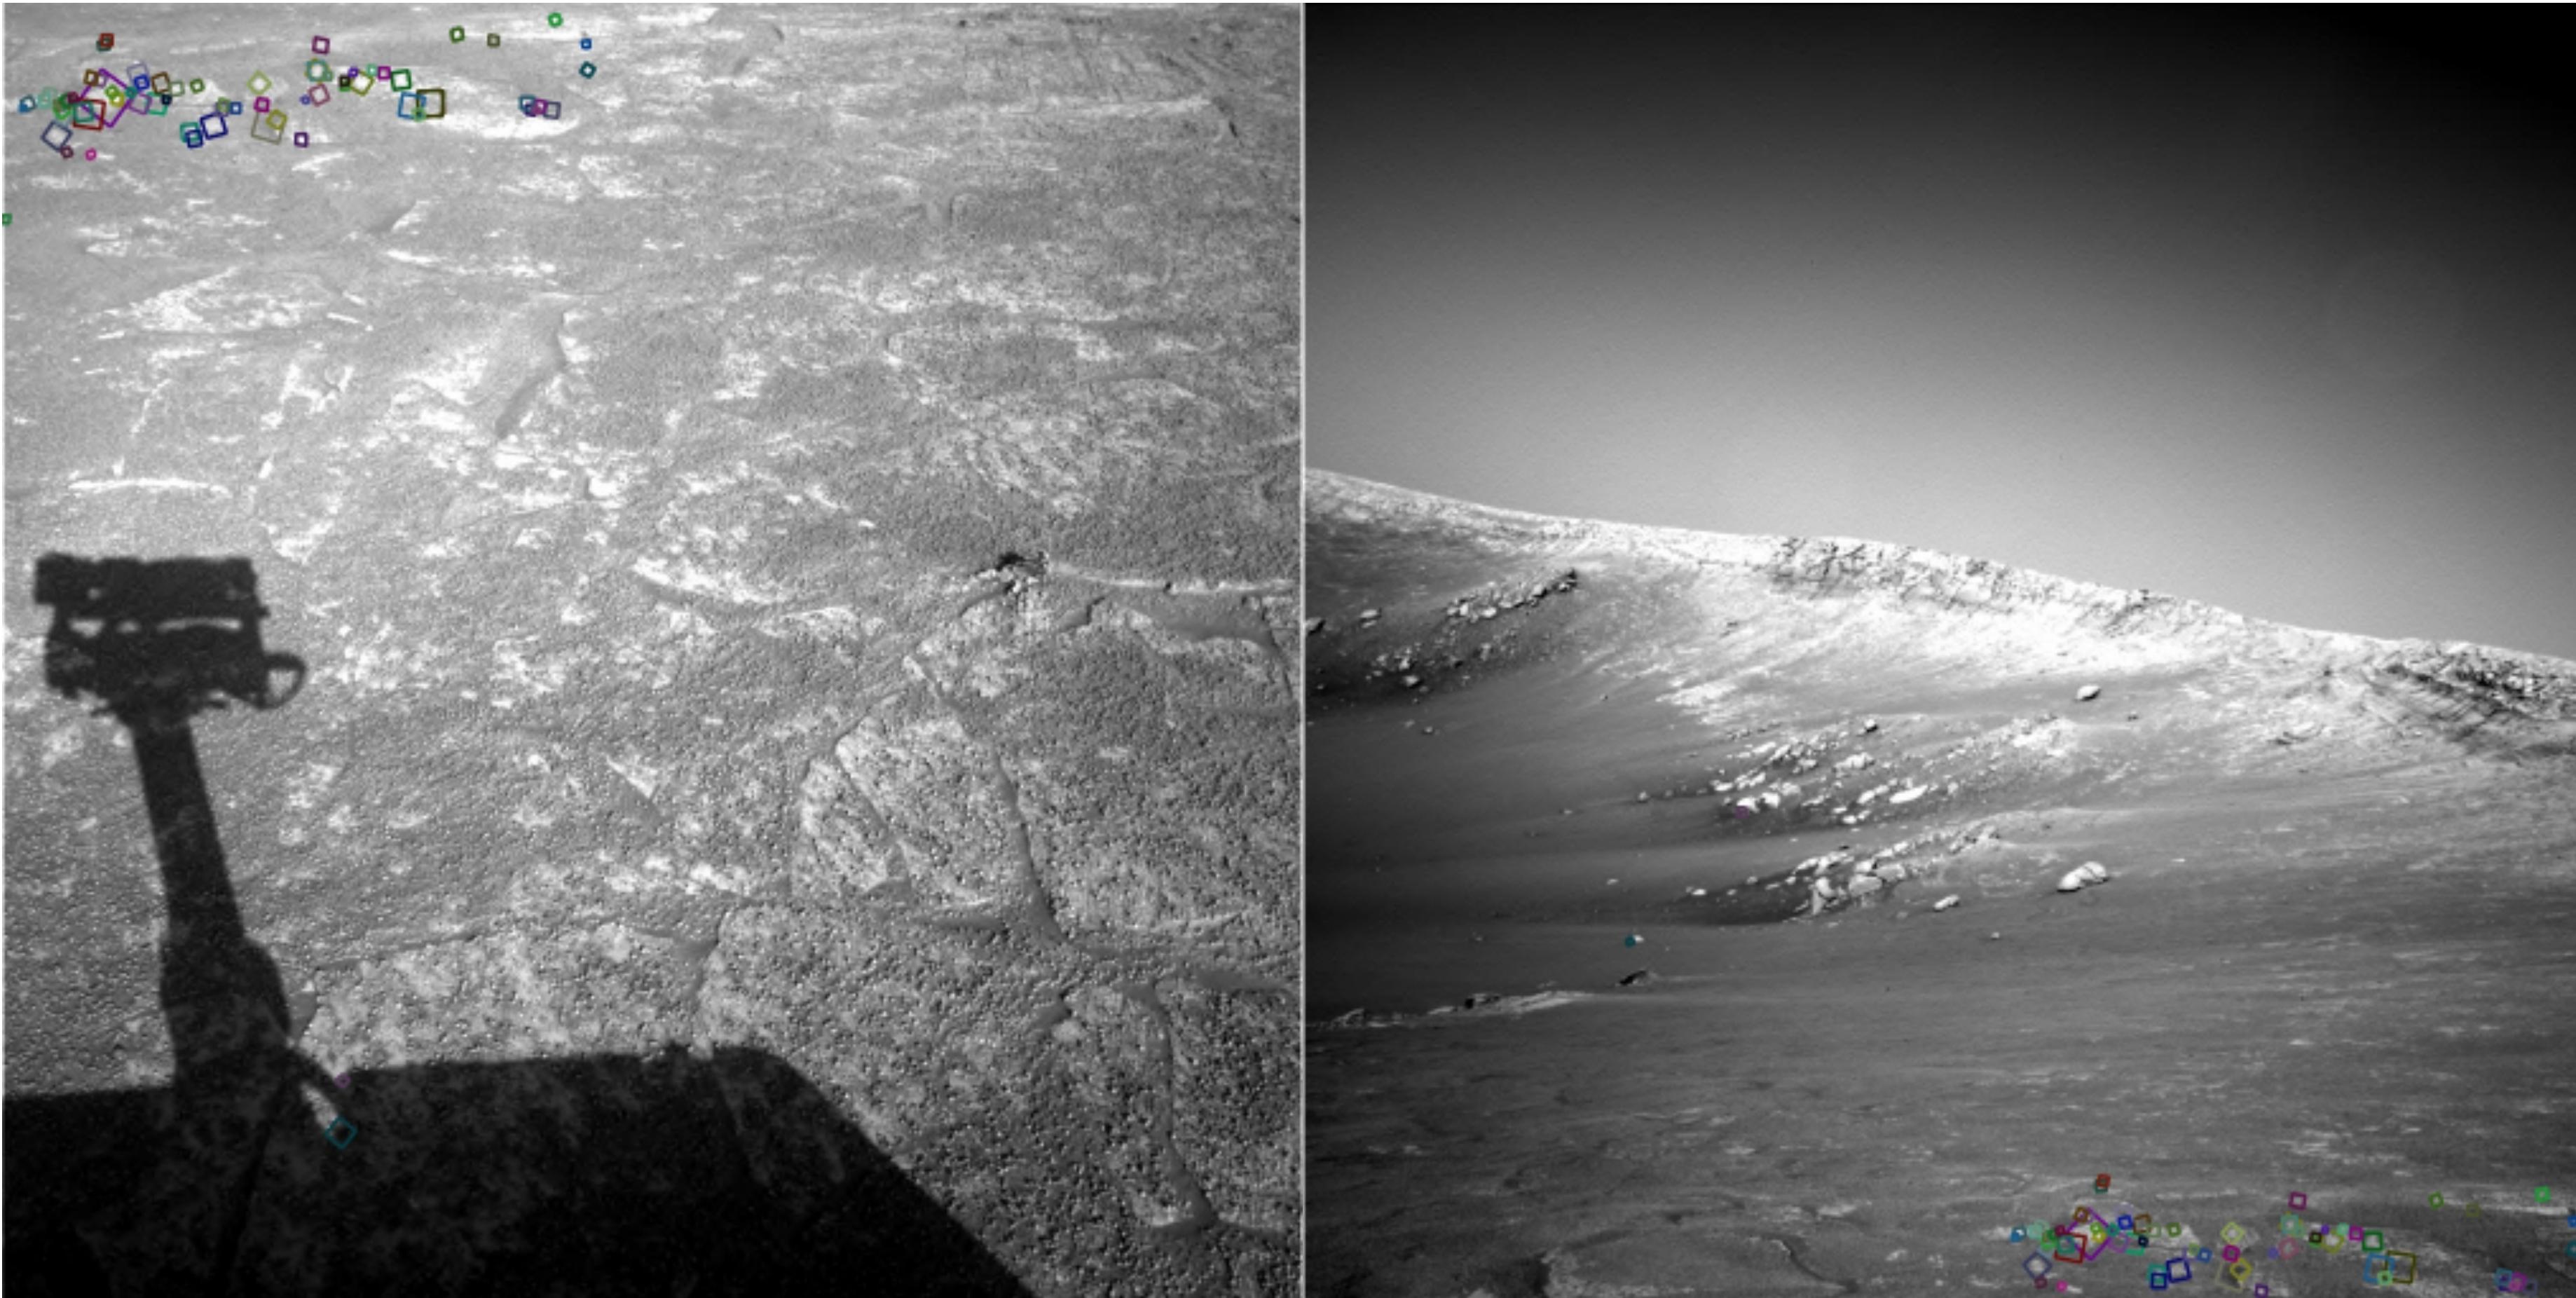


A hard feature matching problem



NASA Mars Rover images

Answer below (look for tiny colored squares...)



NASA Mars Rover images
with SIFT feature matches
Figure by Noah Snavely

Feature extraction: Corners



Corner detection: Attempt one

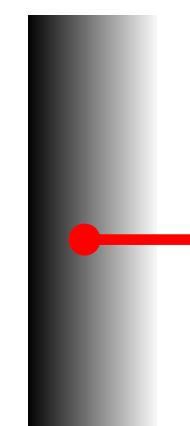
A corner is the intersection of two edges

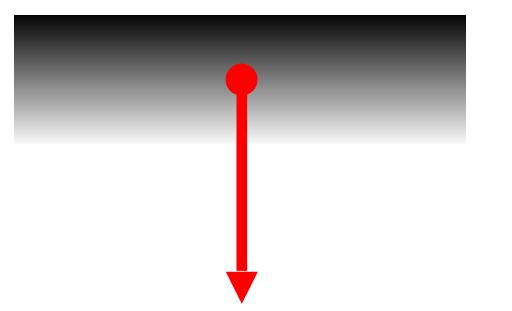
We know how to detect edges

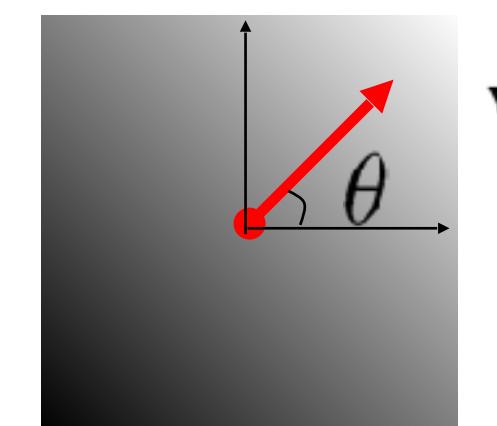
Corner detector (attempt #1)

- Detect edges in images (G_x and G_y)
- Find places where both G_x and G_y are high

Problem: also finds slanted edges!


$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$

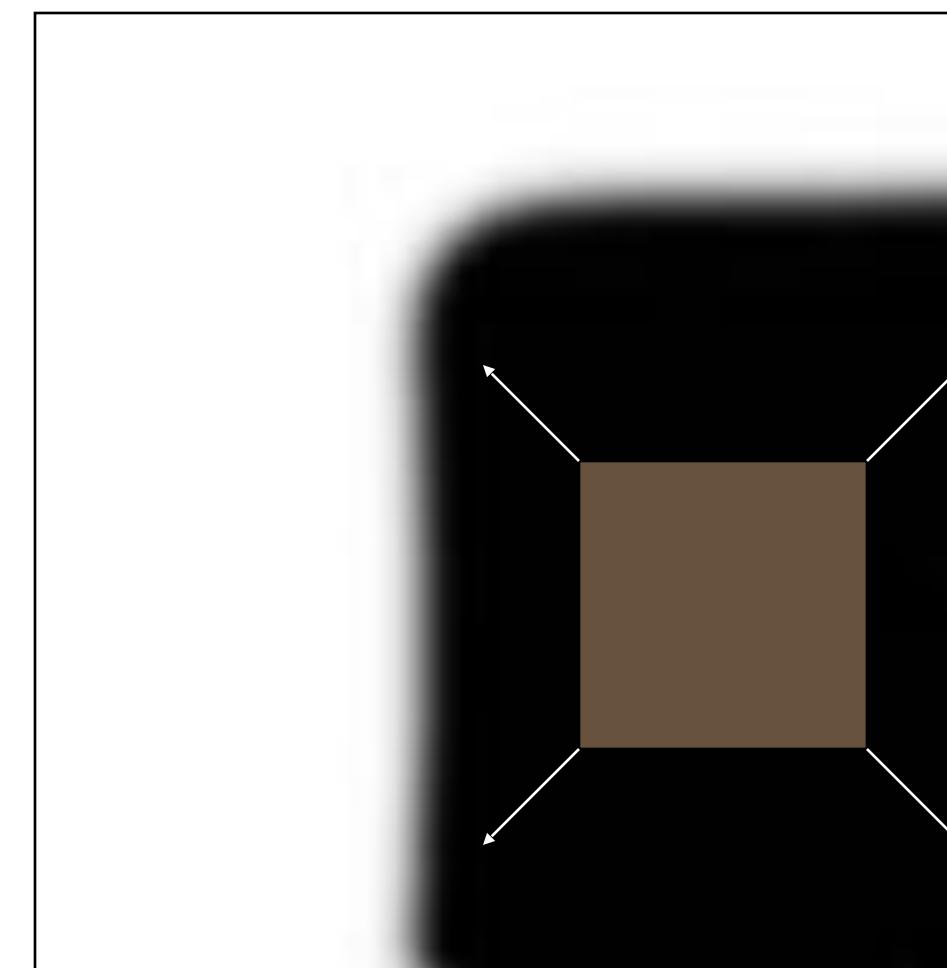

$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

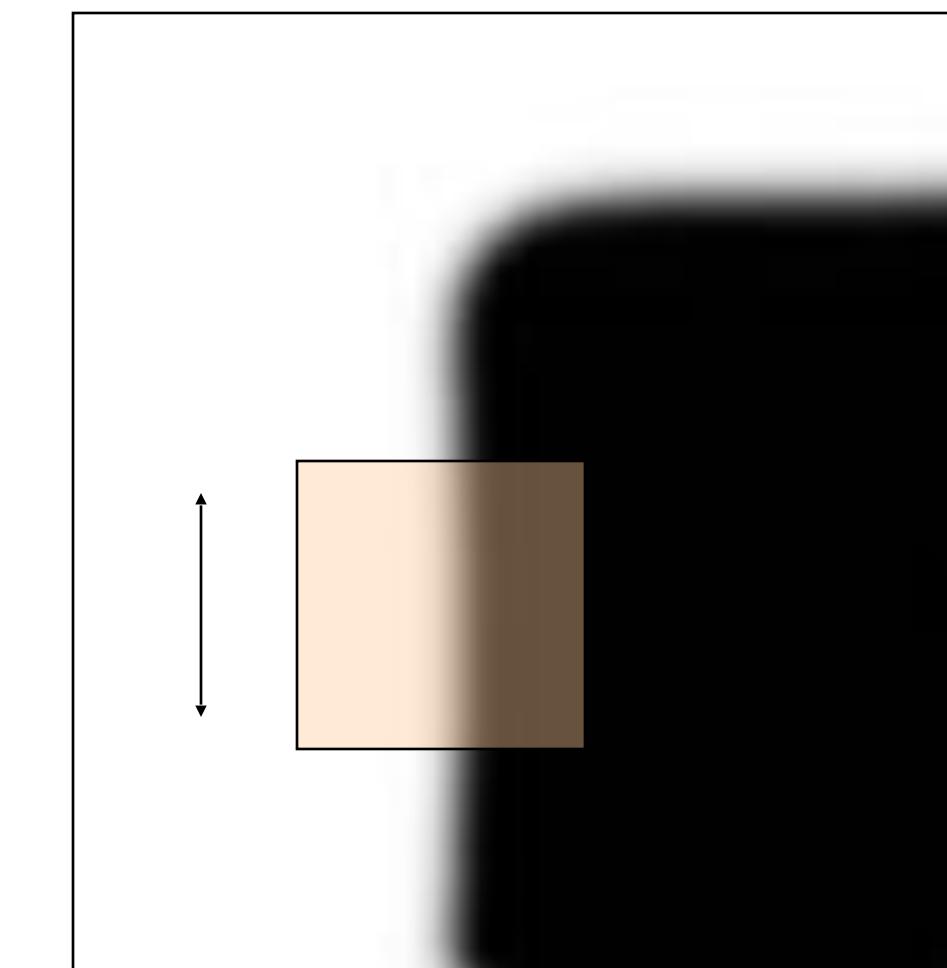
Corner detection: Attempt two

We should easily recognize the corners by looking through a small window

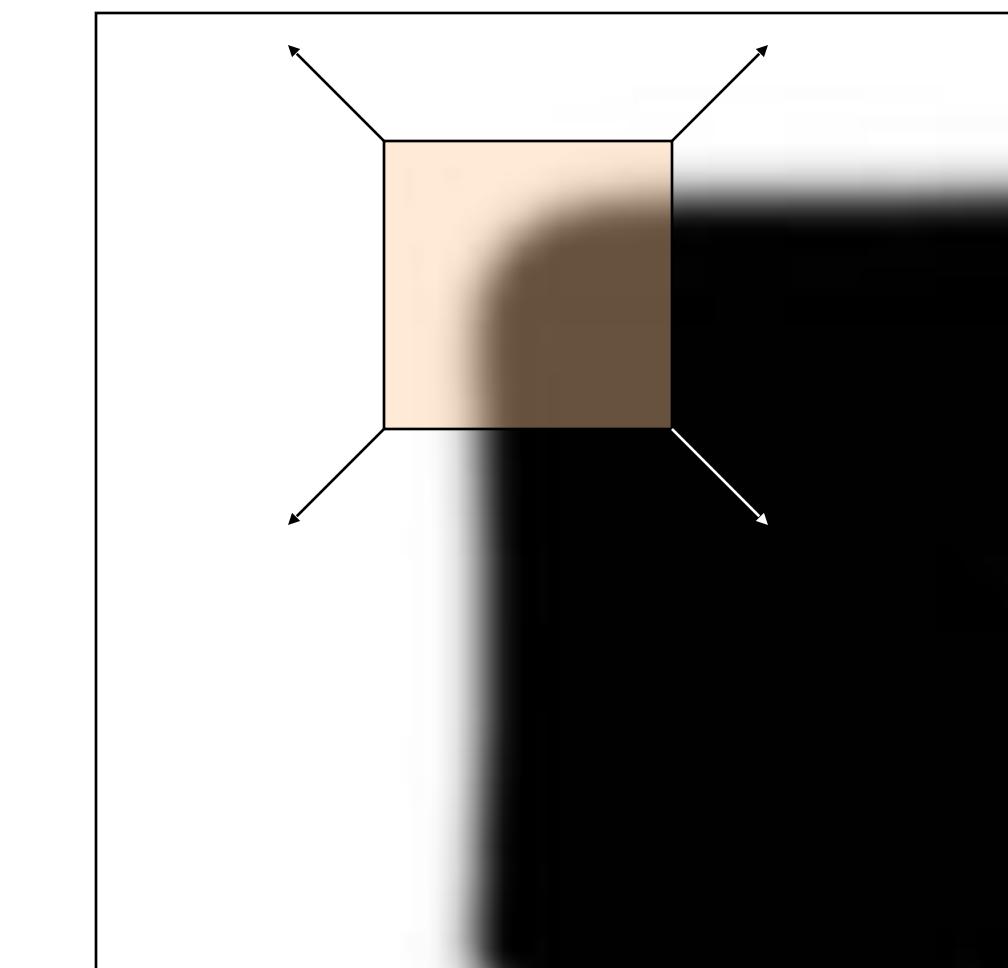
Shifting a window in *any direction* should give a *large change* in intensity at a corner



“flat” region:
no change in
all directions



“edge”:
no change along
the edge direction

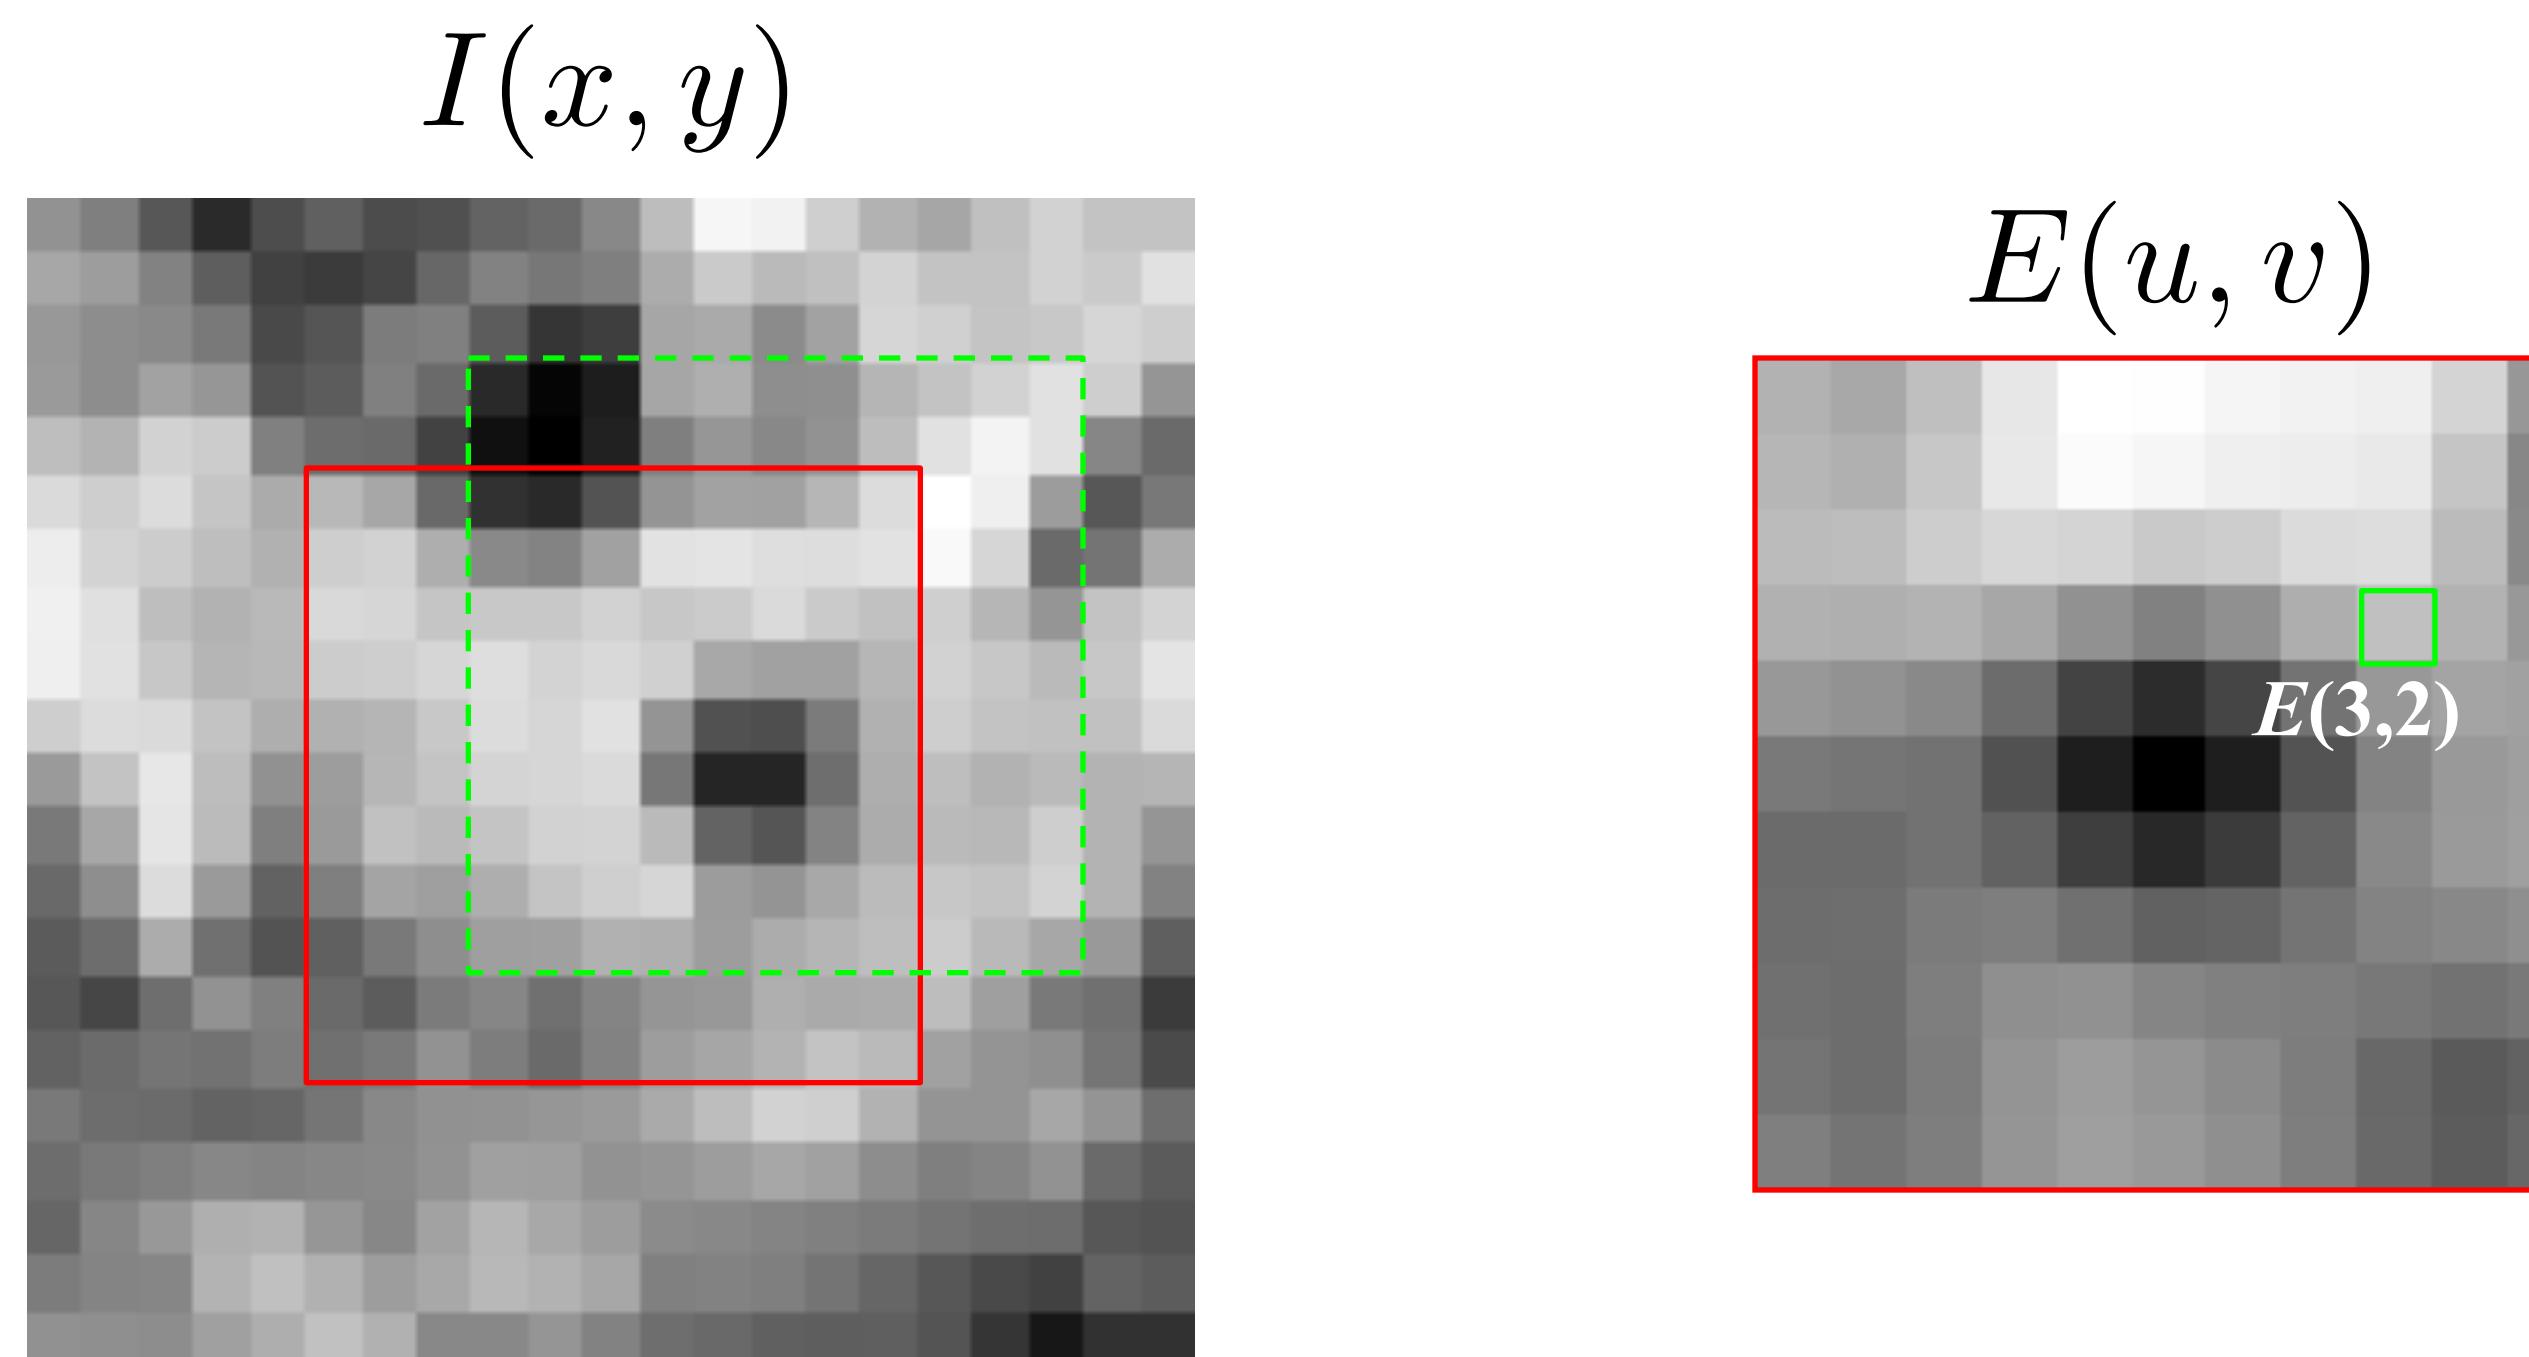


“corner”:
significant change
in all directions

Corner detection: Mathematics

The change in appearance of window W for the shift $[u, v]$:

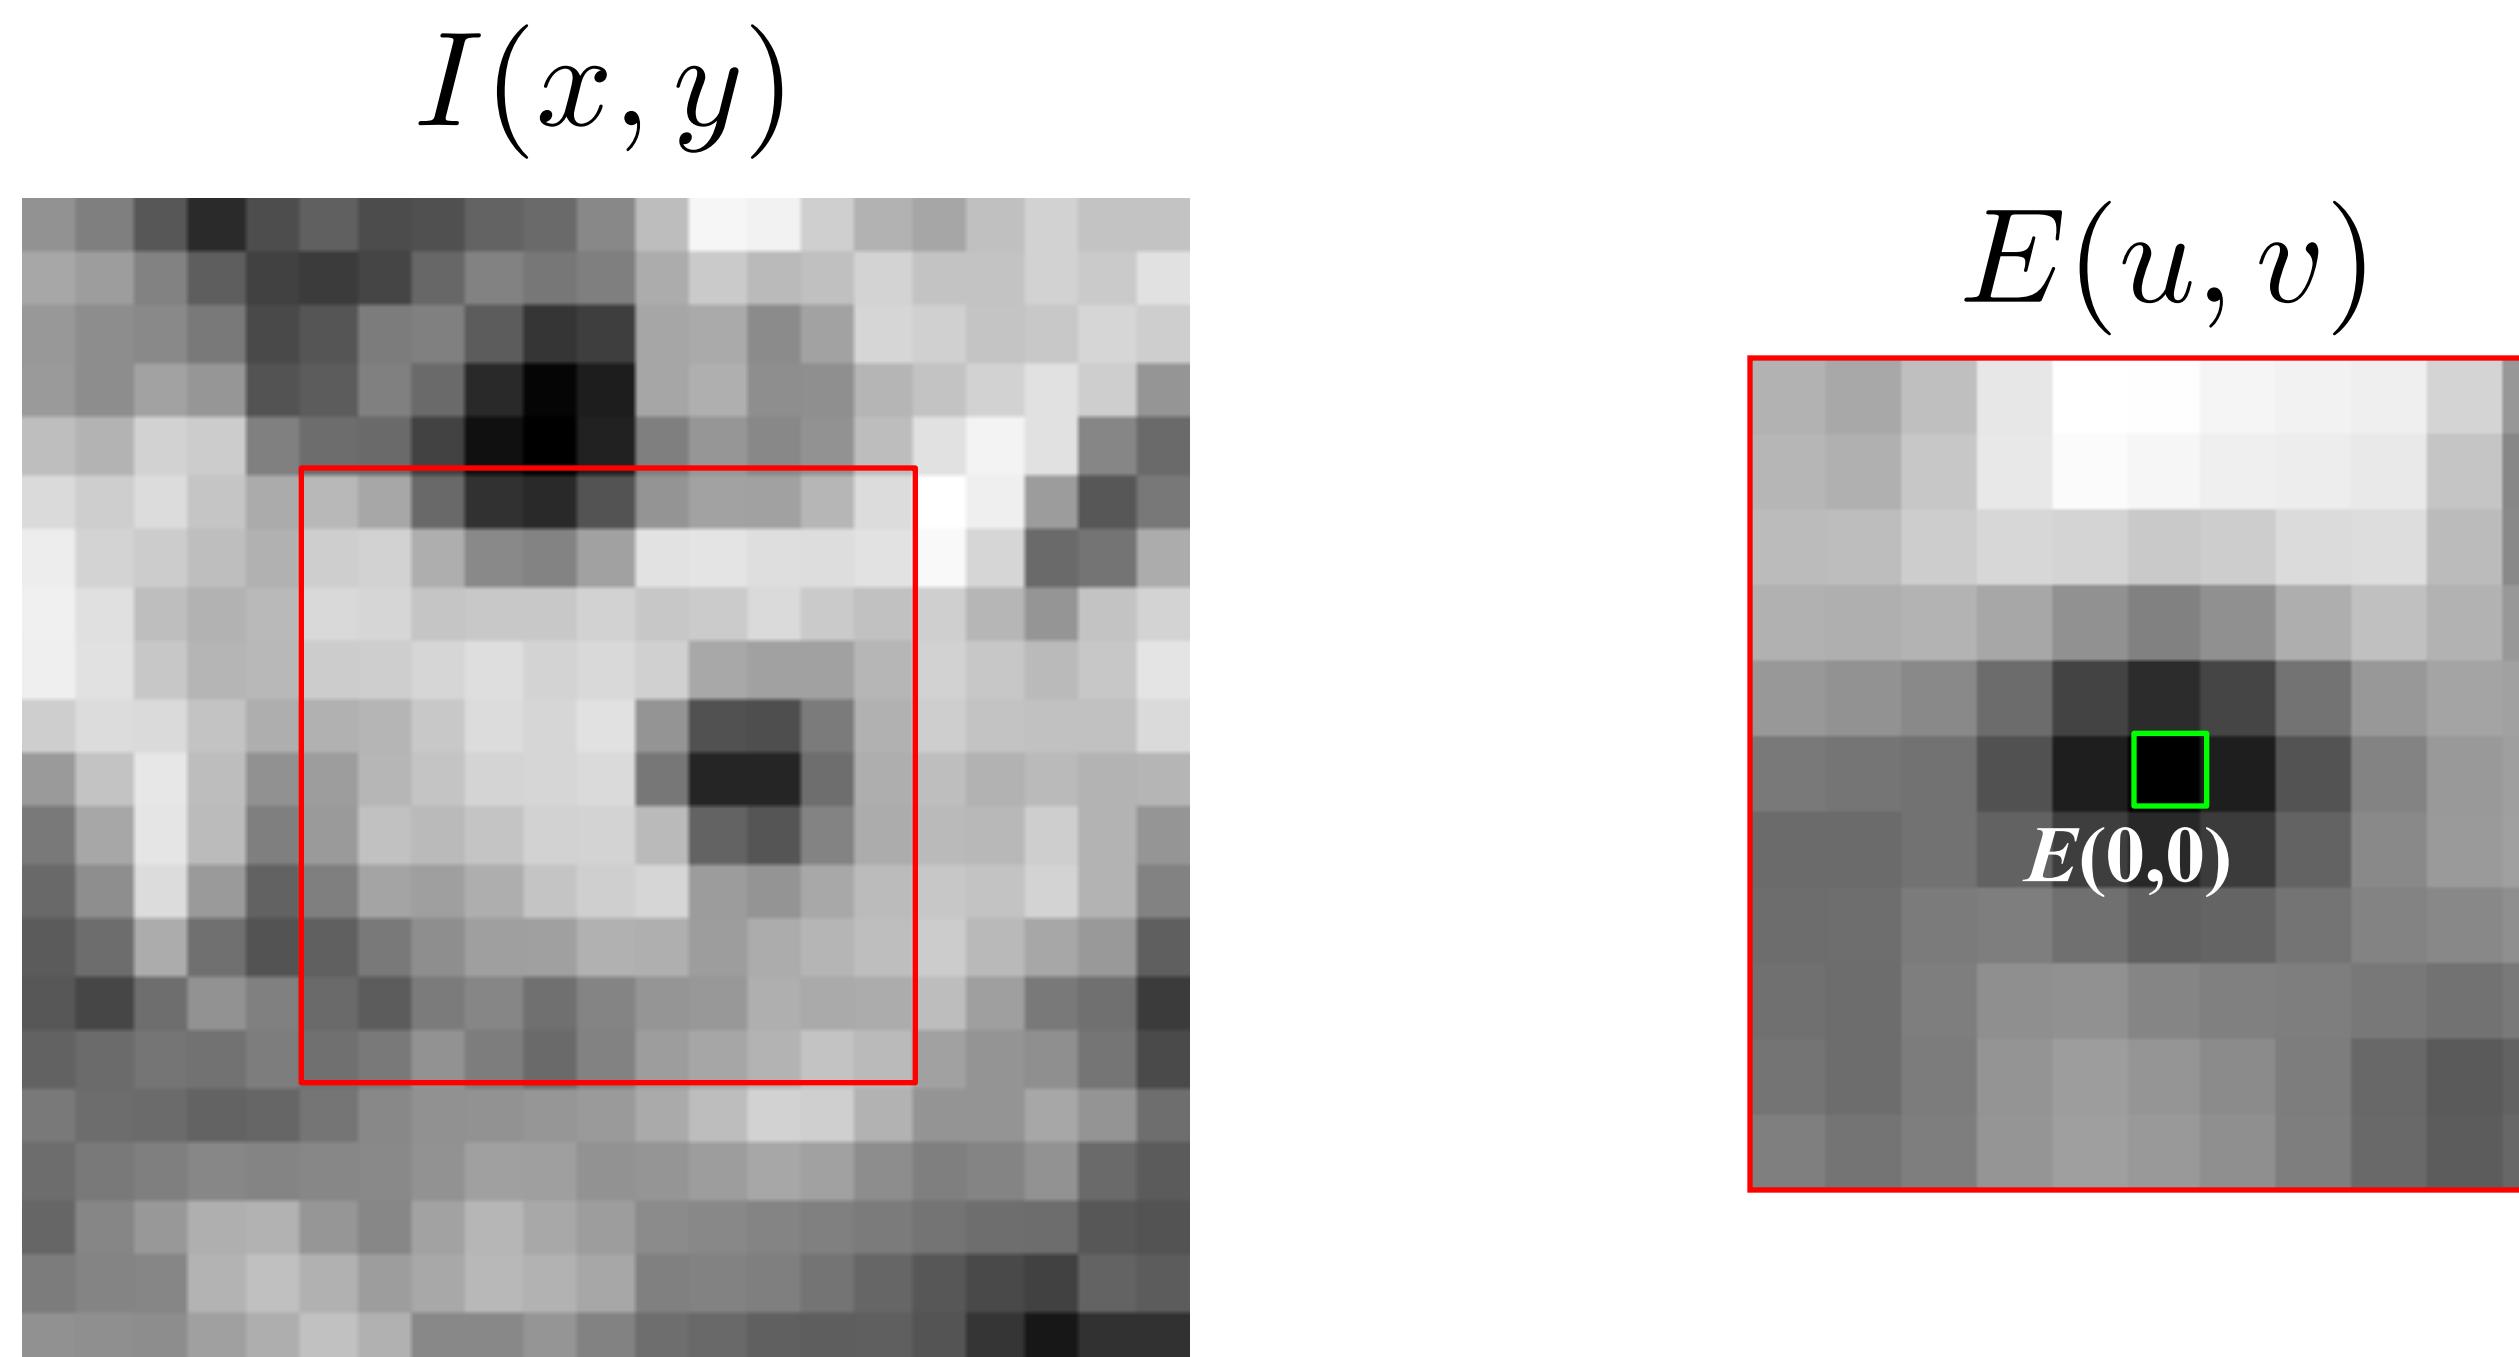
$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$



Corner detection: Mathematics

The change in appearance of window W for the shift $[u, v]$:

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$



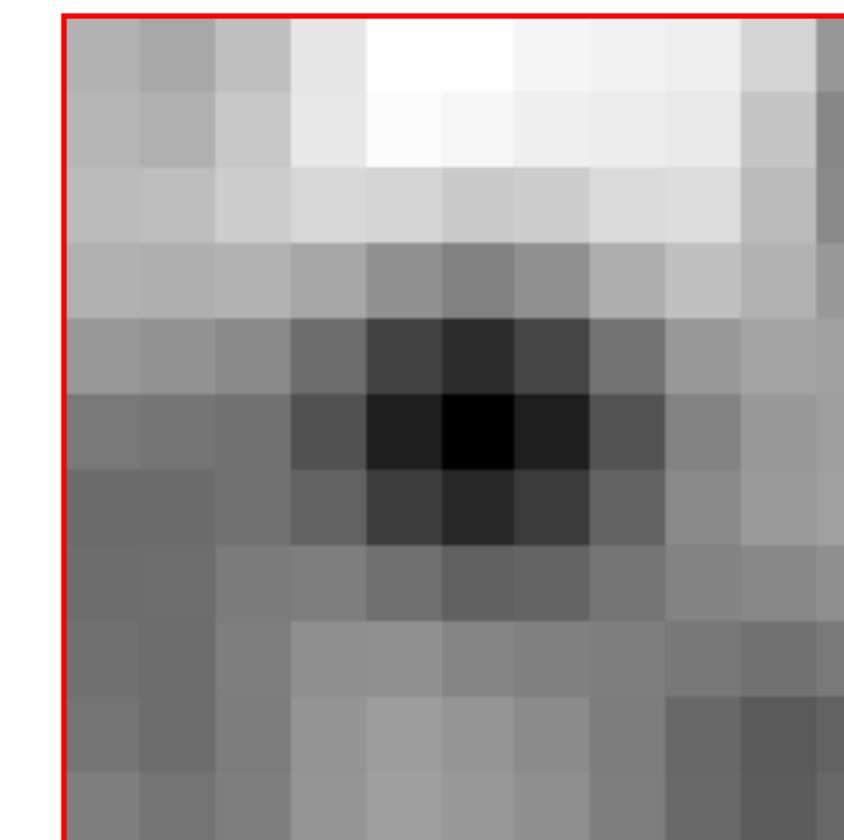
Corner detection: Mathematics

The change in appearance of window W for the shift $[u, v]$:

$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$

We want to find out how this function behaves for small shifts

$$E(u, v)$$



Corner detection: Mathematics

First-order Taylor approximation for small motions $[u, v]$:

$$I(x + u, y + v) = I(x, y) + I_x u + I_y v$$

Let's plug this into $E(u, v)$

$$\begin{aligned} E(u, v) &= \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2 \\ &\simeq \sum_{(x,y) \in W} [I(x, y) + I_x u + I_y v - I(x, y)]^2 \\ &= \sum_{(x,y) \in W} [I_x u + I_y v]^2 \\ &= \sum_{(x,y) \in W} [I_x^2 u^2 + I_x I_y u v + I_y I_x u v + I_y^2 v^2] \end{aligned}$$

Corner Detection: Mathematics

The quadratic approximation can be written as

$$E(u, v) \approx [u \quad v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a *second moment matrix* computed from image derivatives:

$$M = \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix}$$

(the sums are over all the pixels in the window W)

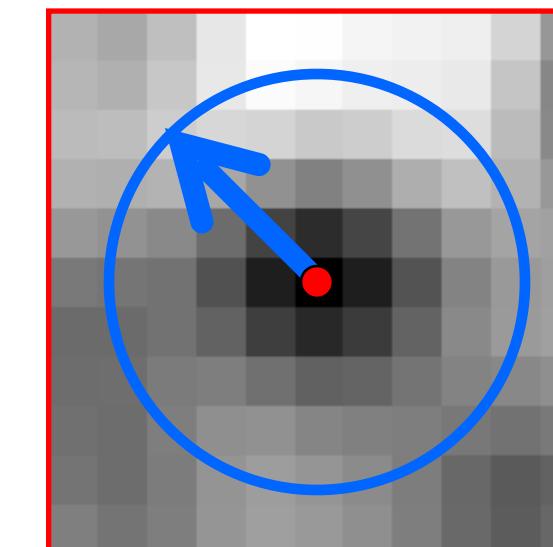
Interpreting the second moment matrix

The surface $E(u, v)$ is locally approximated by a quadratic form. Let's try to understand its shape.

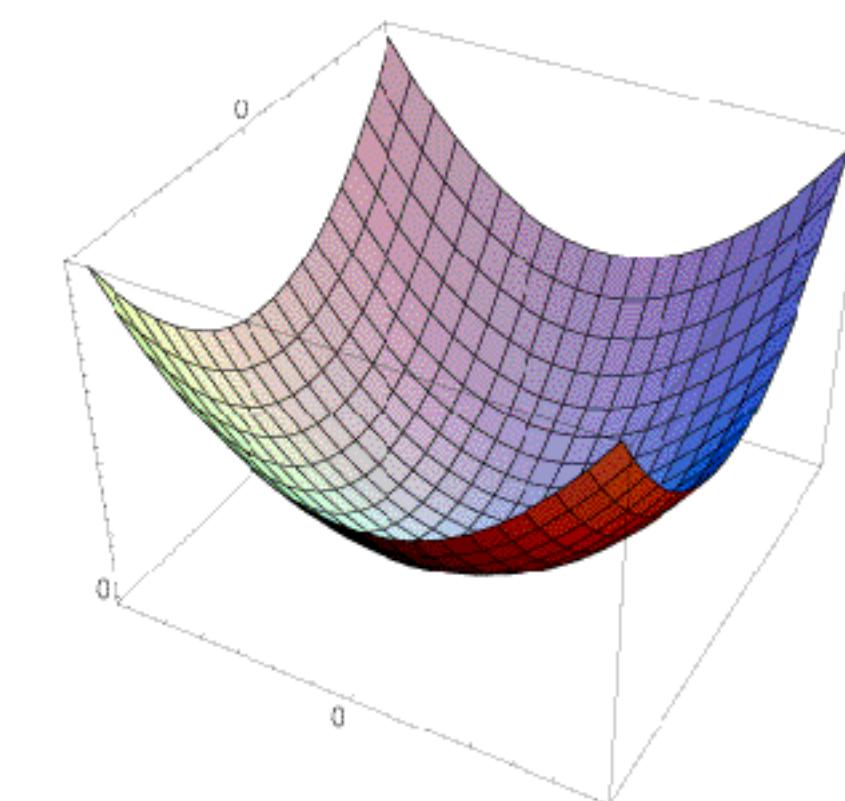
Specifically, in which directions does it have the smallest/greatest change?

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

$E(u, v)$



$$M = \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix}$$



Interpreting the second moment matrix

First, consider the axis-aligned case
(gradients are either horizontal or vertical)

$$M = \begin{bmatrix} \sum_{x,y} I_x^2 & \sum_{x,y} I_x I_y \\ \sum_{x,y} I_x I_y & \sum_{x,y} I_y^2 \end{bmatrix} = \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$$

If either a or b is close to 0, then this is **not** a corner, so look for locations where both are large.

Interpreting the second moment matrix

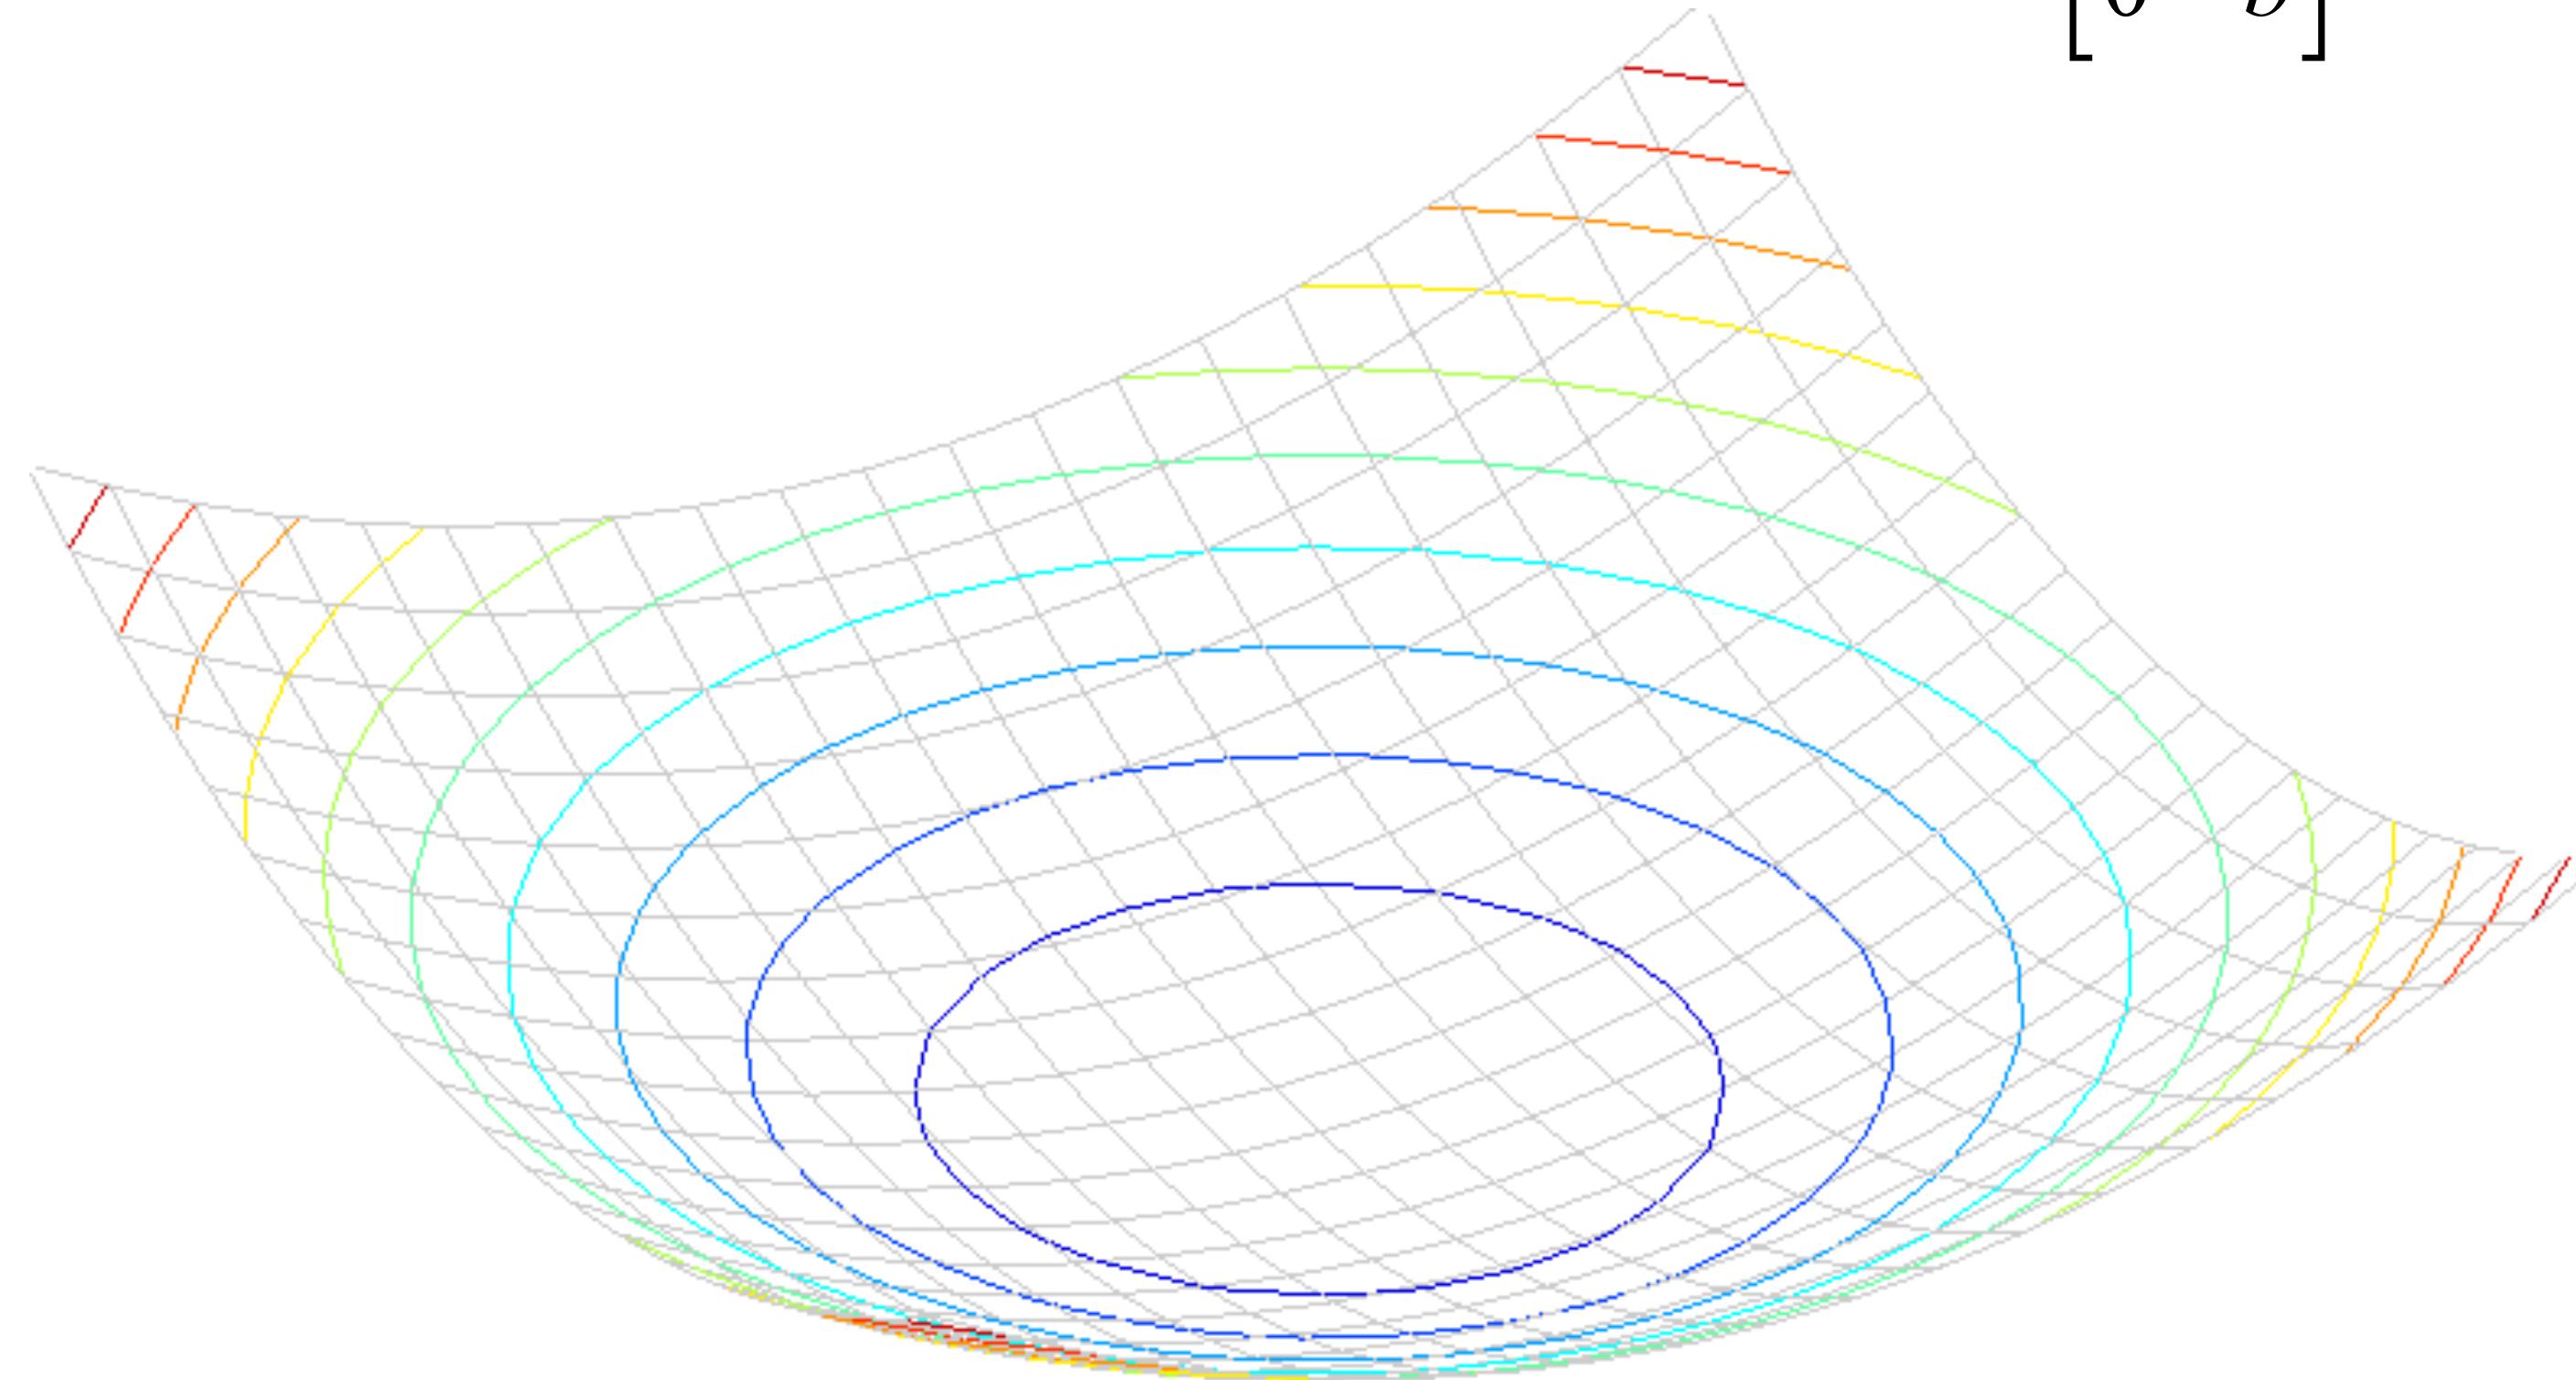
Consider a horizontal “slice” of $E(u, v)$:

$$[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$$

\downarrow

$$\begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$$

This is the equation of an ellipse.



Interpreting the second moment matrix

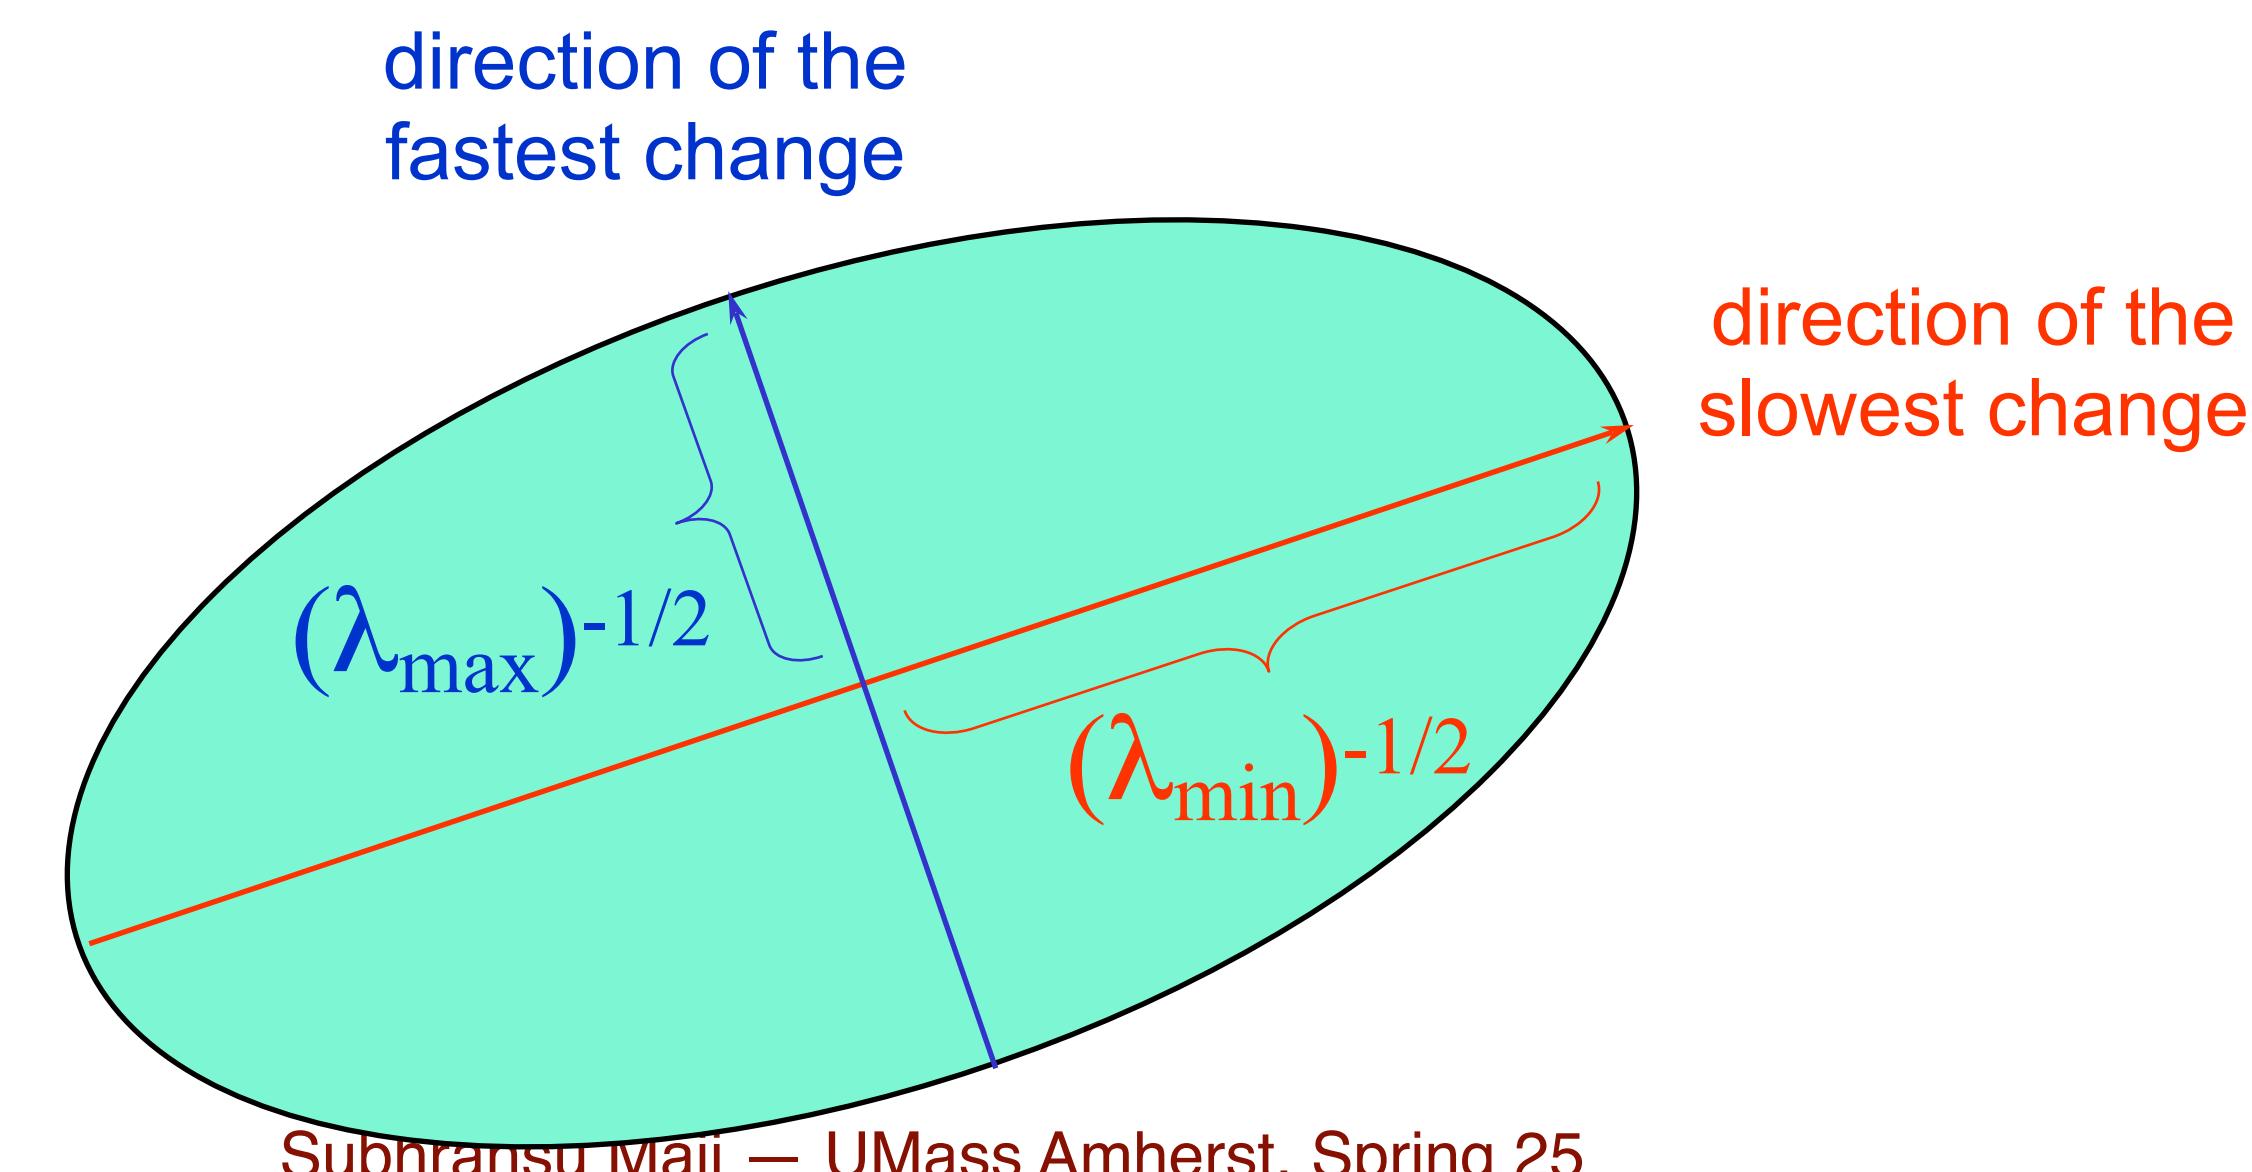
Consider a horizontal “slice” of $E(u, v)$: $[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} = \text{const}$

This is the equation of an ellipse.

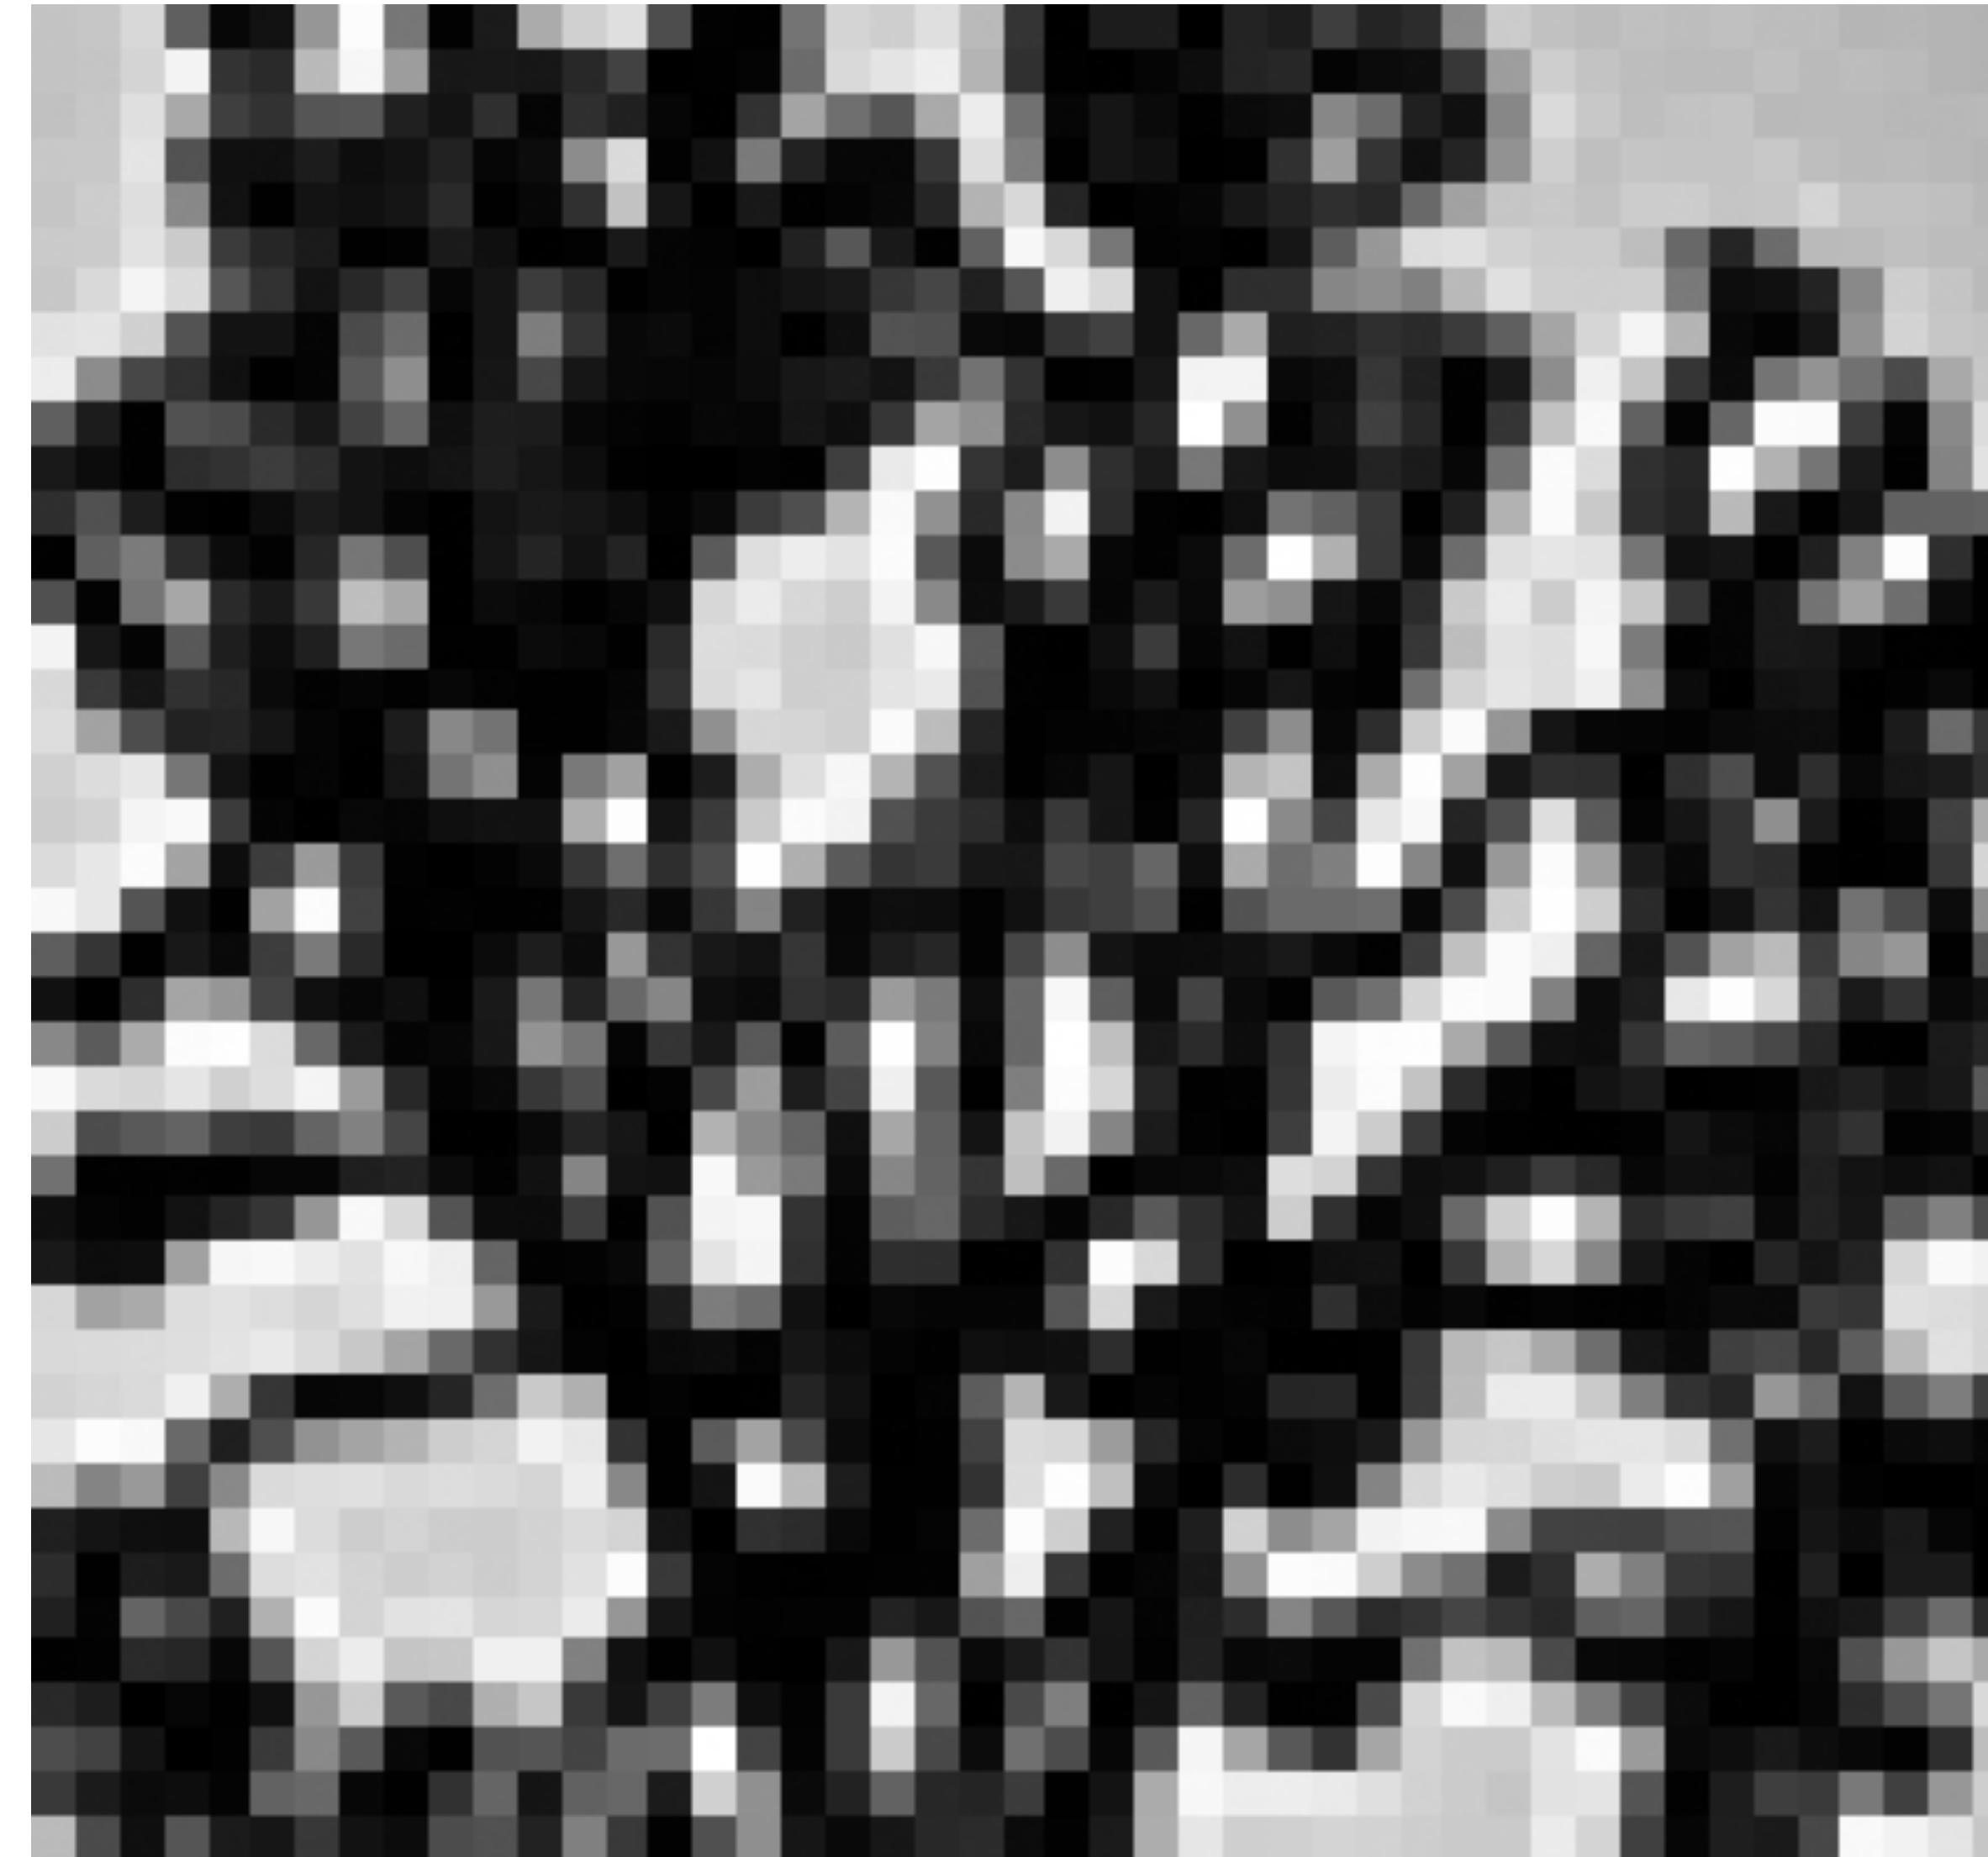
Diagonalization of M :

$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$

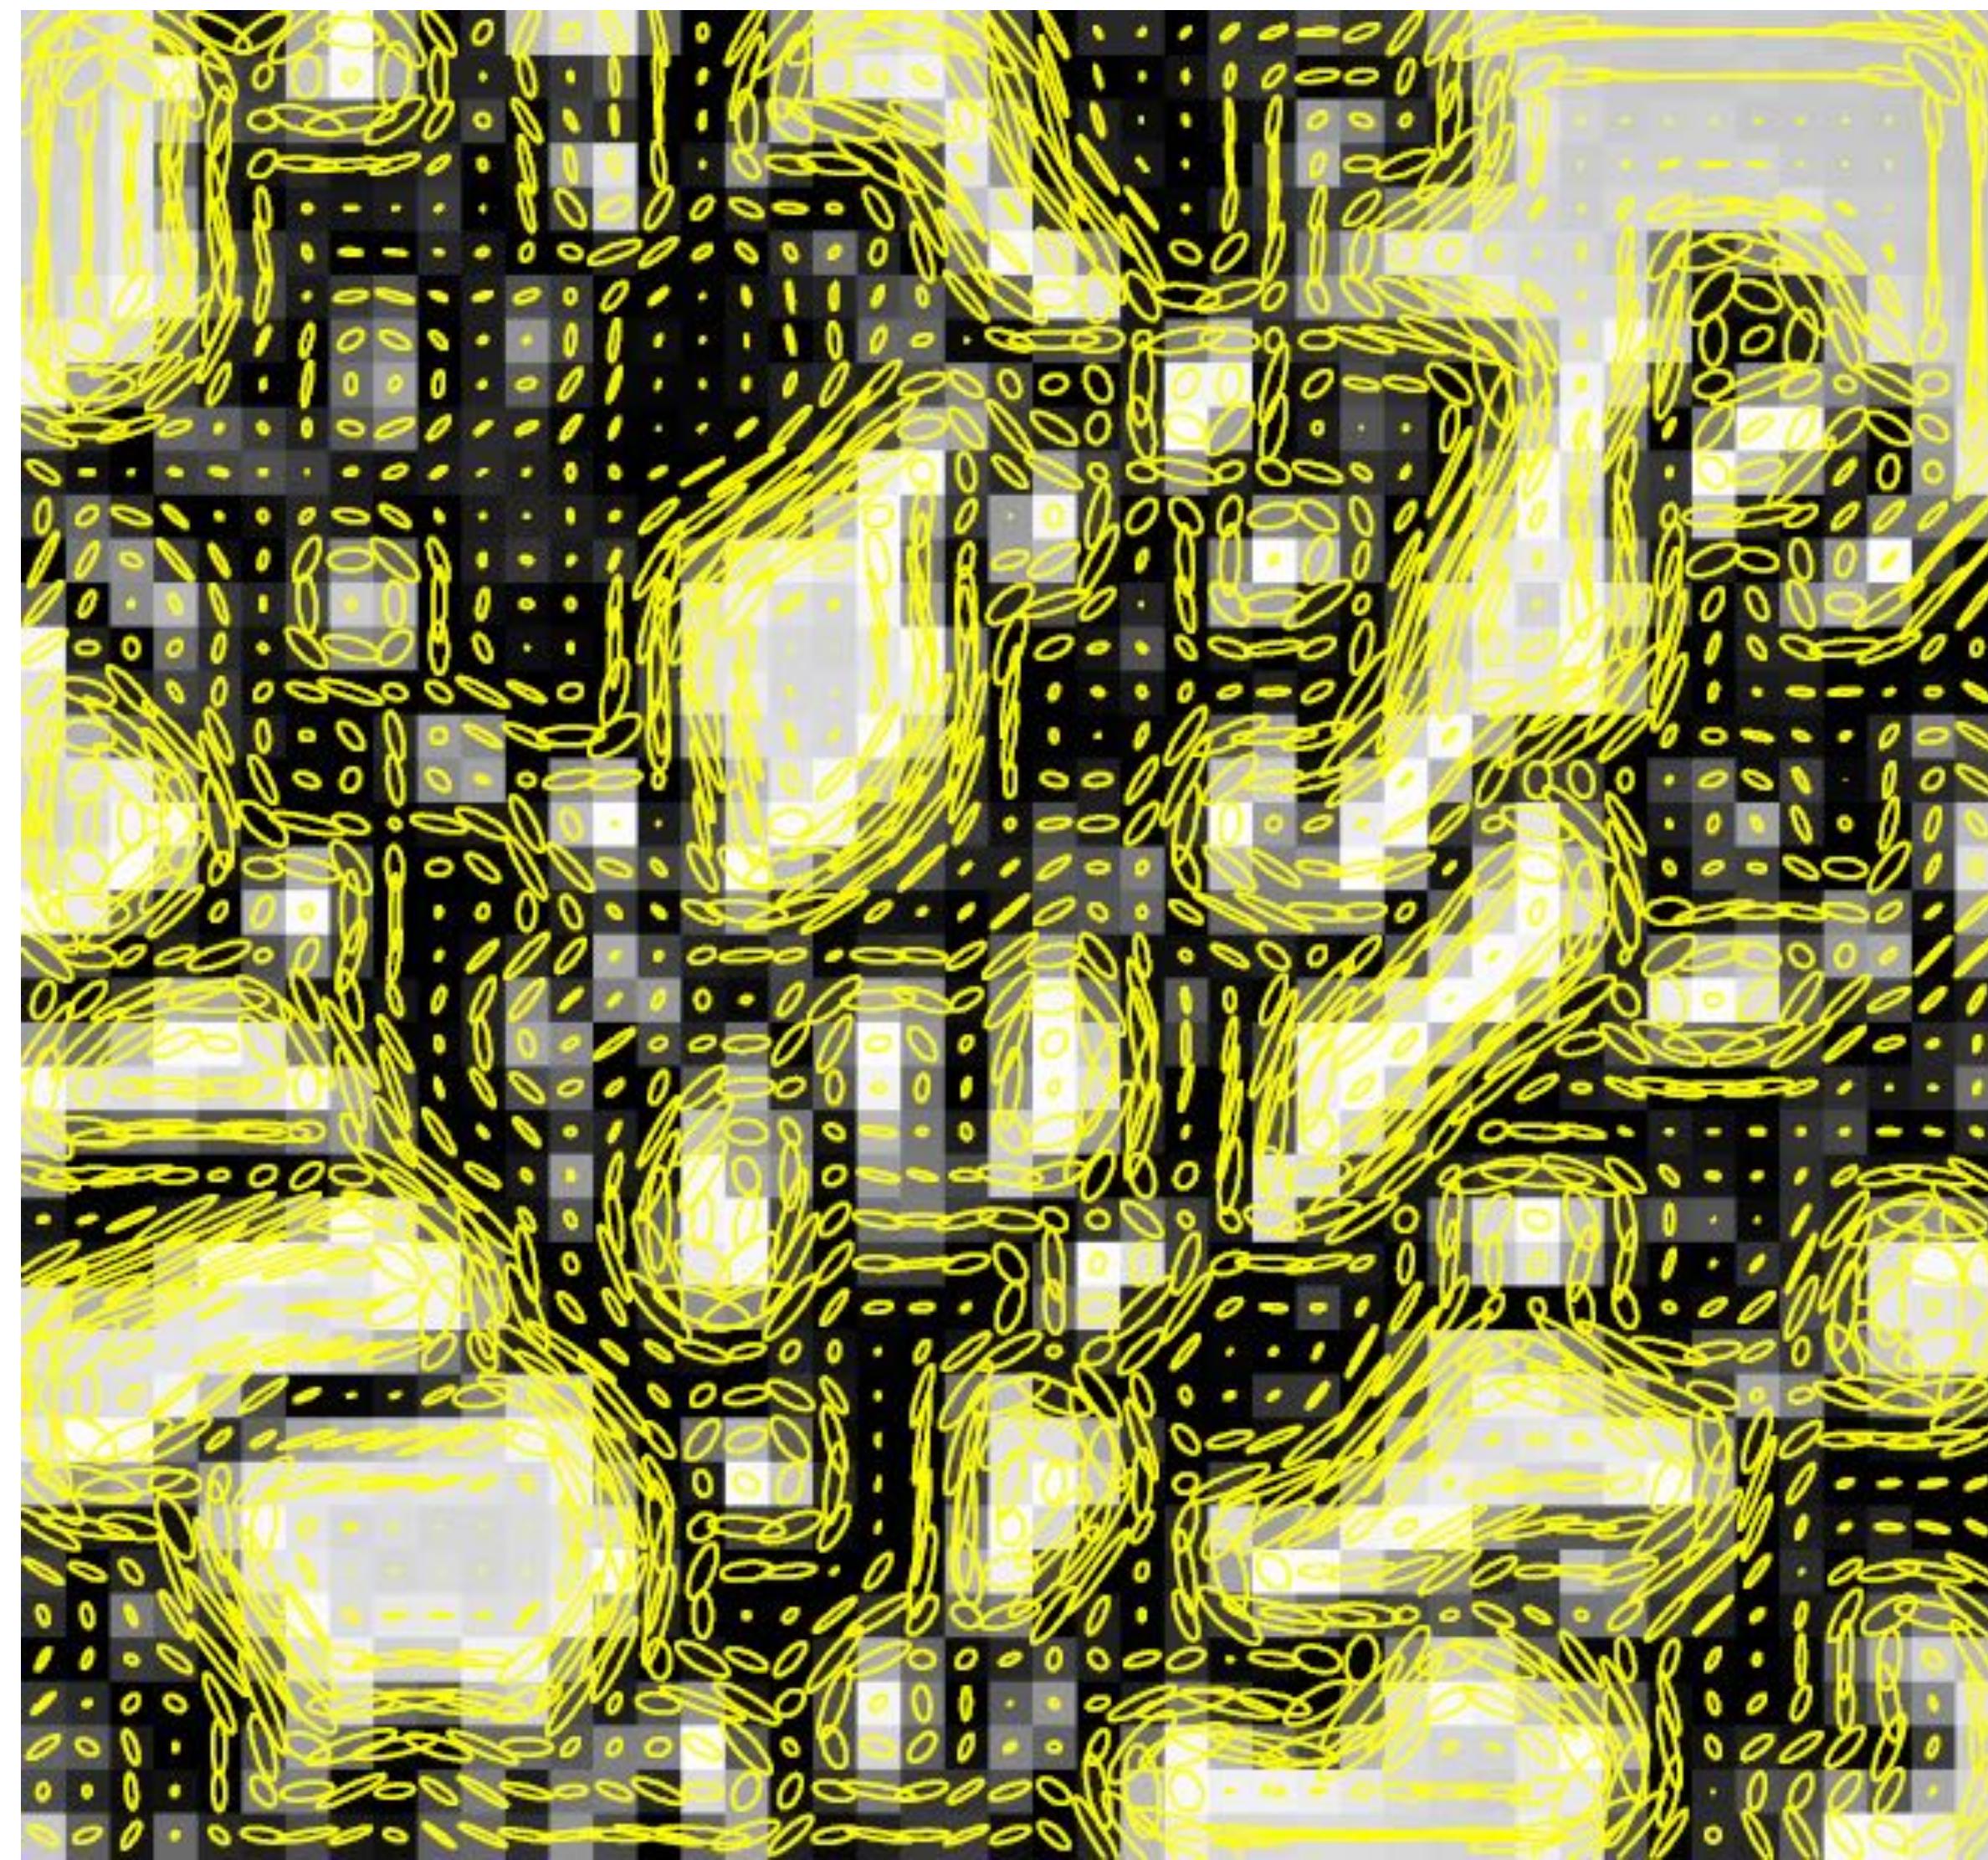
The axis lengths of the ellipse are determined by the eigenvalues and the orientation is determined by R



Visualization of second moment matrices

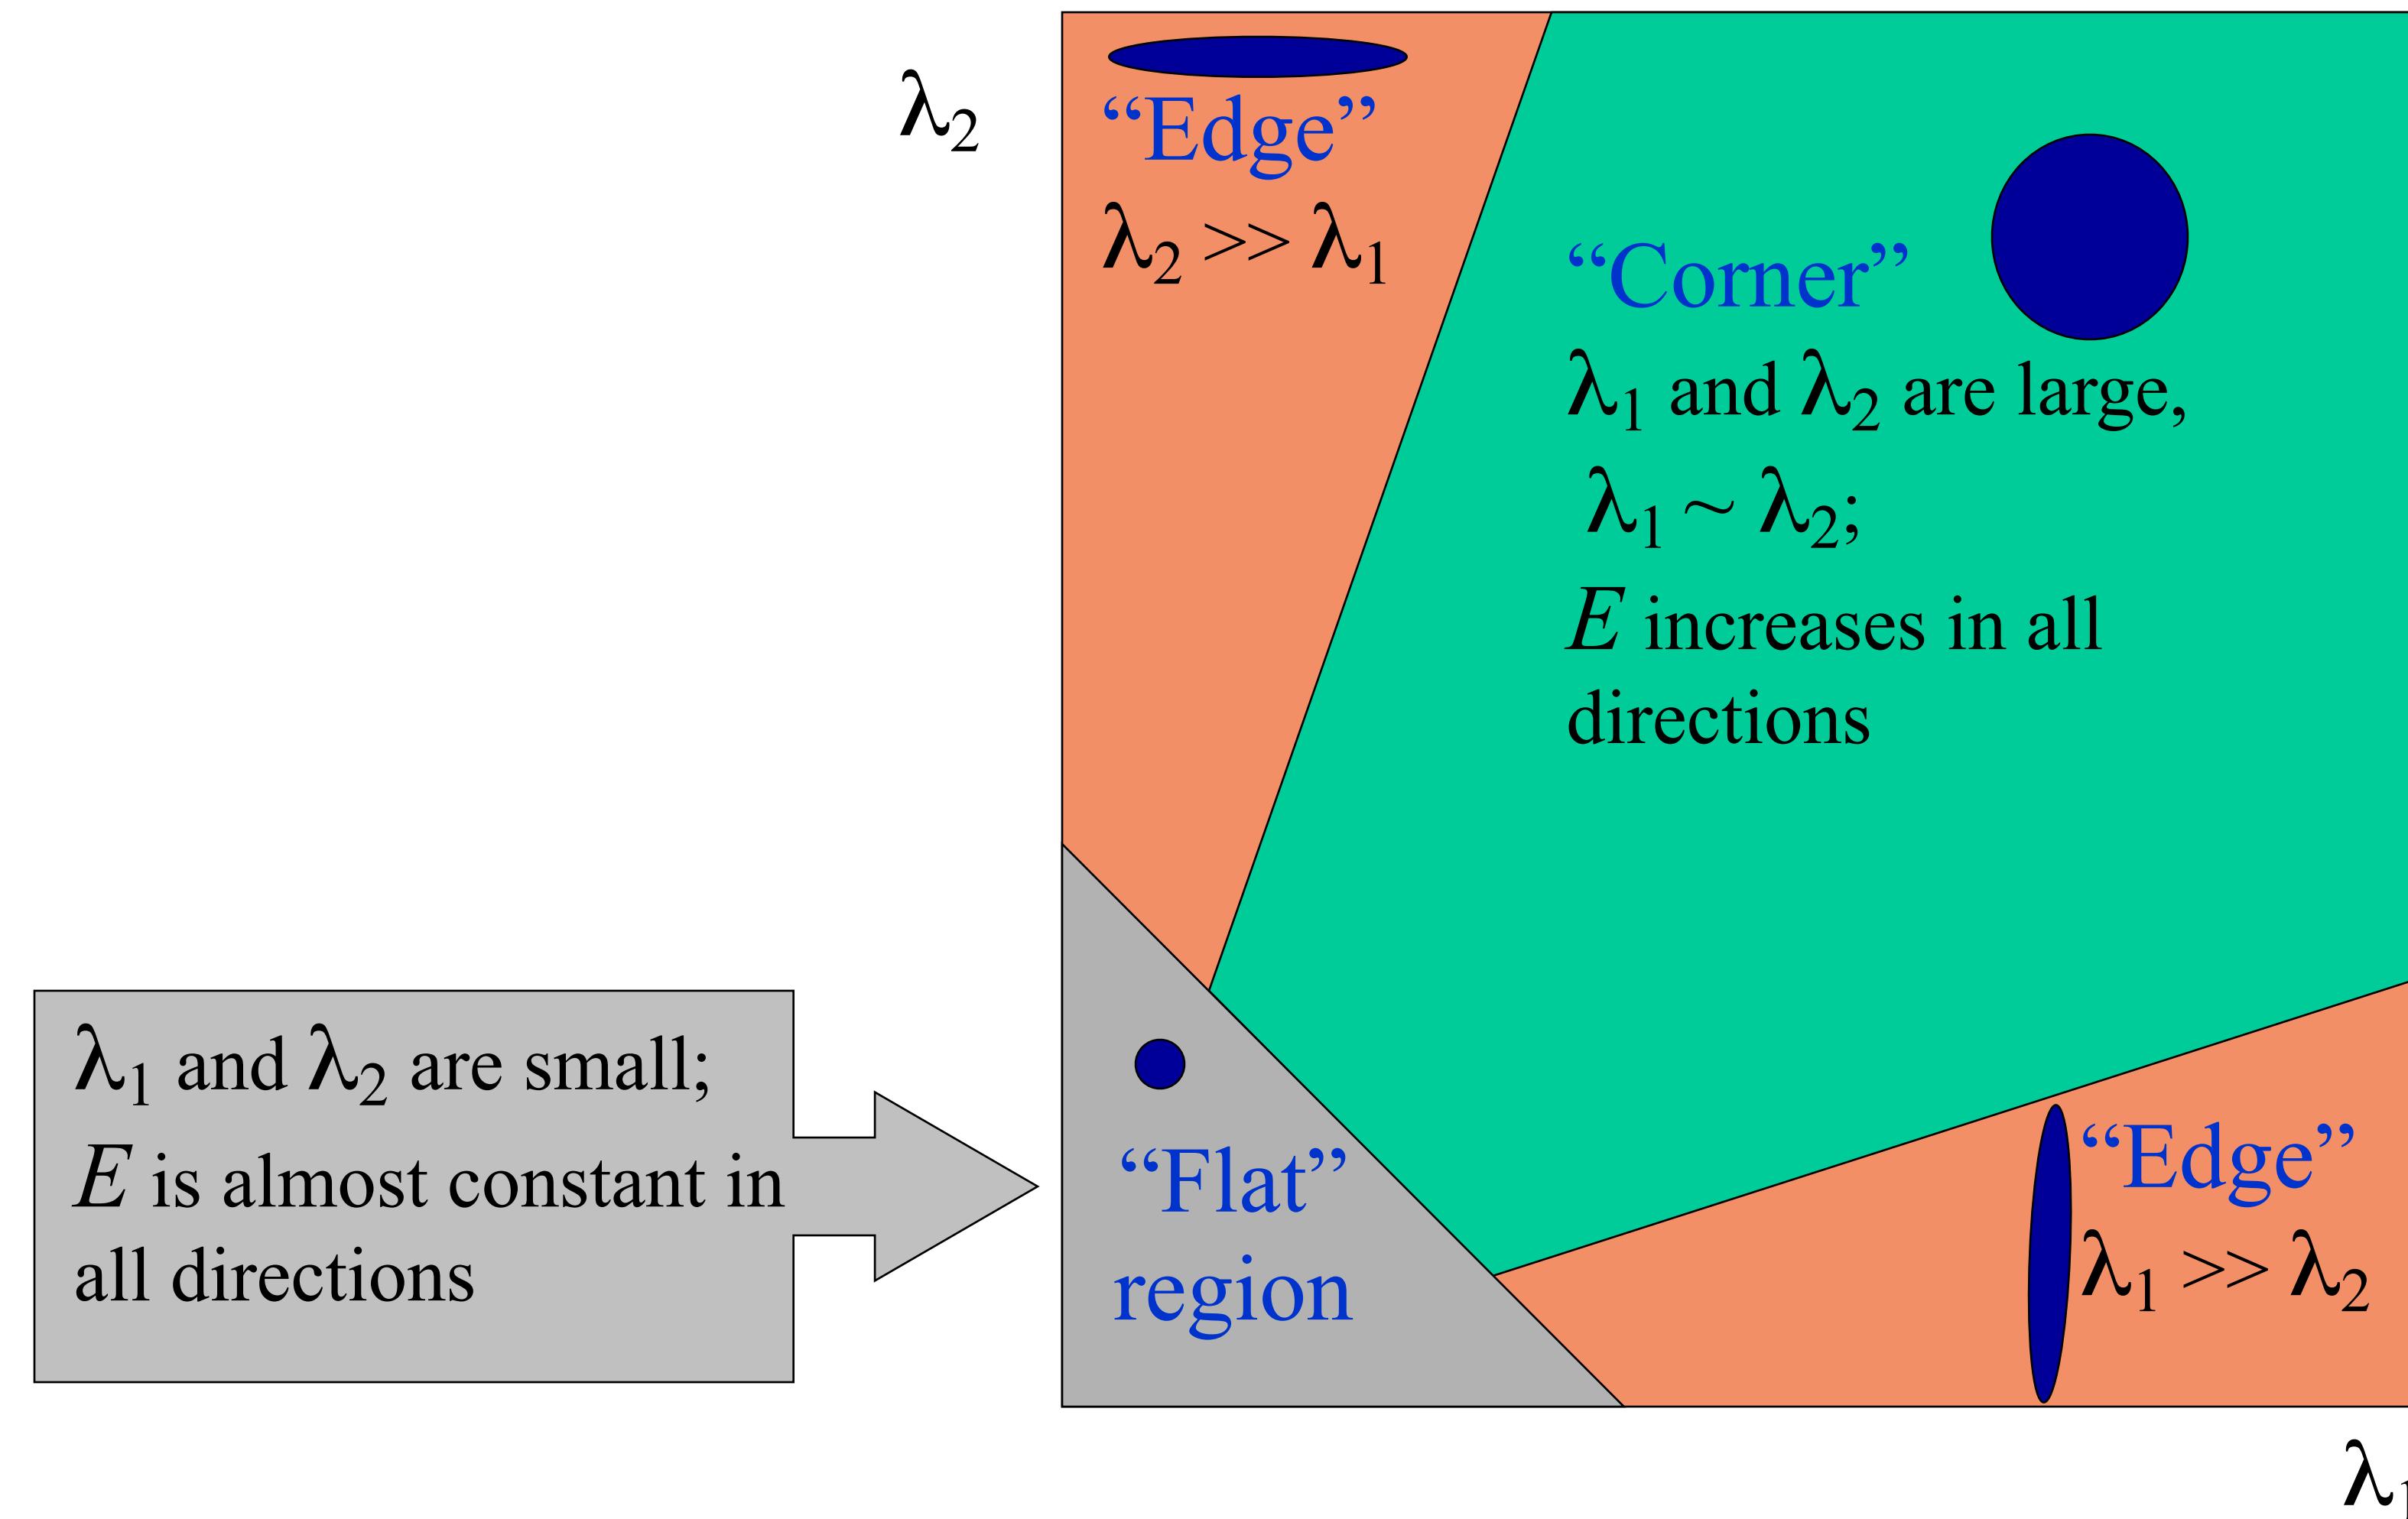


Visualization of second moment matrices



Interpreting the eigenvalues

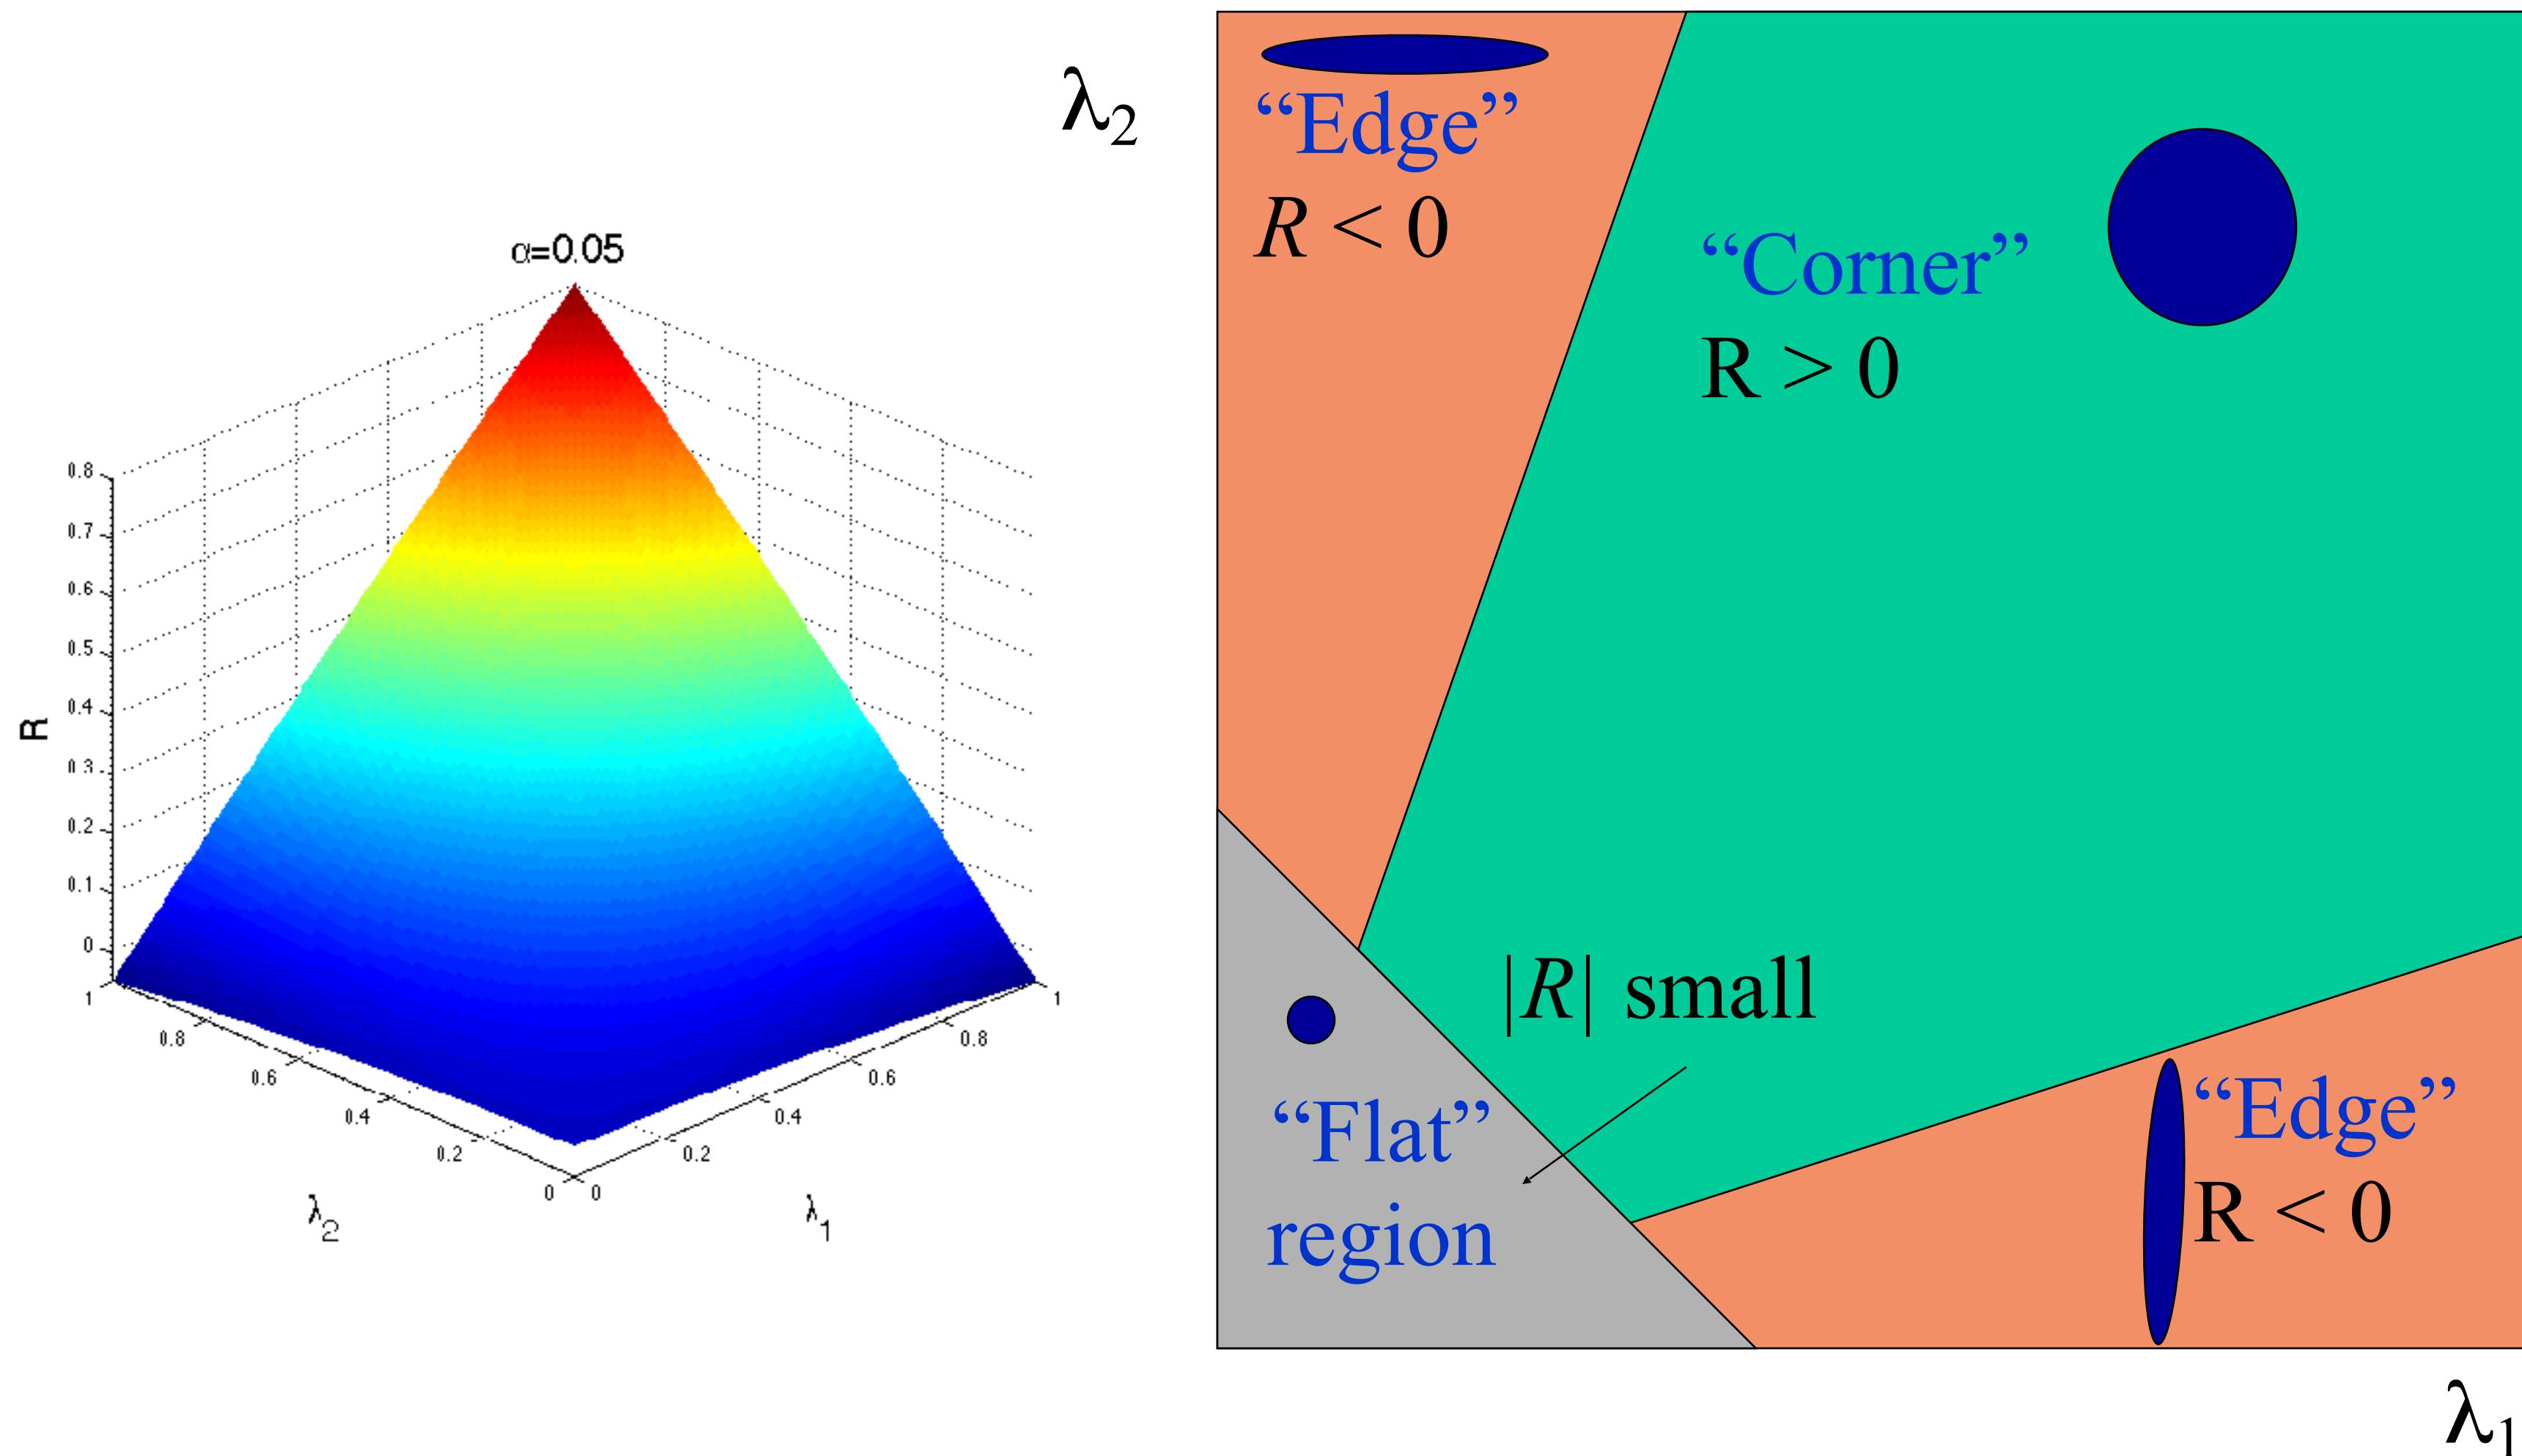
Classification of image points using eigenvalues of M :



Corner response function

$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2$$

α : constant (0.04 to 0.06)



The Harris corner detector

1. Compute partial derivatives at each pixel
2. Compute second moment matrix M in a Gaussian window around each pixel:

$$M = \begin{bmatrix} \sum_{x,y} w(x,y) I_x^2 & \sum_{x,y} w(x,y) I_x I_y \\ \sum_{x,y} w(x,y) I_x I_y & \sum_{x,y} w(x,y) I_y^2 \end{bmatrix}$$

C.Harris and M.Stephens. [**“A Combined Corner and Edge Detector.”**](#)
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

The Harris corner detector

1. Compute partial derivatives at each pixel
2. Compute second moment matrix M in a Gaussian window around each pixel
3. Compute corner response function R

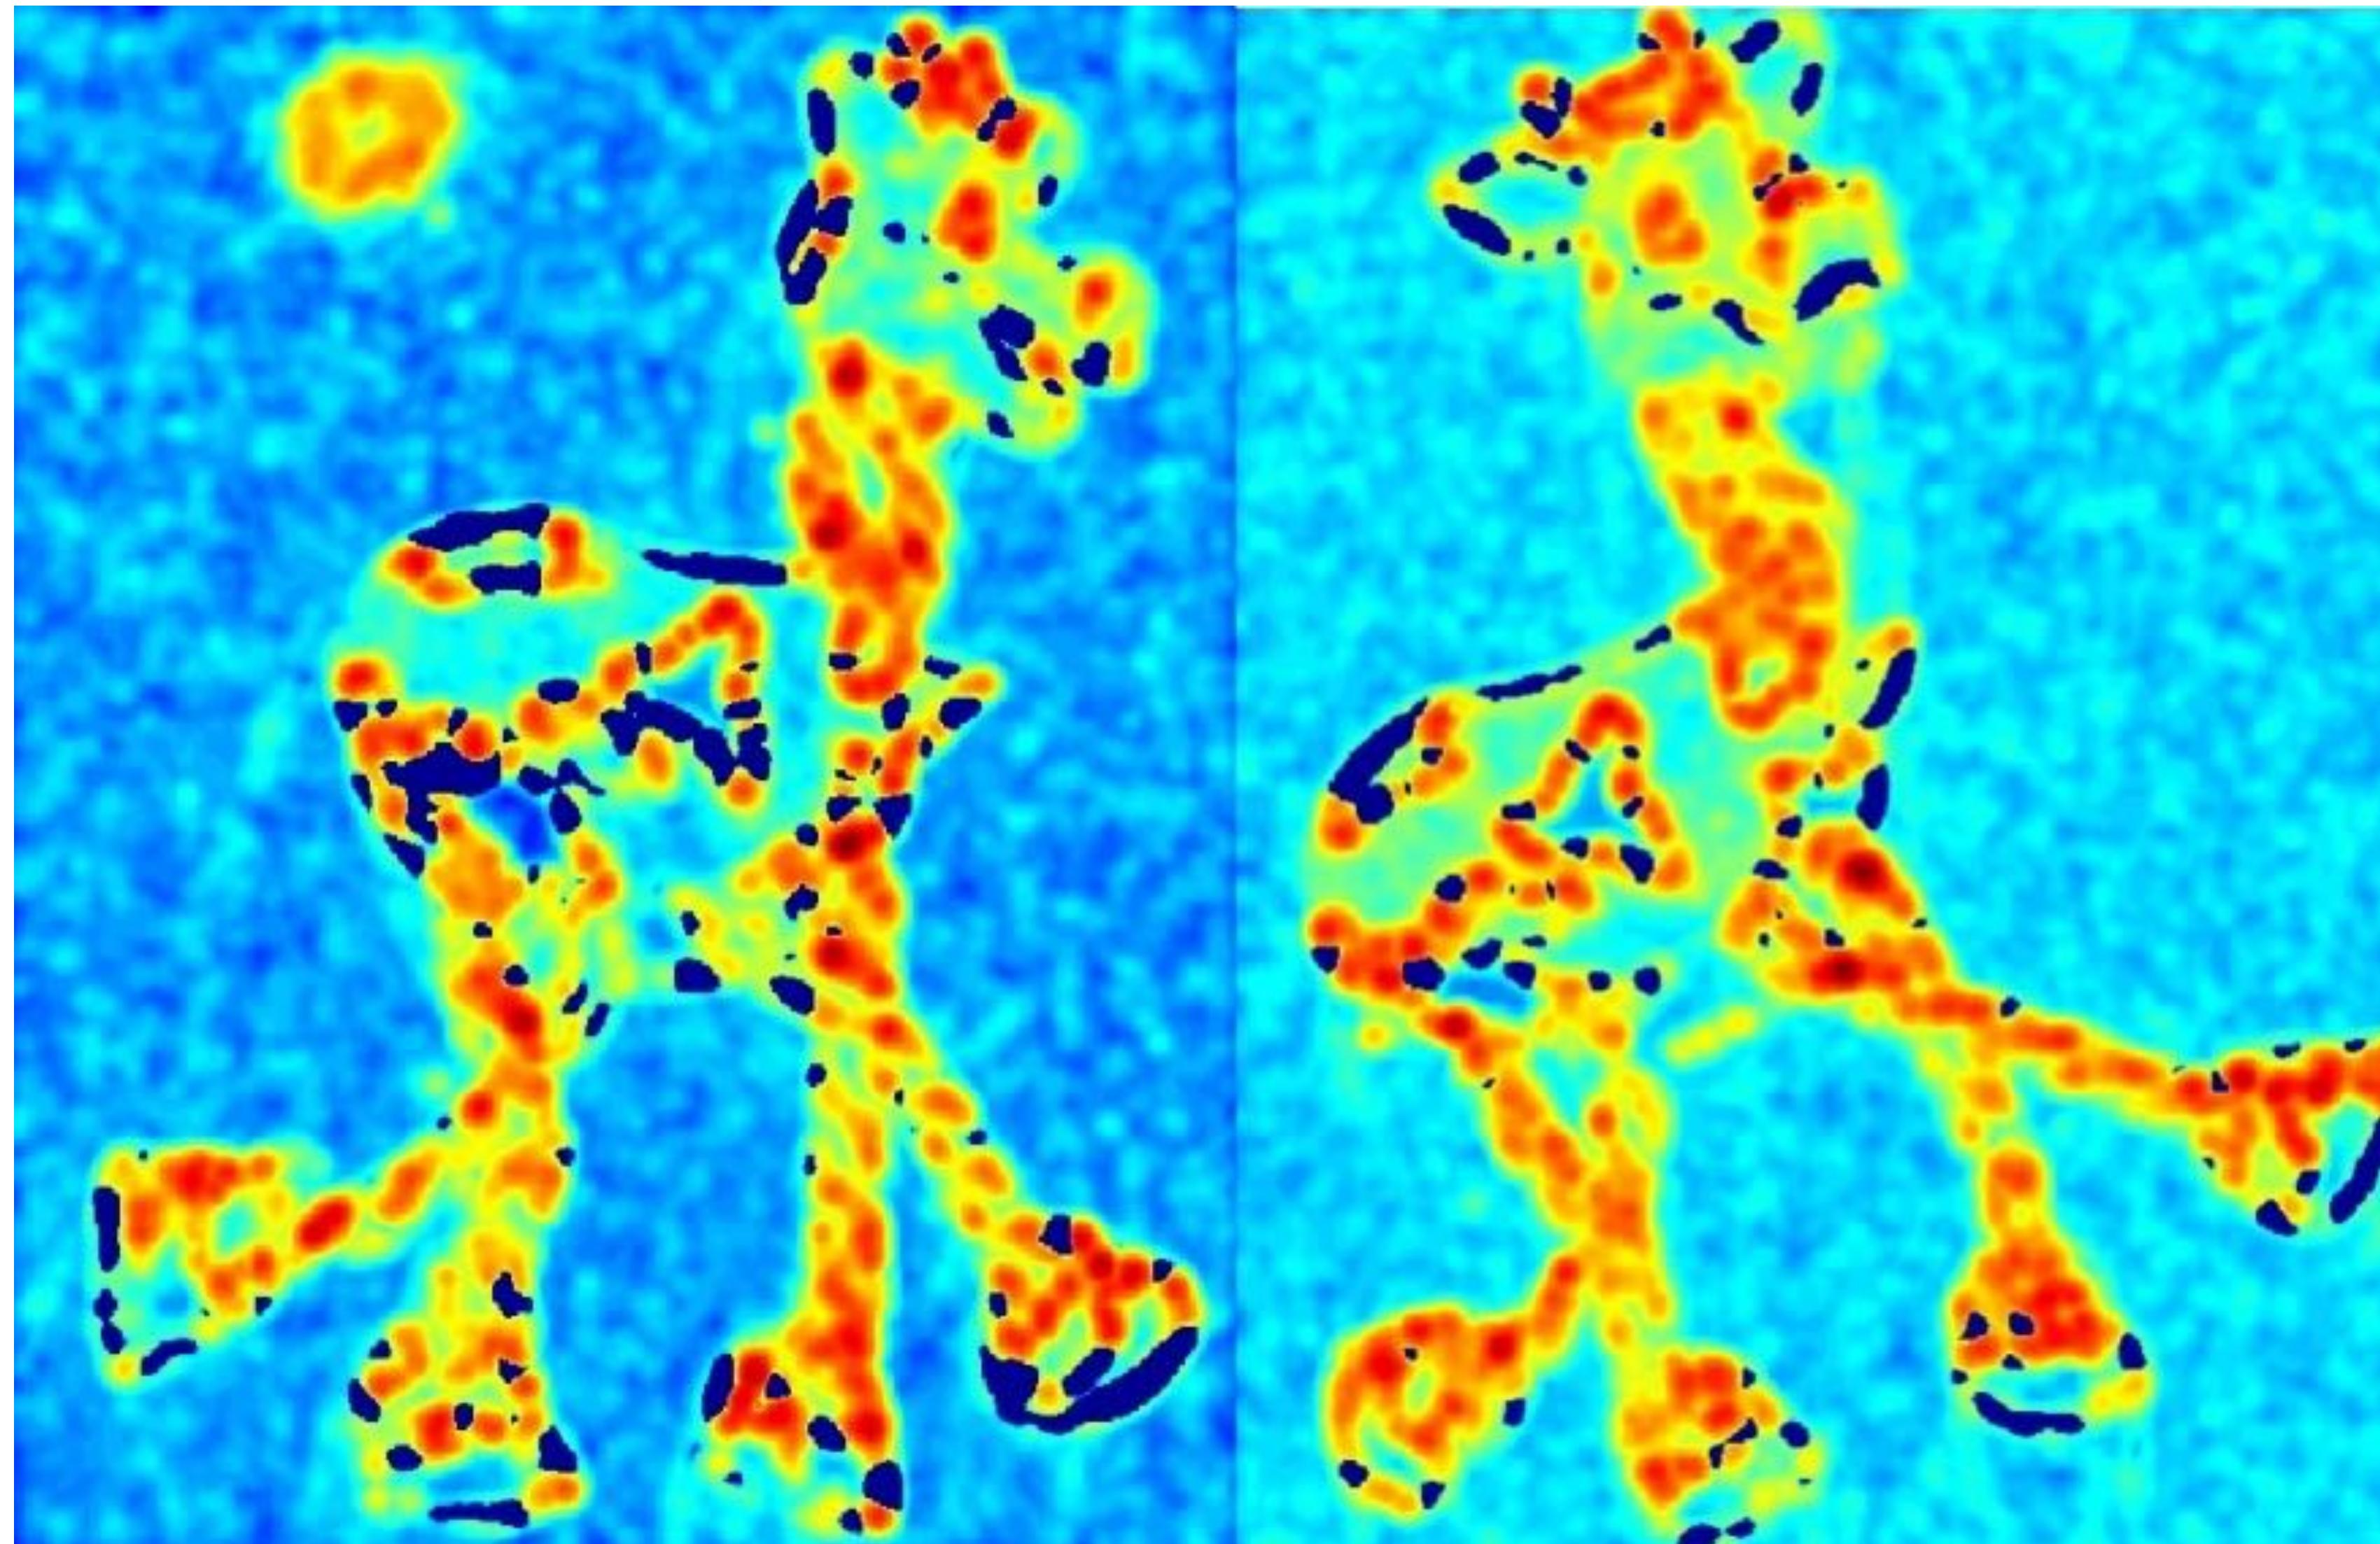
C.Harris and M.Stephens. [**“A Combined Corner and Edge Detector.”**](#)
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

Harris Detector: Steps



Harris Detector: Steps

Compute corner response R



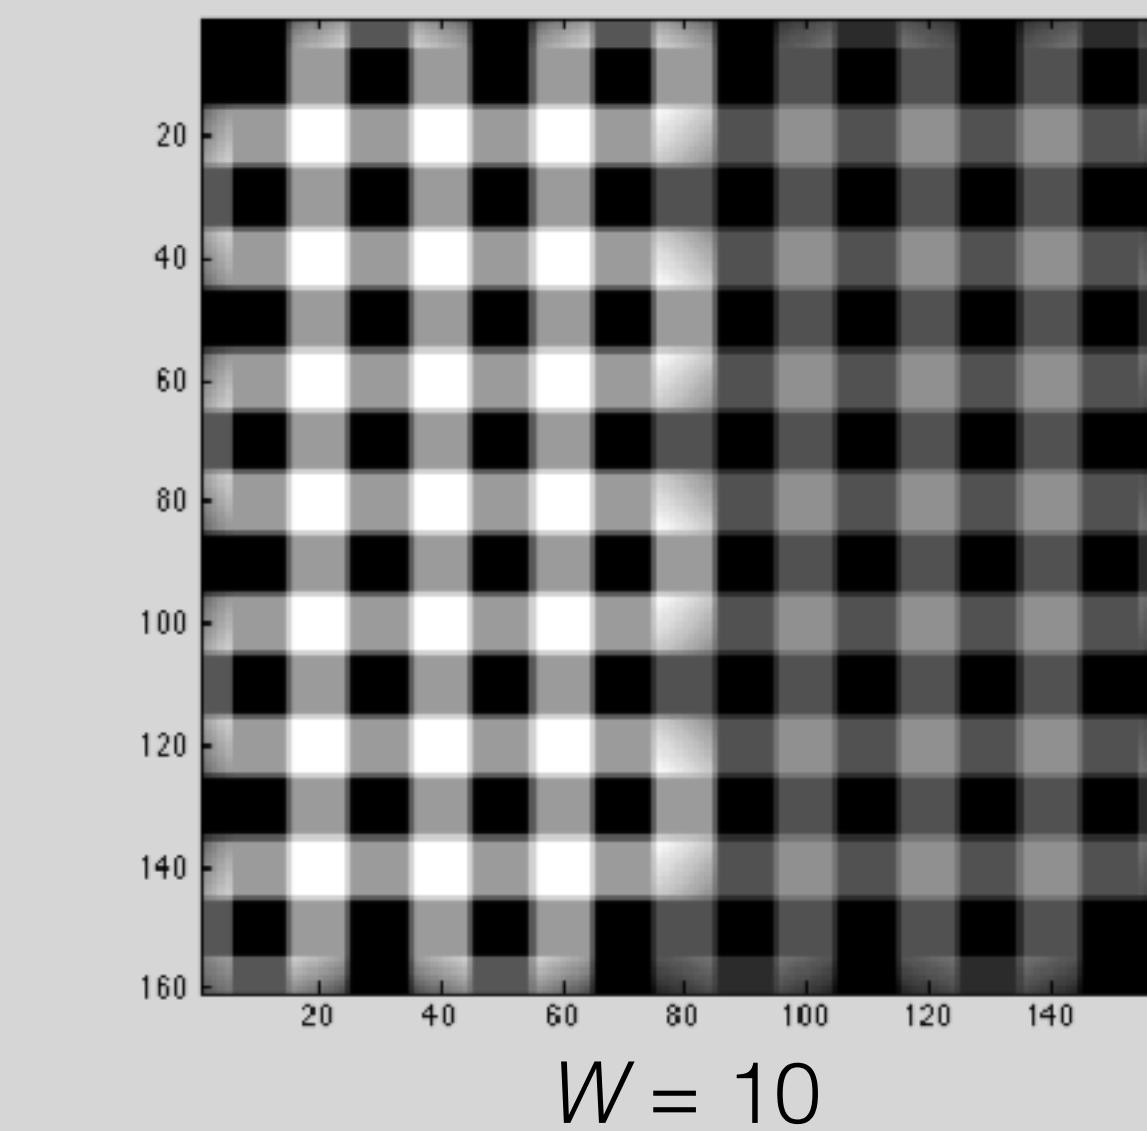
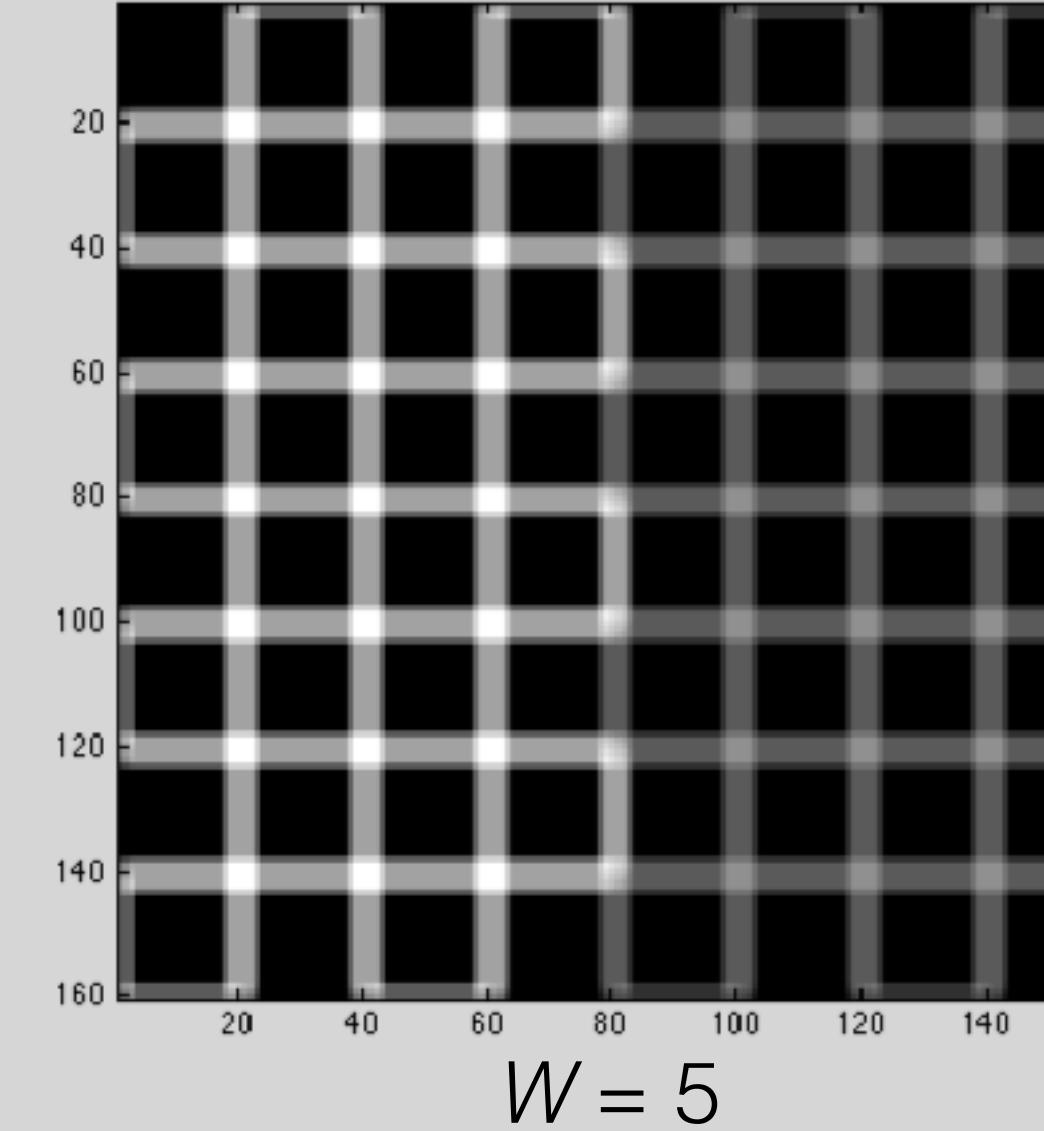
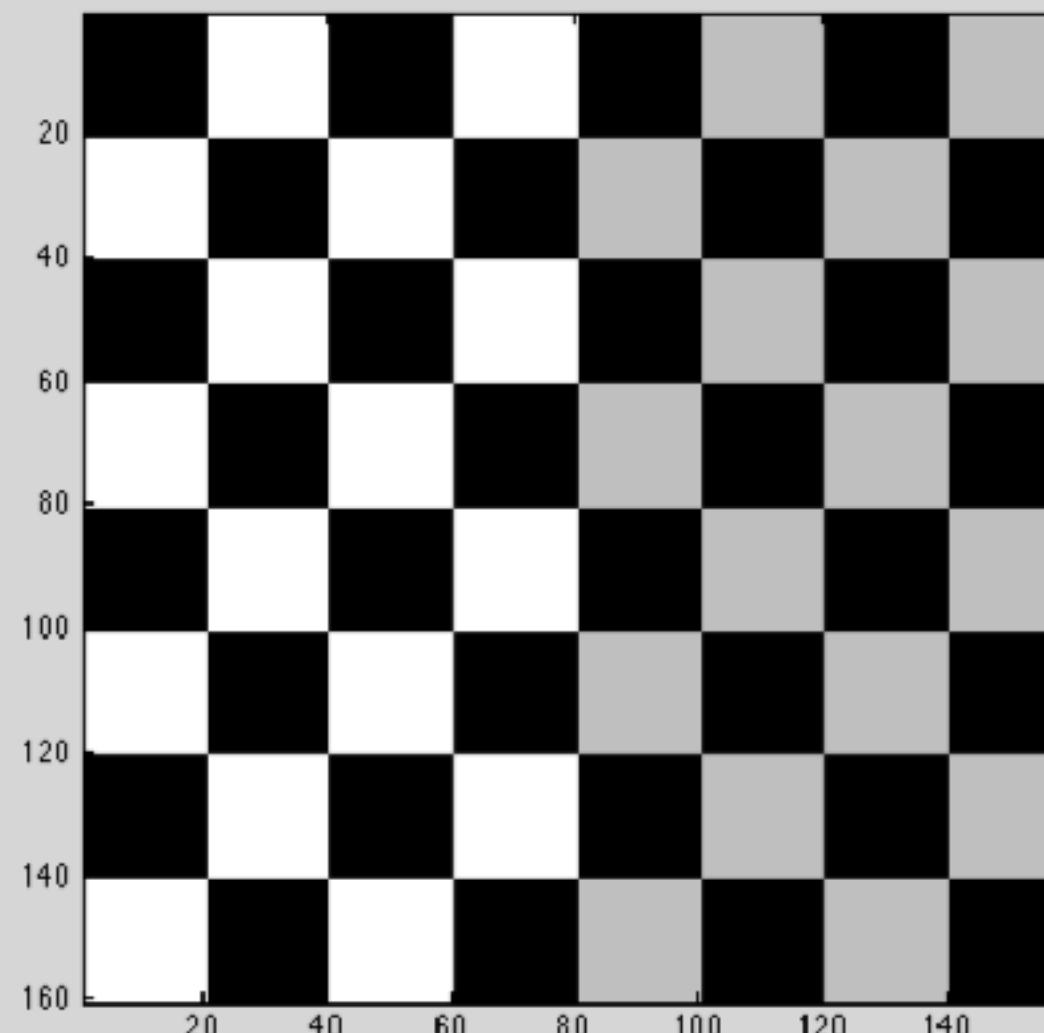
The Harris corner detector

1. Compute partial derivatives at each pixel
2. Compute second moment matrix M in a Gaussian window around each pixel
3. Compute corner response function R
4. Threshold R
5. Find local maxima of response function (non-maximum suppression)

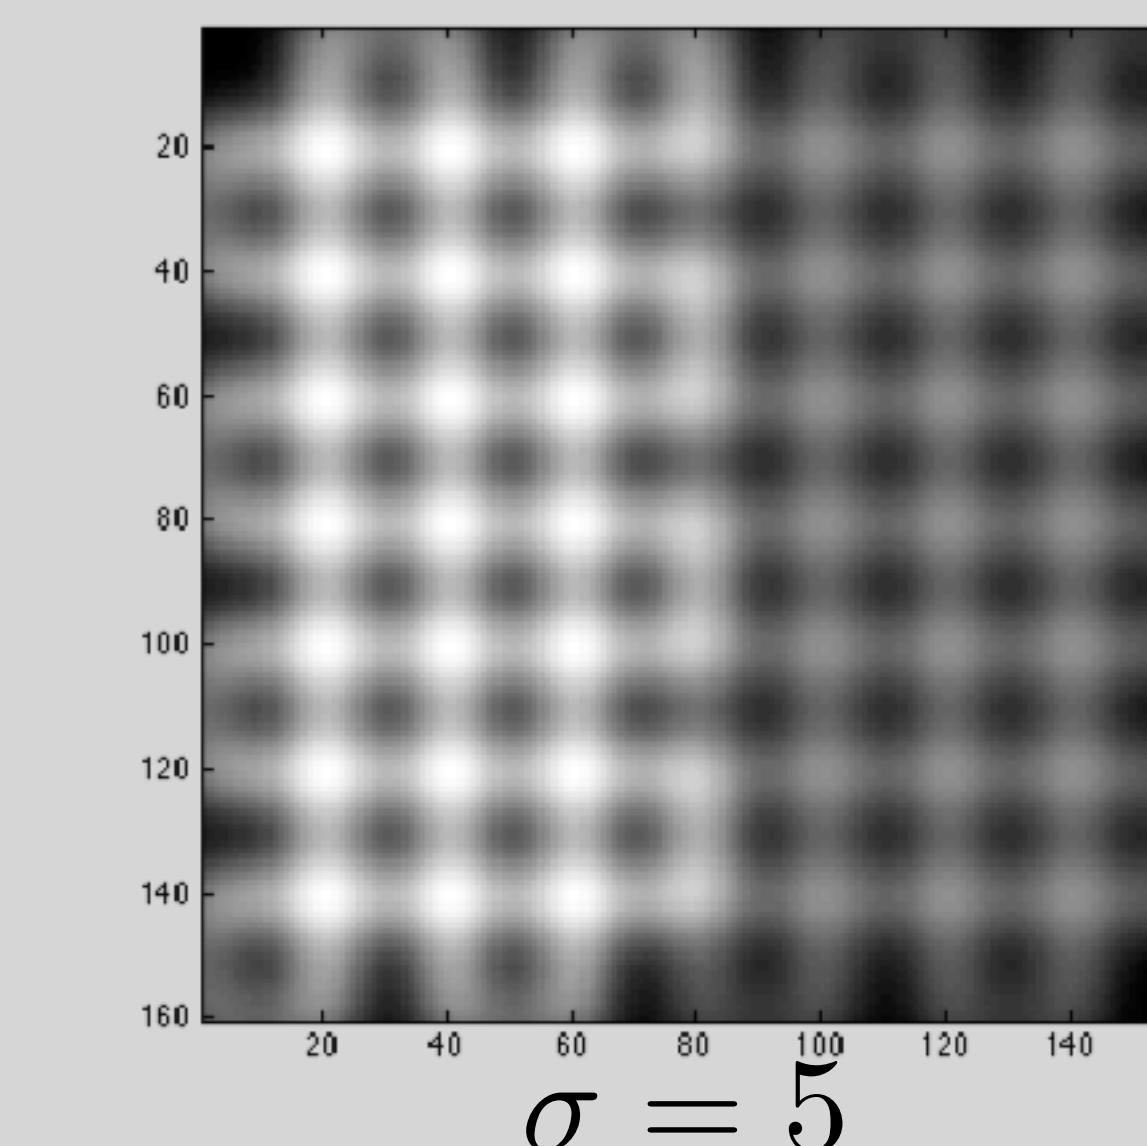
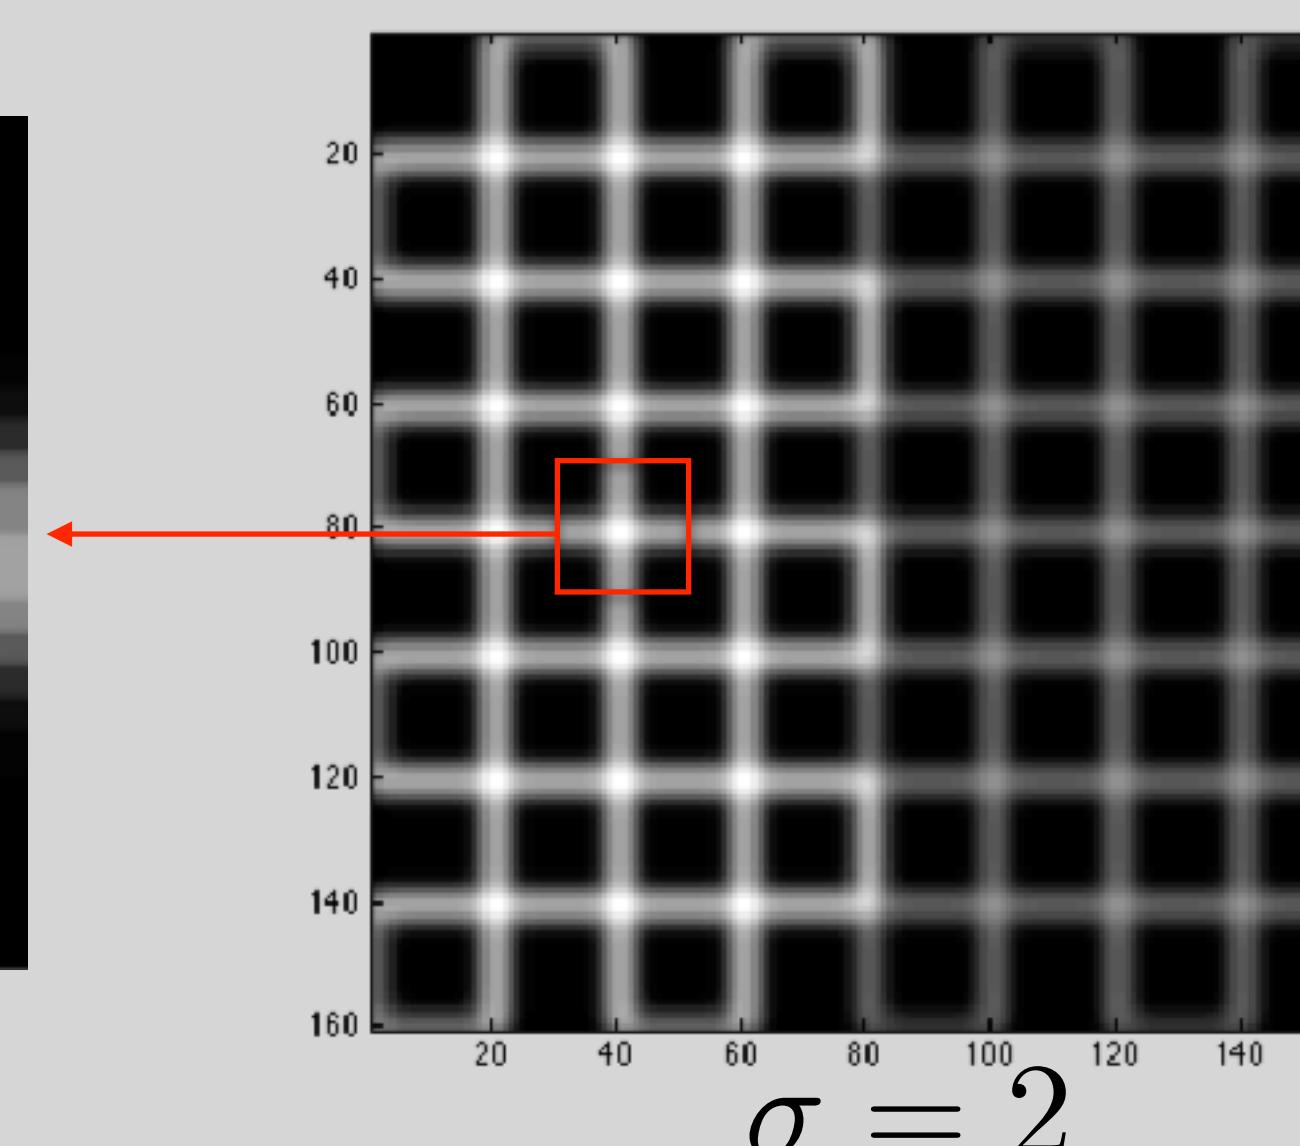
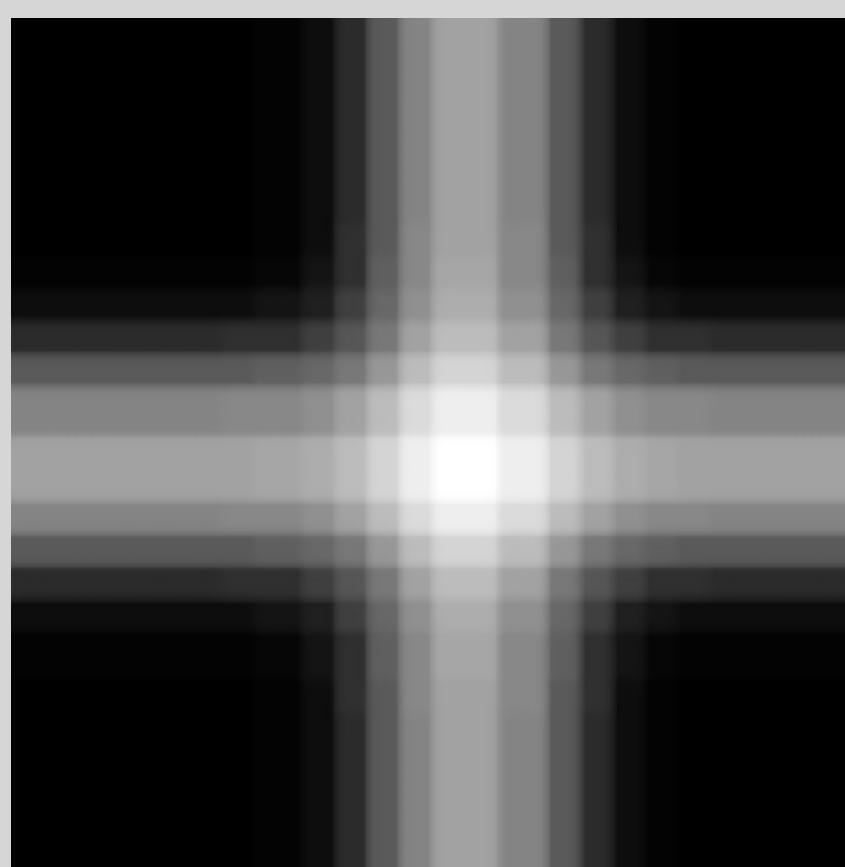
C.Harris and M.Stephens. [**“A Combined Corner and Edge Detector.”**](#)
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

Corner score example

image



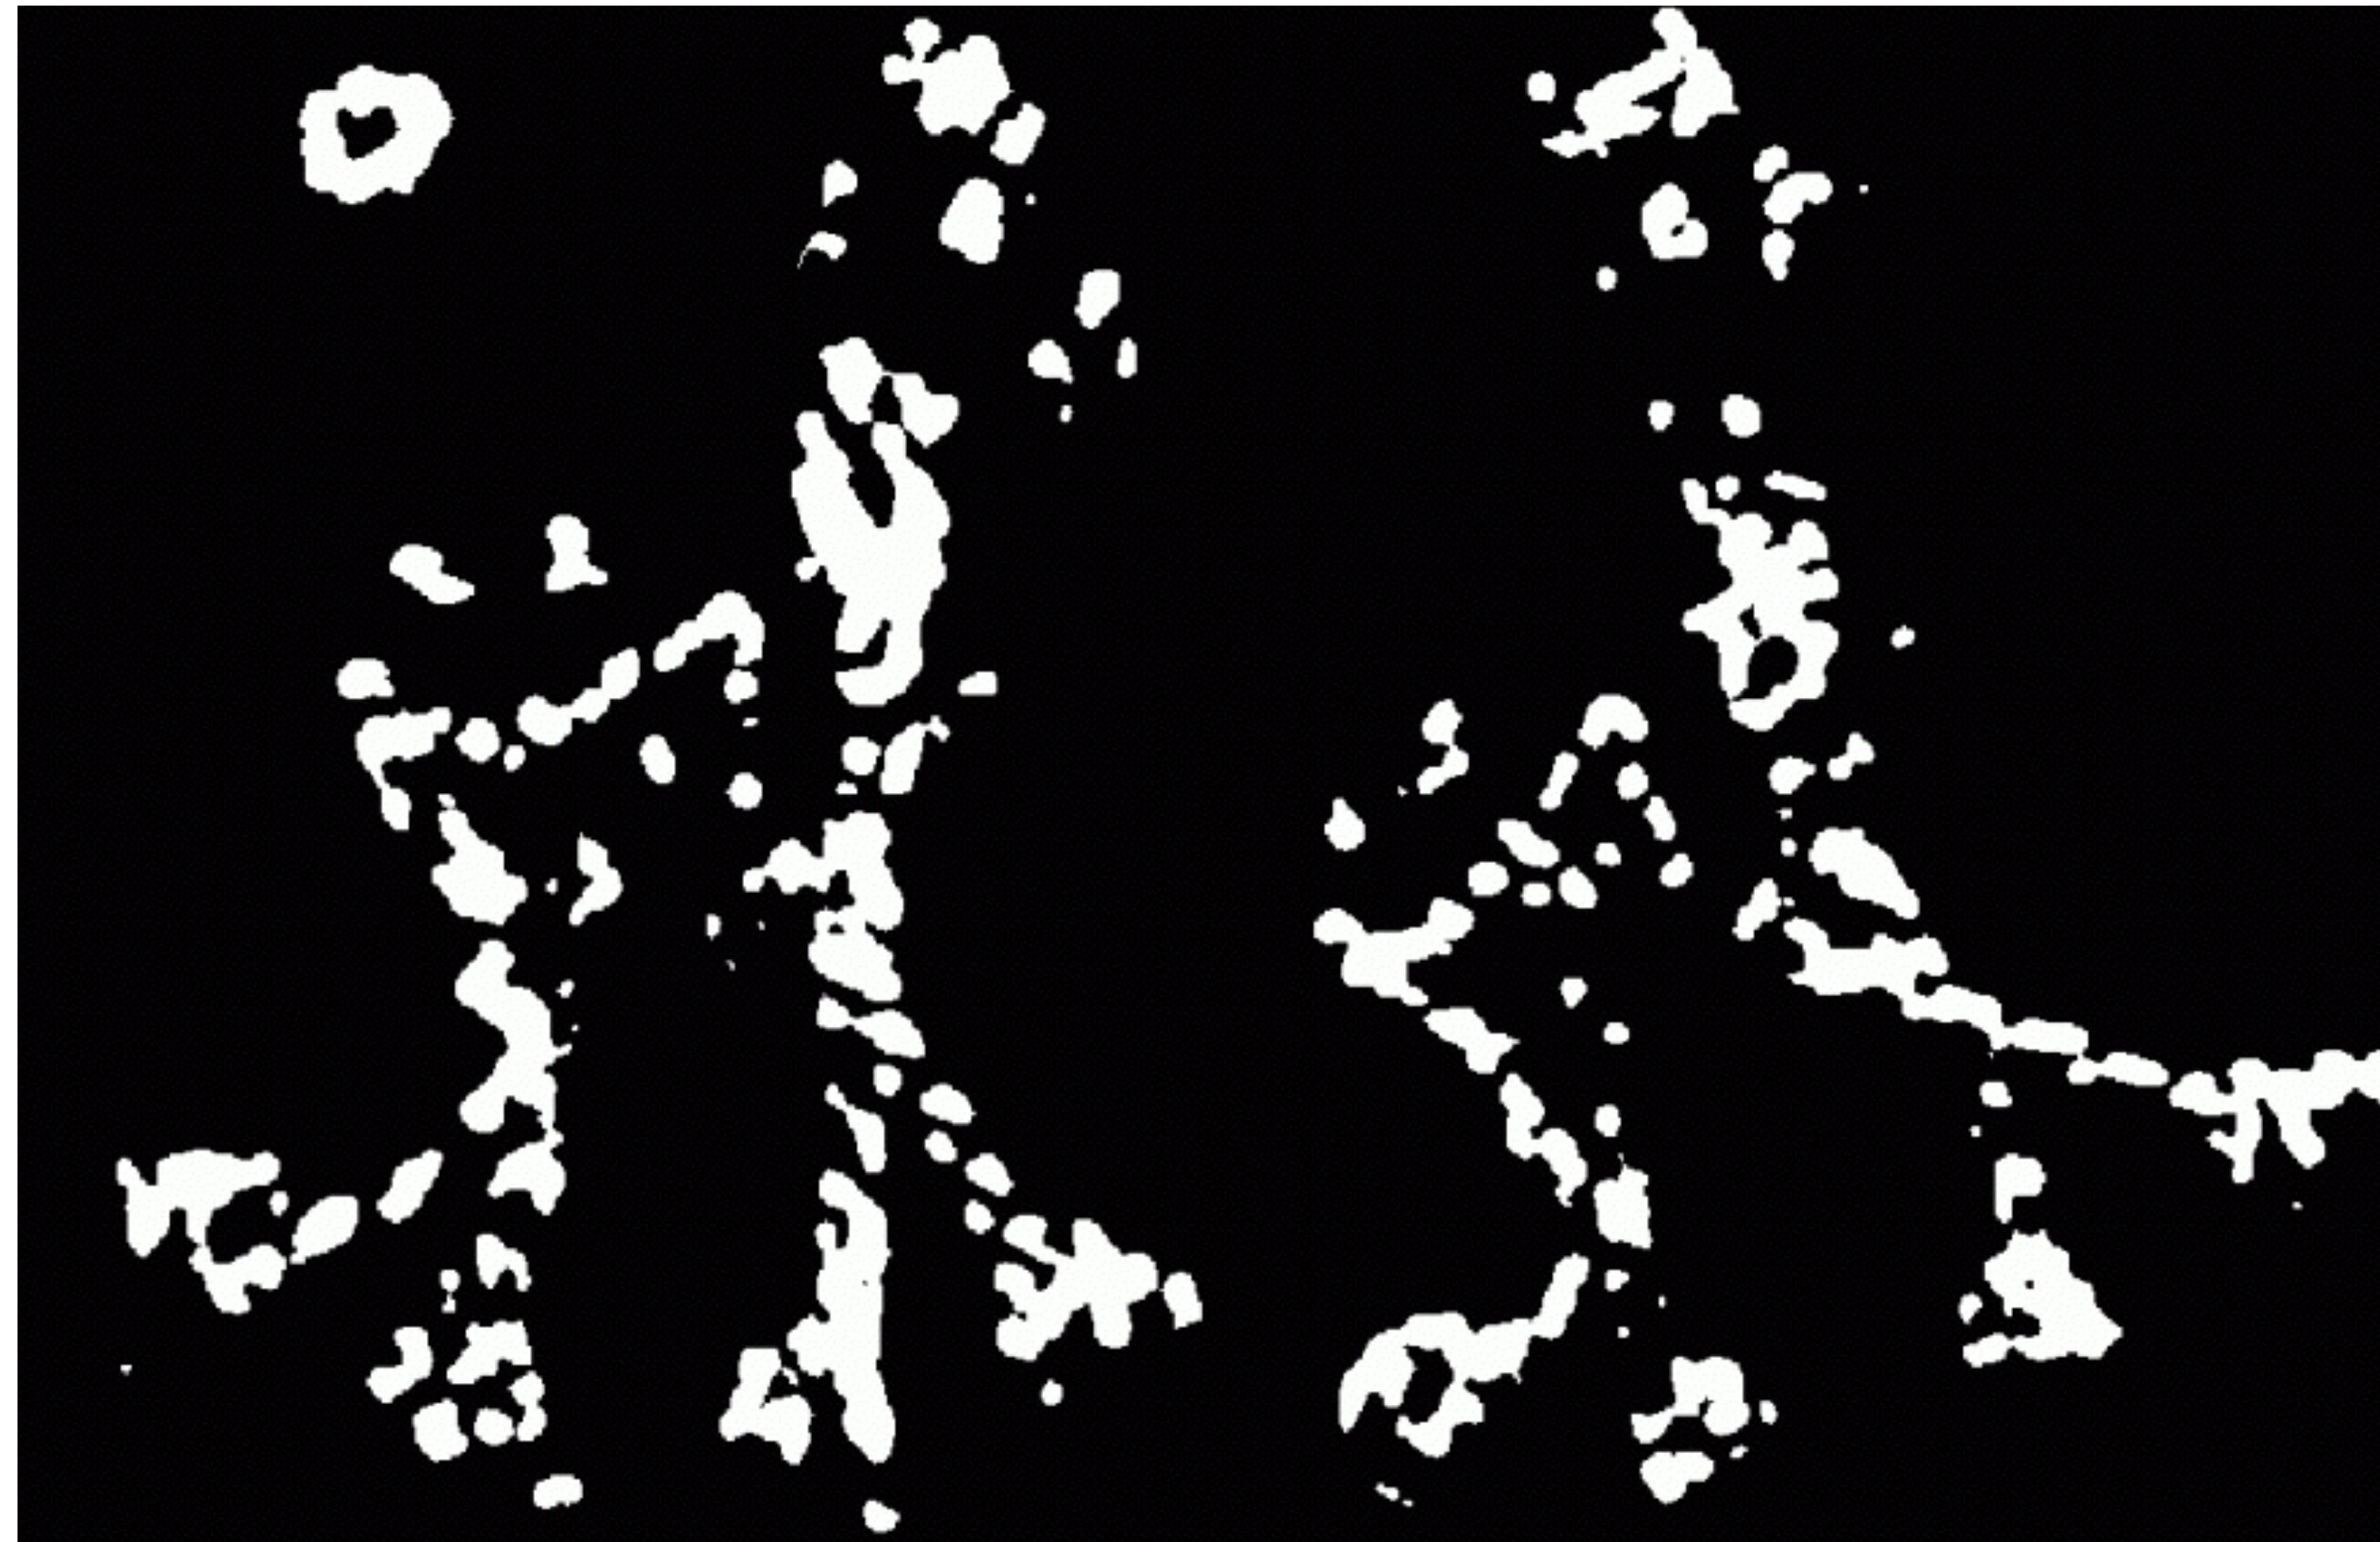
Box filter



Gaussian filter

Harris Detector: Steps

Find points with large corner response: $R > \text{threshold}$



Harris Detector: Steps

Take only the points of local maxima of R



Harris Detector: Steps



Further thoughts and readings...

Original corner detector paper

- C.Harris and M.Stephens, [**“A Combined Corner and Edge Detector.”**](#) Proceedings of the 4th Alvey Vision Conference, 1988

Other corner functions

- Can you think of other $f(\lambda_1, \lambda_2)$ that work for finding corners?

Invariance and covariance

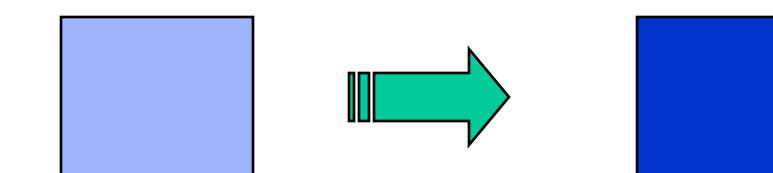
Invariance: transformations *do not change* the corner locations

Covariance or Equivariance: transformations change corner locations *in a predictable way*

We want corners to be *invariant* to photometric transformations and *covariant* to geometric transformations

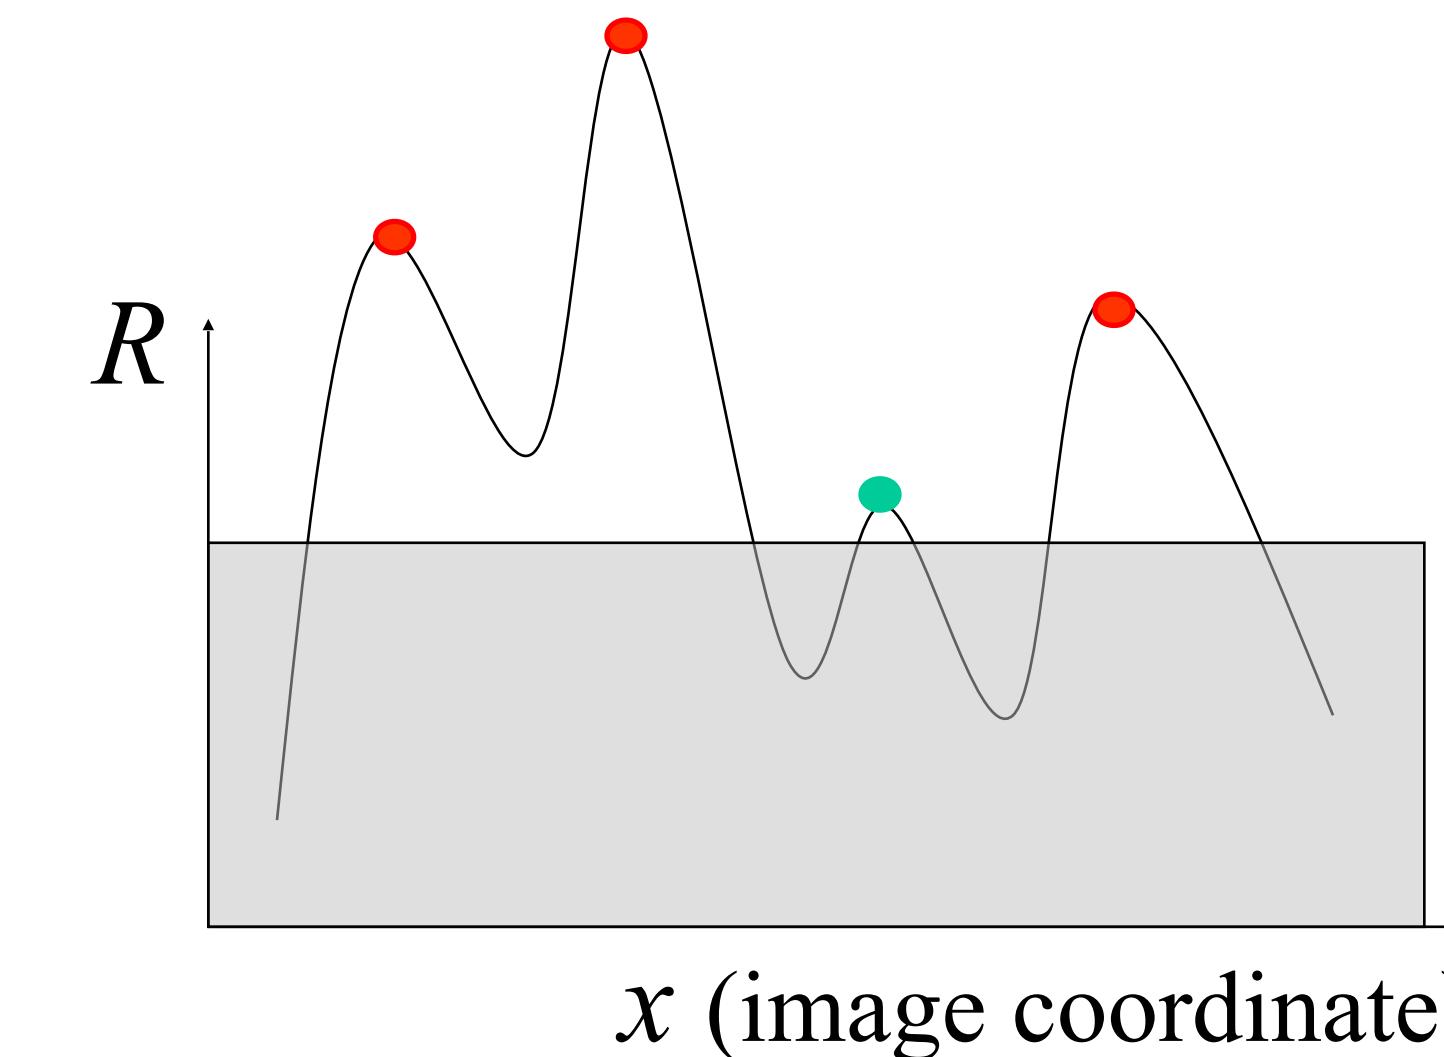
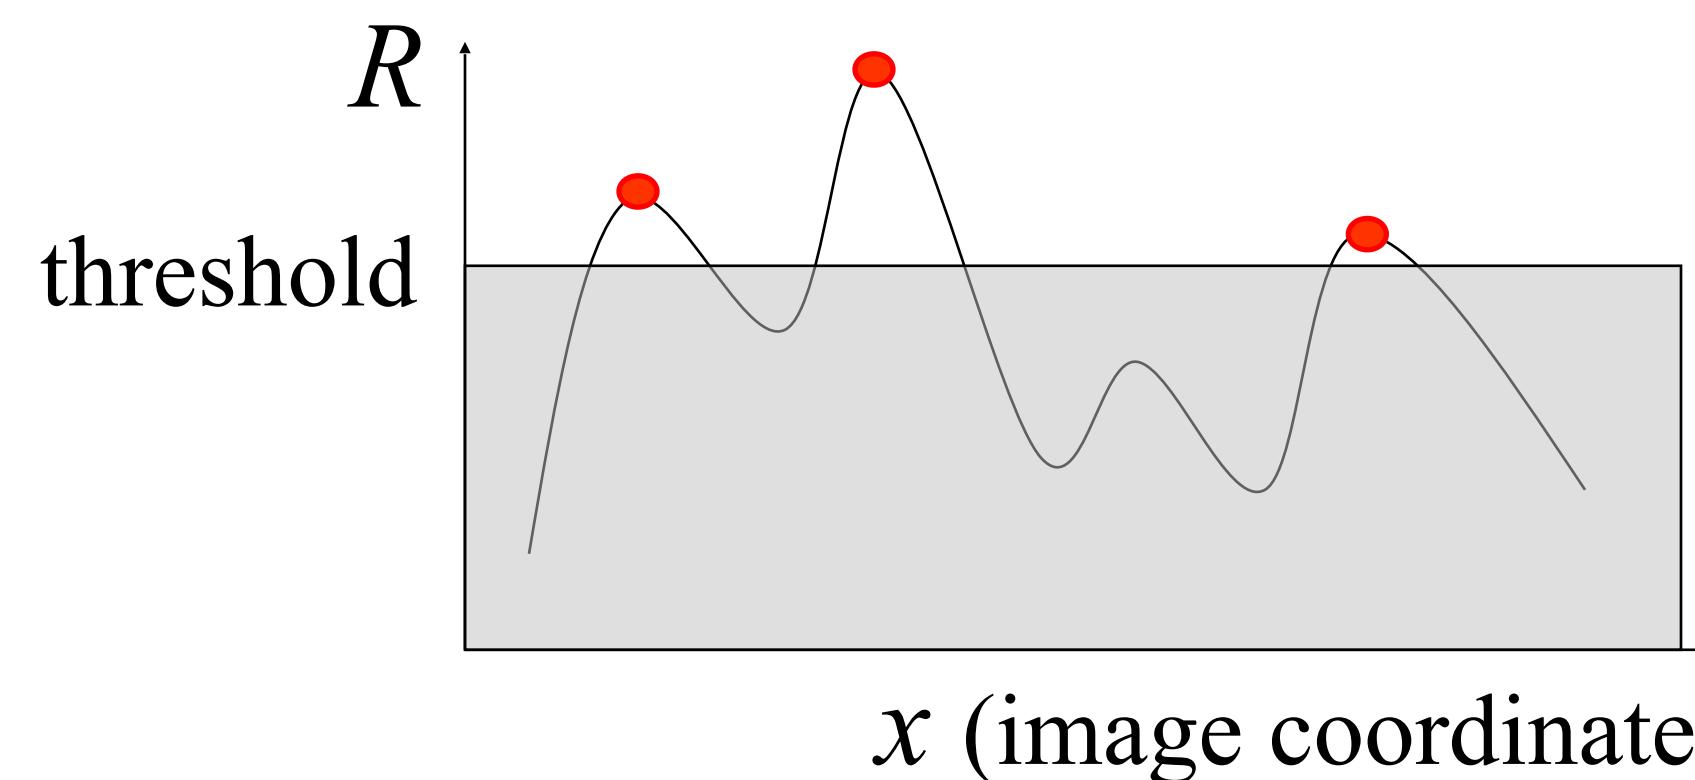


Affine intensity change



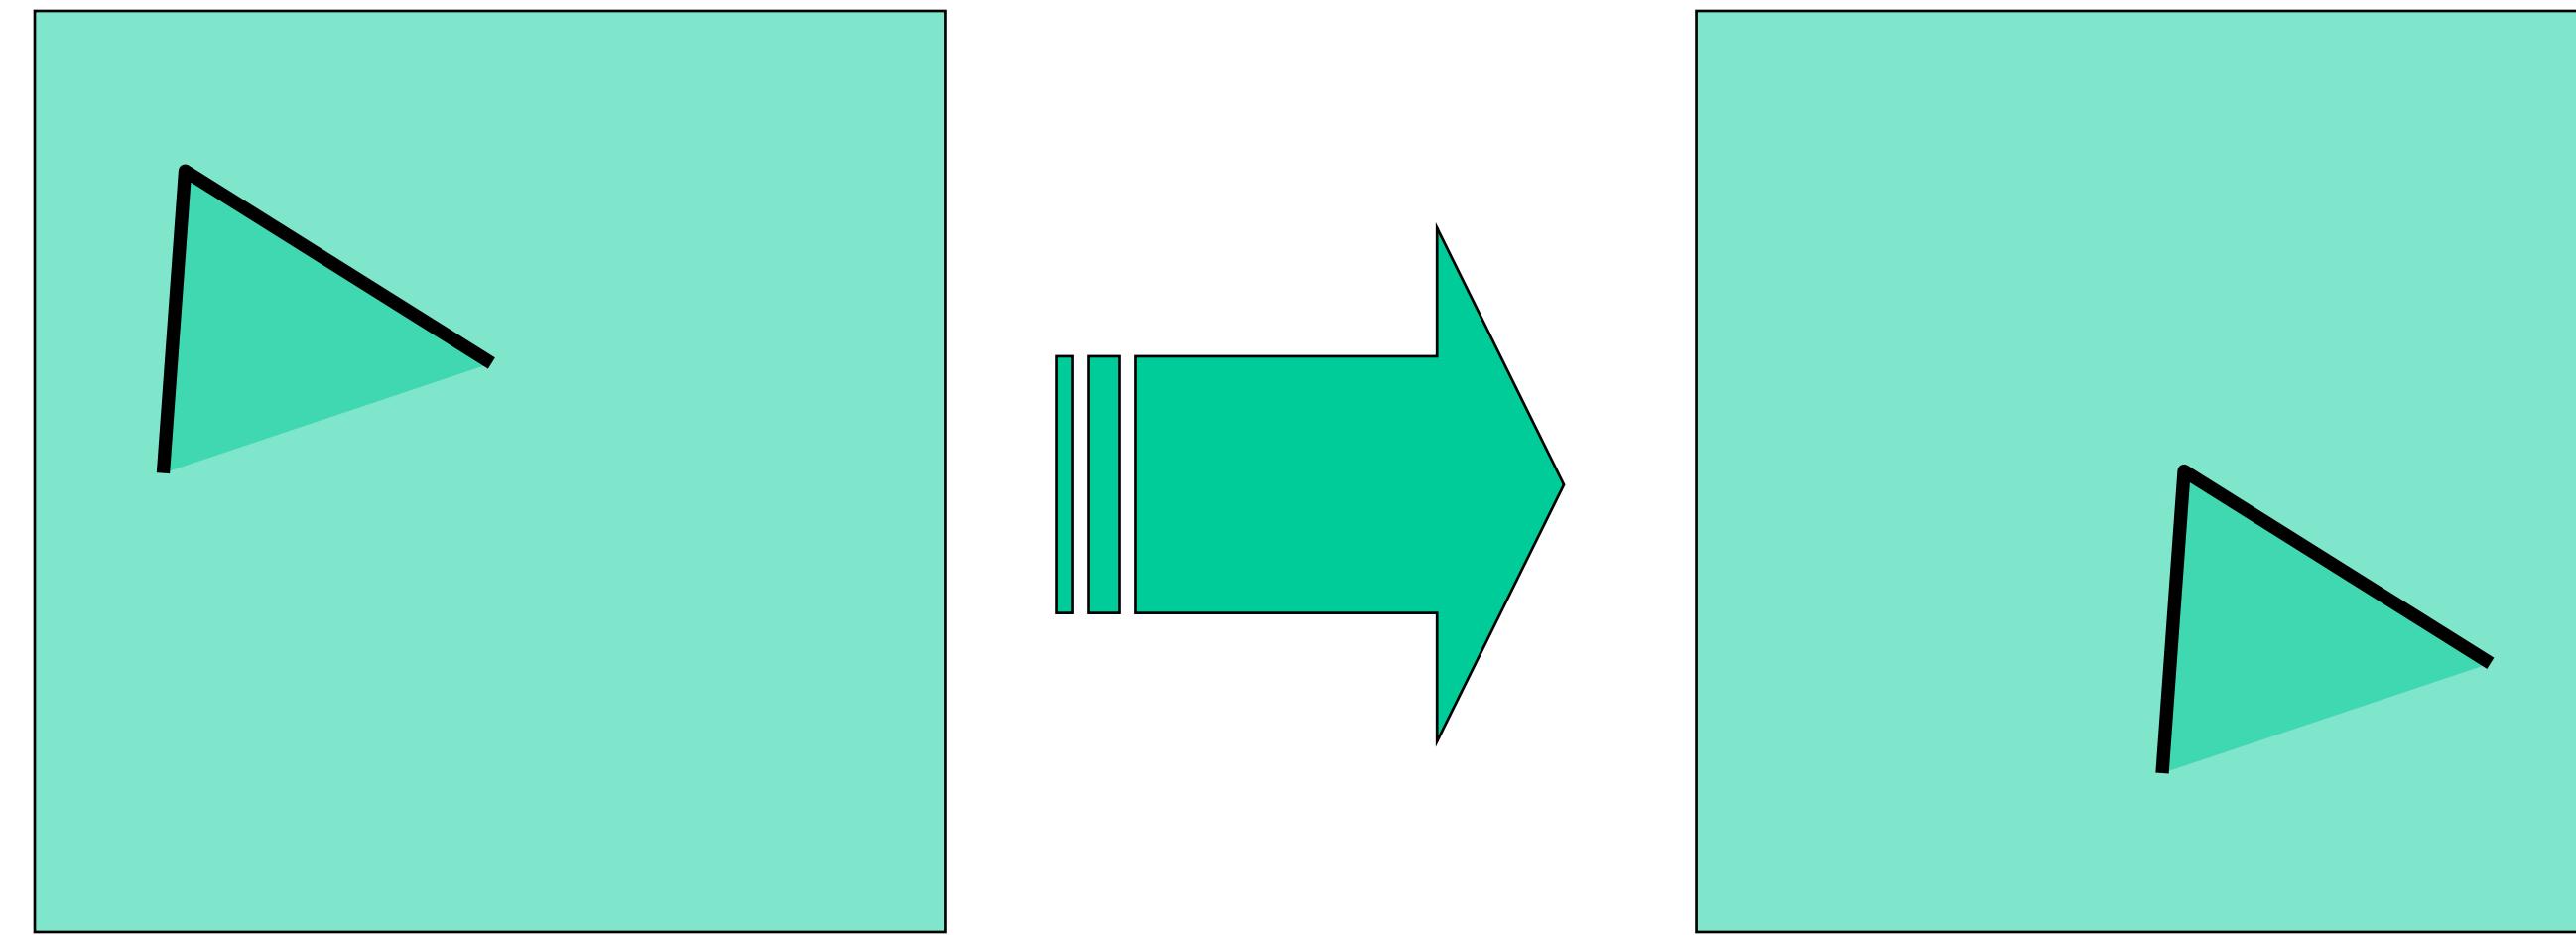
$$I \rightarrow aI + b$$

- Only derivatives are used => invariance to intensity shift $I \rightarrow I + b$
- Intensity scaling: $I \rightarrow aI$



Corner location is partially invariant to affine intensity change

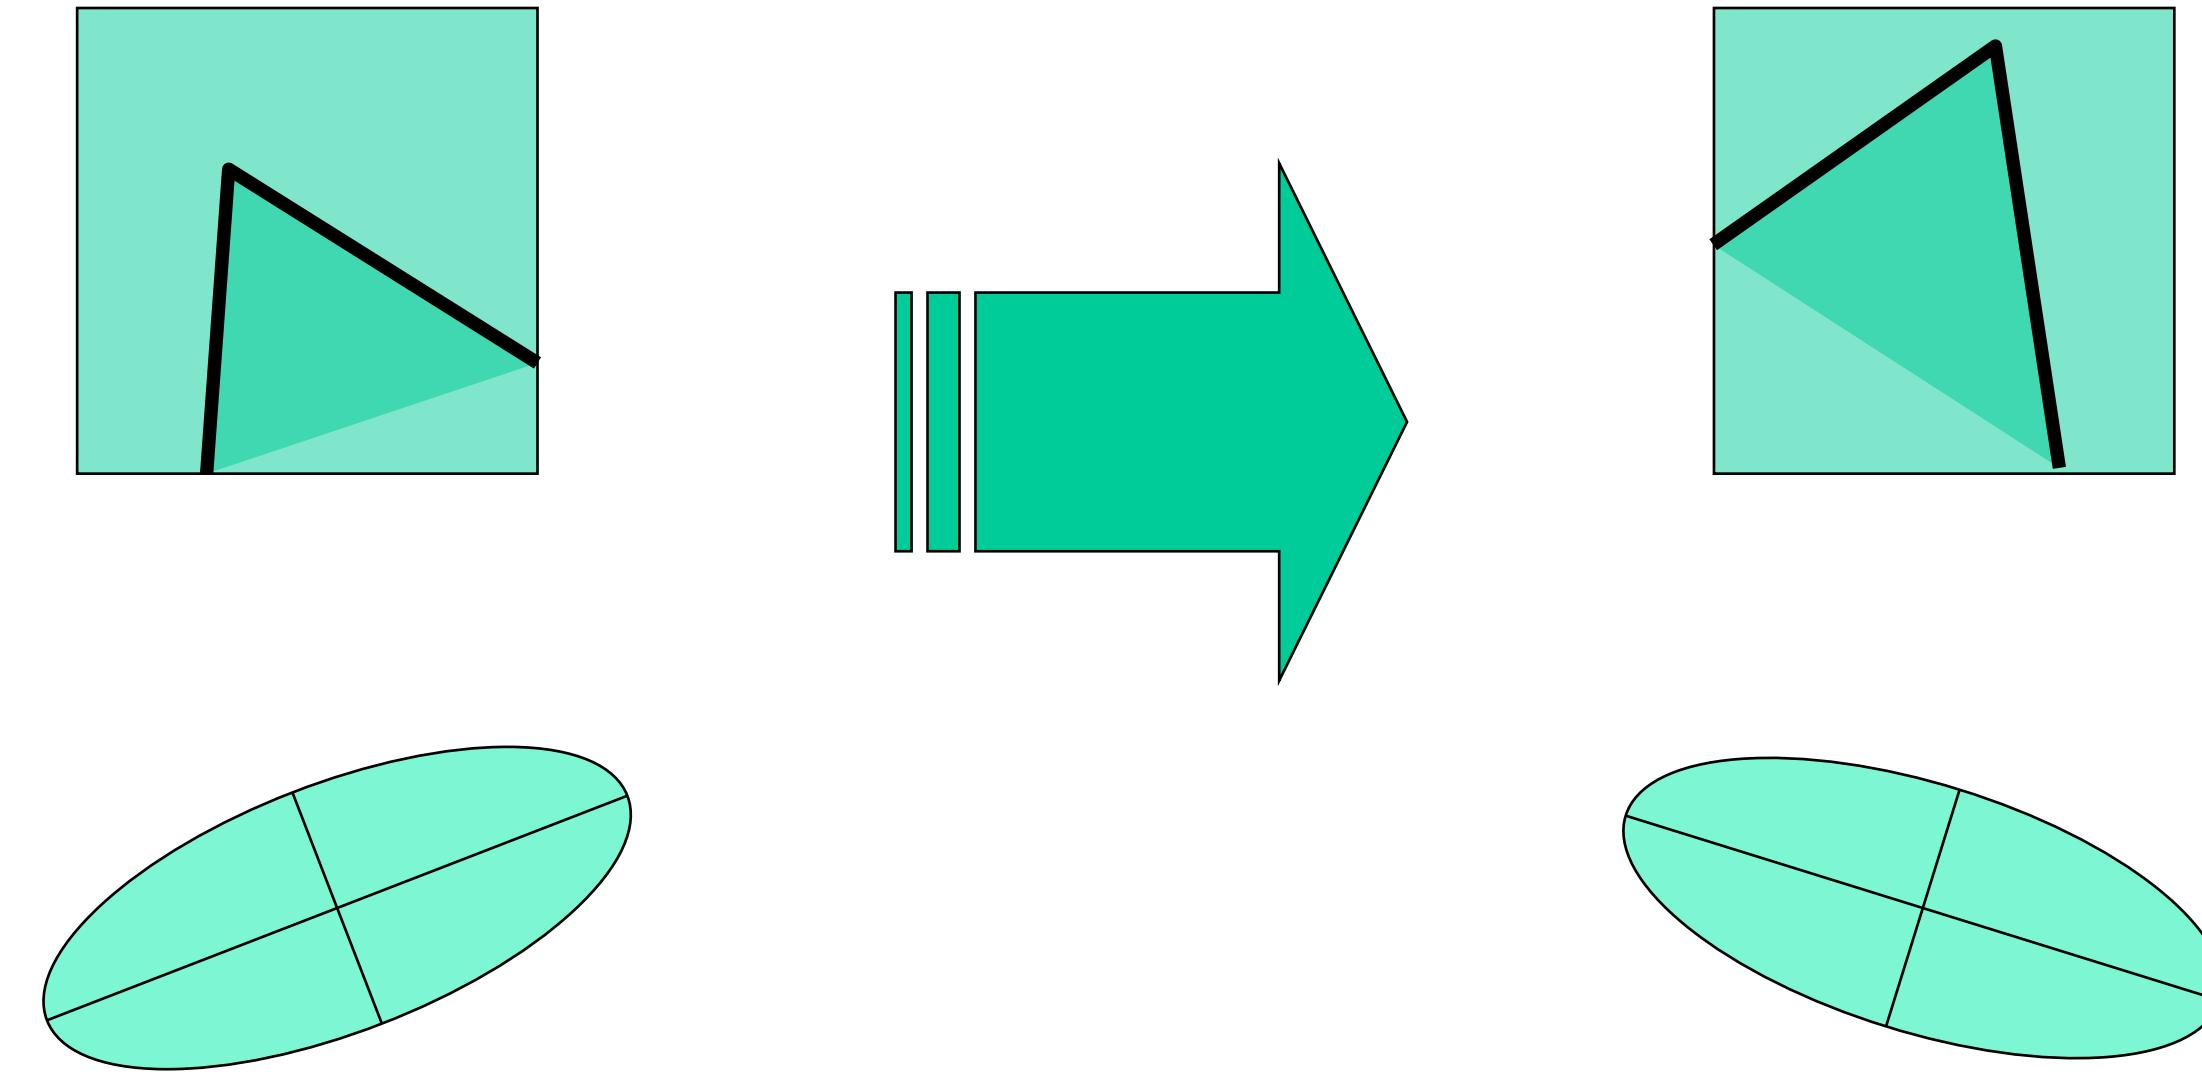
Translation



- Derivatives and window function are shift-invariant

Corner location is covariant w.r.t. translation

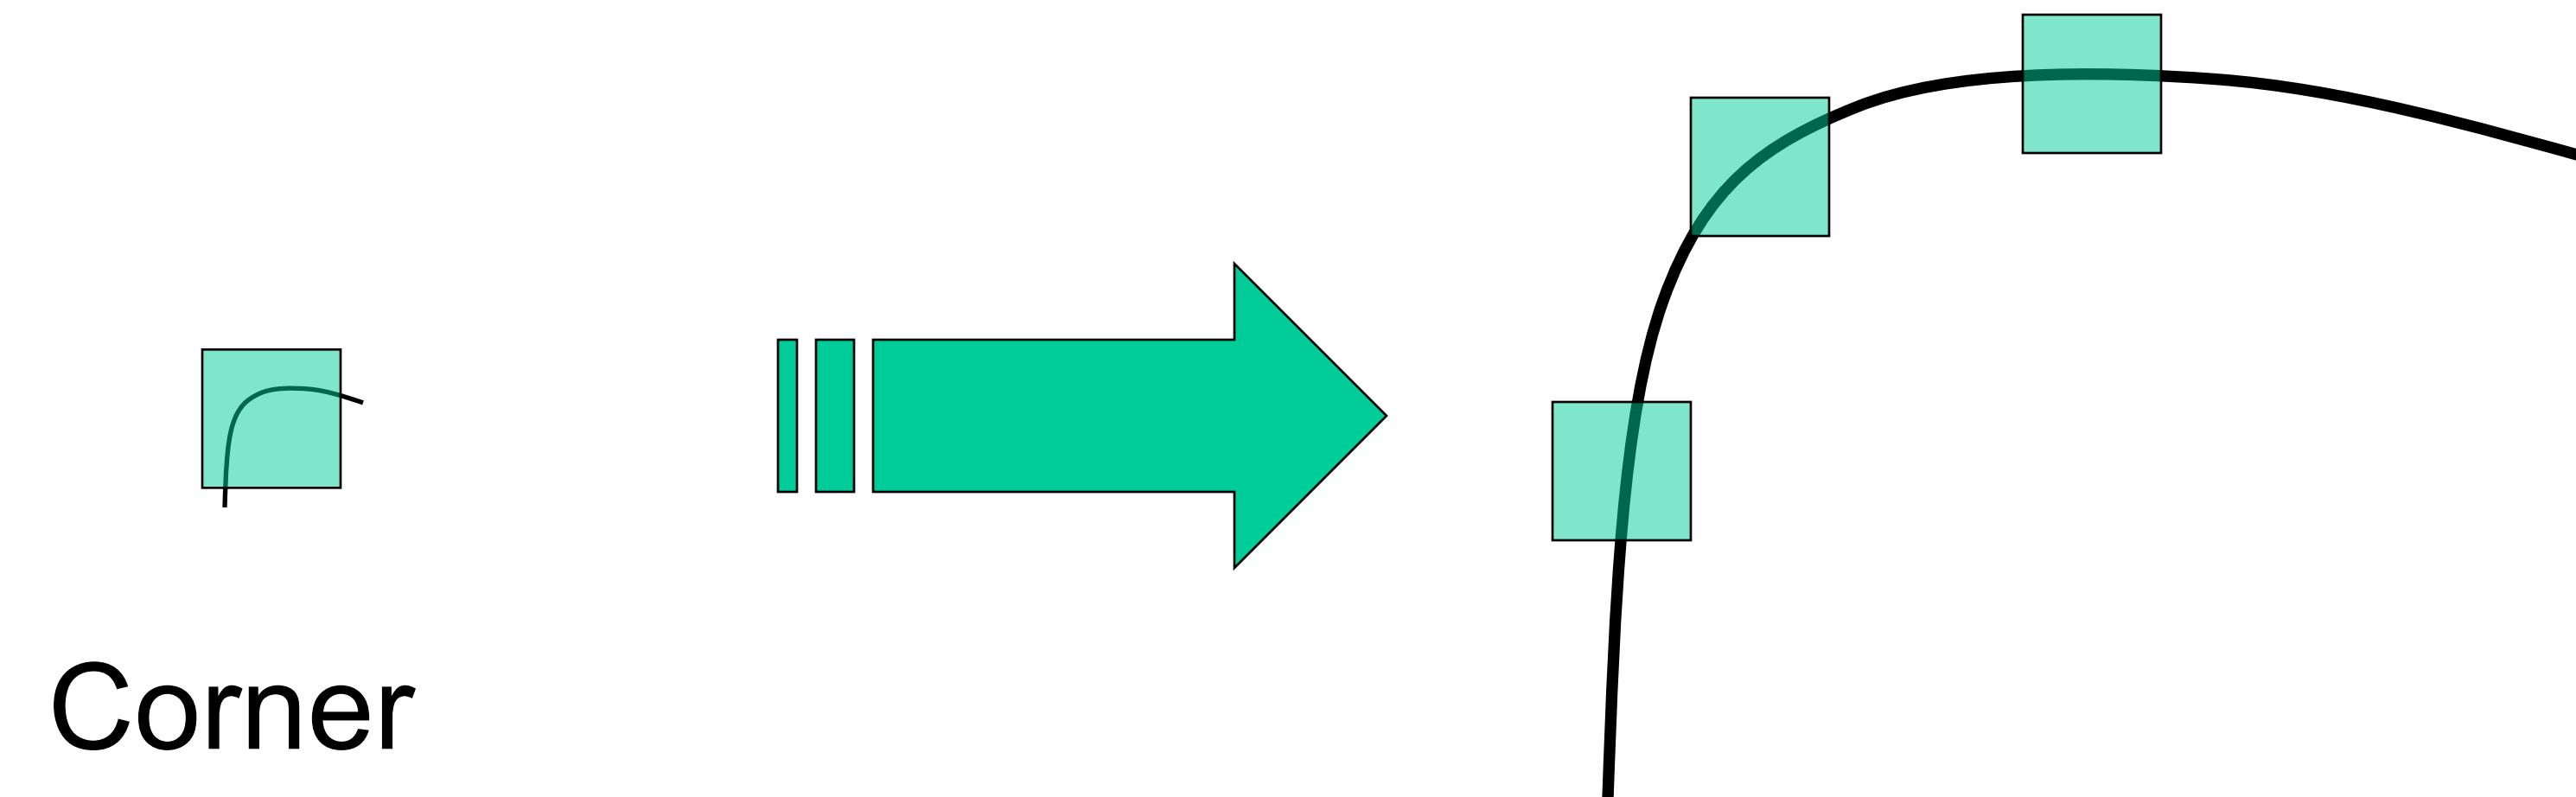
Rotation



Second moment ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner location is covariant w.r.t. rotation

Scaling



All points will be
classified as
edges

Corner detection is sensitive to the image scale!

Feature detection with scale selection

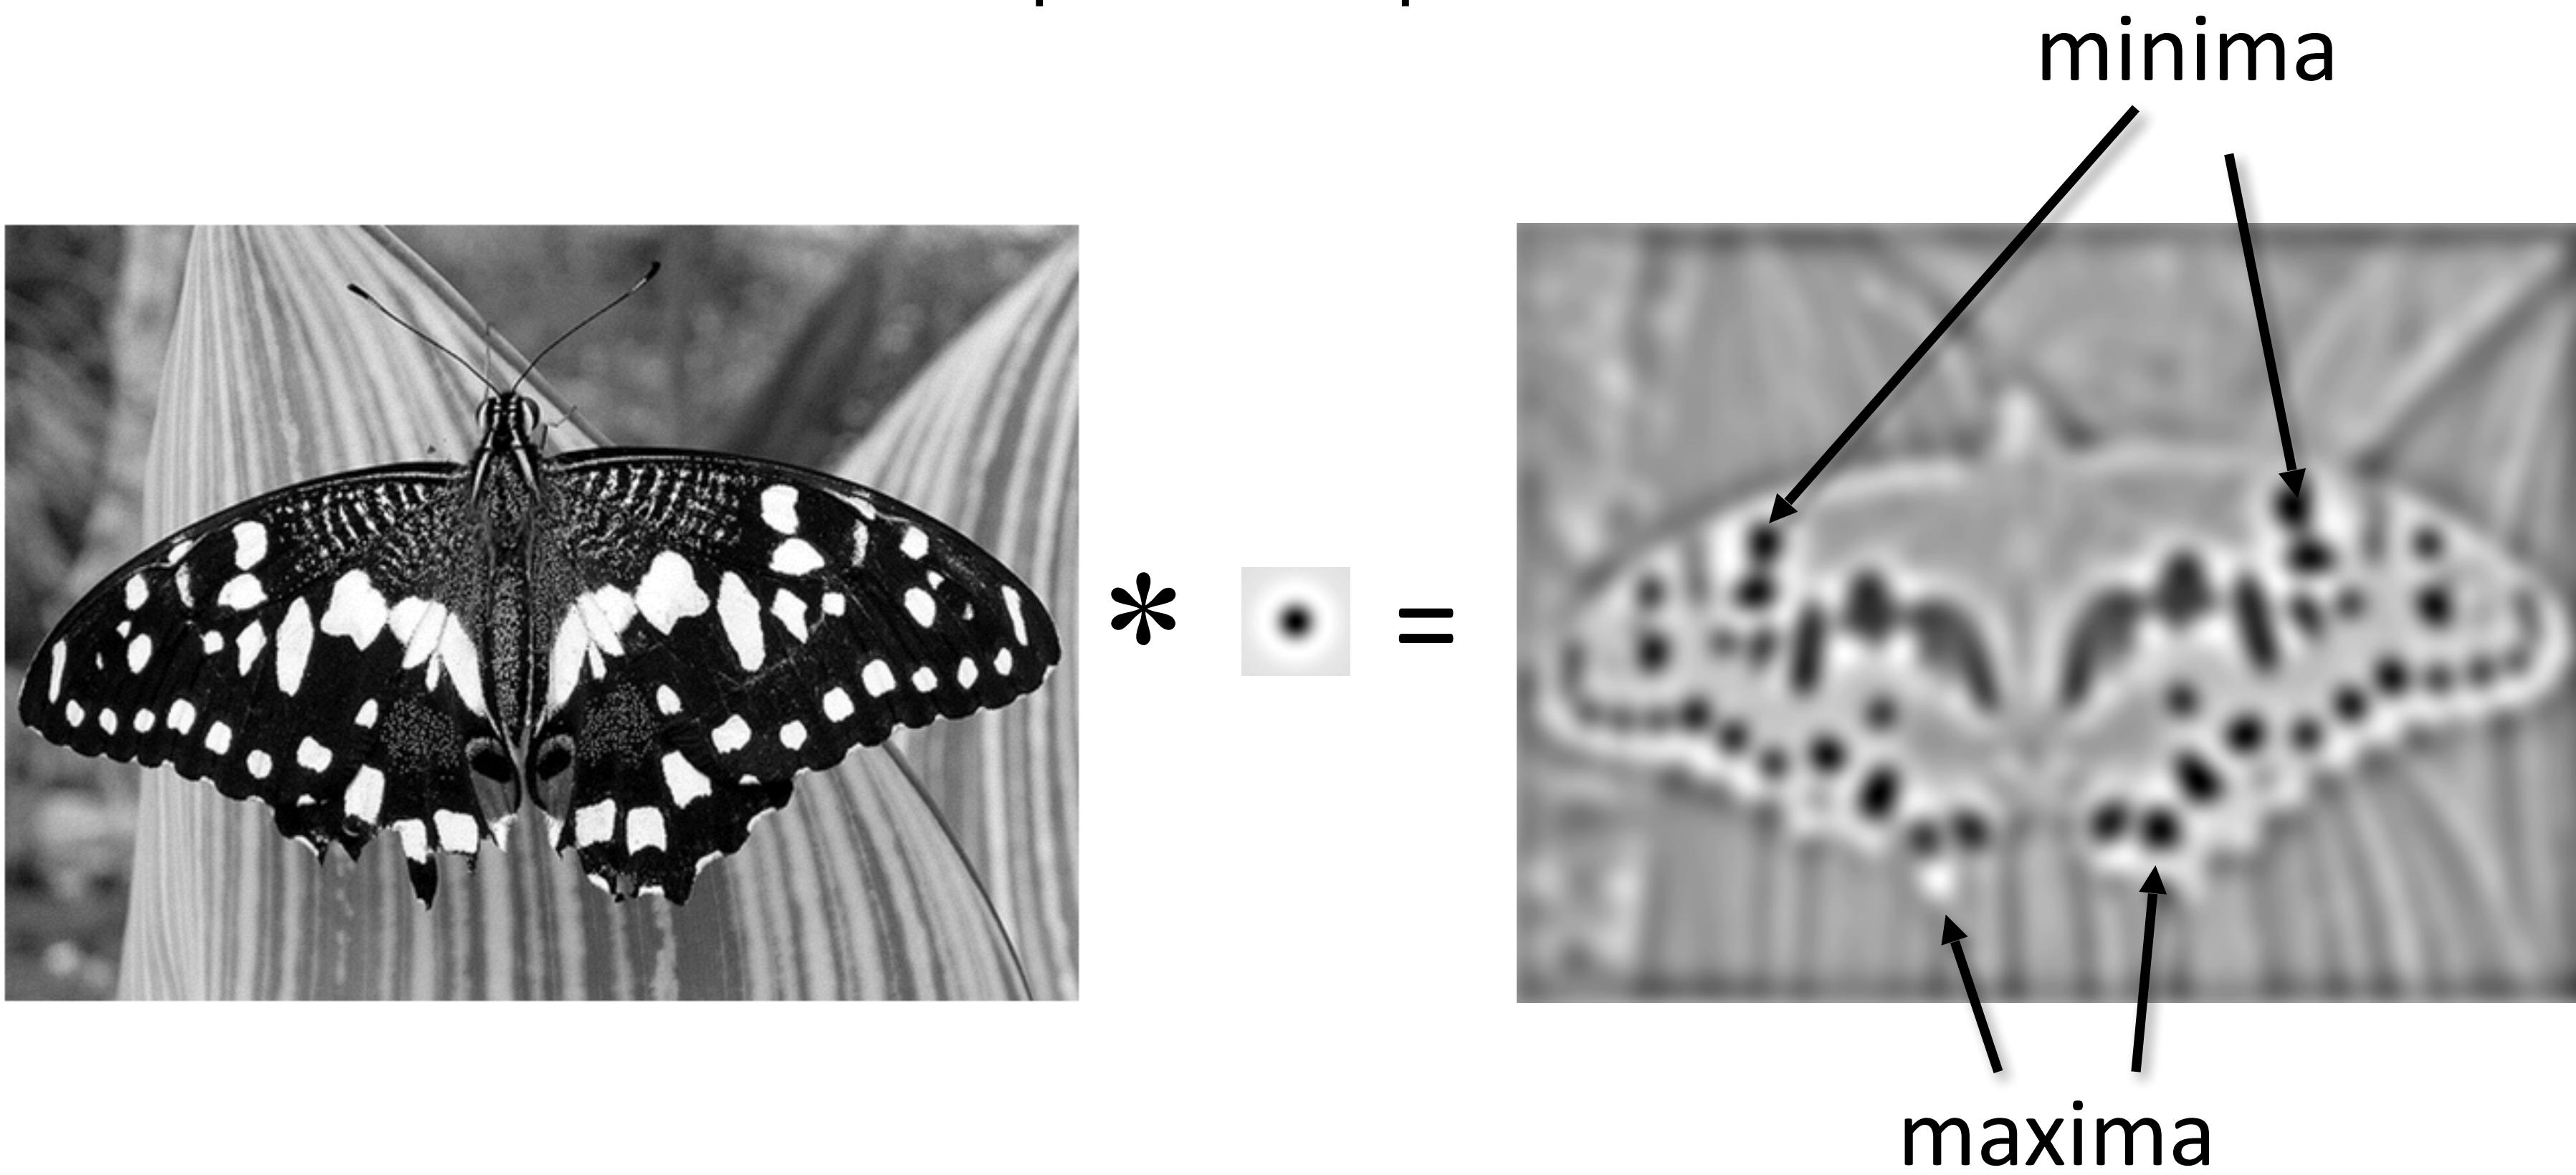
We want to extract features with characteristic scale that matches the image transformation such as **scaling** and **translation**



Matching regions across scales

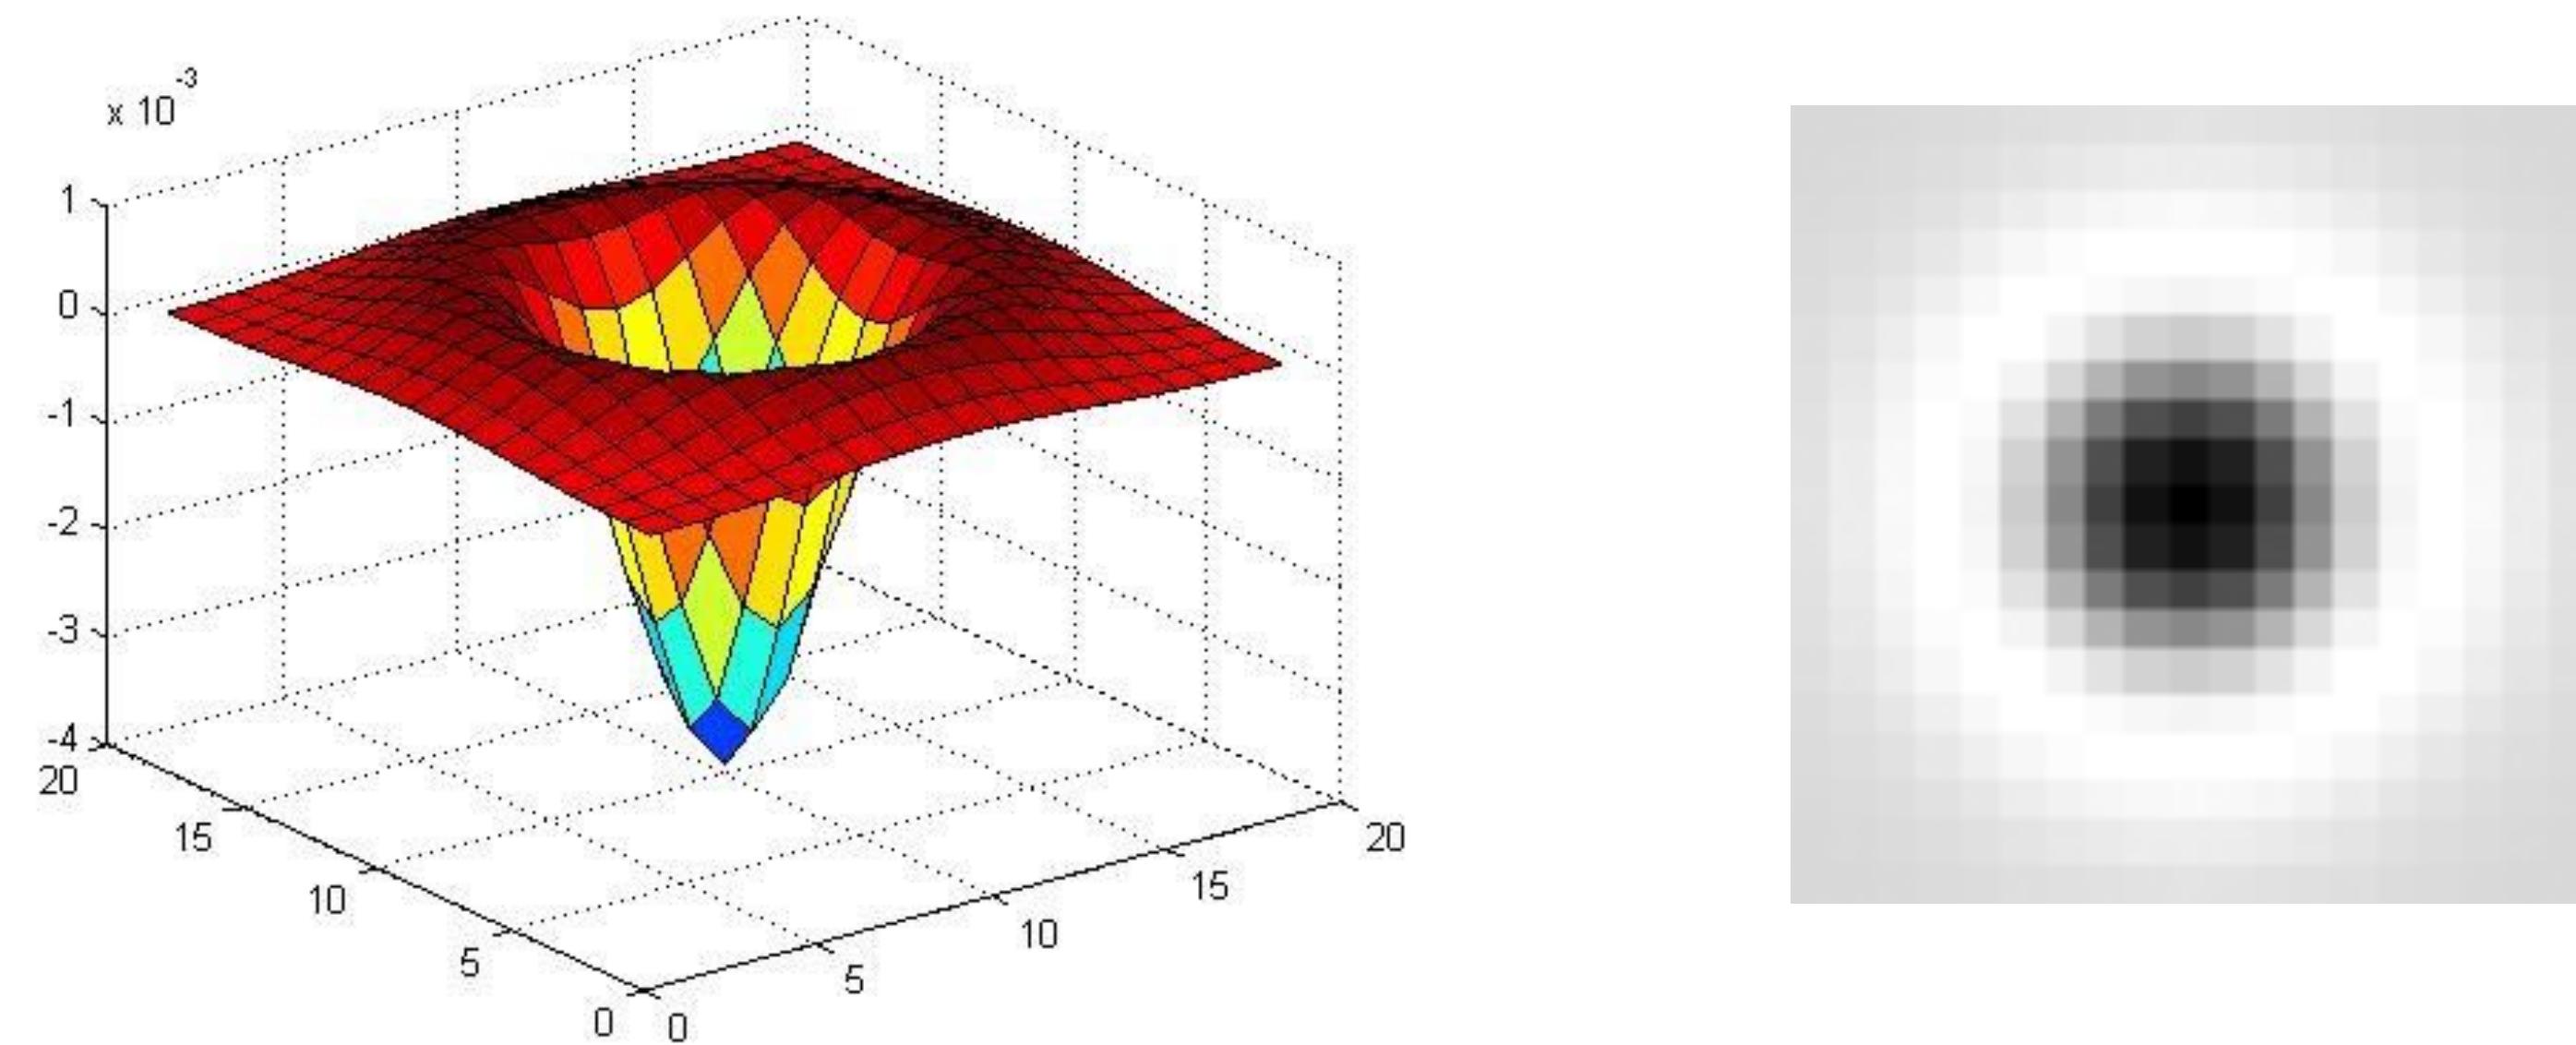
Blob detection: basic idea

Find maxima and minima of blob filter response in space and scale



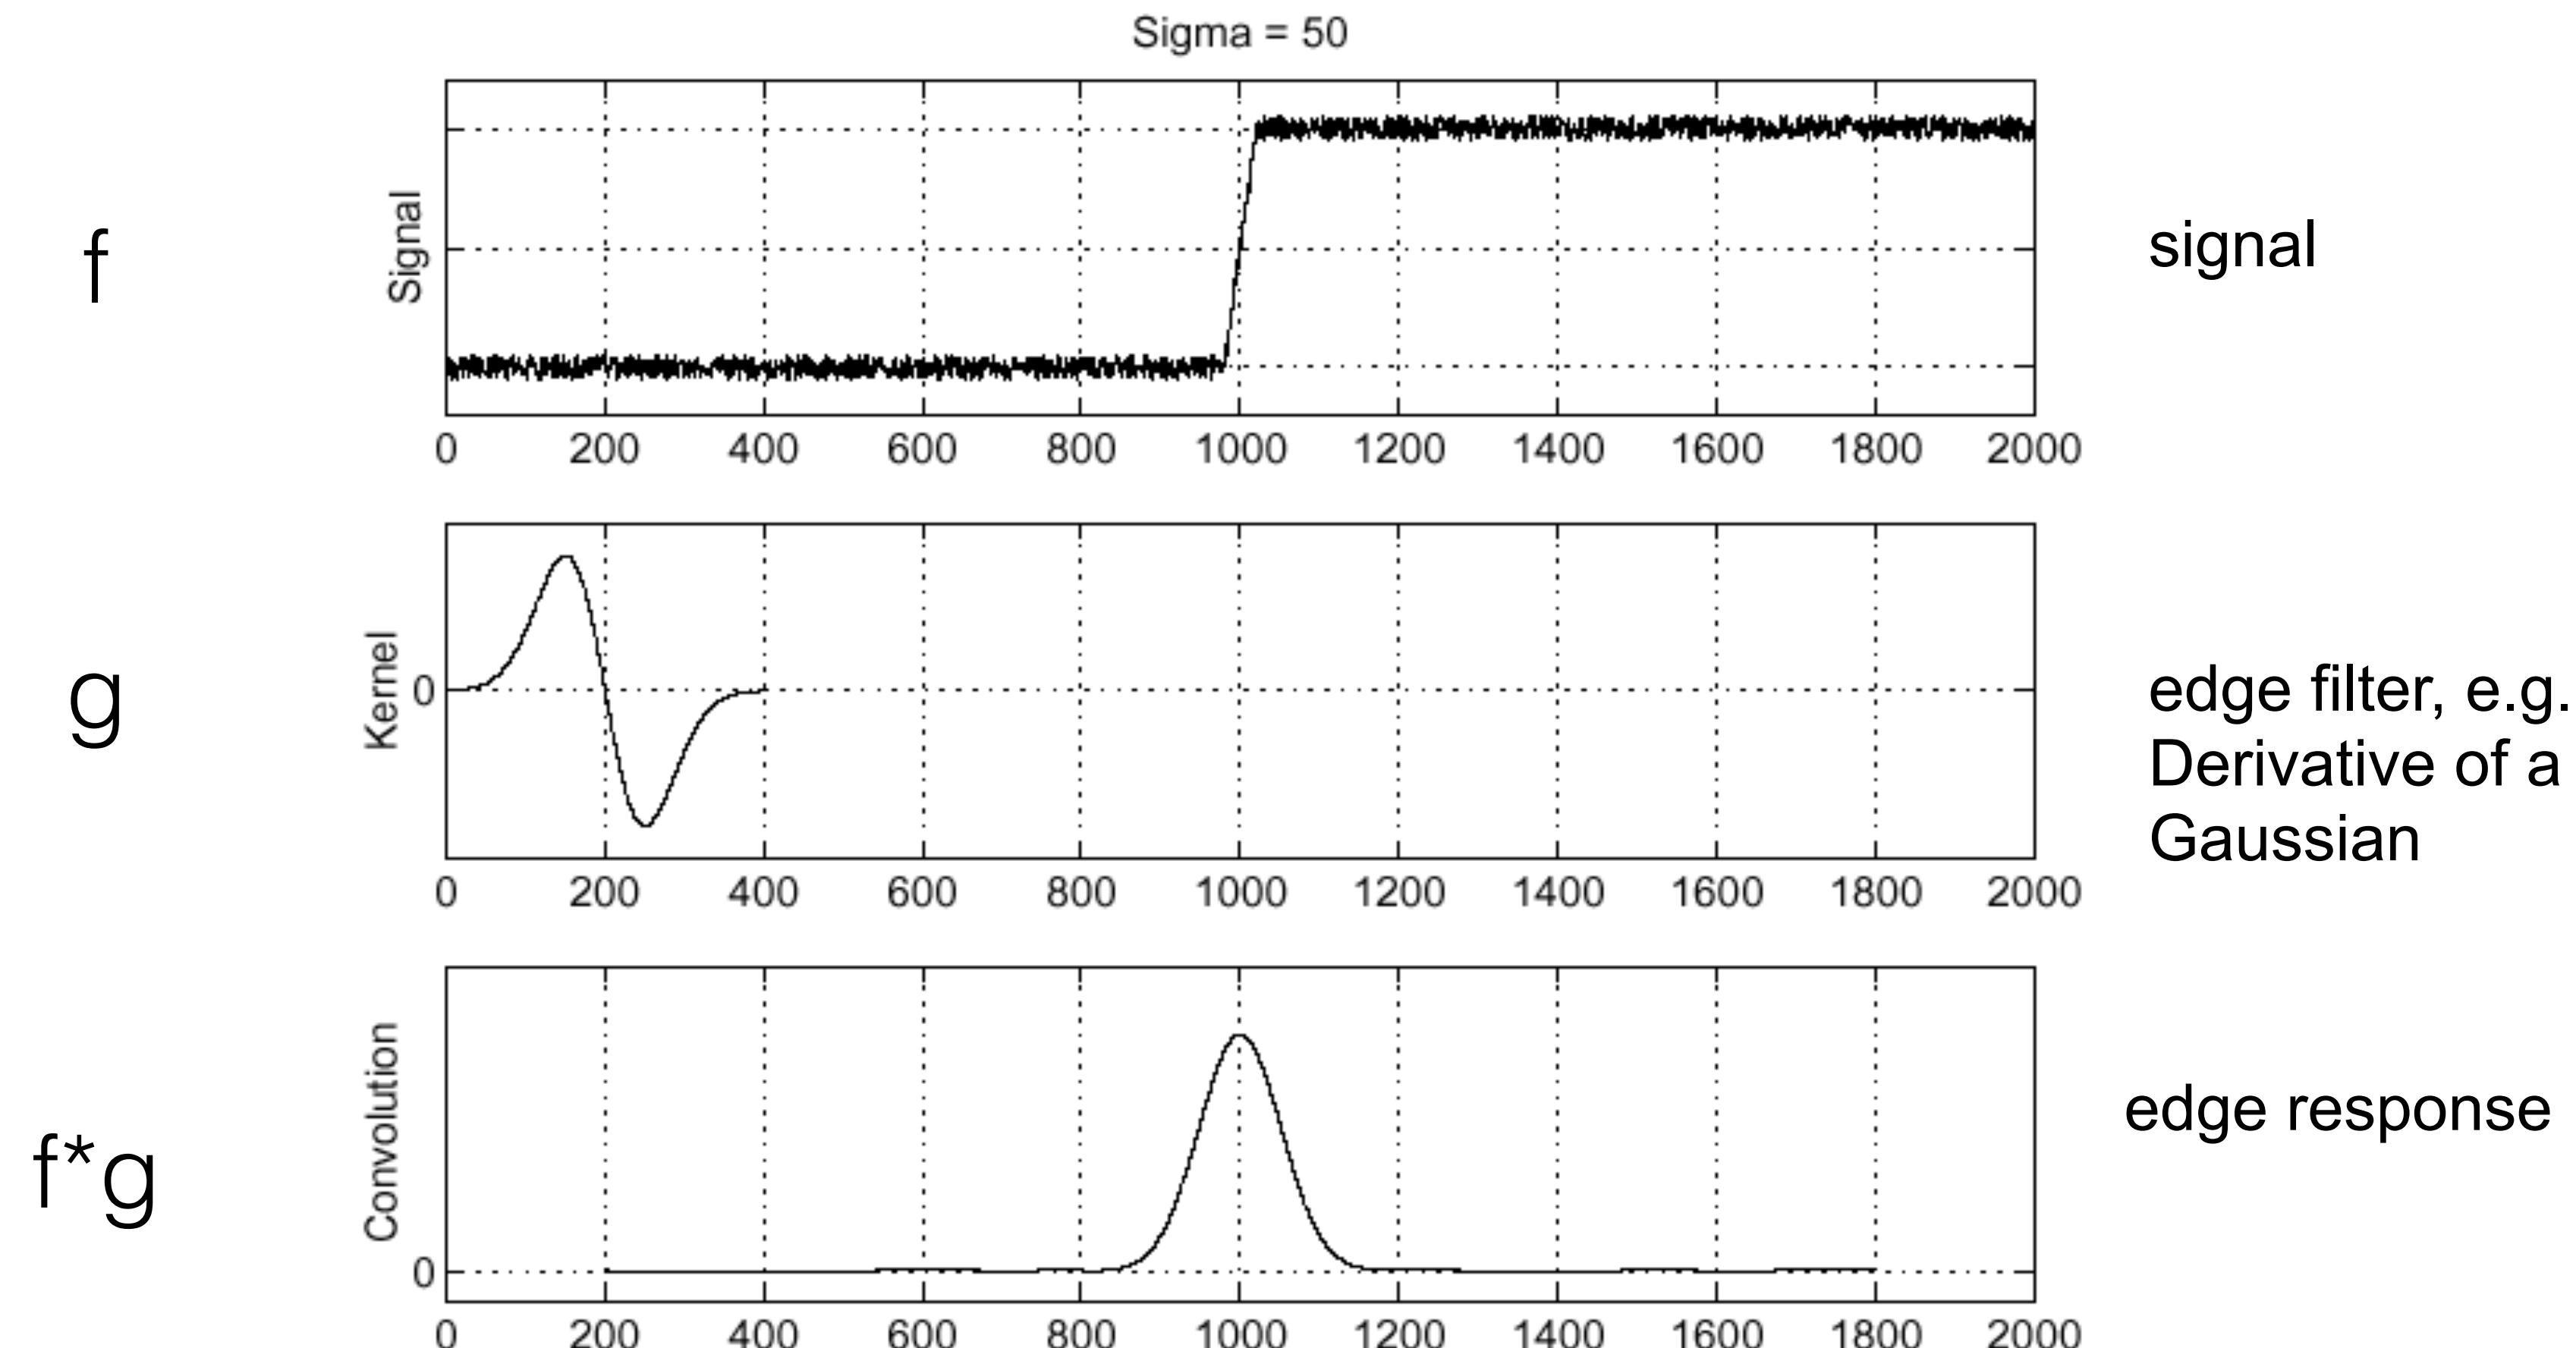
Blob filter

Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D

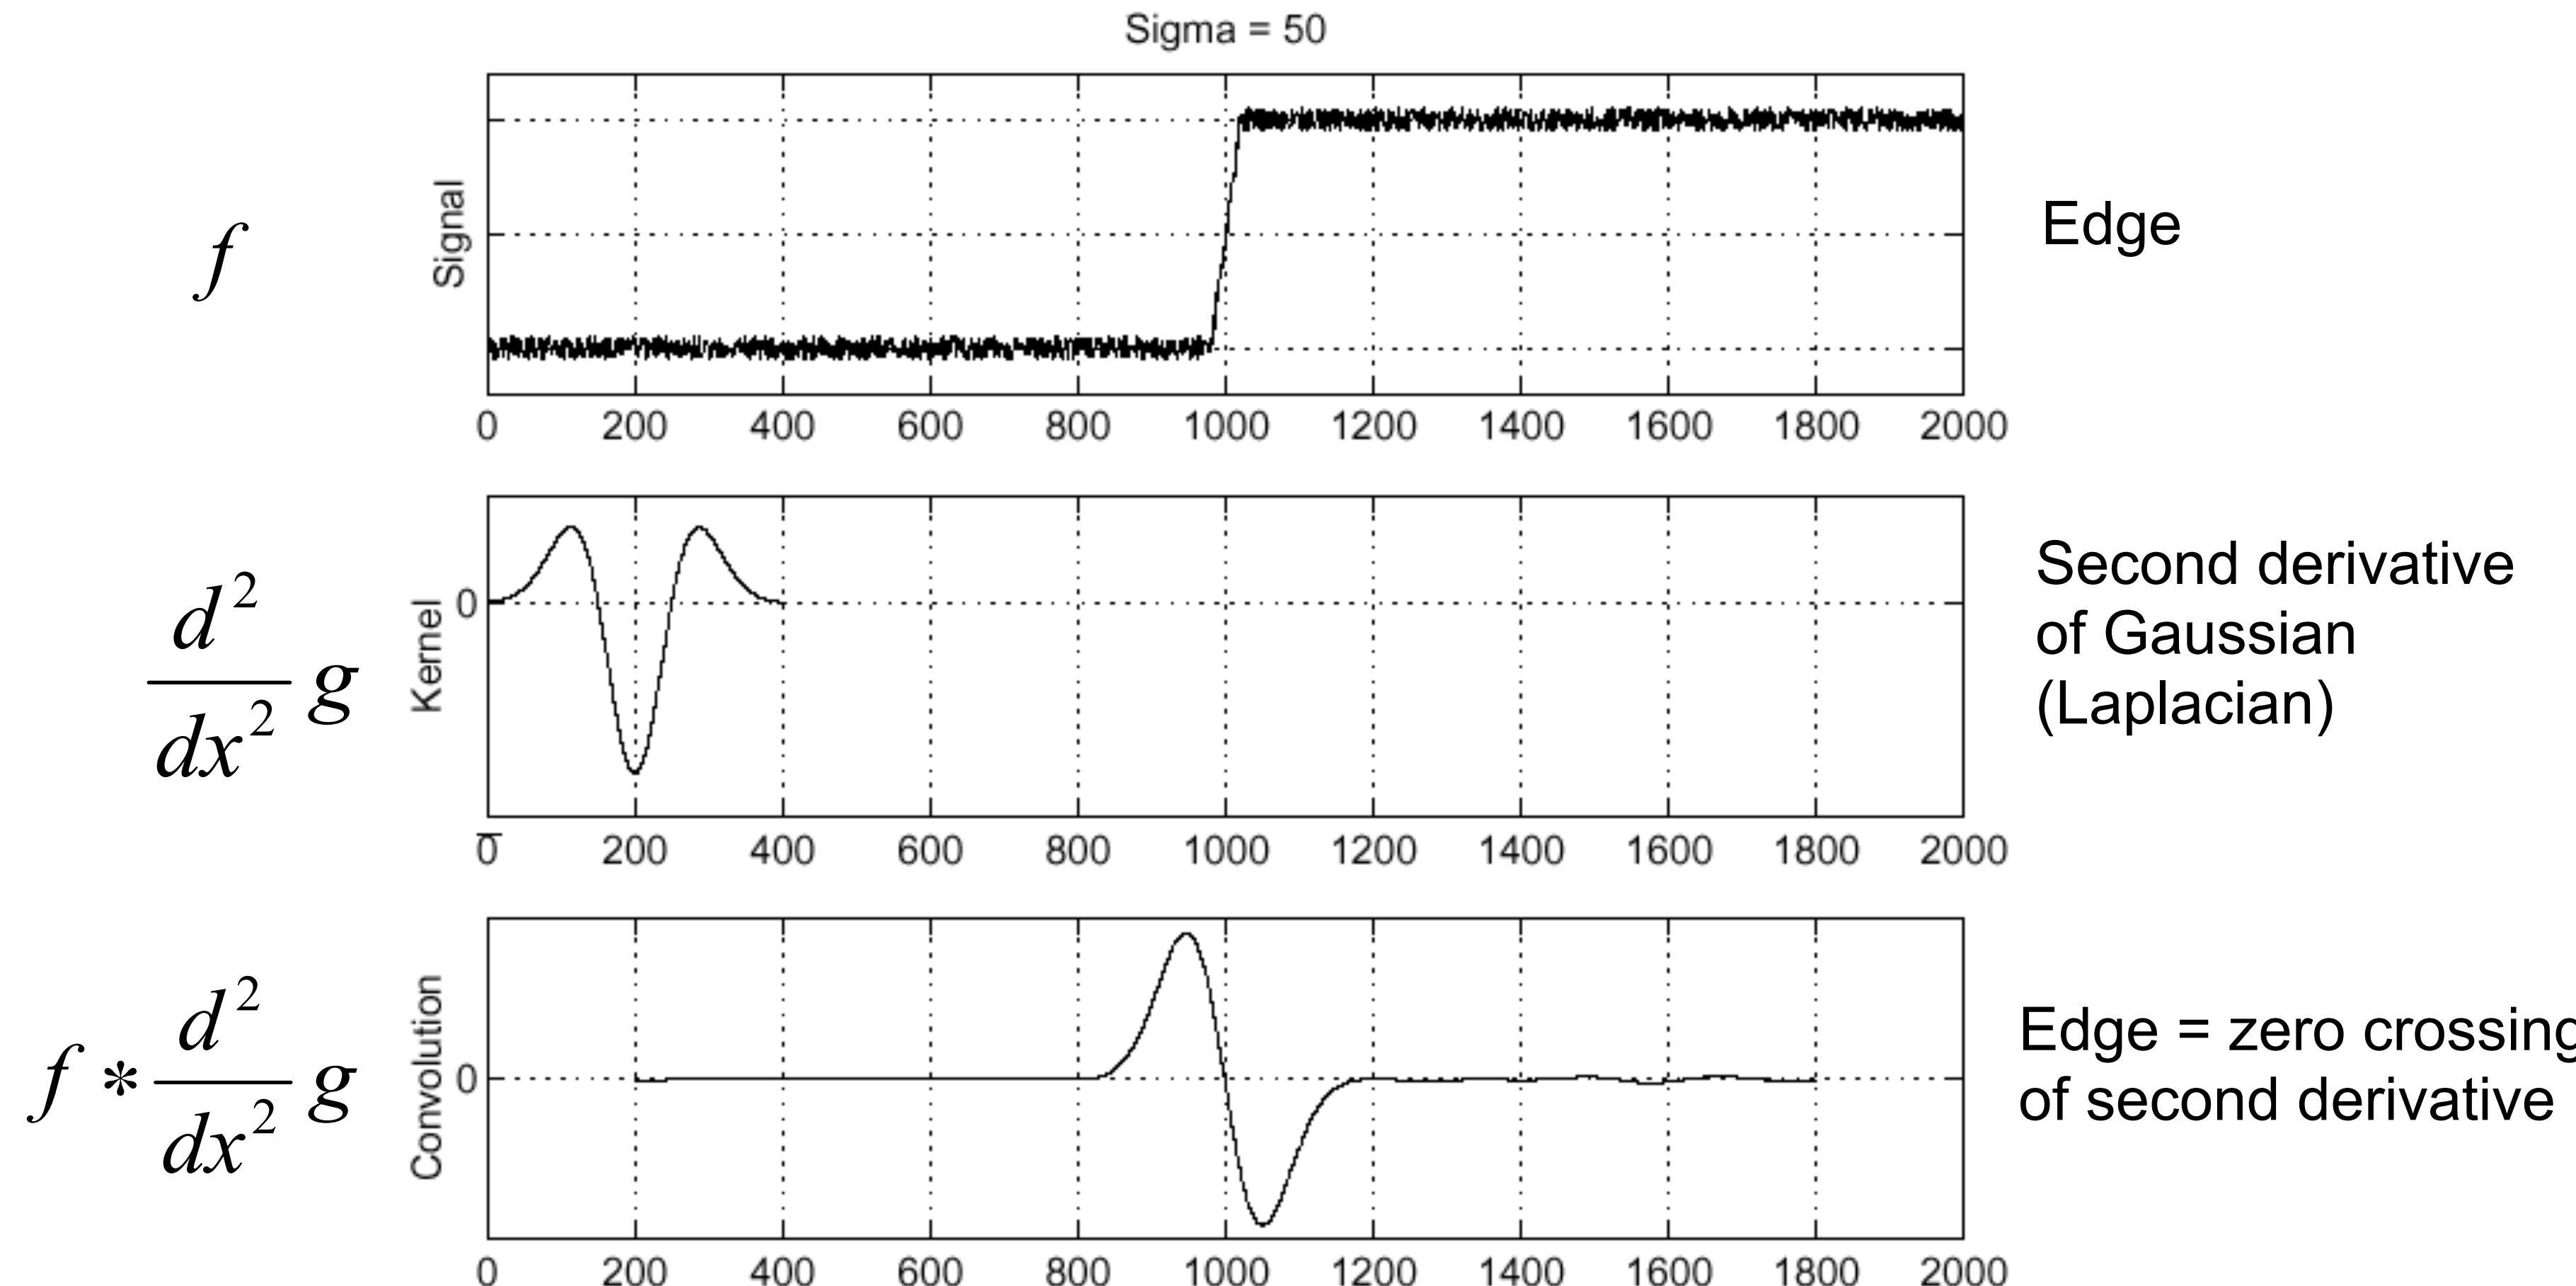


$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

Recall: edge detection



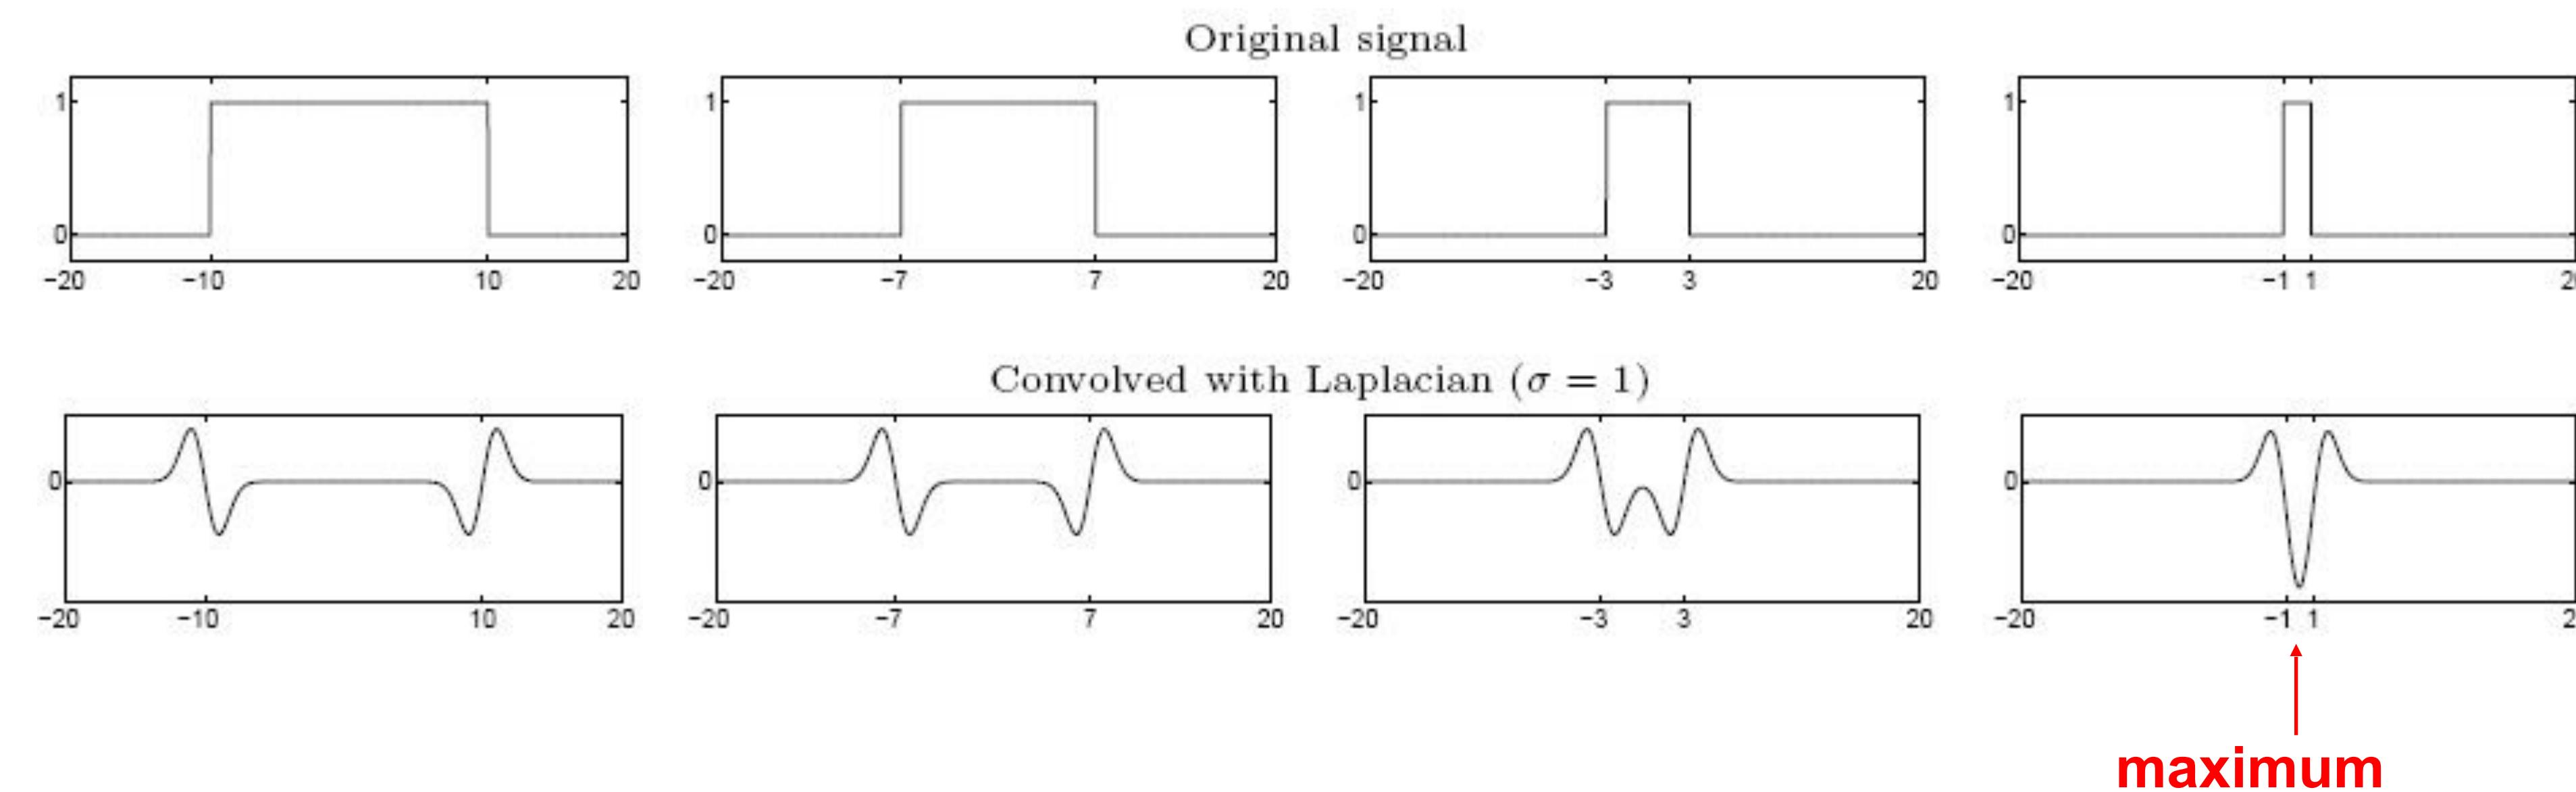
Edge detection using a Laplacian



From edges to blobs

Edge – ripple

Blob – superposition of two ripples

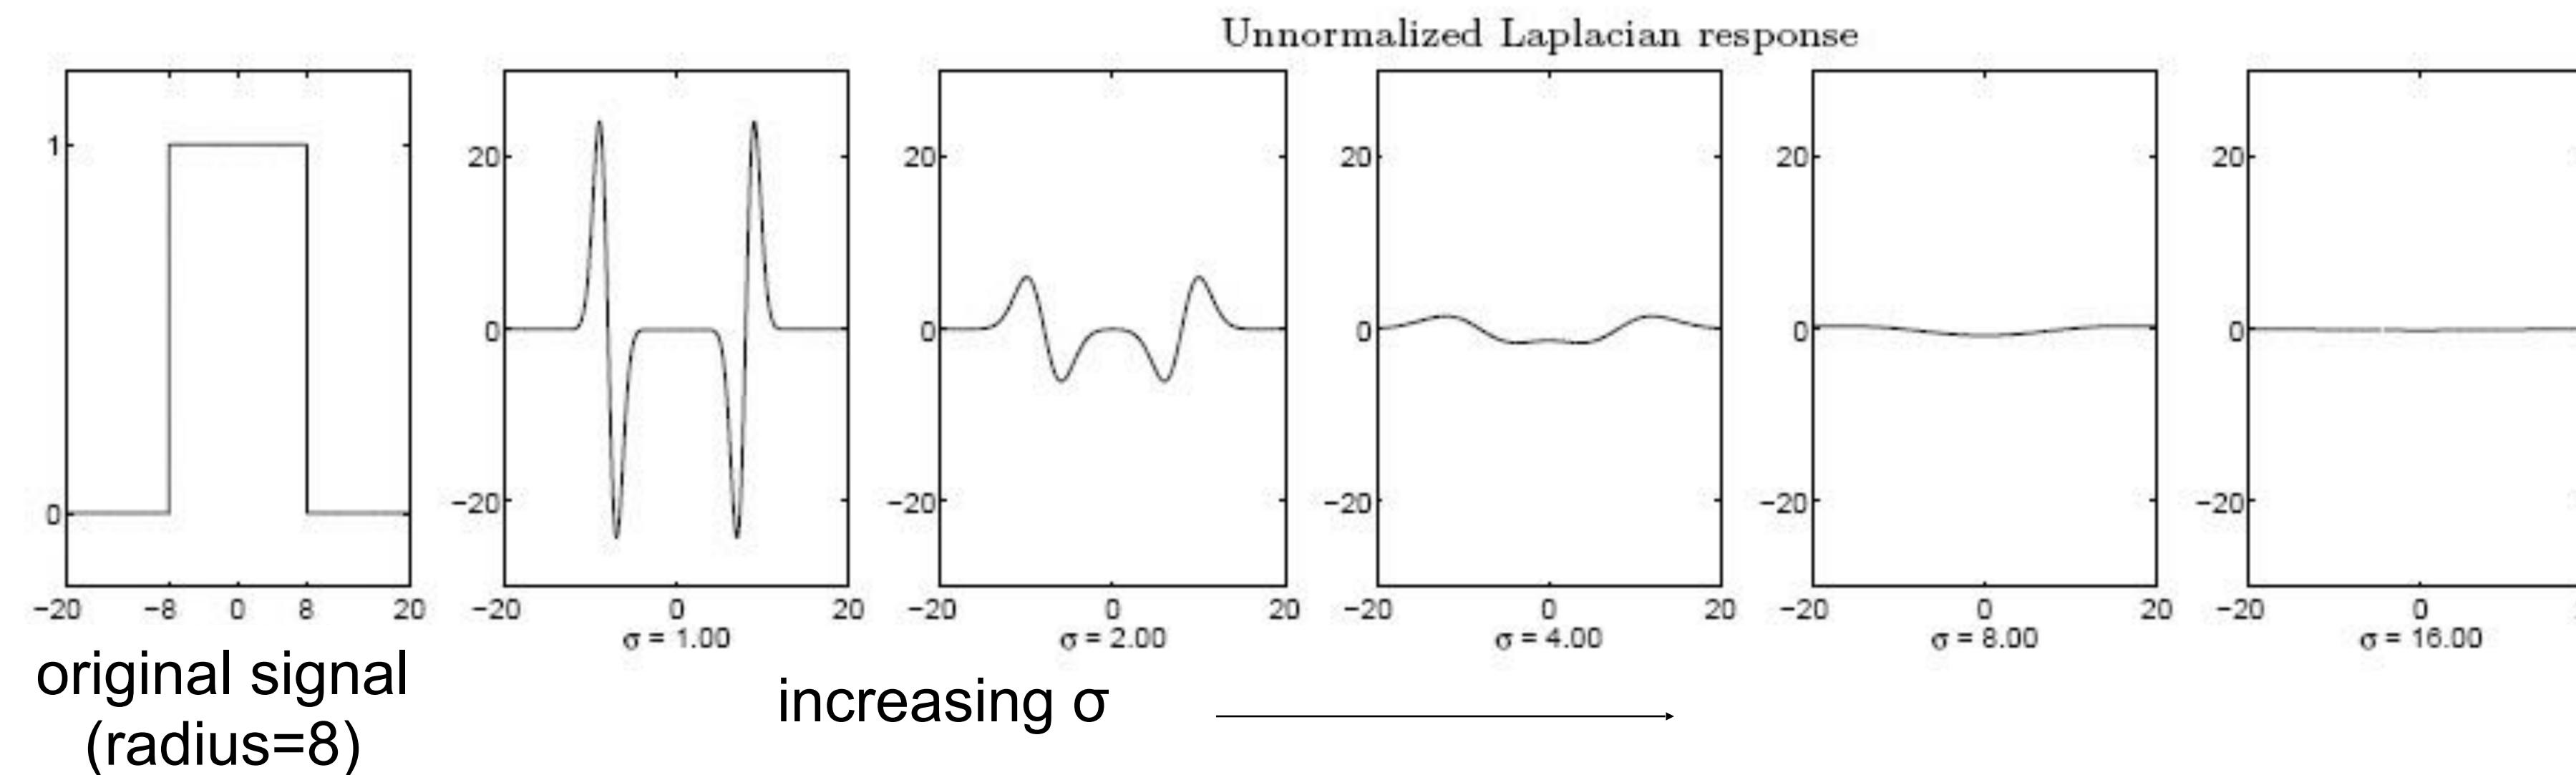


Spatial selection: the magnitude of the Laplacian response will achieve a maximum at the center of the blob, provided the scale of the Laplacian is “matched” to the scale of the blob

Scale selection

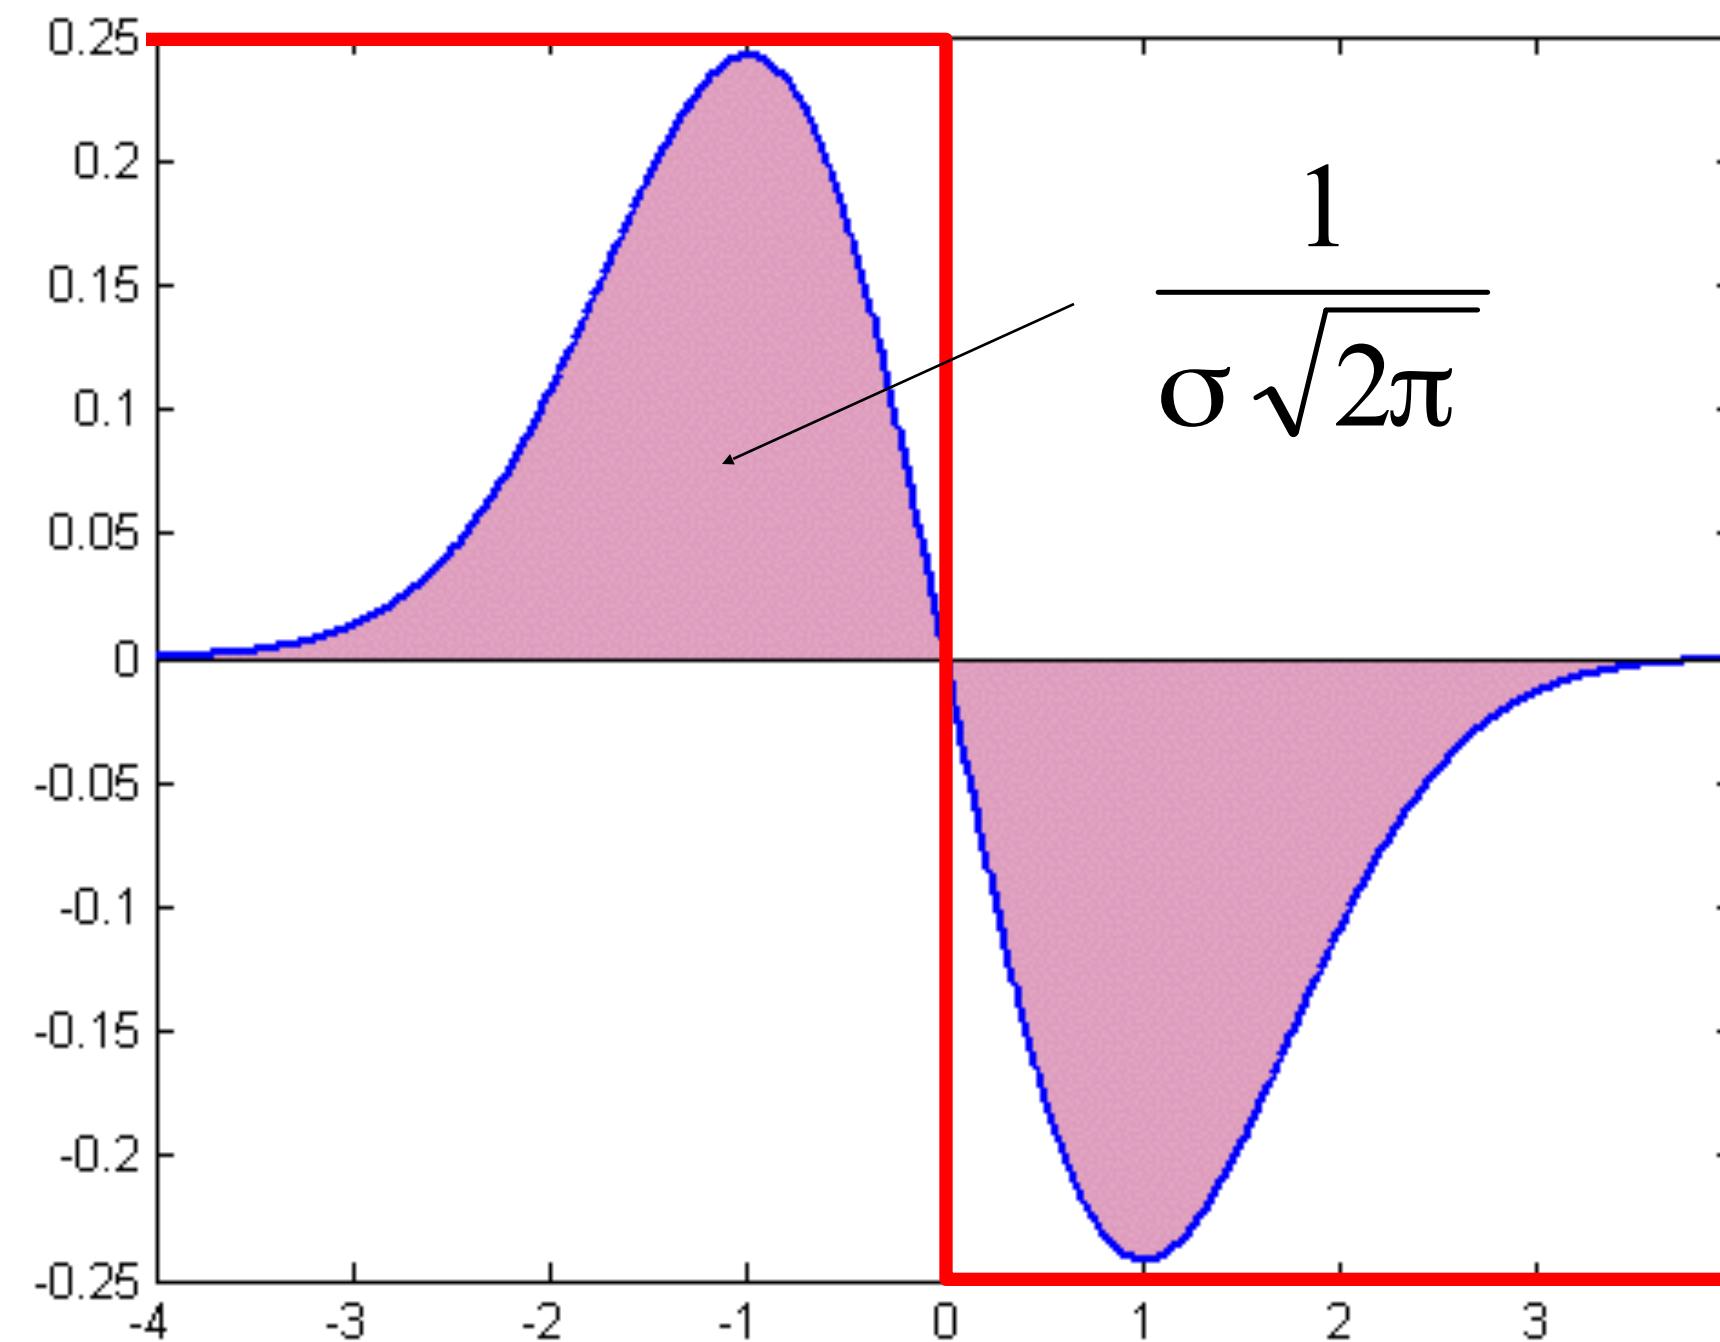
Find the characteristic scale of the blob by convolving it with Laplacians at several scales and looking for the maximum response

However, the Laplacian response decays as scale increases:



Scale normalization

The response of a derivative of Gaussian filter to a perfect step edge decreases as σ increases

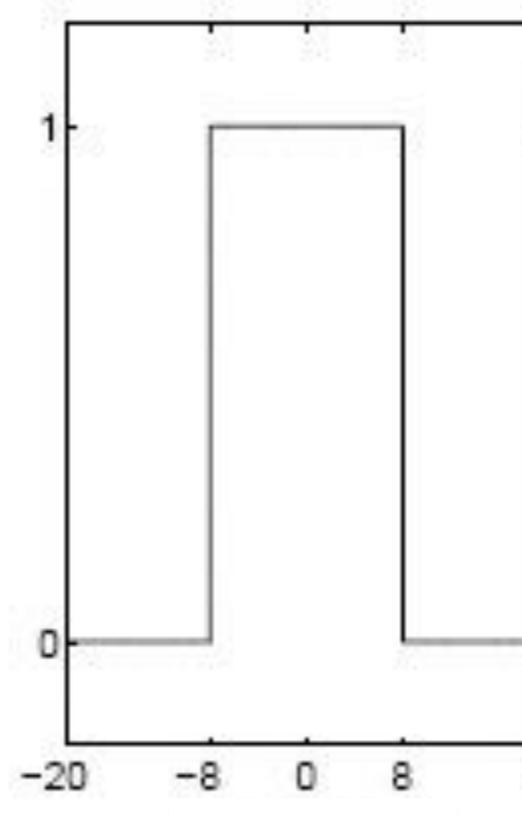


Scale normalization

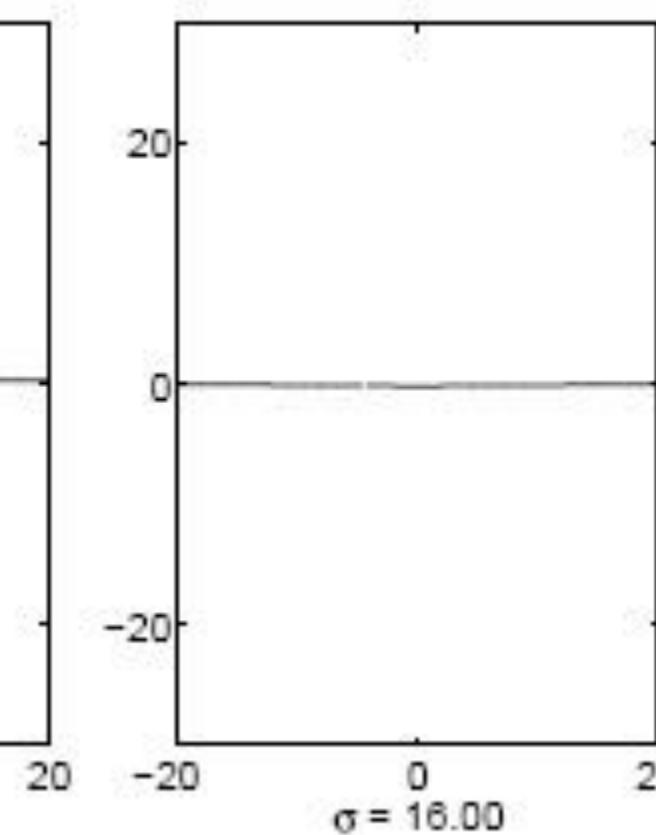
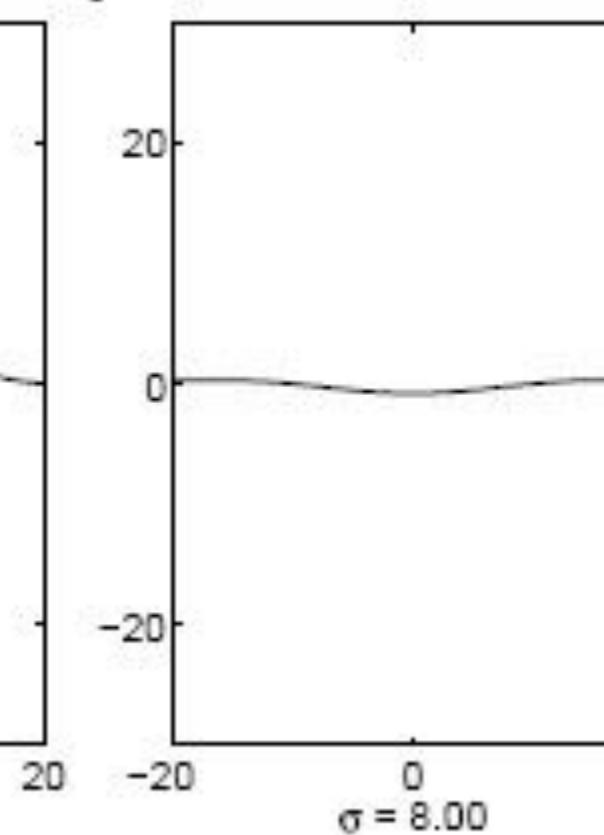
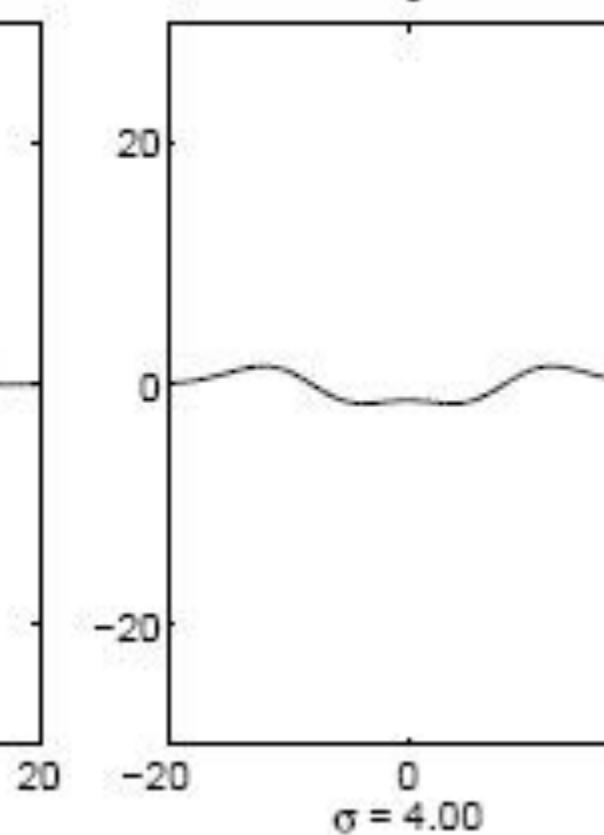
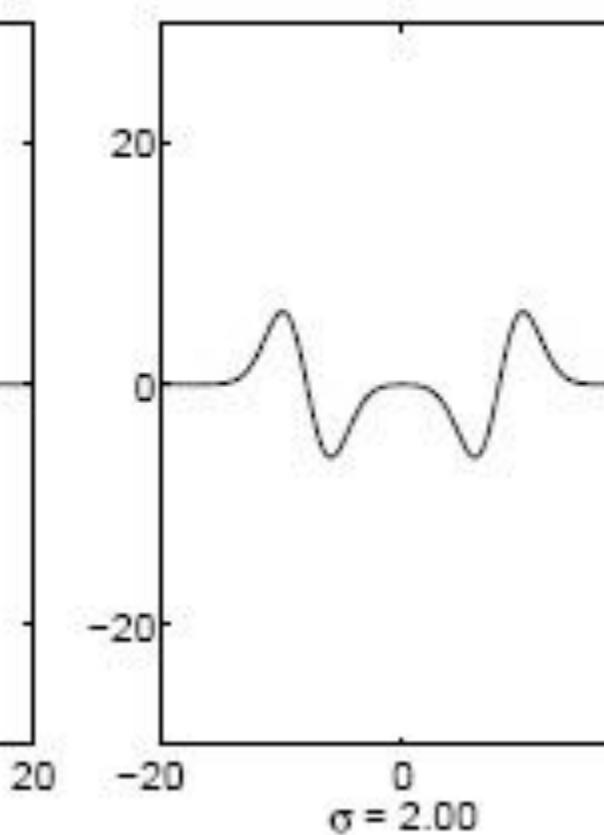
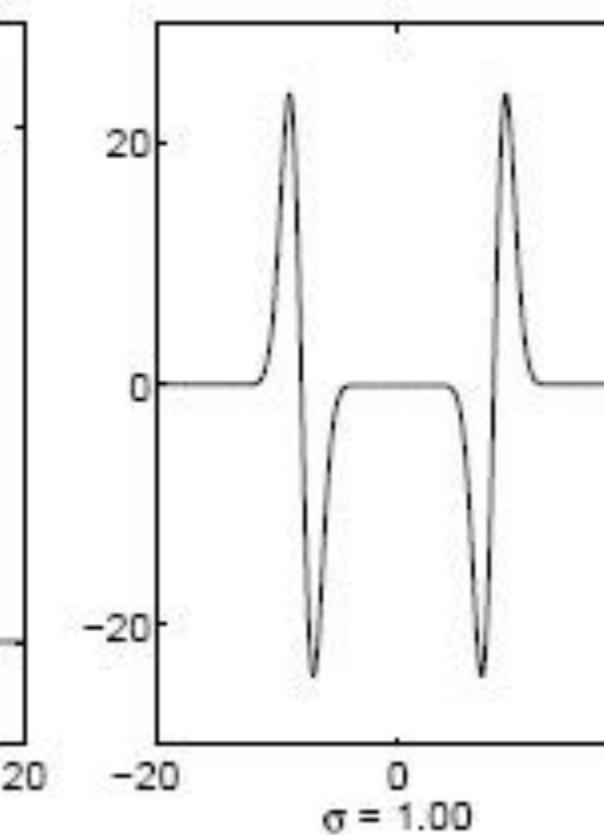
The response of a derivative of Gaussian filter to a perfect step edge decreases as σ increases
To keep response the same (scale-invariant), must multiply Gaussian derivative by σ
Laplacian is the second Gaussian derivative, so it must be multiplied by σ^2

Effect of scale normalization

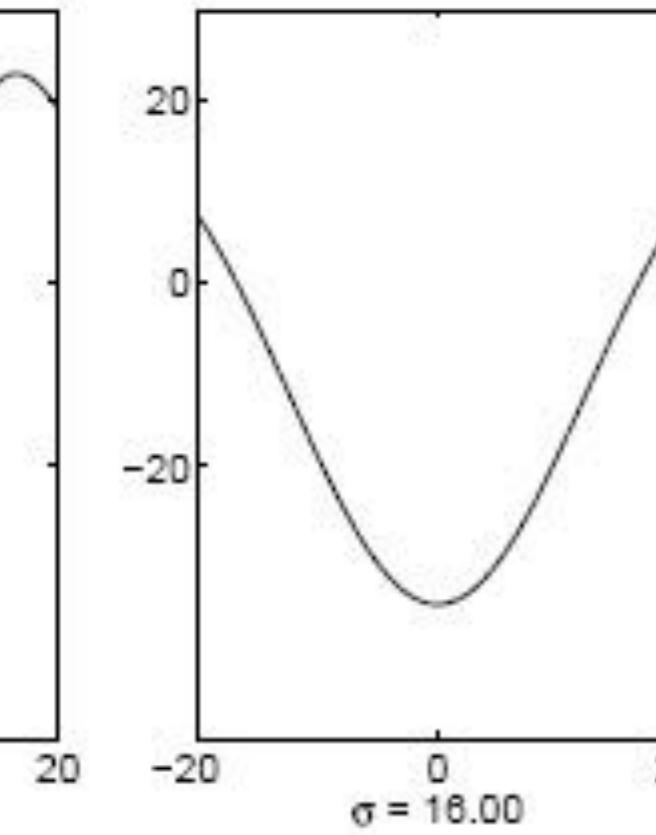
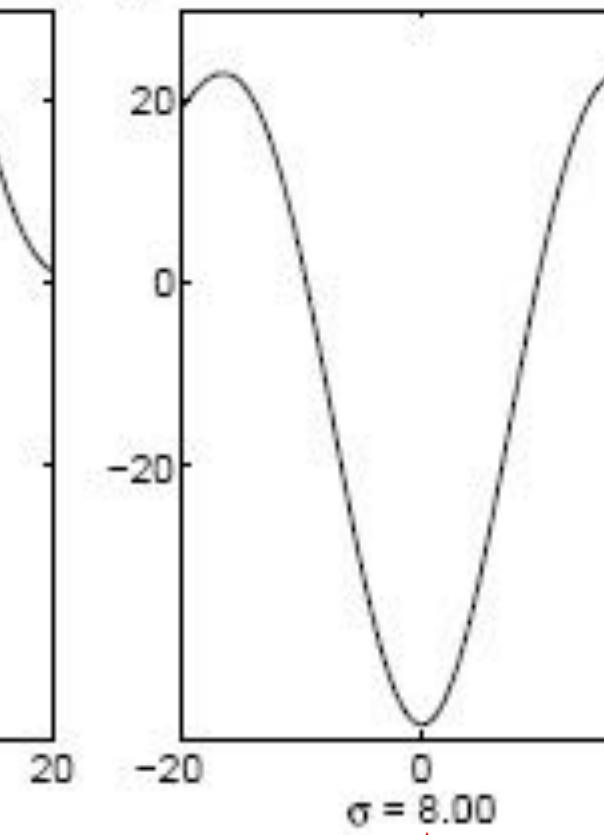
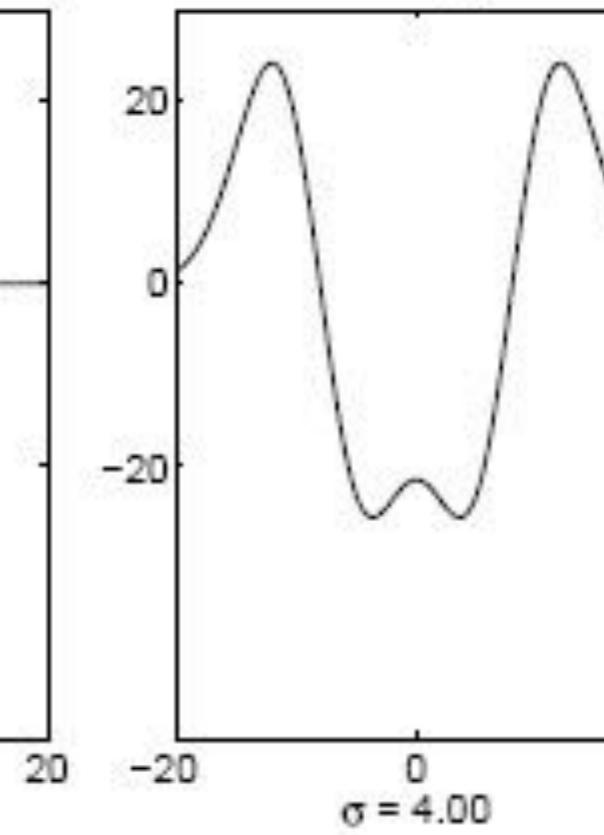
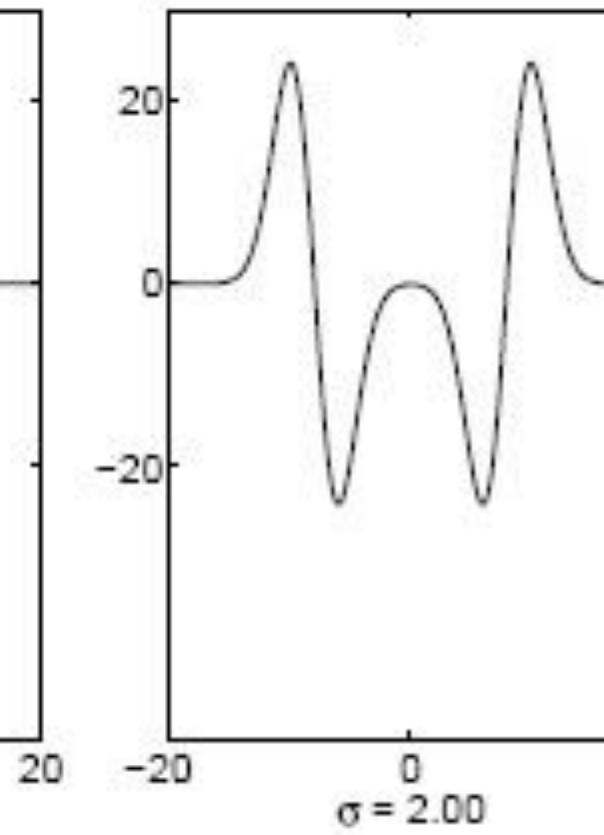
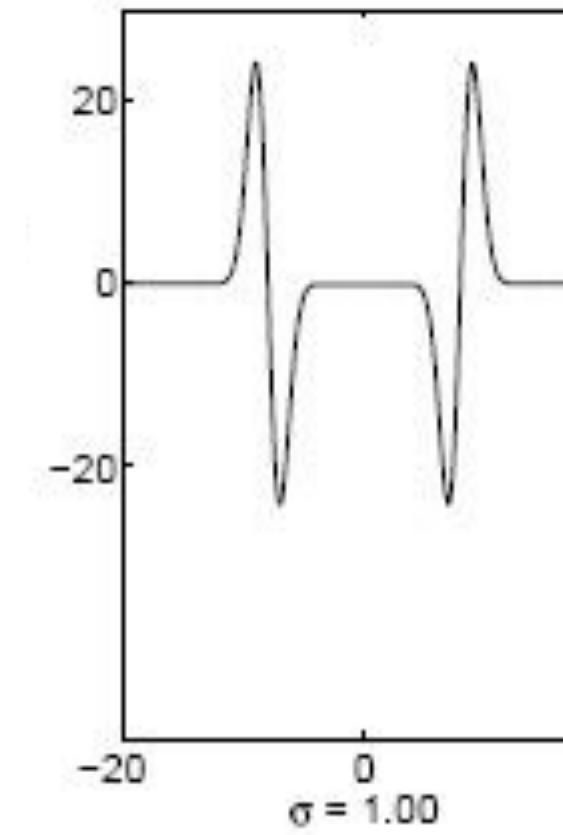
Original signal



Unnormalized Laplacian response



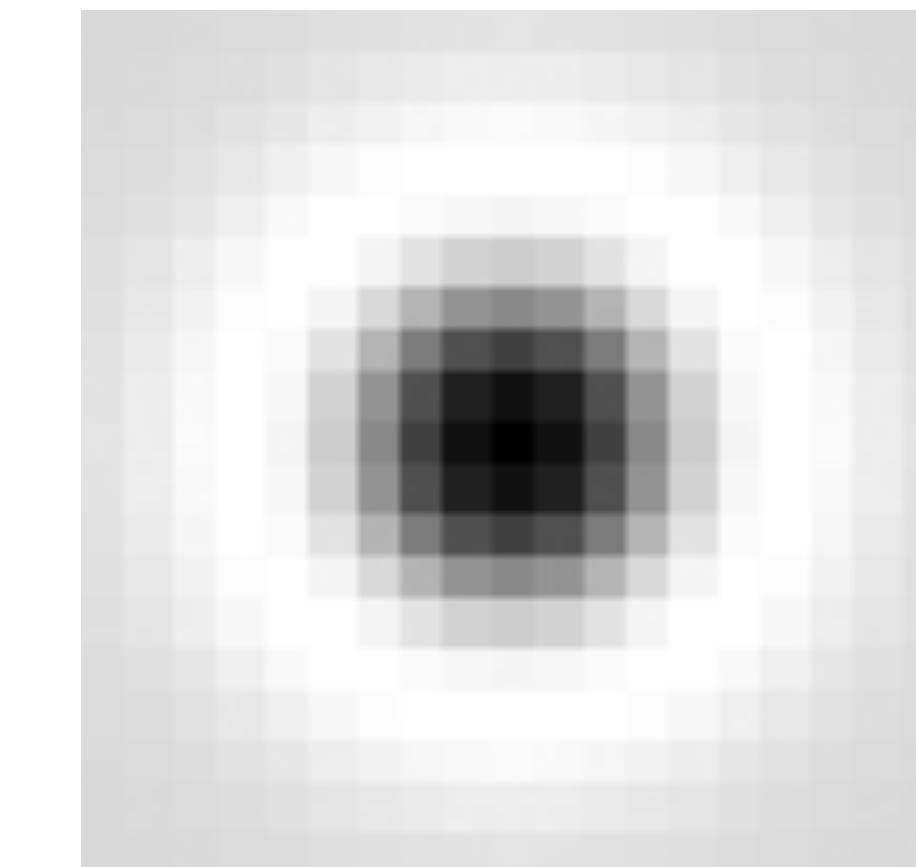
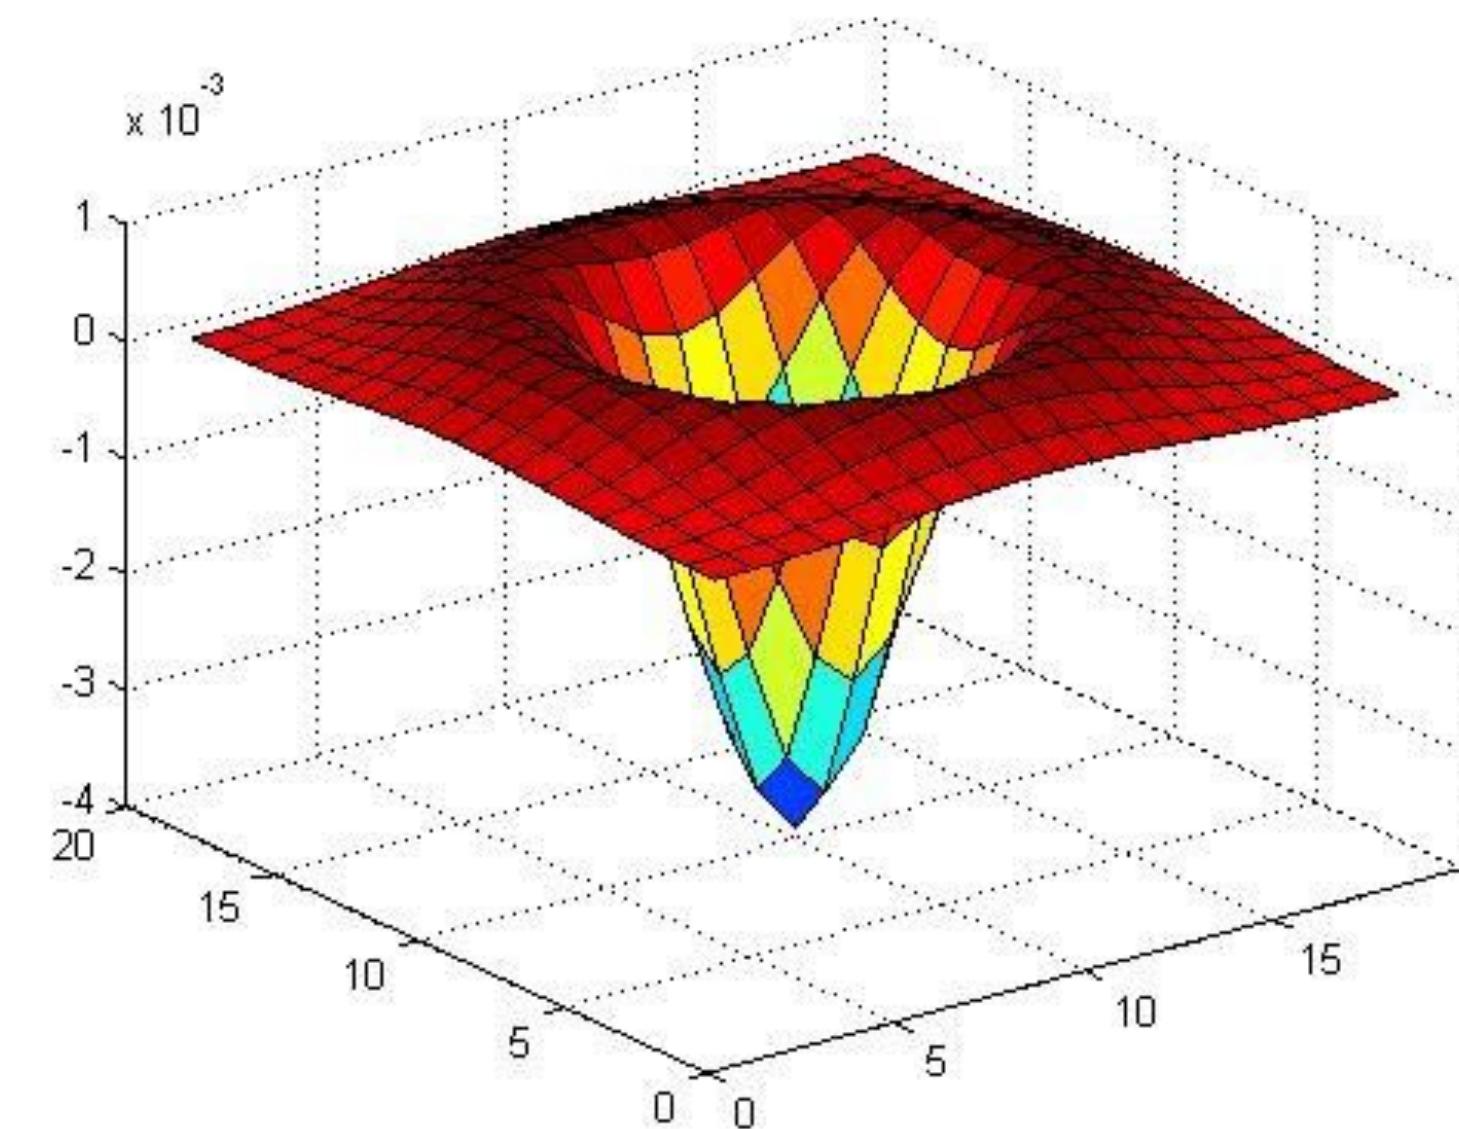
Scale-normalized Laplacian response



maximum

Blob detection in 2D

Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D

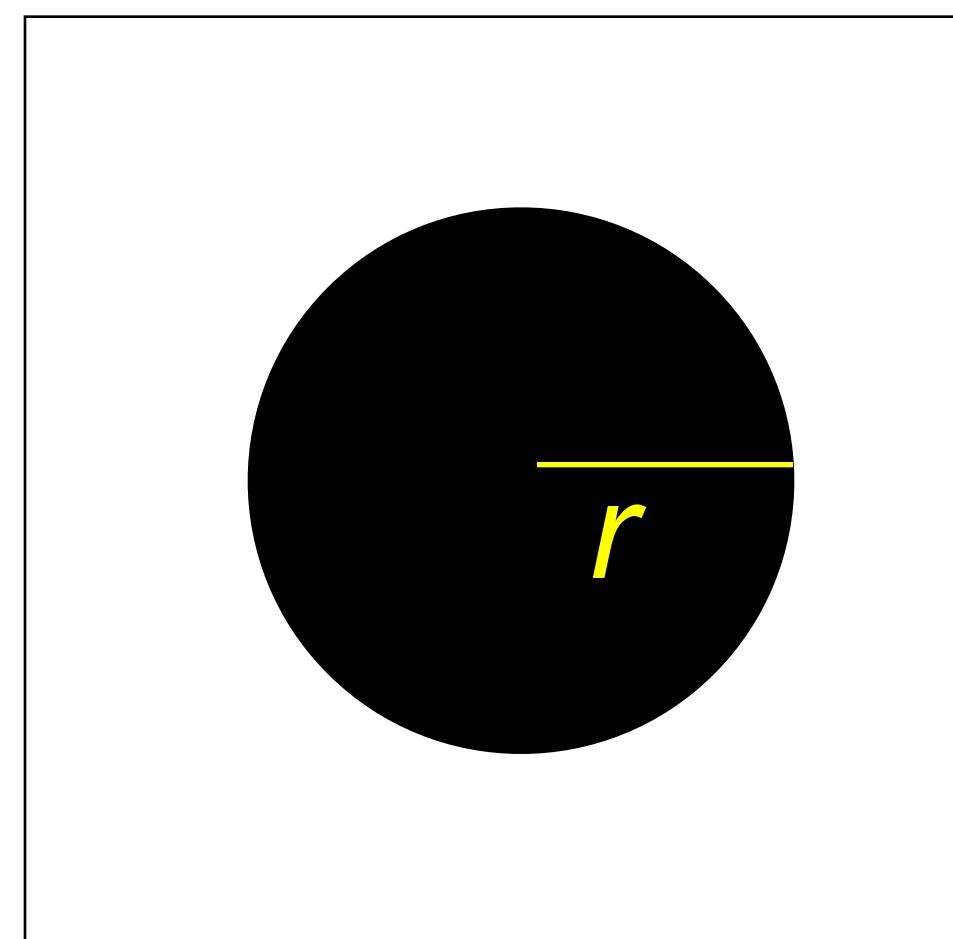


Scale-normalized:

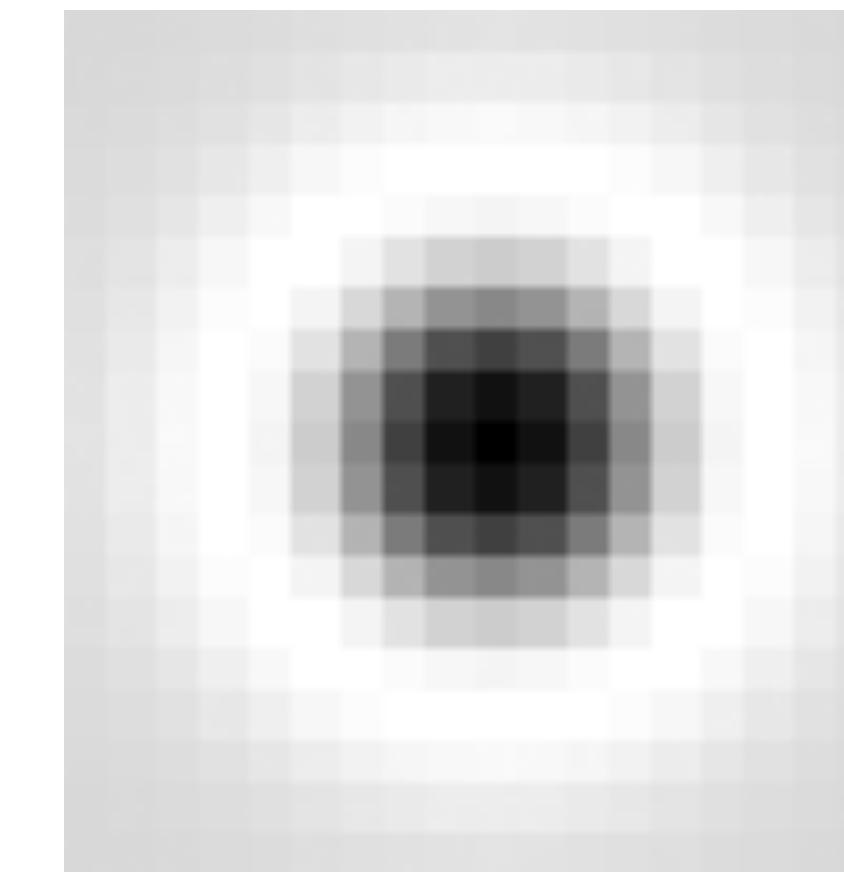
$$\nabla_{\text{norm}}^2 g = \sigma^2 \left(\frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$

Scale selection

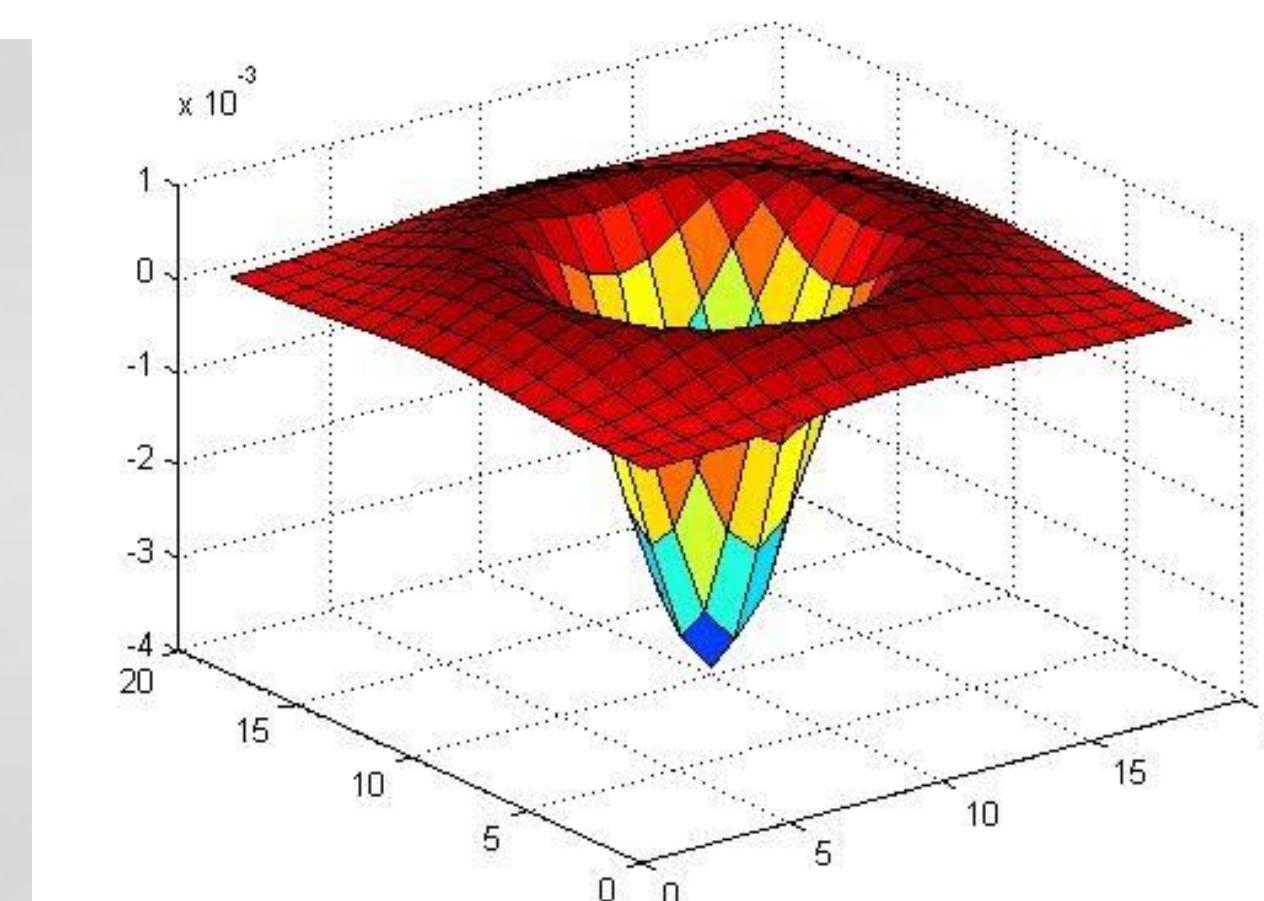
At what scale does the Laplacian achieve a maximum response to a binary circle of radius r ?



image



Laplacian



Scale selection

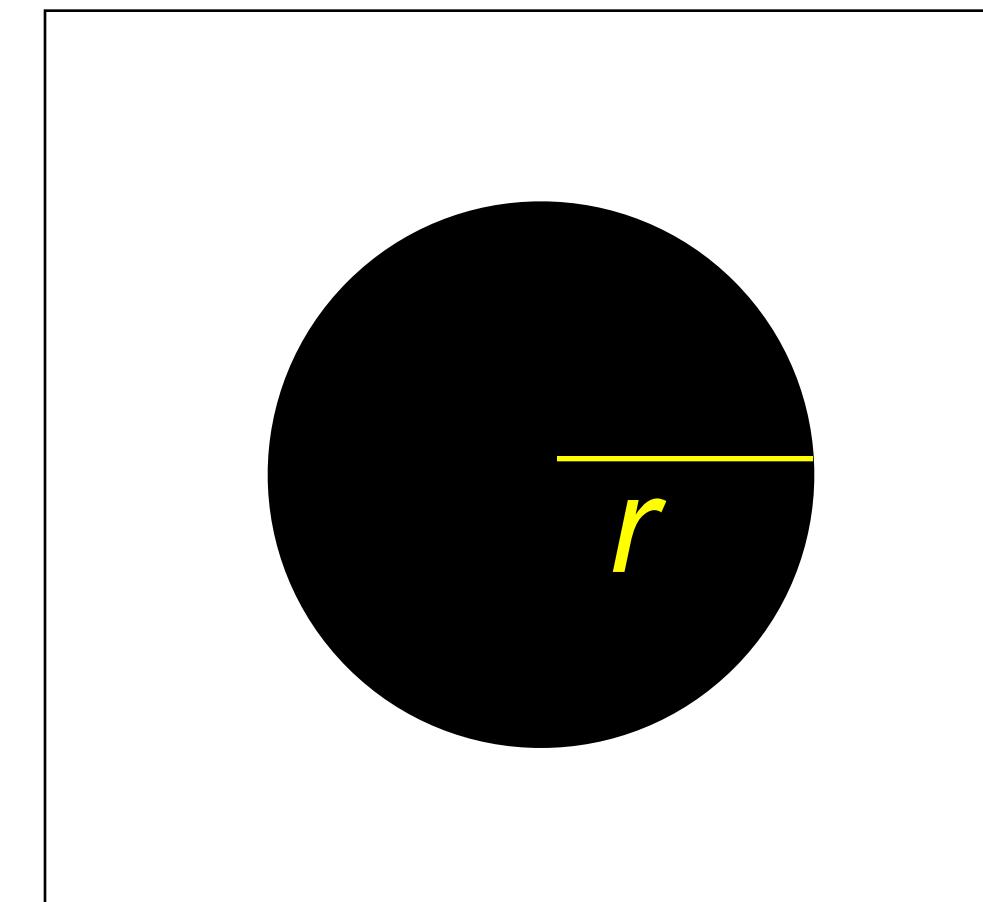
At what scale does the Laplacian achieve a maximum response to a binary circle of radius r ?

To get maximum response, the zeros of the Laplacian have to be aligned with the circle
The Laplacian is given by (up to scale):

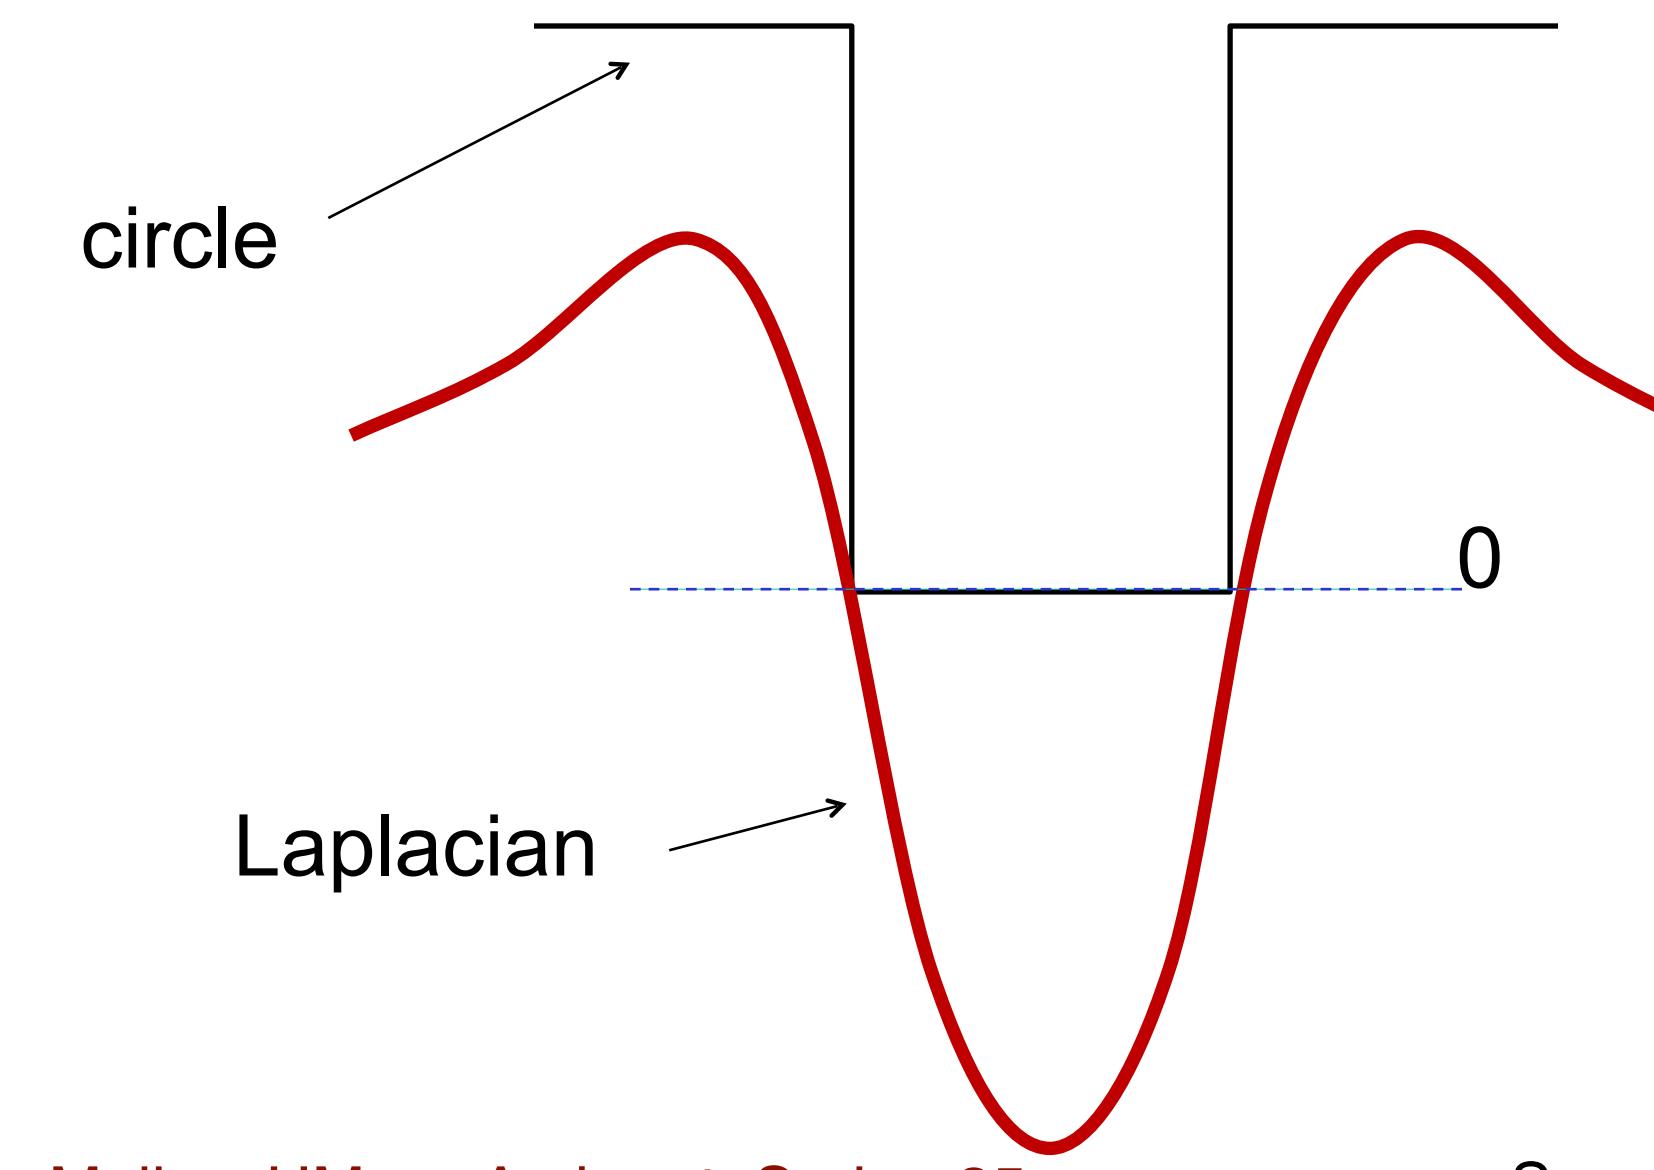
$$(x^2 + y^2 - 2\sigma^2) e^{-(x^2+y^2)/2\sigma^2}$$

Therefore, the maximum response occurs at

$$\sigma = r / \sqrt{2}.$$

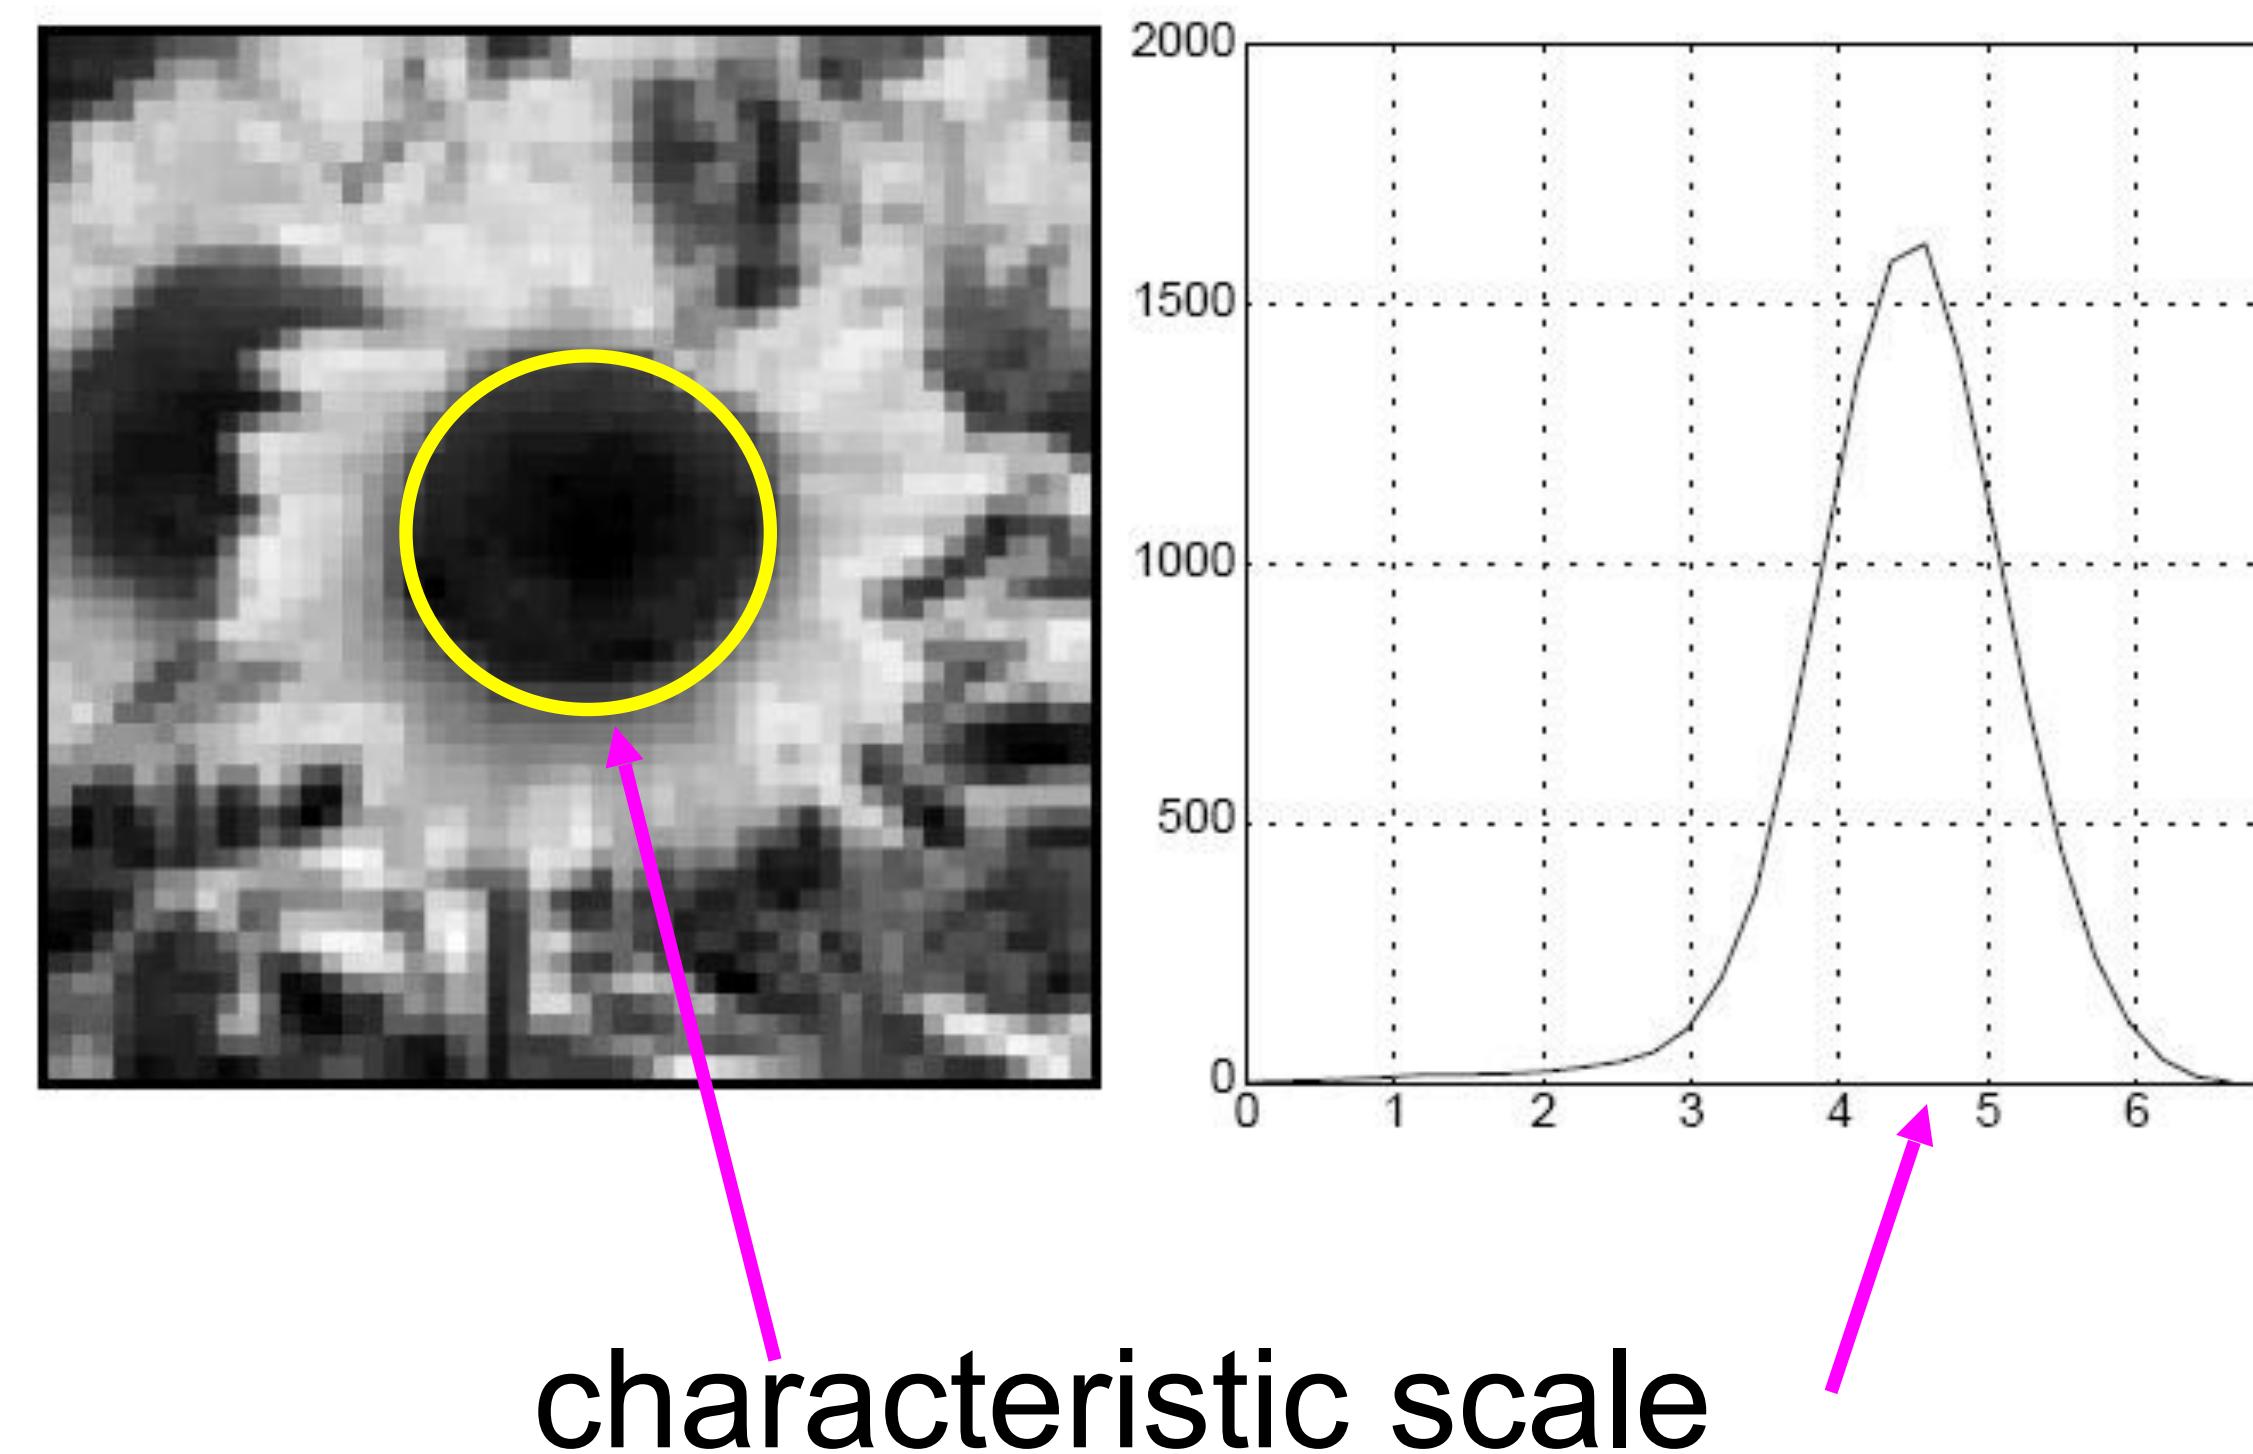


image



Characteristic scale

We define the characteristic scale of a blob as the scale that produces peak of Laplacian response in the blob center



T. Lindeberg (1998). ["Feature detection with automatic scale selection."](#)
International Journal of Computer Vision 30 (2): pp 77--116.

Scale-space blob detector

1. Convolve image with scale-normalized Laplacian at several scales

Scale-space blob detector: Example



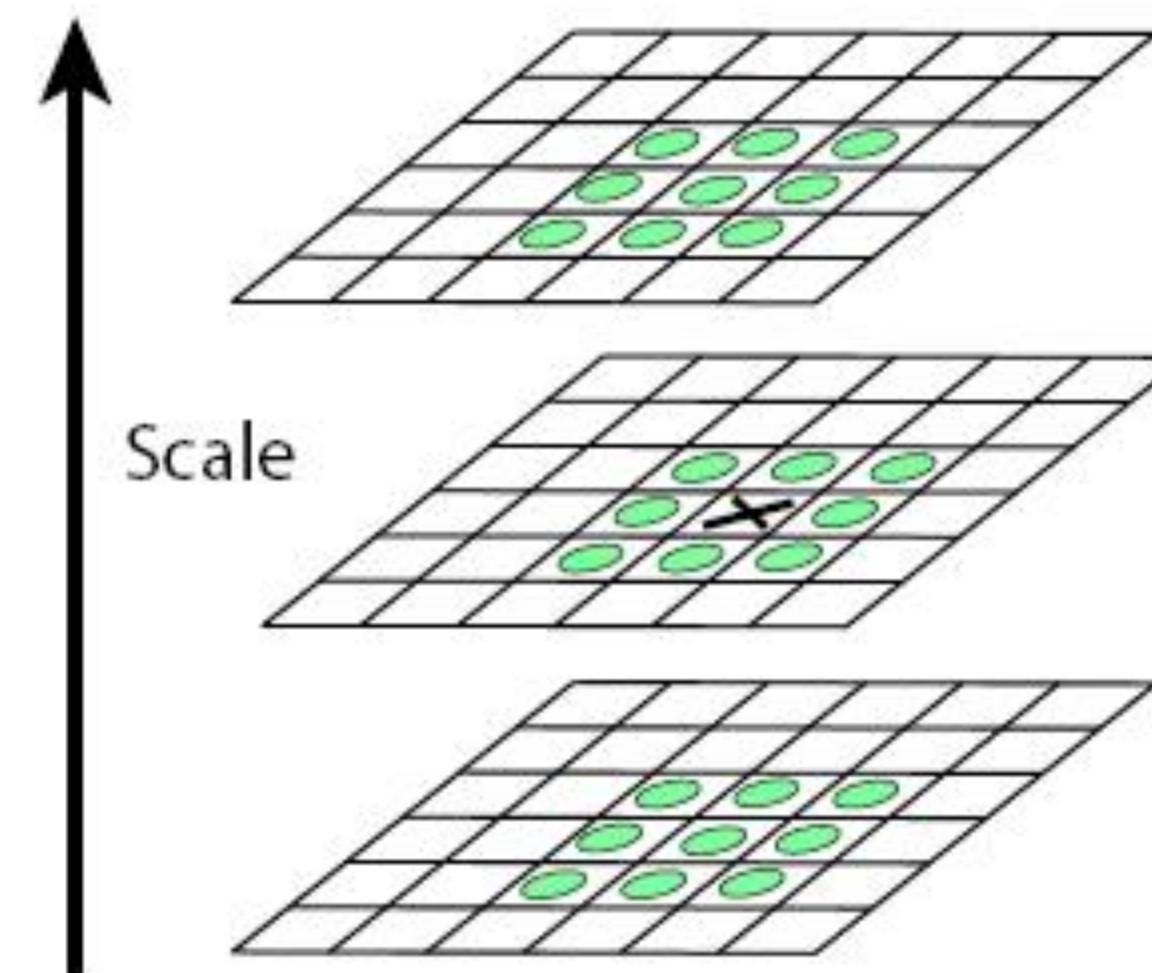
Scale-space blob detector: Example



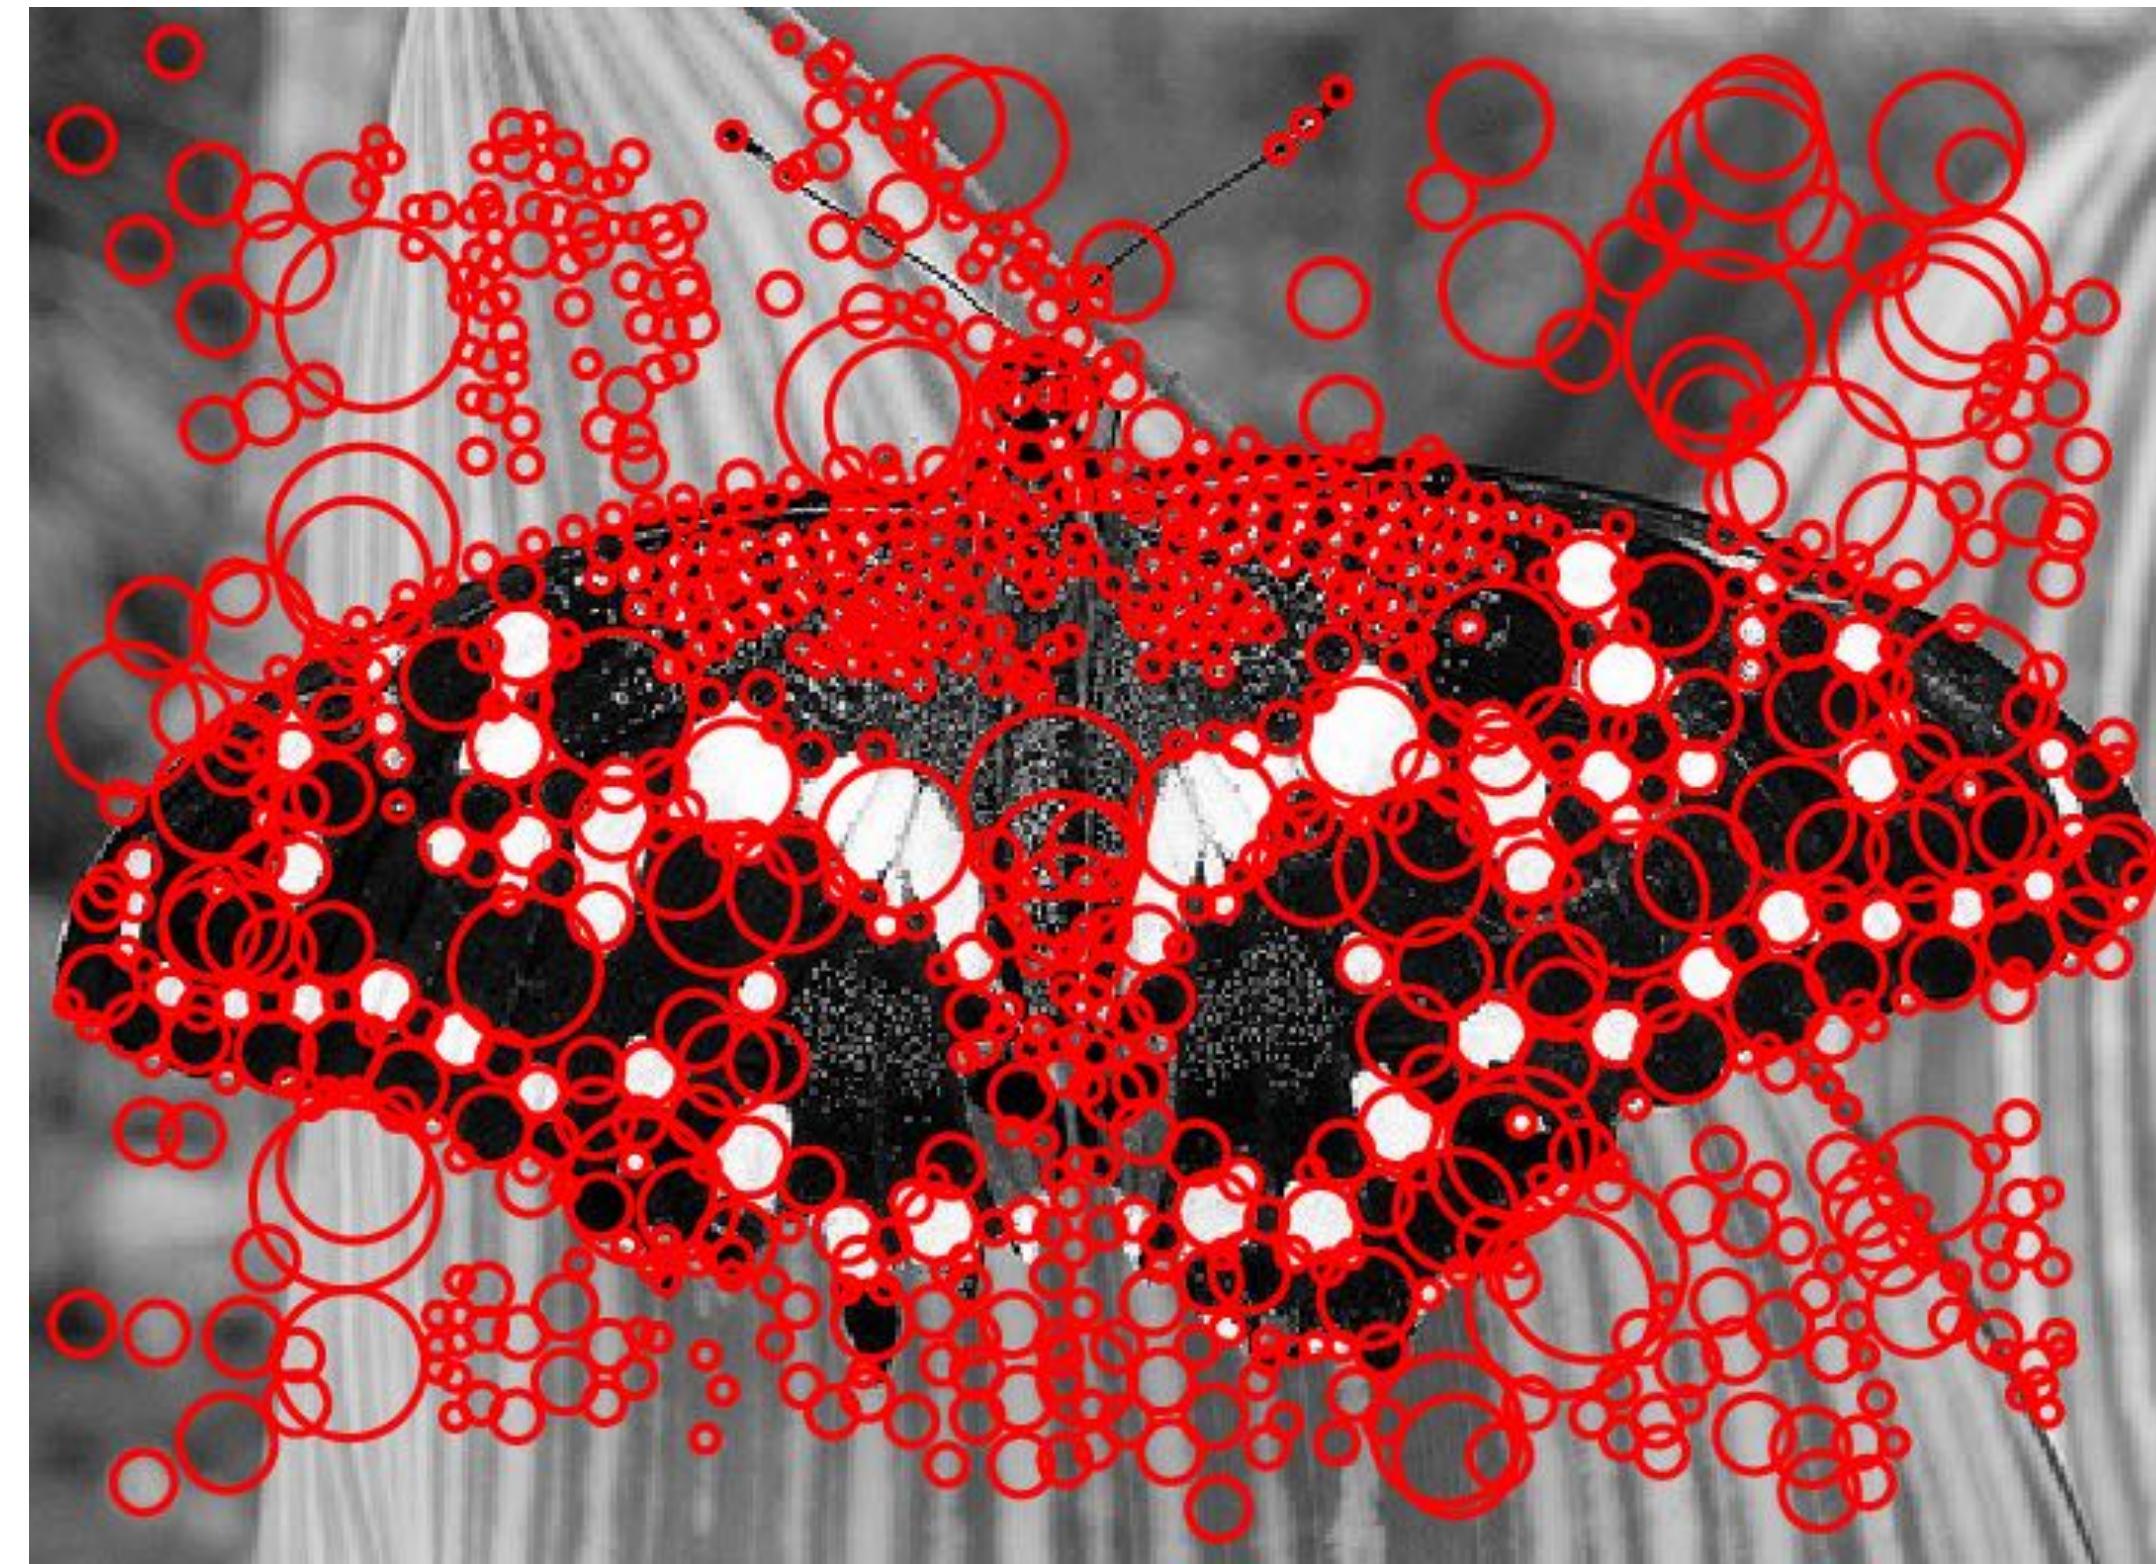
$\sigma = 11.9912$

Scale-space blob detector

1. Convolve image with scale-normalized Laplacian at several scales
2. Find maxima of squared Laplacian response in scale-space



Scale-space blob detector: Example



Efficient implementation

Is the Laplacian separable?

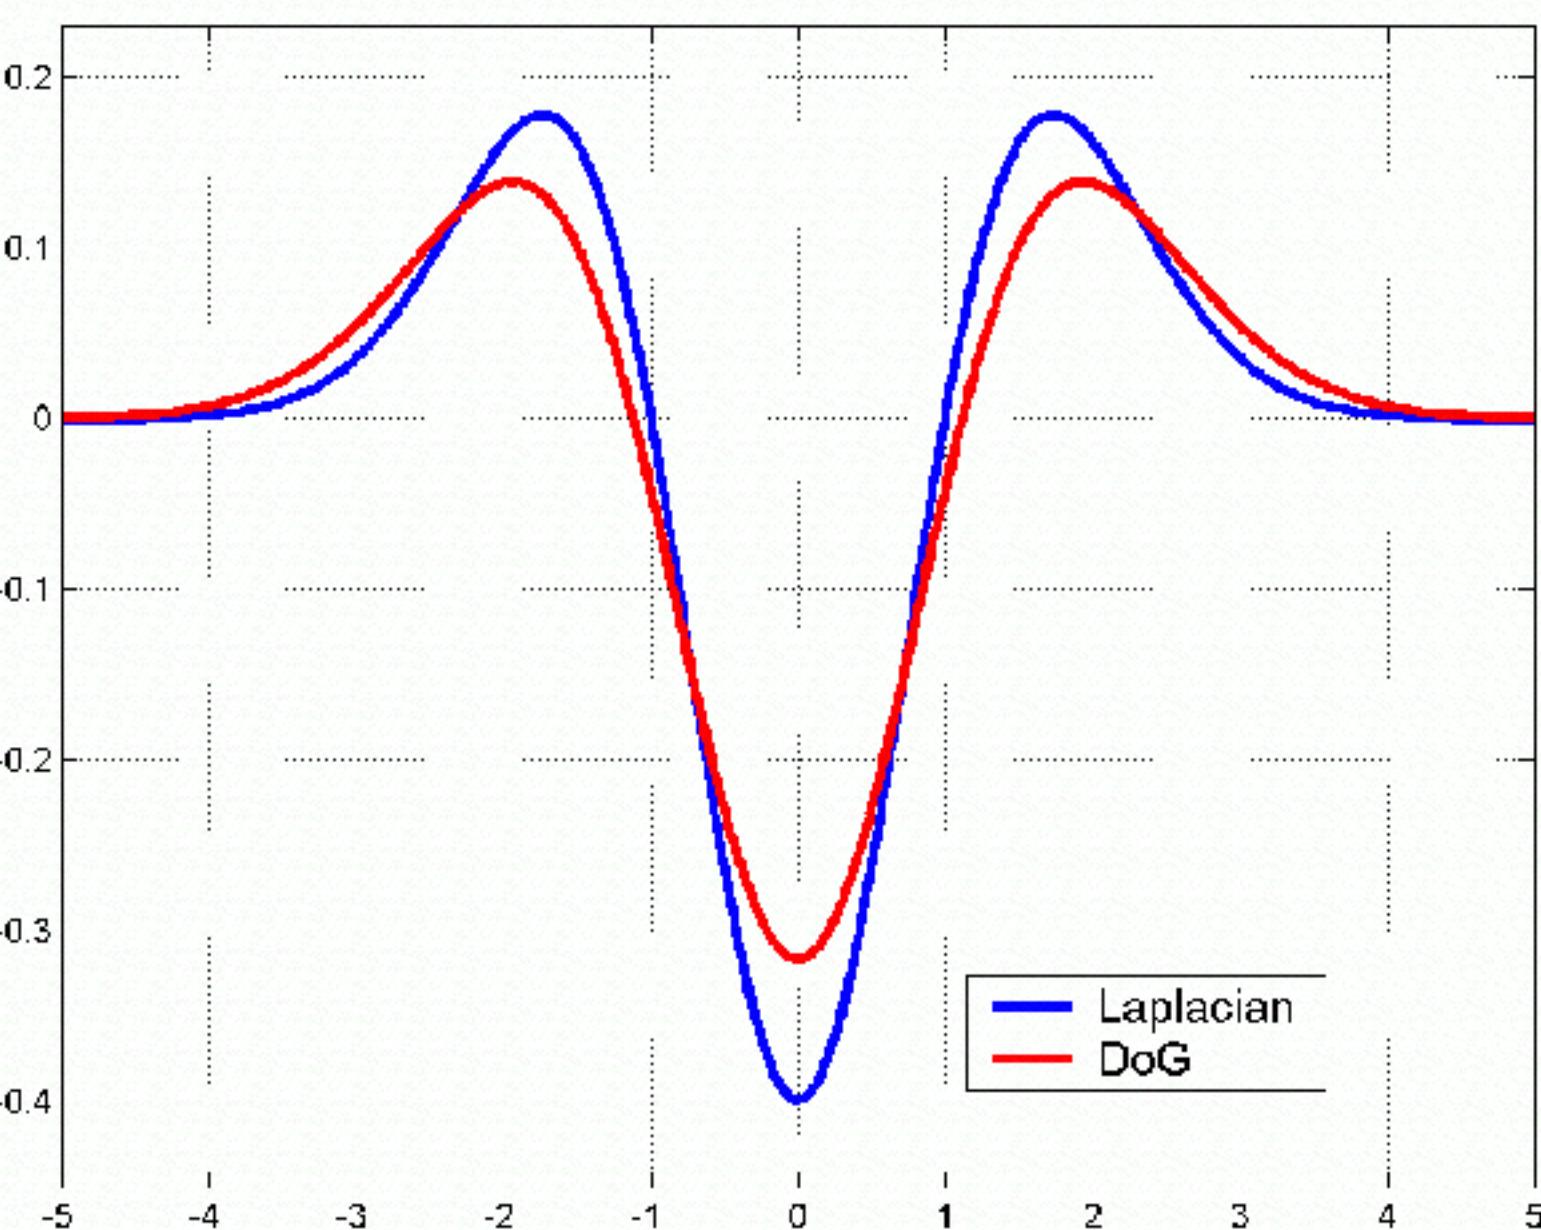
$$L = \sigma^2 (G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma))$$

(Laplacian)

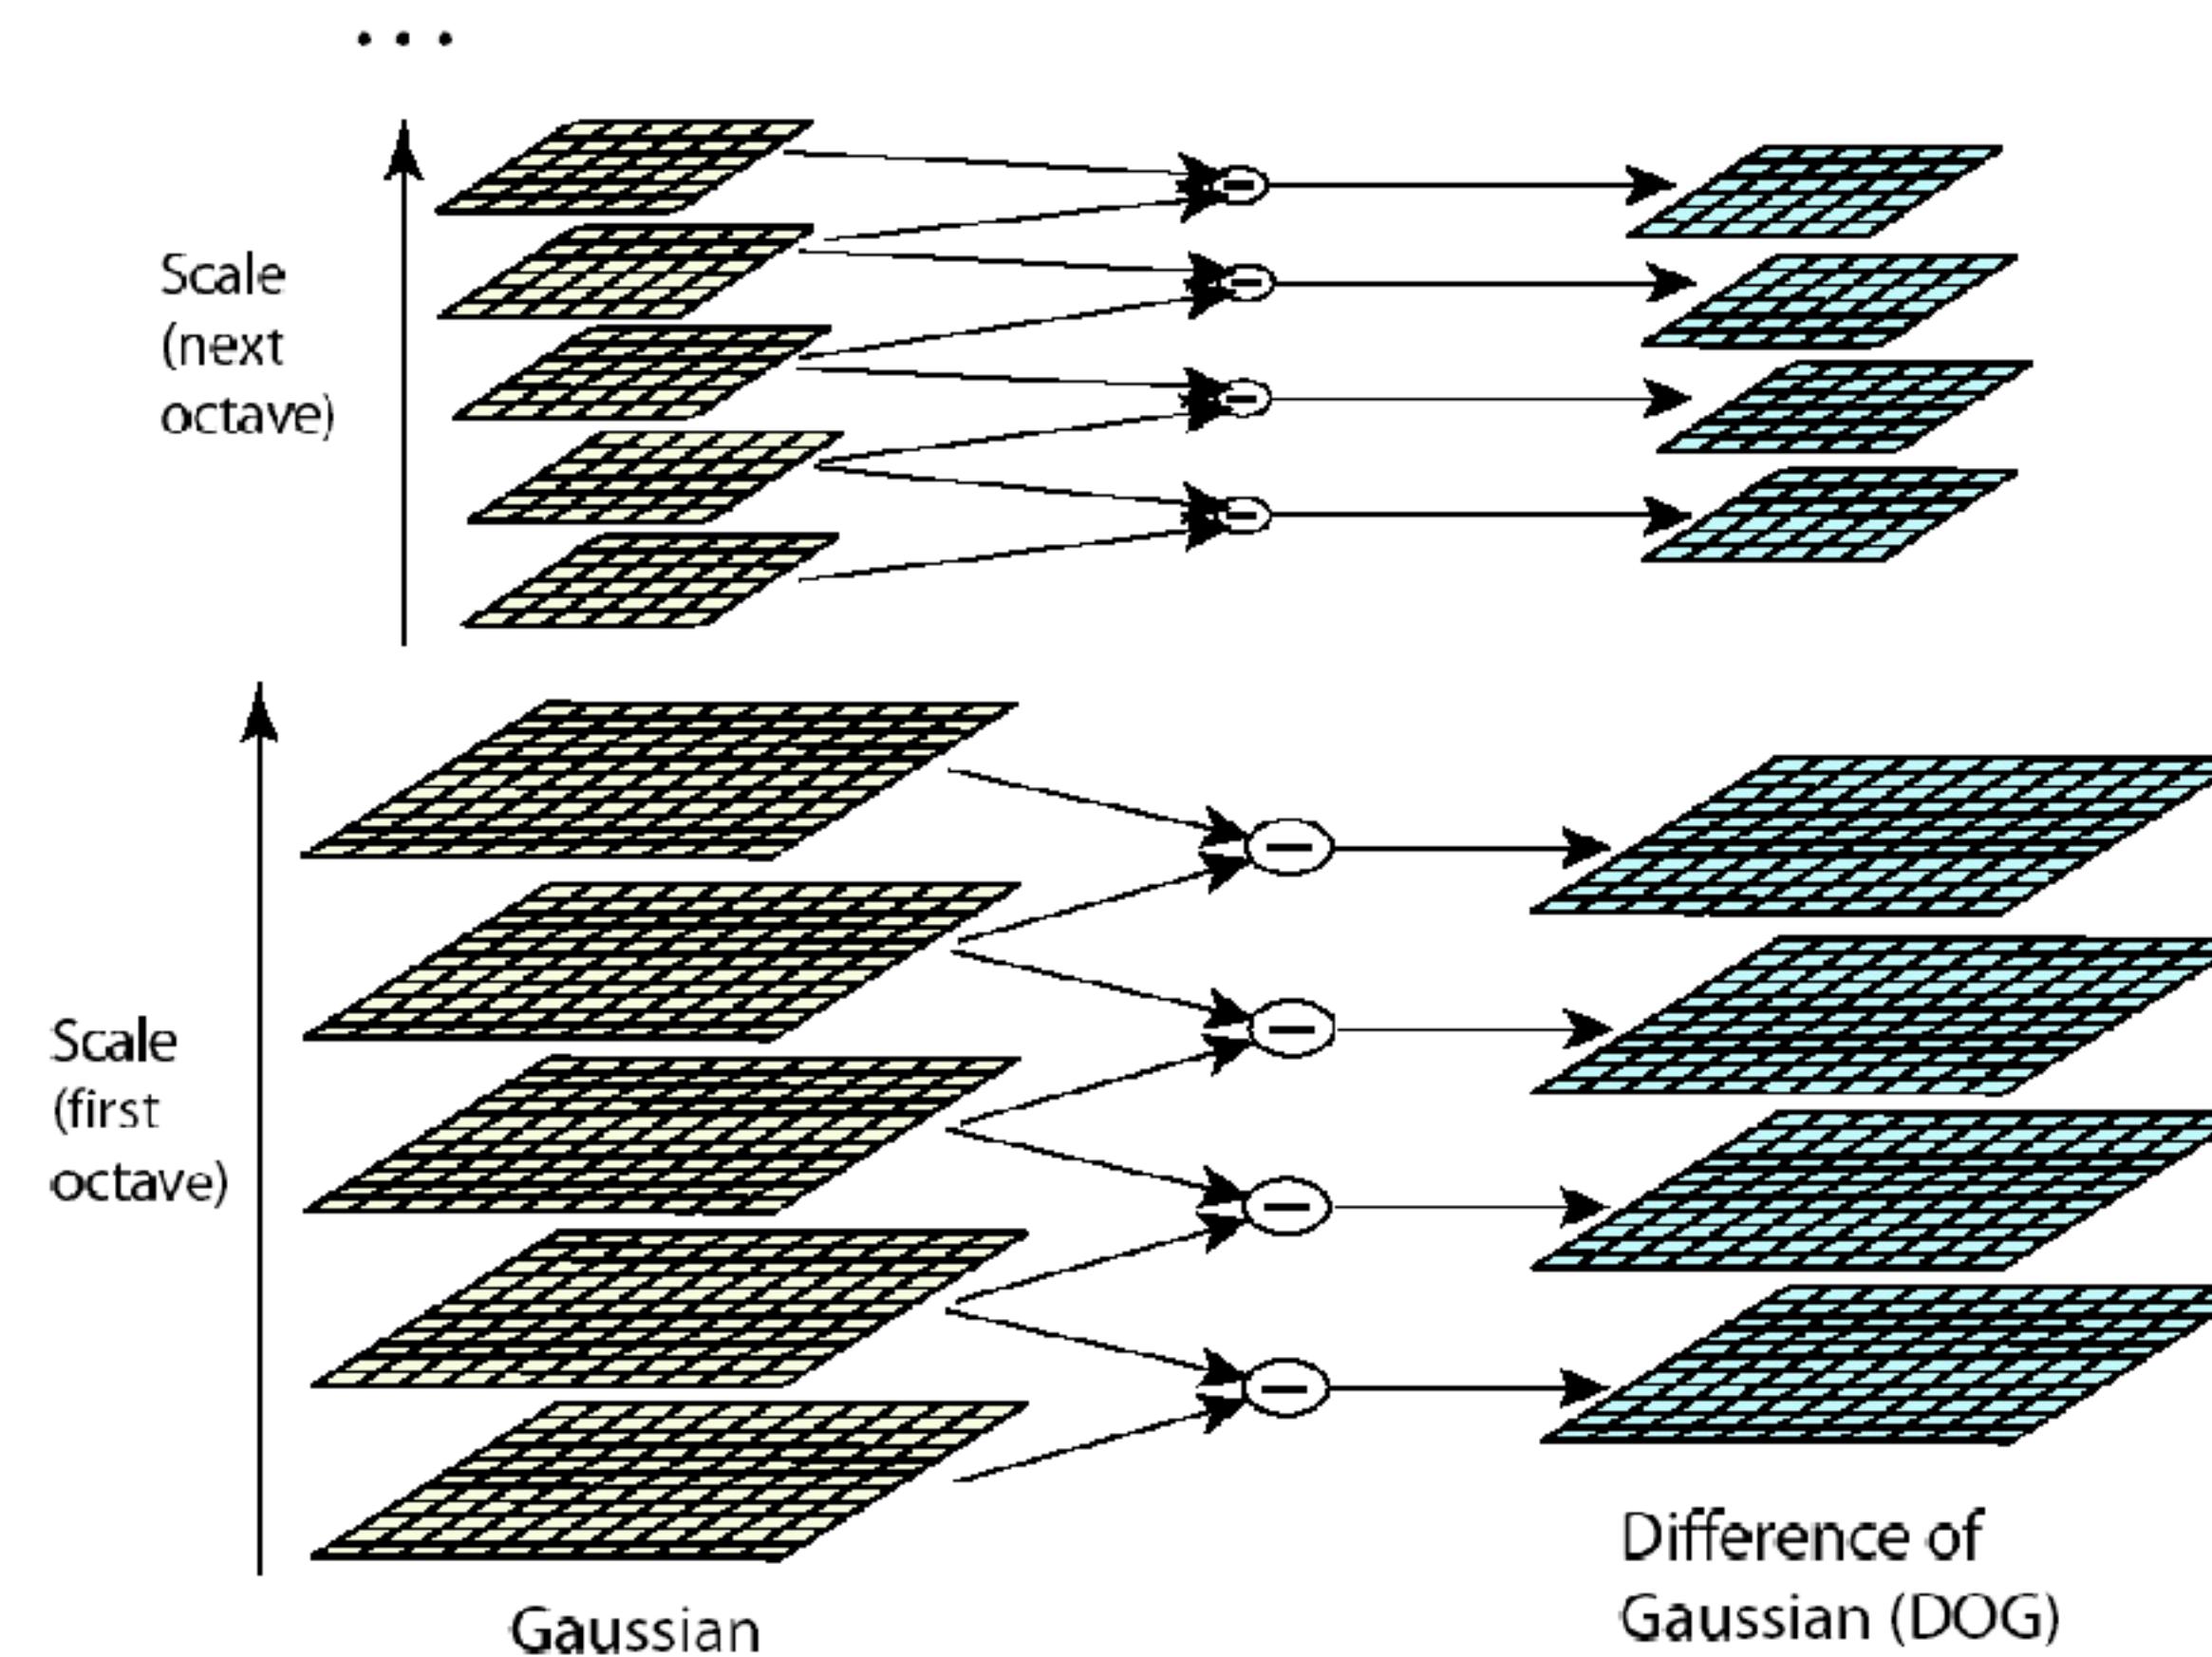
Approximating the Laplacian with a difference of Gaussians:

$$DoG = G(x, y, k\sigma) - G(x, y, \sigma)$$

(Difference of Gaussians)



Efficient implementation



David G. Lowe. [**"Distinctive image features from scale-invariant keypoints."**](#)
IJCV 60 (2), pp. 91-110, 2004.